

# OPERATING SYSTEMS



# IPC: INTER-PROCESS COMMUNICATION


- Two primary methods:
  - Shared Memory
  - Messaging



# SHARED MEMORY VS MESSAGE PASSING

## (a) Shared memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Advantage: very fast and efficient




(a)



# SHARED MEMORY VS MESSAGE PASSING

## (a) Shared memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Advantage: very fast and efficient
- Disadvantage?




(a)



# SHARED MEMORY VS MESSAGE PASSING

## (a) Shared memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Advantage: very fast and efficient
- Disadvantage?
- OS needs to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in later chapters.




(a)



# SHARED MEMORY VS MESSAGE PASSING

## (b) Message passing




(b)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(message)
  - **receive**(message)



# SHARED MEMORY VS MESSAGE PASSING

## (b) Message passing




(b)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(message)
  - **receive**(message)
- The *message* size is either fixed or variable



# SHARED MEMORY VS MESSAGE PASSING

## (b) Message passing



(b)


- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(message)`
  - `receive(message)`
- The *message* size is either fixed or variable
- Advantage: Can be easily synchronized or even used for synchronization, which will be discussed later.





# SHARED MEMORY VS MESSAGE PASSING

## (b) Message passing




(b)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(message)`
  - `receive(message)`
- The *message* size is either fixed or variable
- Advantage: Can be easily synchronized or even used for synchronization, which will be discussed later.
- Disadvantage?



# SHARED MEMORY VS MESSAGE PASSING

## (b) Message passing




(b)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(message)`
  - `receive(message)`
- The *message* size is either fixed or variable
- Advantage: Can be easily synchronized or even used for synchronization, which will be discussed later.
- Disadvantage: Requires more operations and more read/writes than shared memory.




# SHARED MEMORY VS MESSAGE PASSING

(a) Shared memory



(b) Message passing



---

# POSIX SHARED MEMORY

# SHARED MEMORY EXAMPLES: POSIX SHARED MEMORY

- POSIX: Portable Operating System Interface
- Standards specified by the IEEE Computer Society
- Objective: OS have unified API allowing portability of software.



# POSIX SHARED MEMORY

- POSIX: Portable Operating System Interface
- Standards specified by the IEEE Computer Society
- Objective: OS have unified API allowing portability of software.

Most Linux systems are partially or fully compliant to the POSIX standards making it fairly easy to port code.



# POSIX SHARED MEMORY

- POSIX Shared Memory

- Process first creates shared memory segment

- ```
shm_fd = shm_open(name, O_CREAT);
```

- Also used to open an existing segment.

- Set the size of the object

- ```
ftruncate(shm_fd, 4096);
```

- Map shared memory object to the process's address space

- ```
void* addr = mmap(NULL, 1024, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);
```



# IPC - MESSAGE IMPLEMENTATIONS

- Pipes
- Sockets
- Local/Remote Procedure Calls





# PIPES

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)



# PIPES

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes
- Windows calls these **anonymous pipes**



# BLOCKING VS UNBLOCKING

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** -- the sender sends the message and continue
  - **Non-blocking receive** -- the receiver receives:
    - A valid message, or
    - Null message



# BLOCKING VS UNBLOCKING

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available

Blocked send/receive can result in indefinite wait or even infinite wait ...  
This can freeze the whole program ...

Worksheet Q1: What would be possible solution to this problem?



# BLOCKING VS UNBLOCKING

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available

Blocked send/receive can result in indefinite wait or even infinite wait ...

This can freeze the whole program ...

Worksheet Q1: What would be possible solution to this problem?

- Timeout: set max wait time and return with “null” or error.



# BLOCKING VS UNBLOCKING

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available

Blocked send/receive can result in indefinite wait or even infinite wait ...  
This can freeze the whole program ...

Worksheet Q1: What would be possible solution to this problem?

- Timeout: set max wait time and return with “null” or error.
- Interrupt blocked read/write thread and wake it up.



# BLOCKING VS UNBLOCKING

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available

Blocked send/receive can result in indefinite wait or even infinite wait ...  
This can freeze the whole program ...

Worksheet Q1: What would be possible solution to this problem?

- Timeout: set max wait time and return with “null” or error.
- Interrupt blocked read/write thread and wake it up.
- Create a worker thread just for the read/write. The program can continue running in main thread.



# PIPES IN UNIX

- The vertical bar | is the pipe operator in unix shell.
- The transferred data is never saved in a file, it is simply communicated to the other process.
- Unnamed or “ordinary” pipes are destroyed after the process completes execution.

**Syntax :**

```
command_1 | command_2 | command_3 | .... | command_N
```

```
$ ls -l | more
```





# PIPES IN UNIX

- The vertical bar | is the pipe operator in unix shell.
- The transferred data is never saved in a file, it is simply communicated to the other process.
- Unnamed or “ordinary” pipes are destroyed after the process completes execution.
- Named pipes can be created by the `mknod()` system call with the ‘FIFO’ option:  
`mknod ("mypipe", SIFIFO, 0)`
- You can also use `mkfifo("name", 0666)`
- Named pipes allow communication between any two processes

## Syntax :

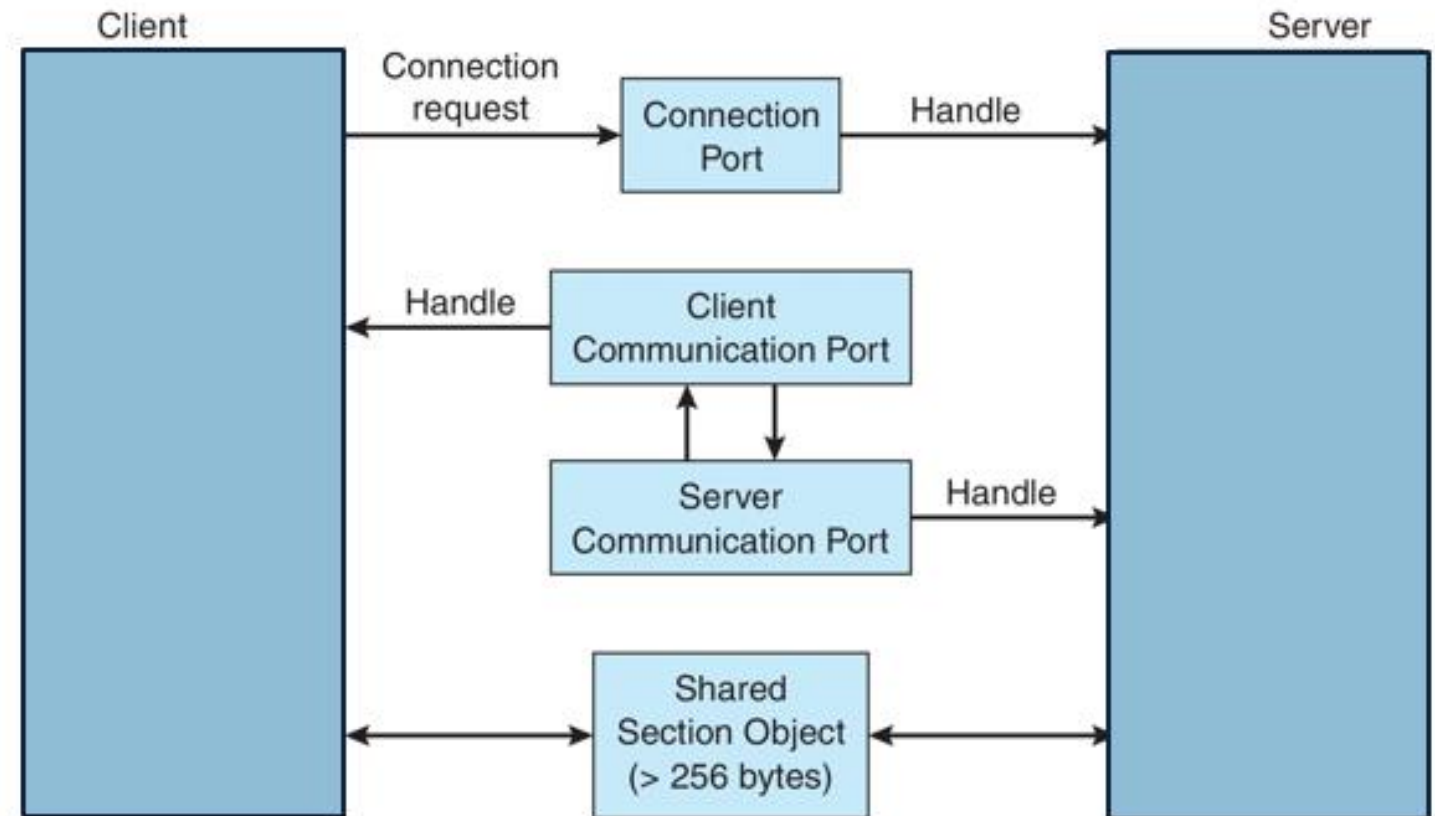
```
command_1 | command_2 | command_3 | .... | command_N
```

```
$ ls -l | more
```



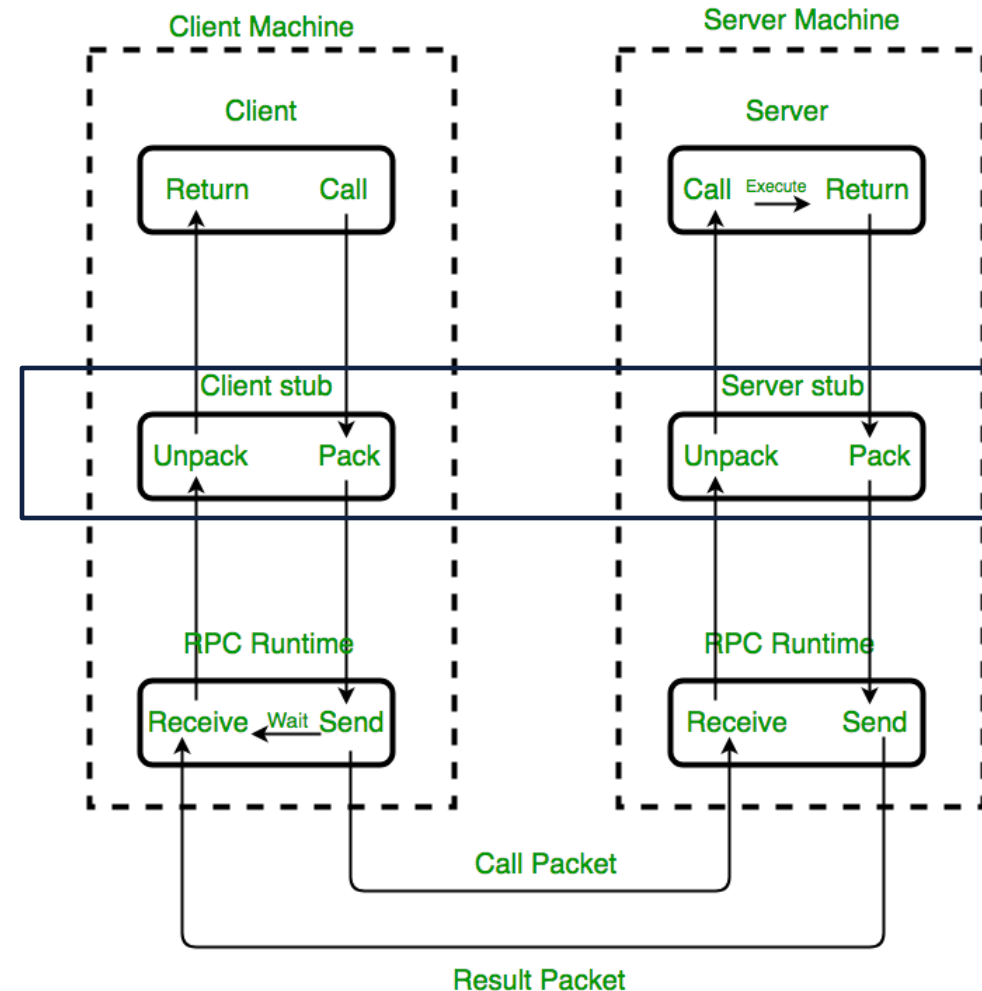
# LOCAL PROCEDURE CALLS IN WINDOWS

- Message-passing centric via **advanced local procedure call (LPC)** facility
  - Only works between processes on the same system
  - Uses ports to establish and maintain communication channels



# REMOTE PROCEDURE CALLS

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
  - Again uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and **marshalls** (packs) the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server



Implementation of RPC mechanism

# SOCKETS

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running

