OPERATING SYSTEMS INTERNALS

CSCI 509

PAGE SHARING

- A lot of code is shared by many processed in a system.
- Example: Libraries, routines ..

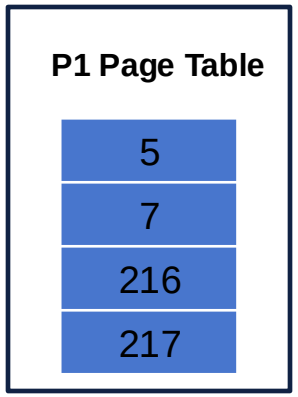

Using OpenGL: Open Graphics Library

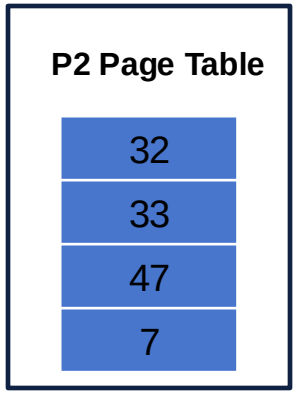

Using OpenGL

PAGE SHARING

- A lot of code is shared by many processed in a system.
- Example: Libraries, routines ..

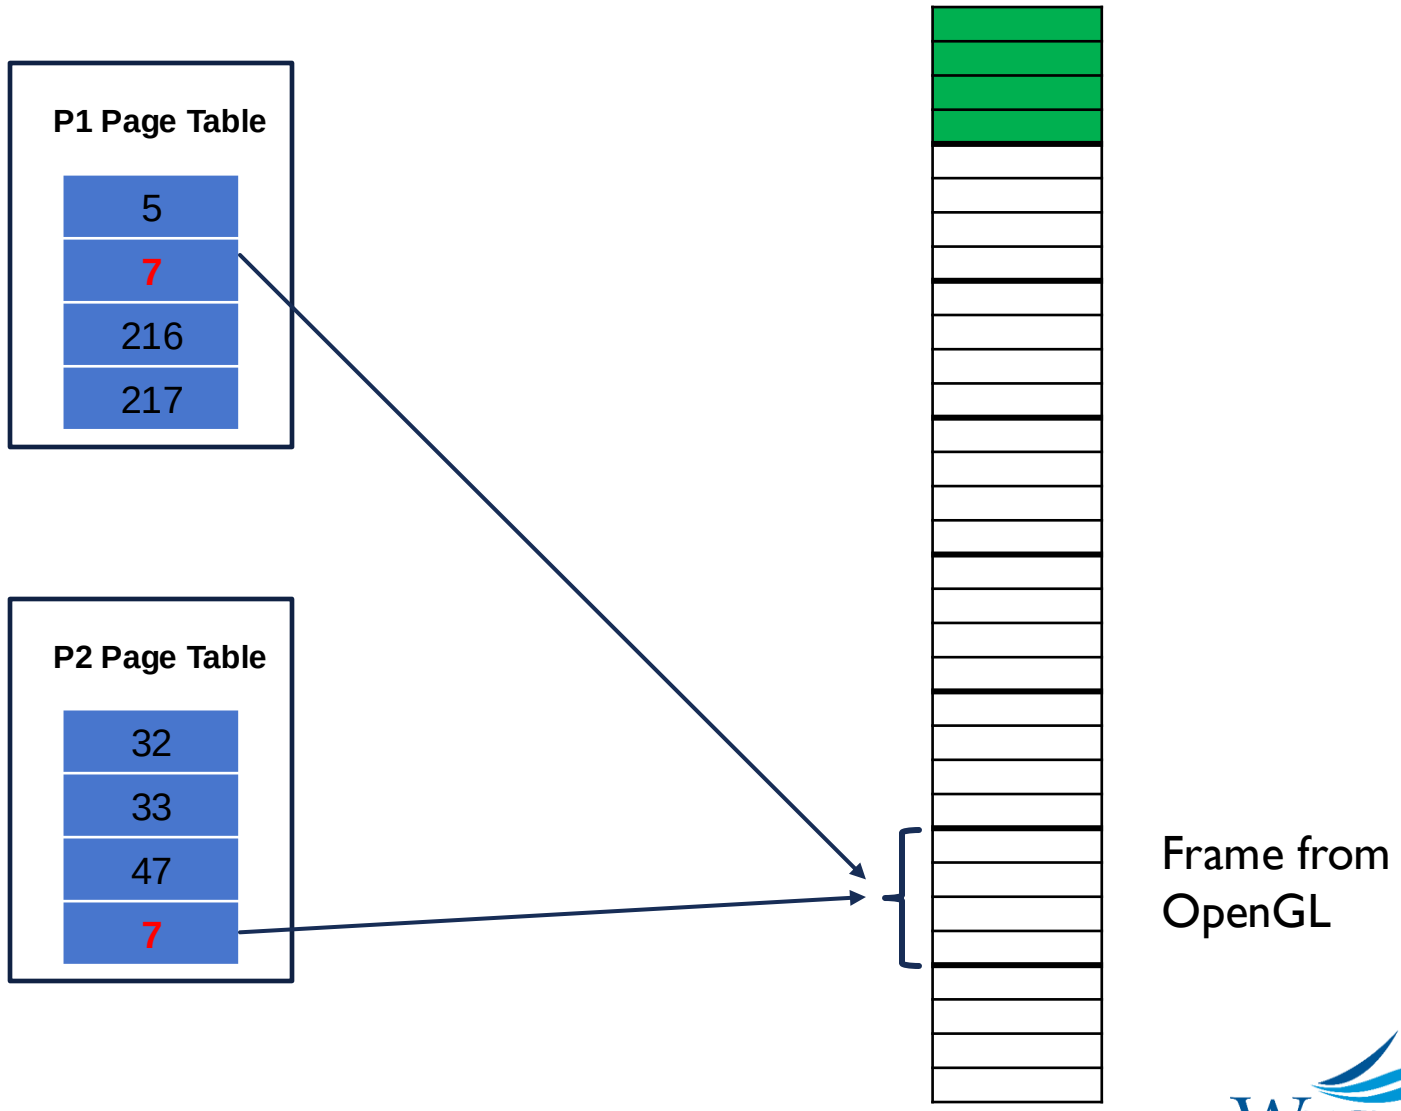

PAGE SHARING

- A lot of code is shared by many processed in a system.
- Example: Libraries, routines ..
- Read-Only: Enforced by OS.

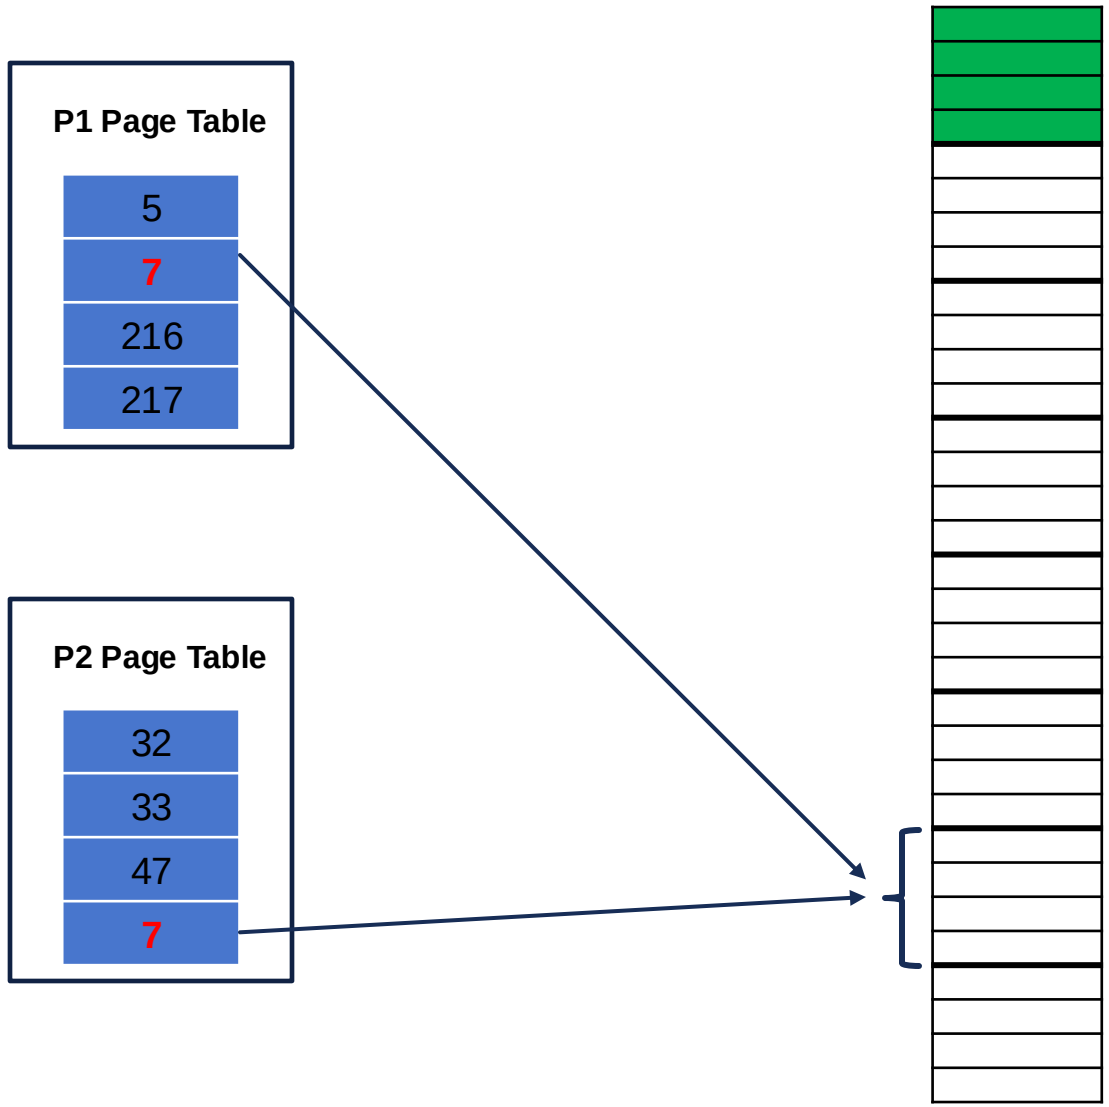

PAGE SHARING

- A lot of code is shared by many processed in a system.
- Example: Libraries, routines ..
- Read-Only: Enforced by OS.
- This is different that shared memory segments and IPC which support read/write on shared memory.

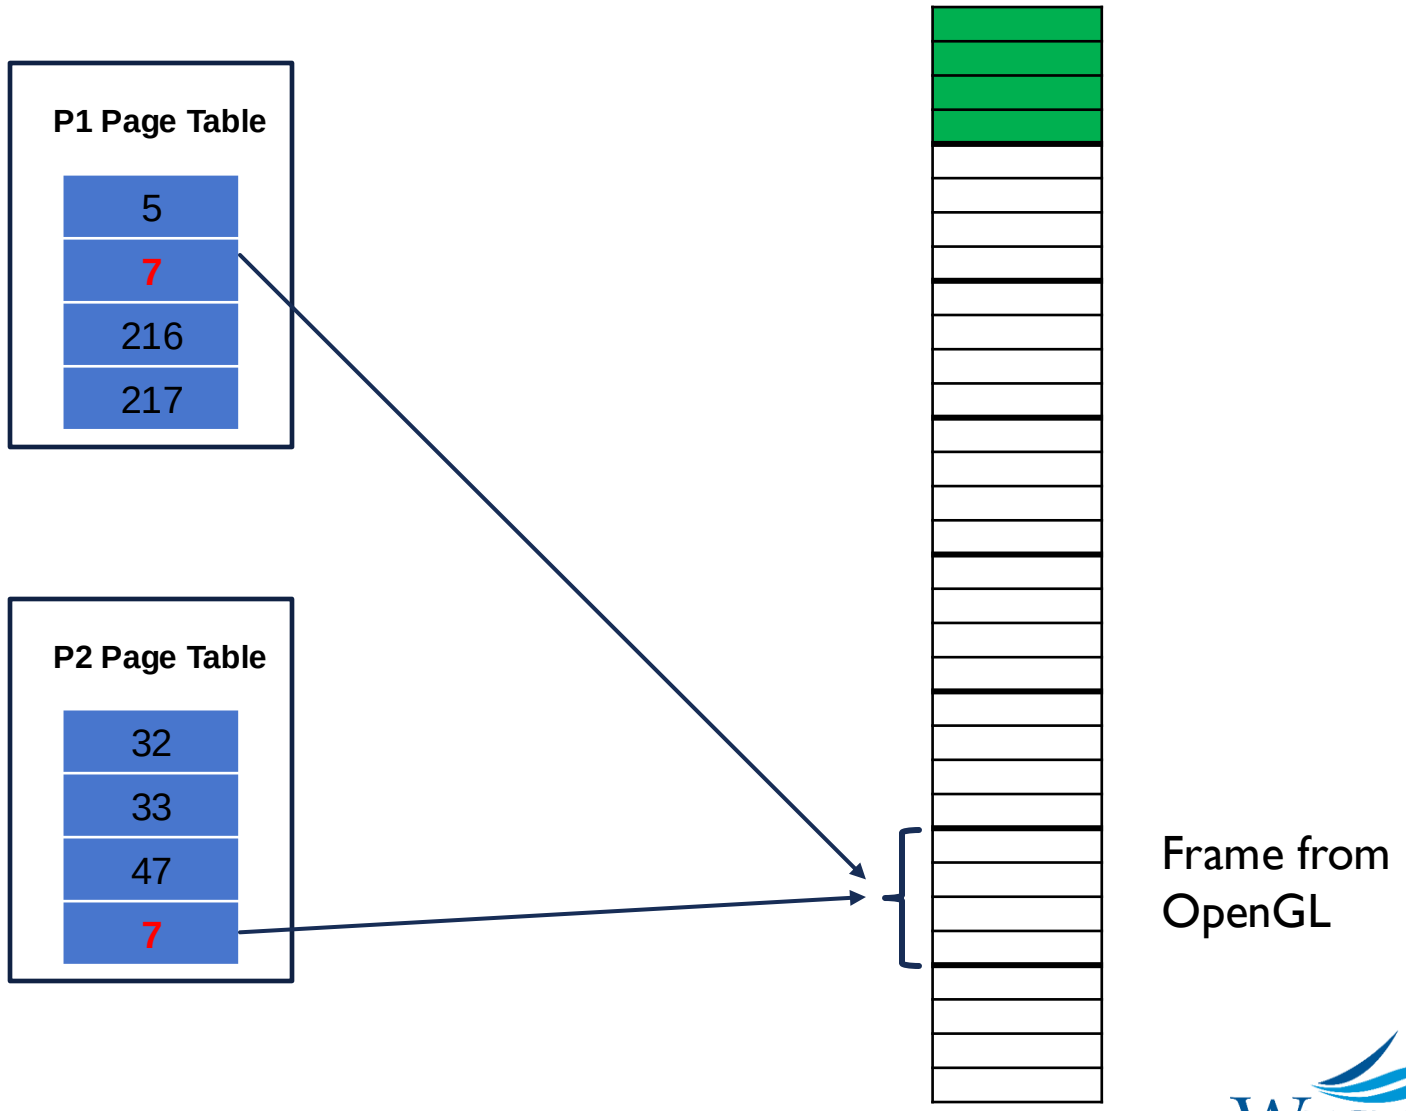

MULTI-LEVEL PAGING

- Assume 32-bit machine.
- If page size is 1 KB, how large would the page table be?
- Size of table: 2^{22}= 4MB for 1 byte per entry.
- For 4 bytes for every entry, that's 16 MB.

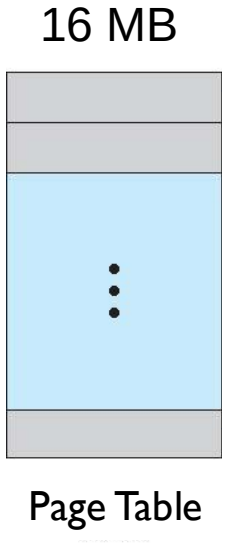

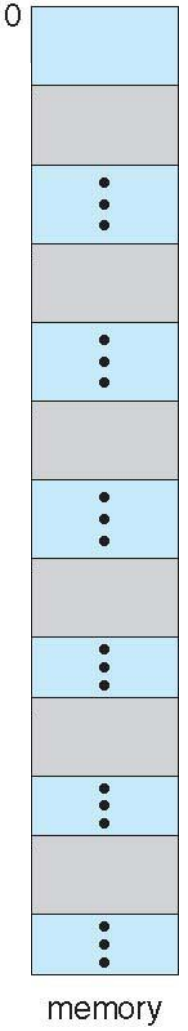

MULTI-LEVEL PAGING

- Assume 32-bit machine.
- If page size is 1 KB, how large would the page table be?
- Size of table: 2^{22}= 4MB for 1 byte per entry.
- For 4 bytes for every entry, that's 16 MB.
- TLB Speeds things up but the page table still needs to be accessed frequently.

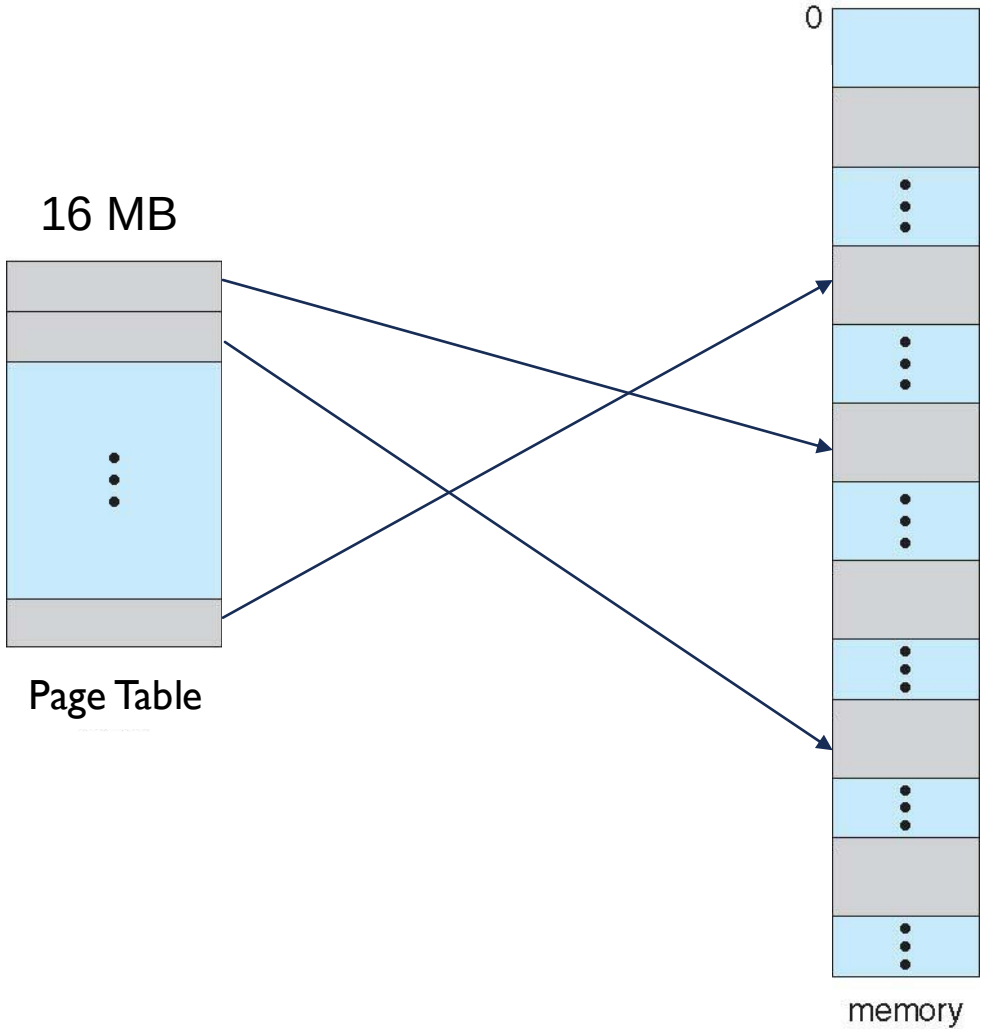

MULTI-LEVEL PAGING

- Solution: Paging the page table.

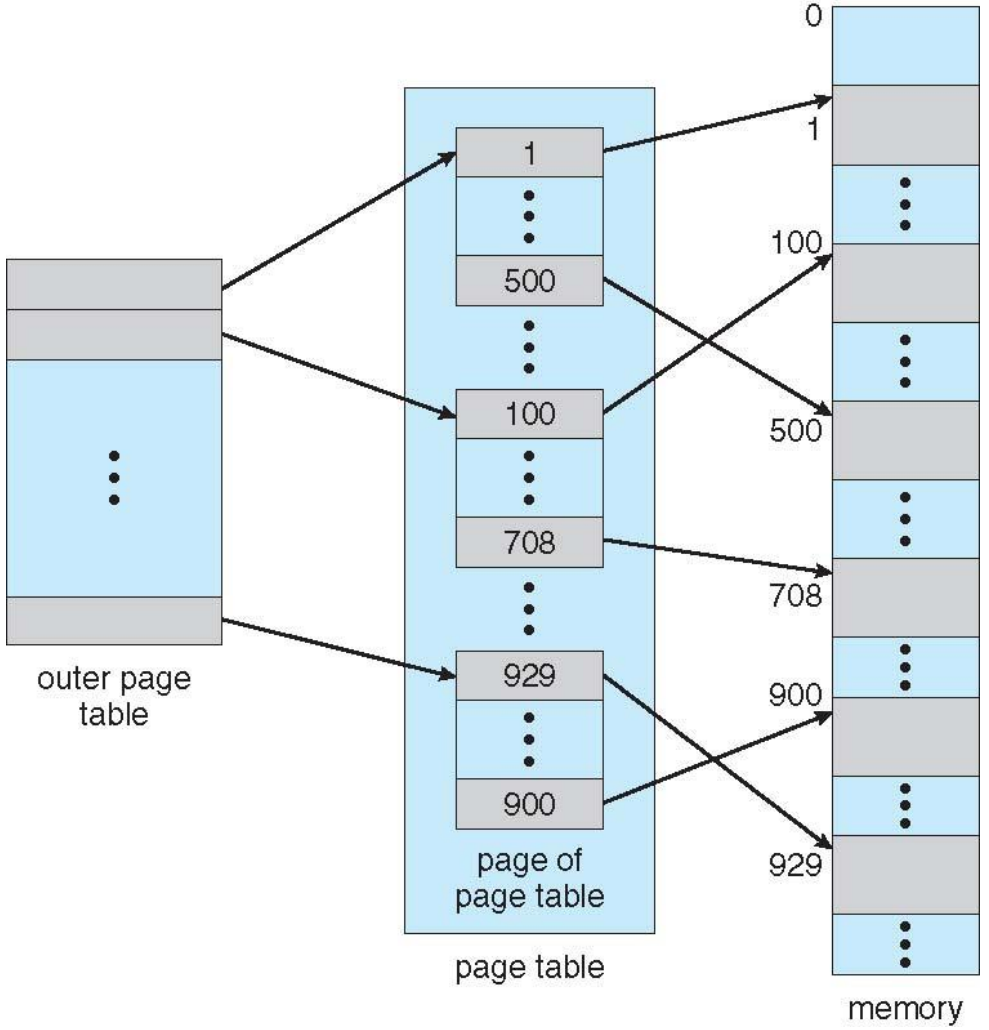

MULTI-LEVEL PAGING

- Solution: Paging the page table.
- The virtual address is now split into outer page, inner page and page offset.

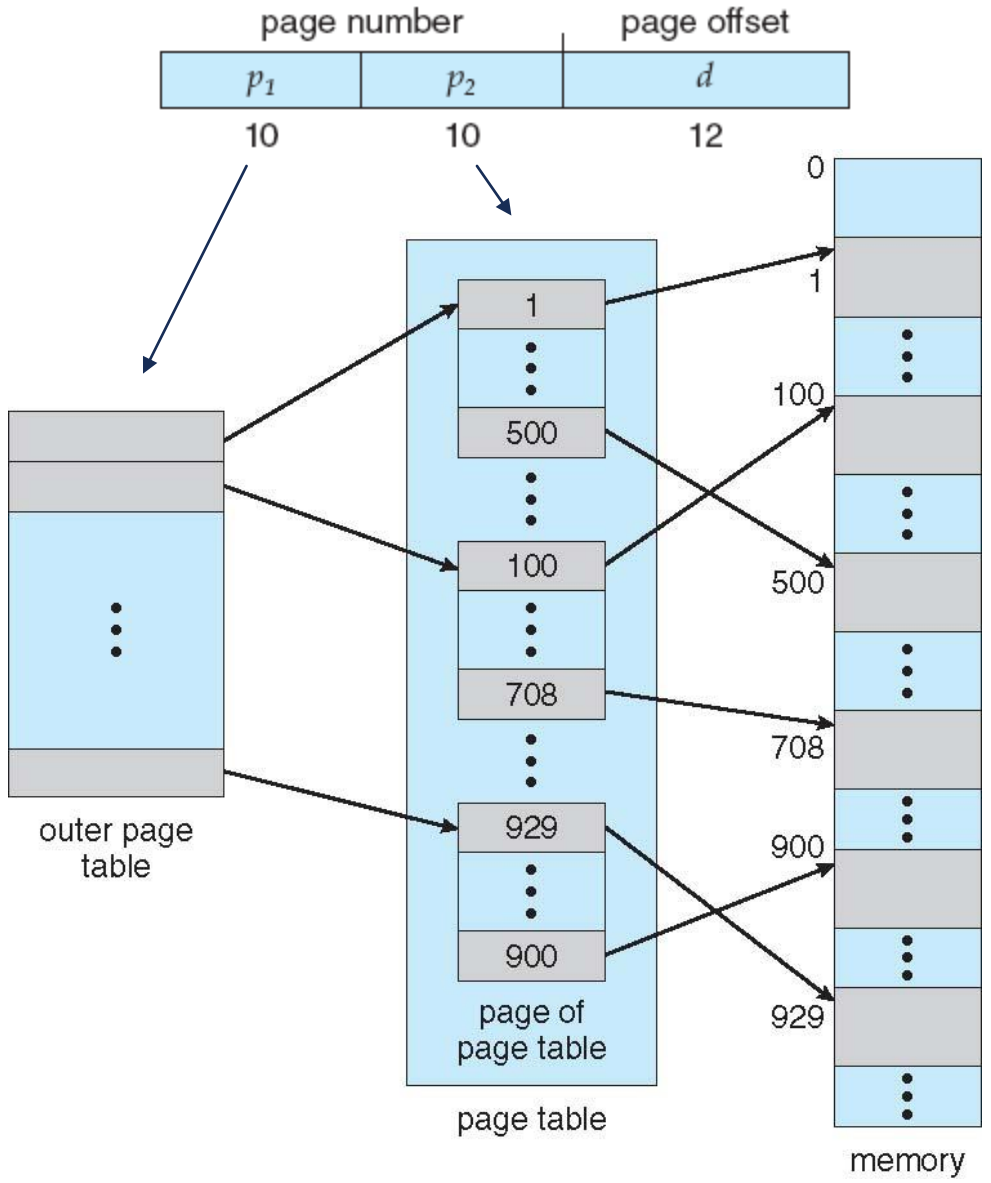

MULTI-LEVEL PAGING

- Solution: Paging the page table.
- The virtual address is now split into outer page, inner page and page offset.

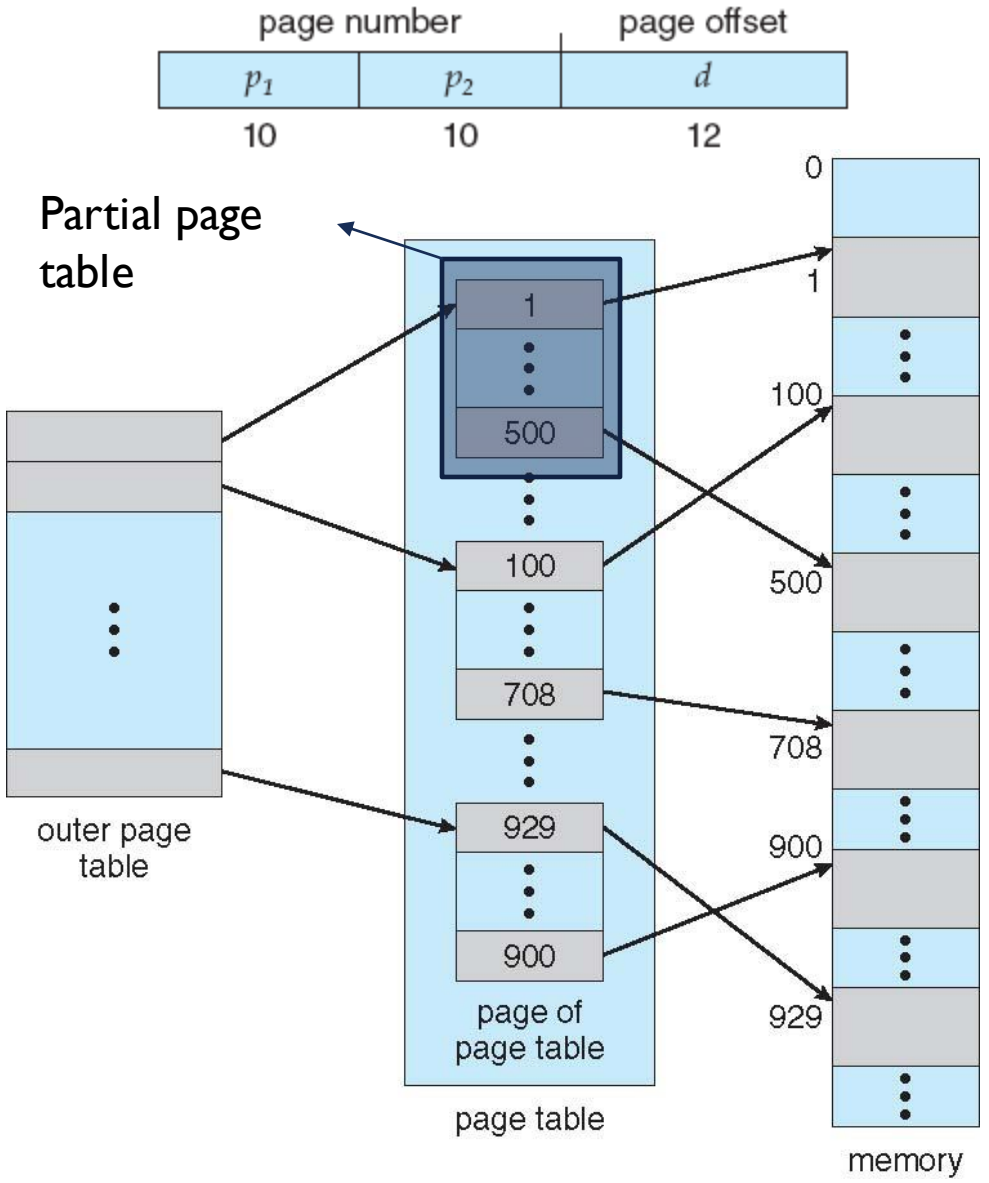

MULTI-LEVEL PAGING

00100011100011011111101101111010

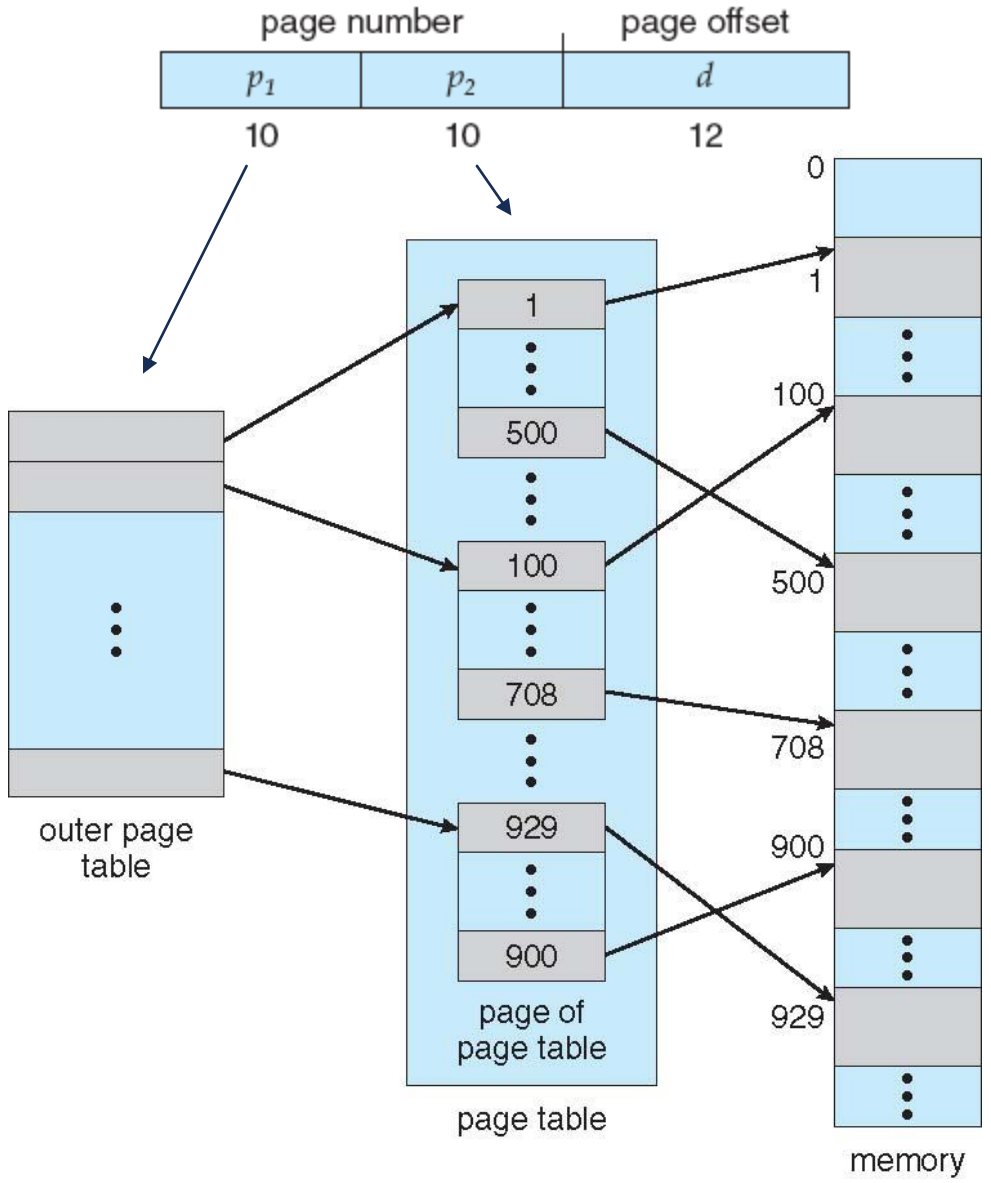

MULTI-LEVEL PAGING

P1	P2	D
0010001110	0011011111	1101111010

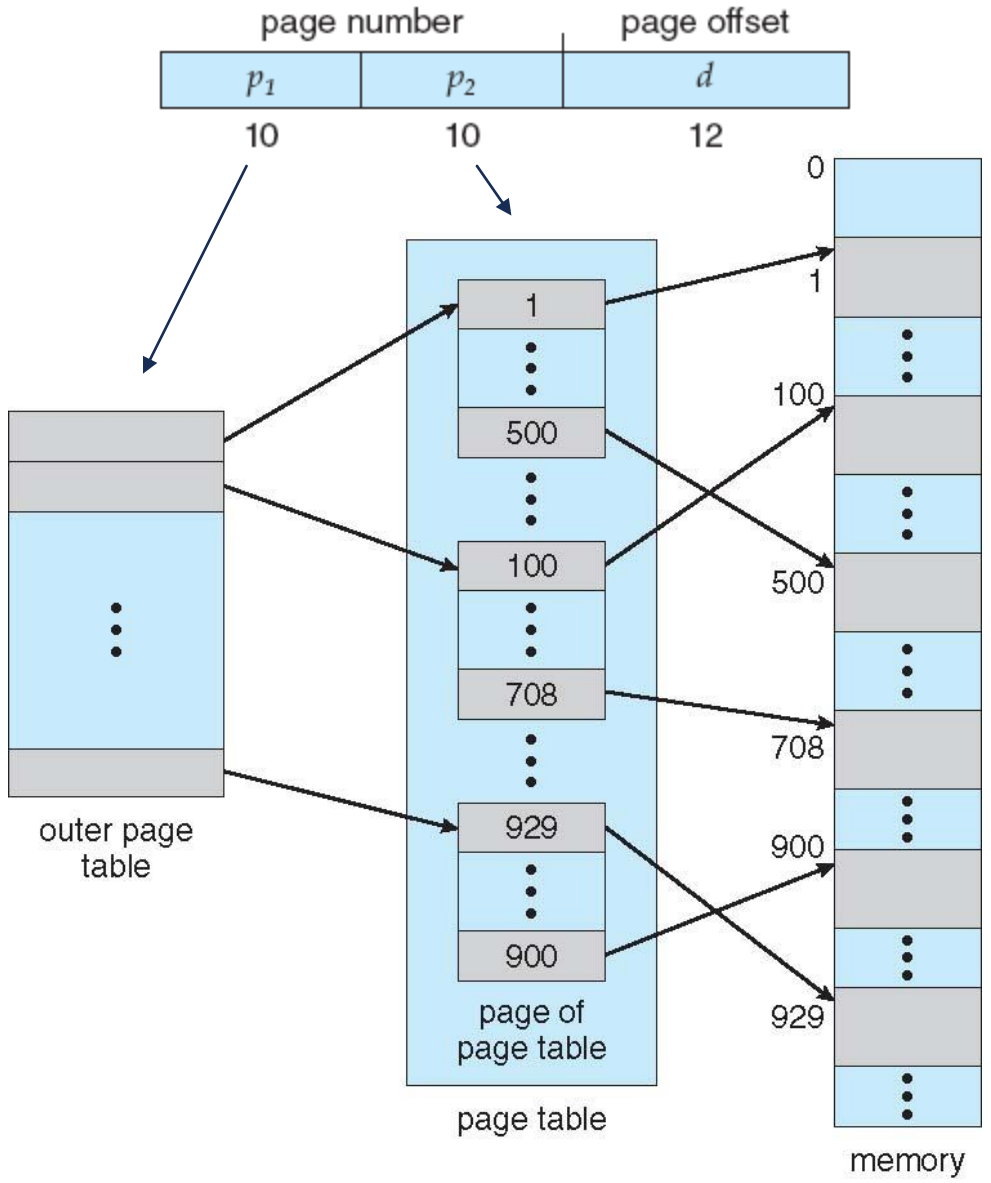

MULTI-LEVEL PAGING

P1	P2	D
0010001110	0011011111	11101111010
142	447	

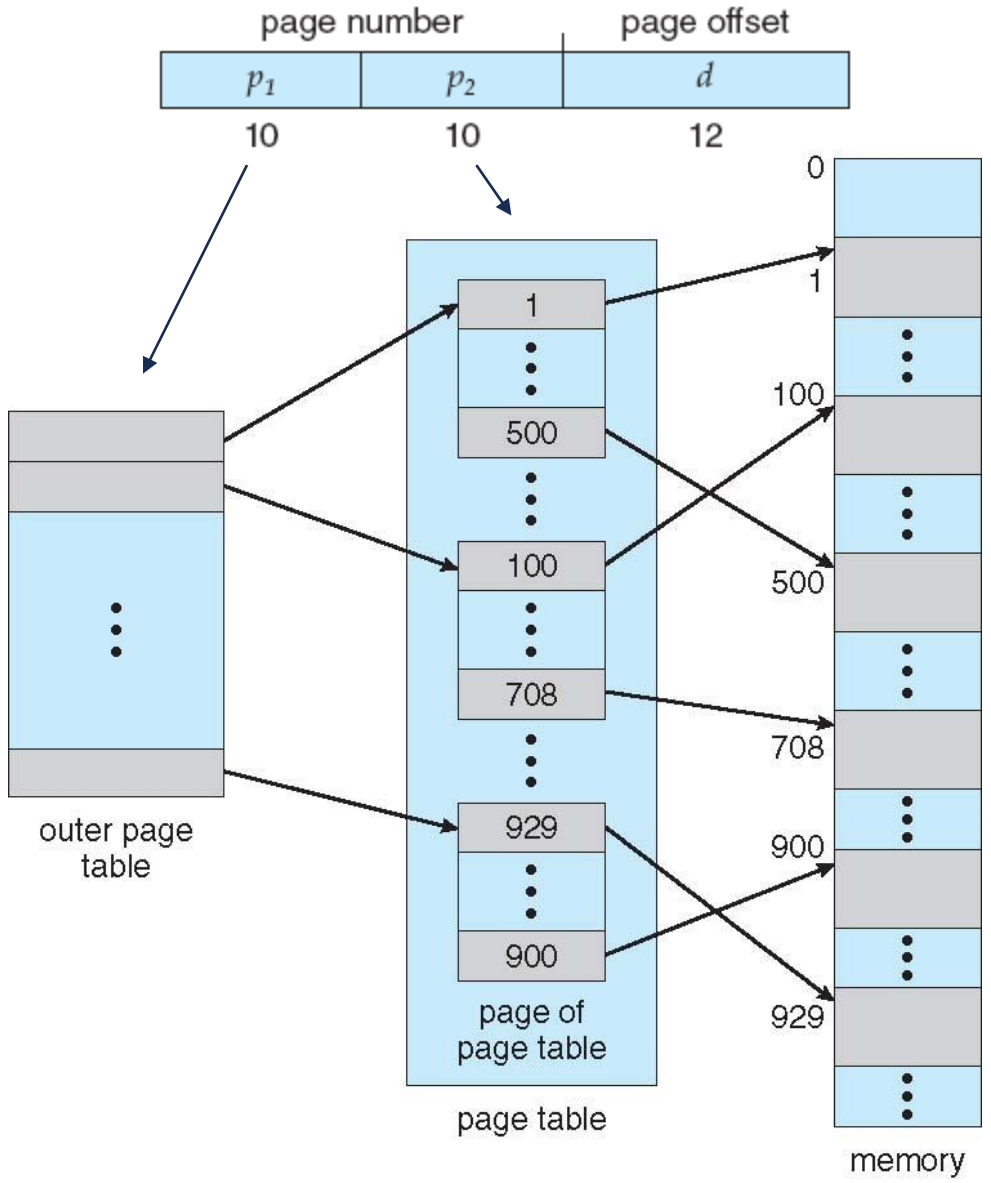

MULTI-LEVEL PAGING

P1	P2	D
0010001110	0011011111	1101111010
142	447	

Use outer page table to find address of the 142nd inner page table

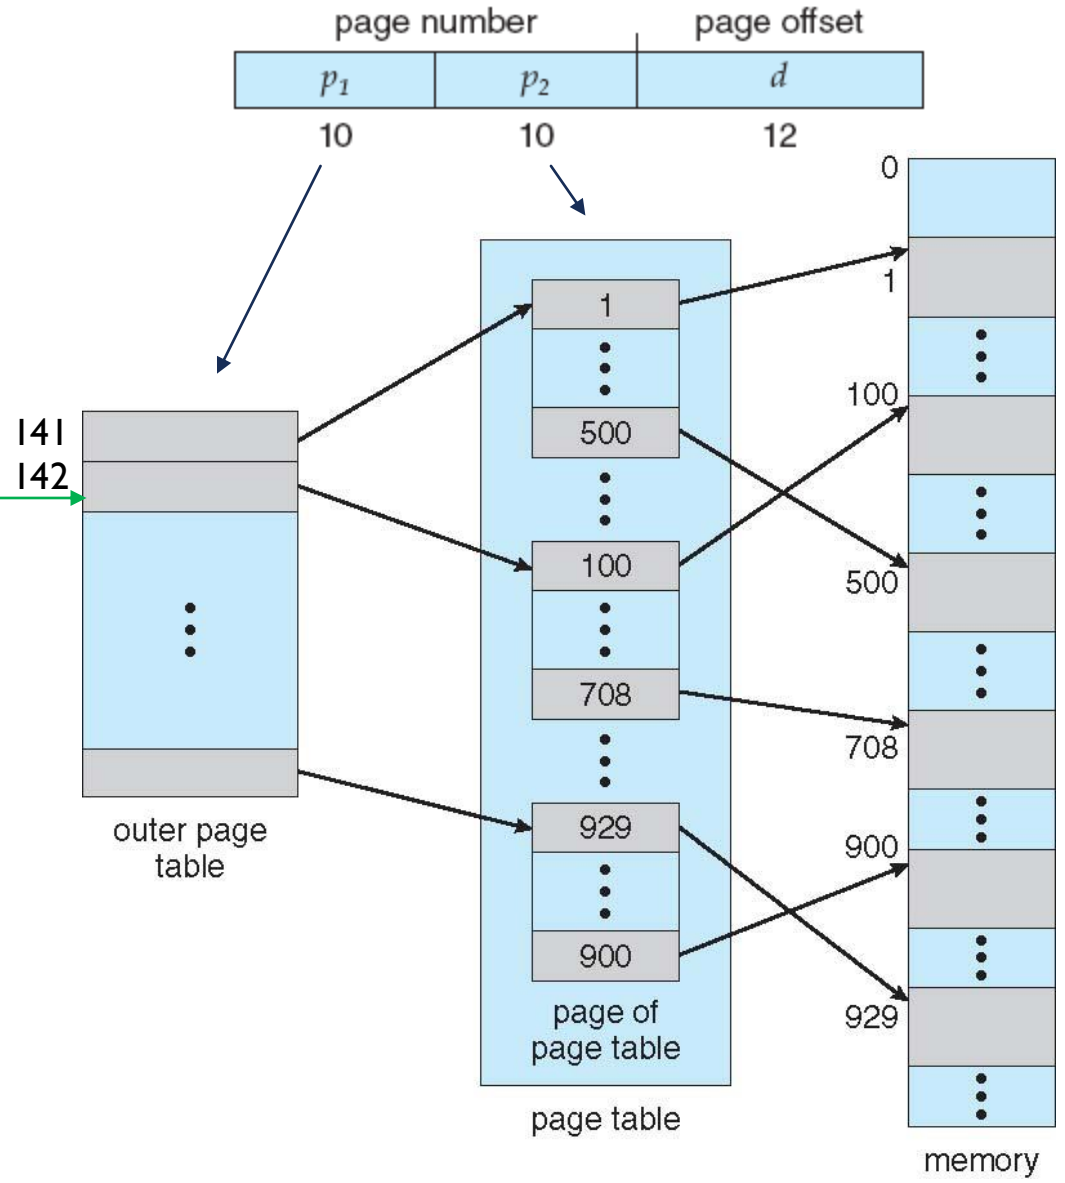

MULTI-LEVEL PAGING

P1 P2 D

0010001110 0011011111 1101111010

142 447

Use inner page table to find address of page 447

Use outer page table to find address of the 142nd inner page table

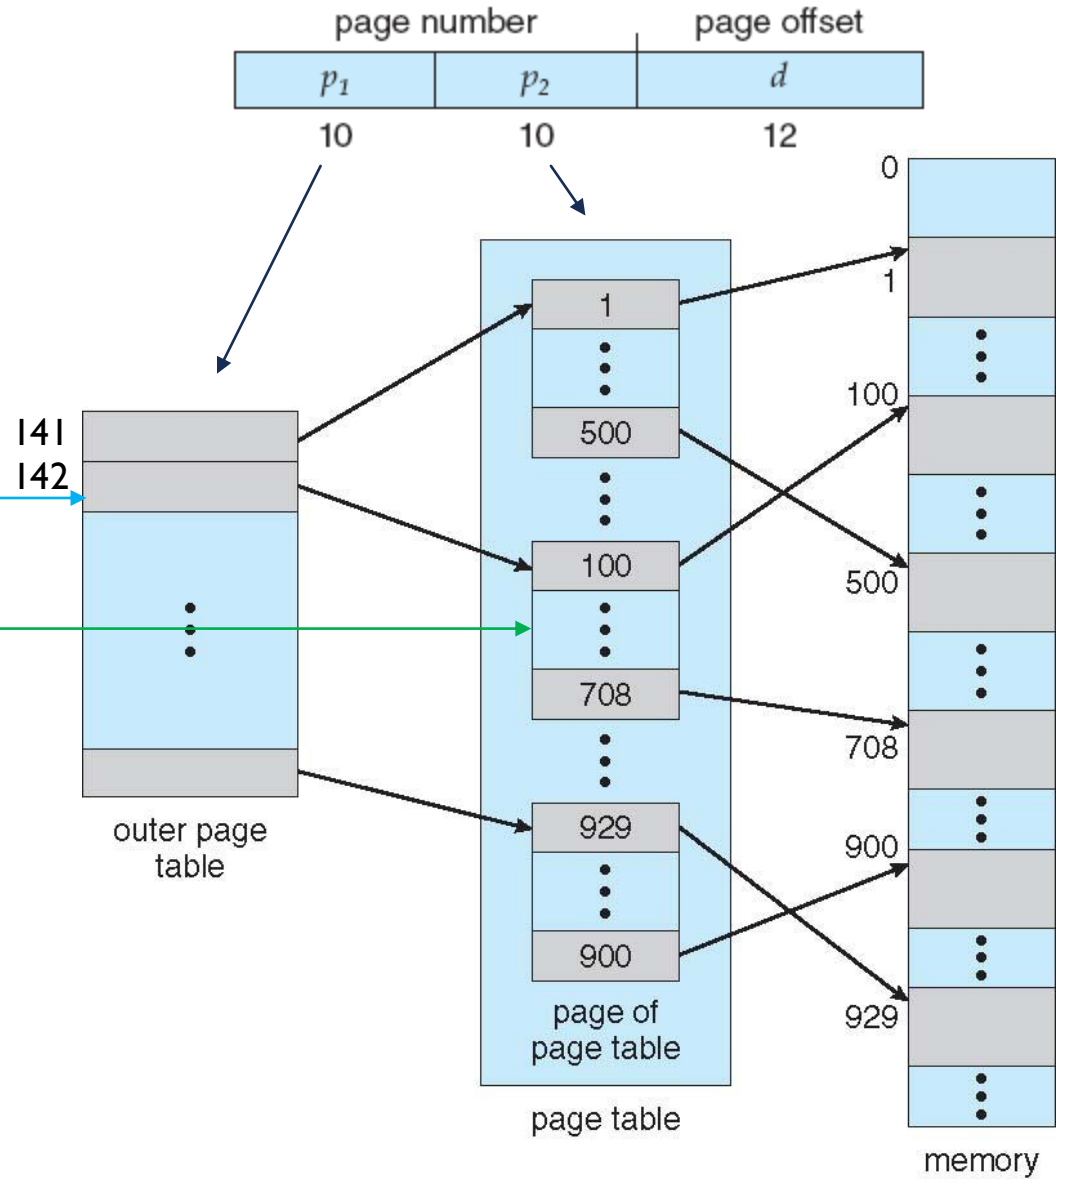

MULTI-LEVEL PAGING

P1 P2 D

00100011100011011111101101111010

142 447

Use inner page table to find address of page 447

Use outer page table to find address of the 142nd inner page table

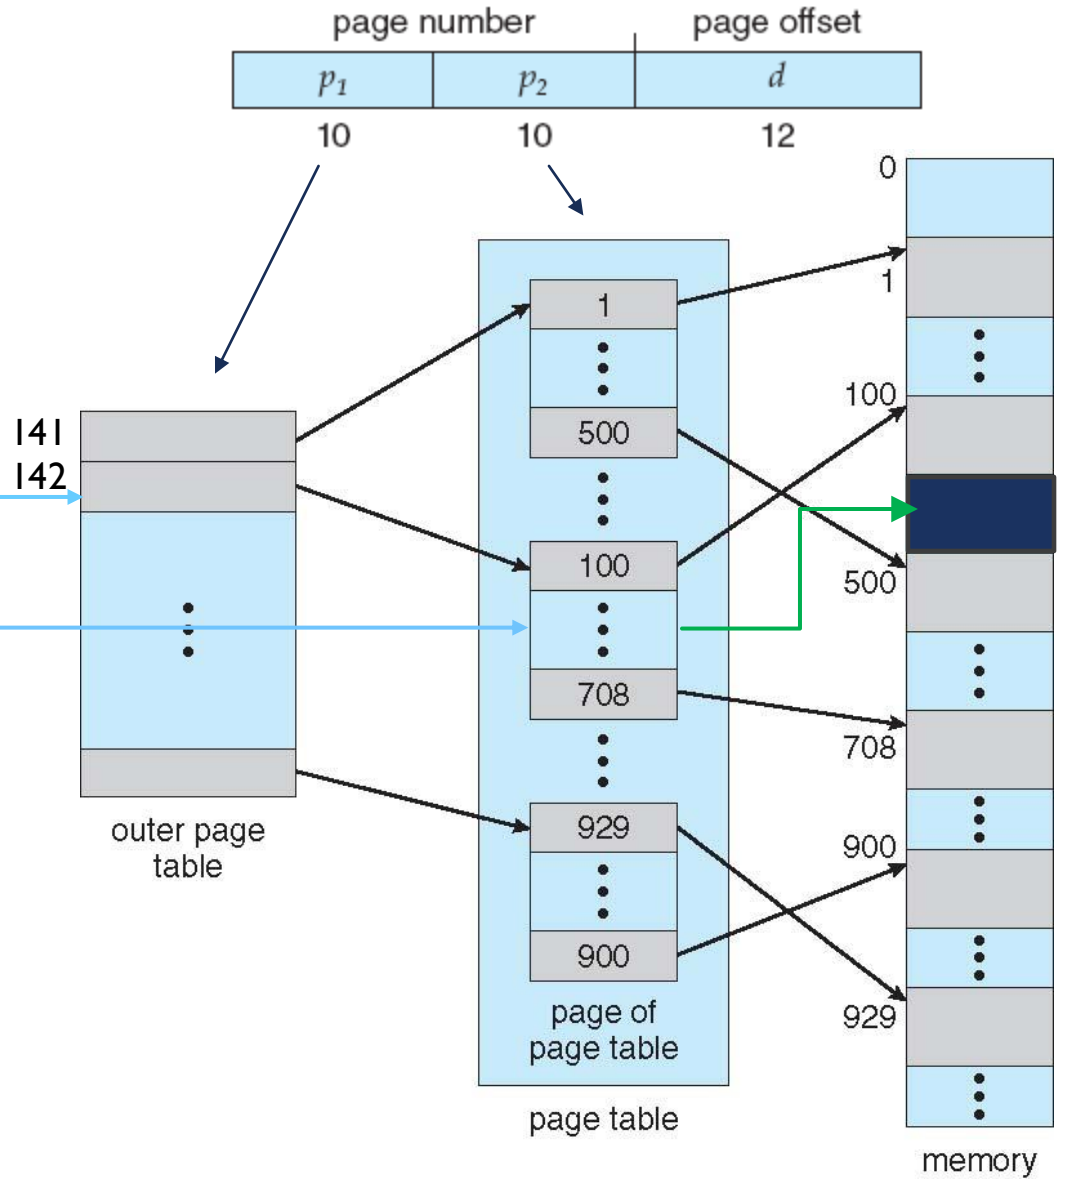

Retrieve the page

MULTI-LEVEL PAGING

2-level paging

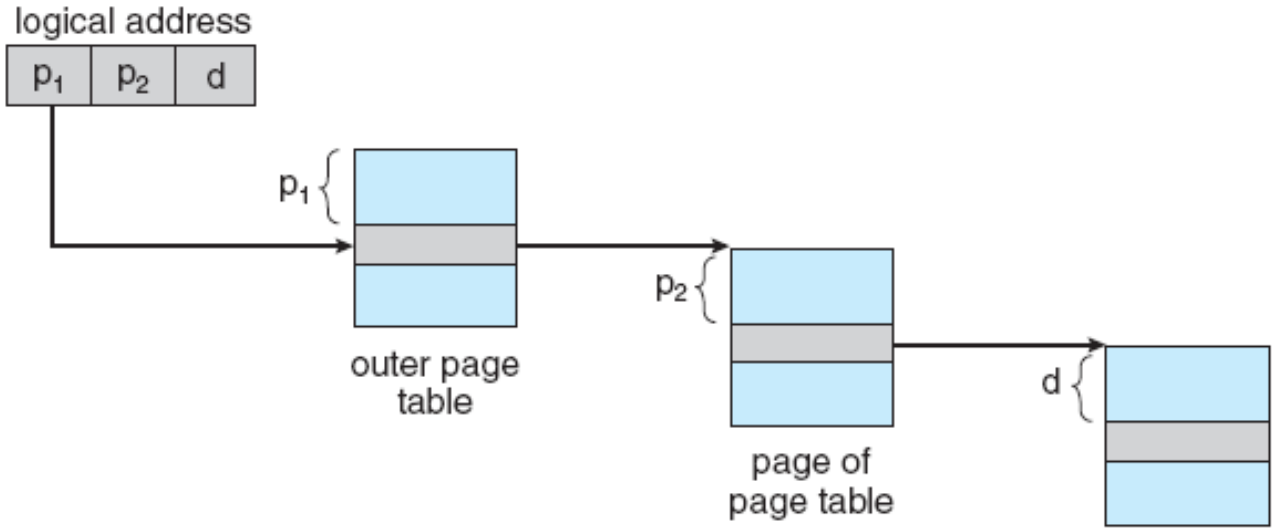

We can have 3 or even more levels of paging …

3-level paging

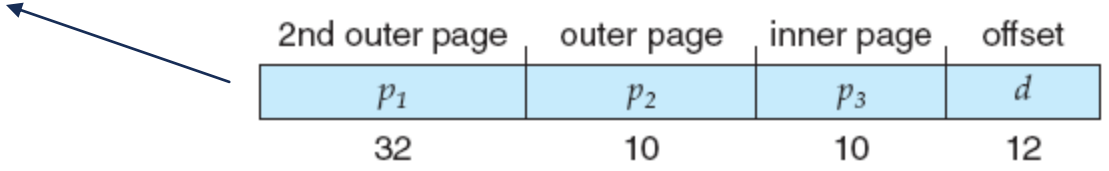

HASHED PAGED TABLE

- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element

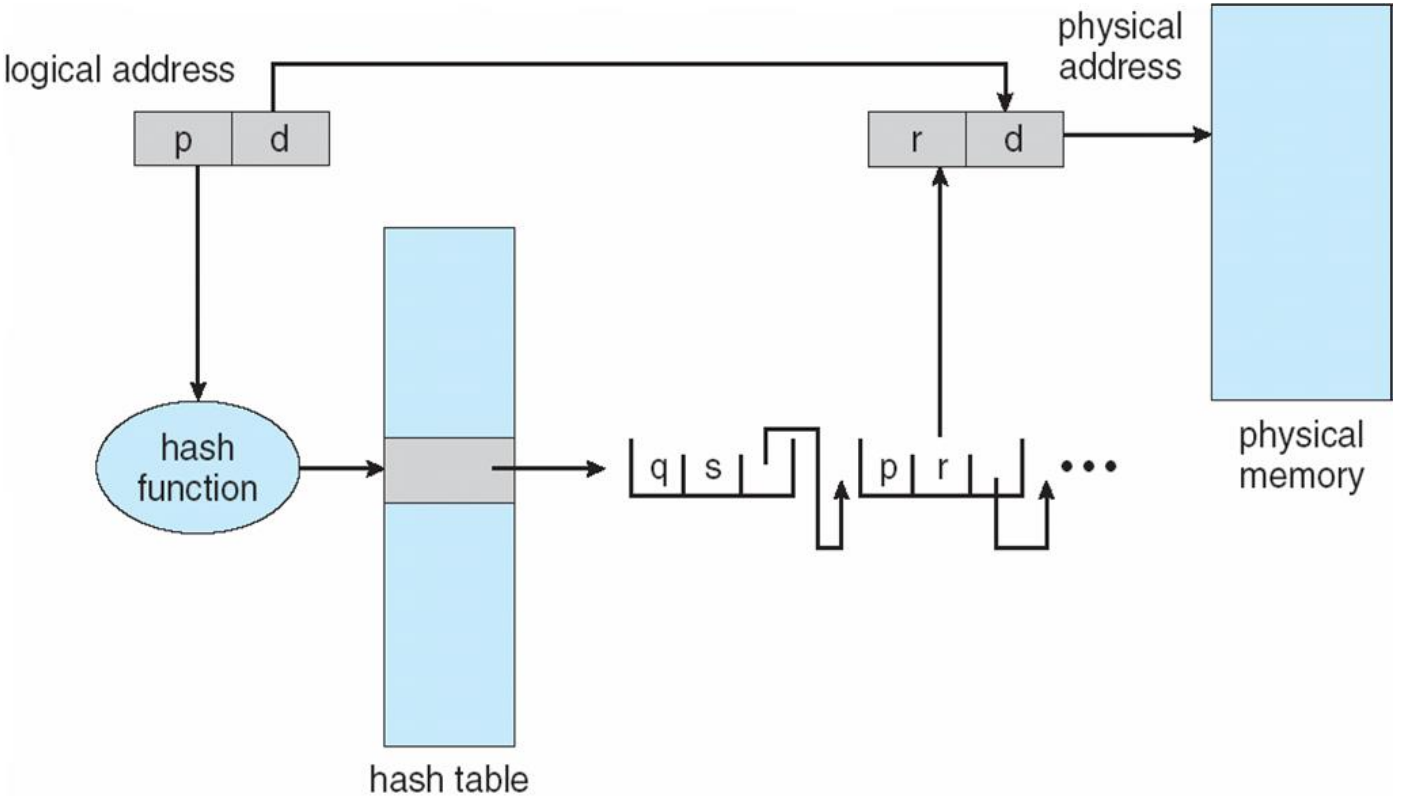

SINGLE PAGE TABLE

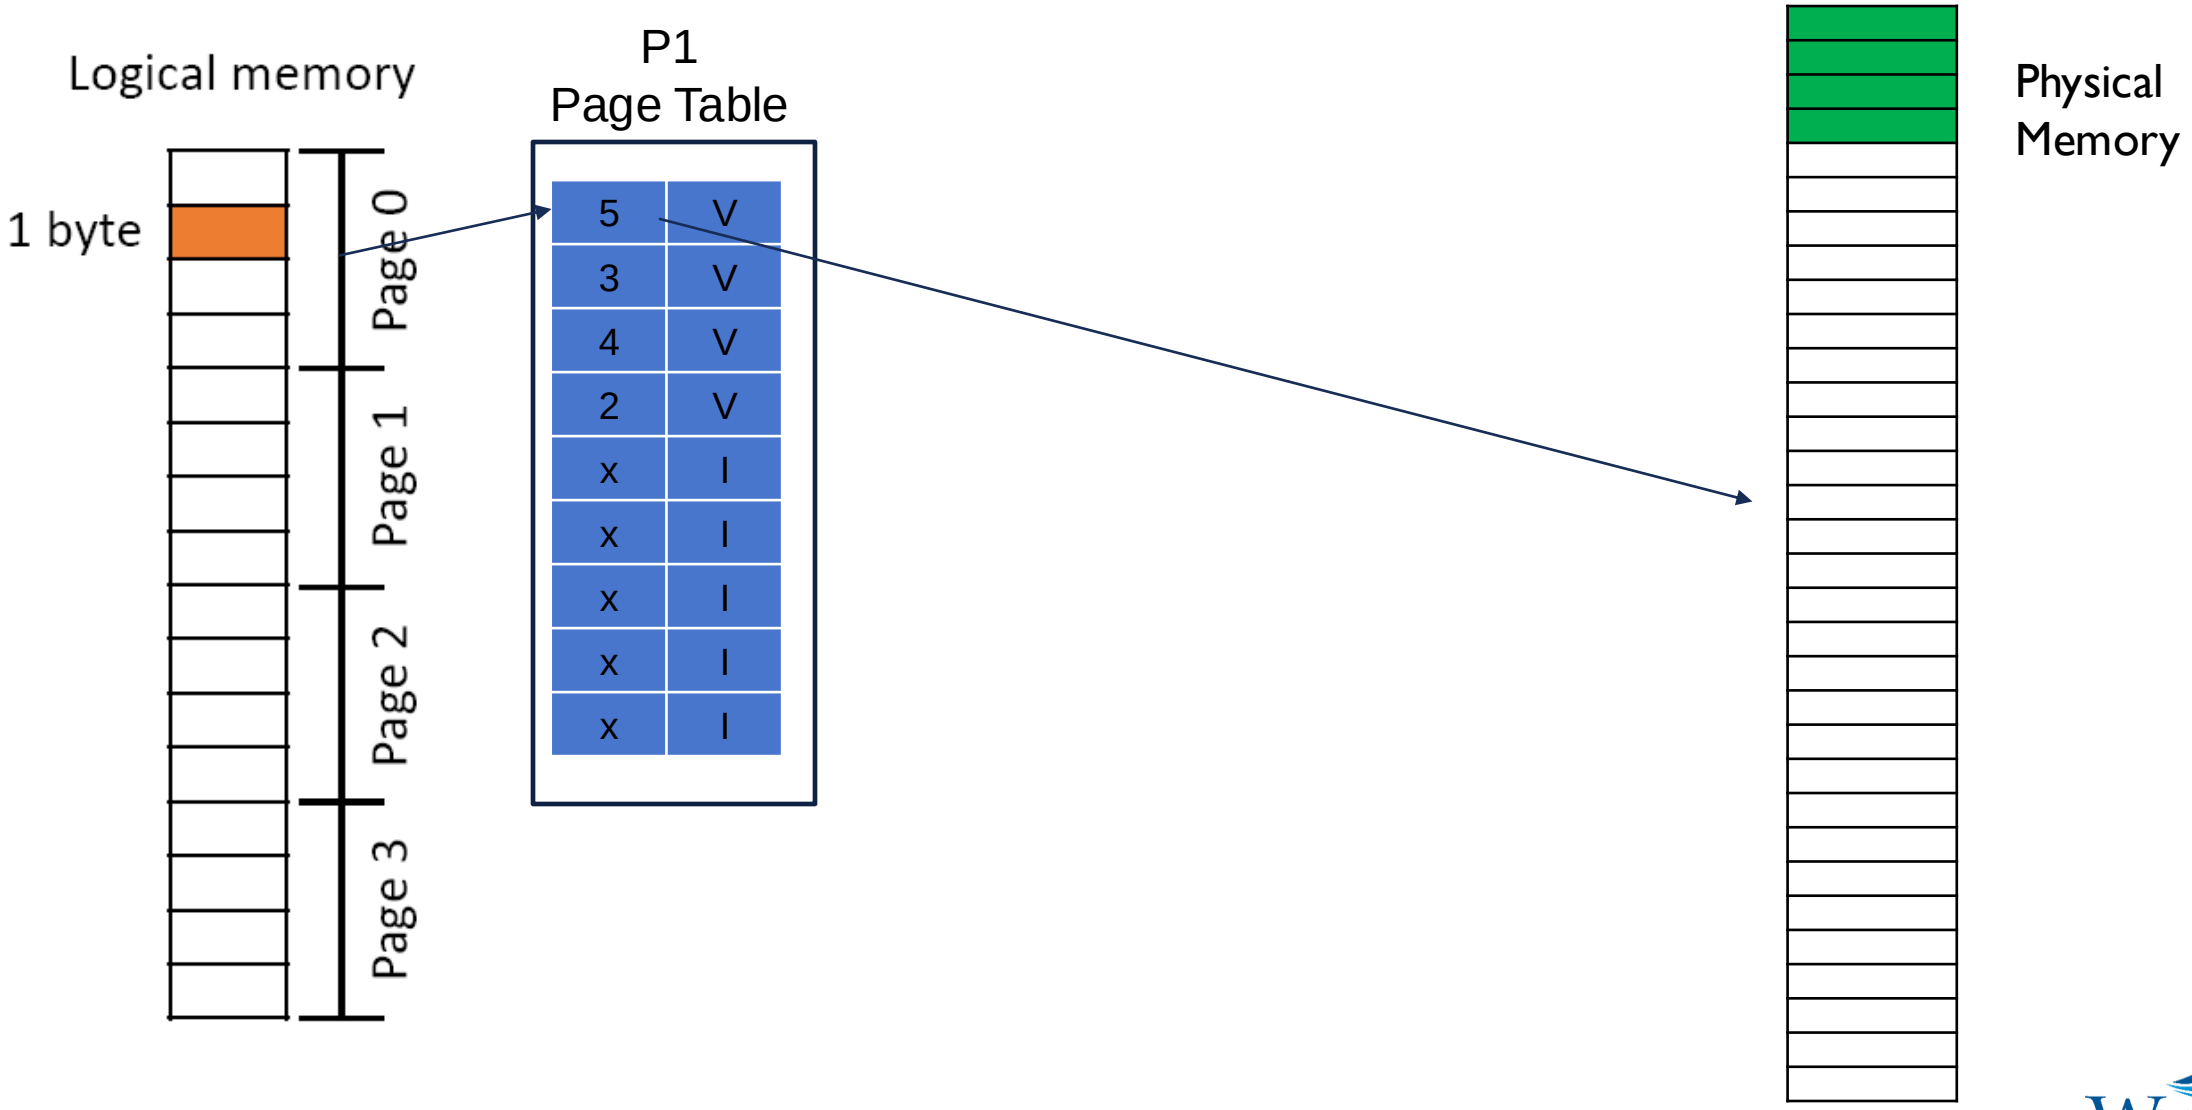

SINGLE PAGE TABLE

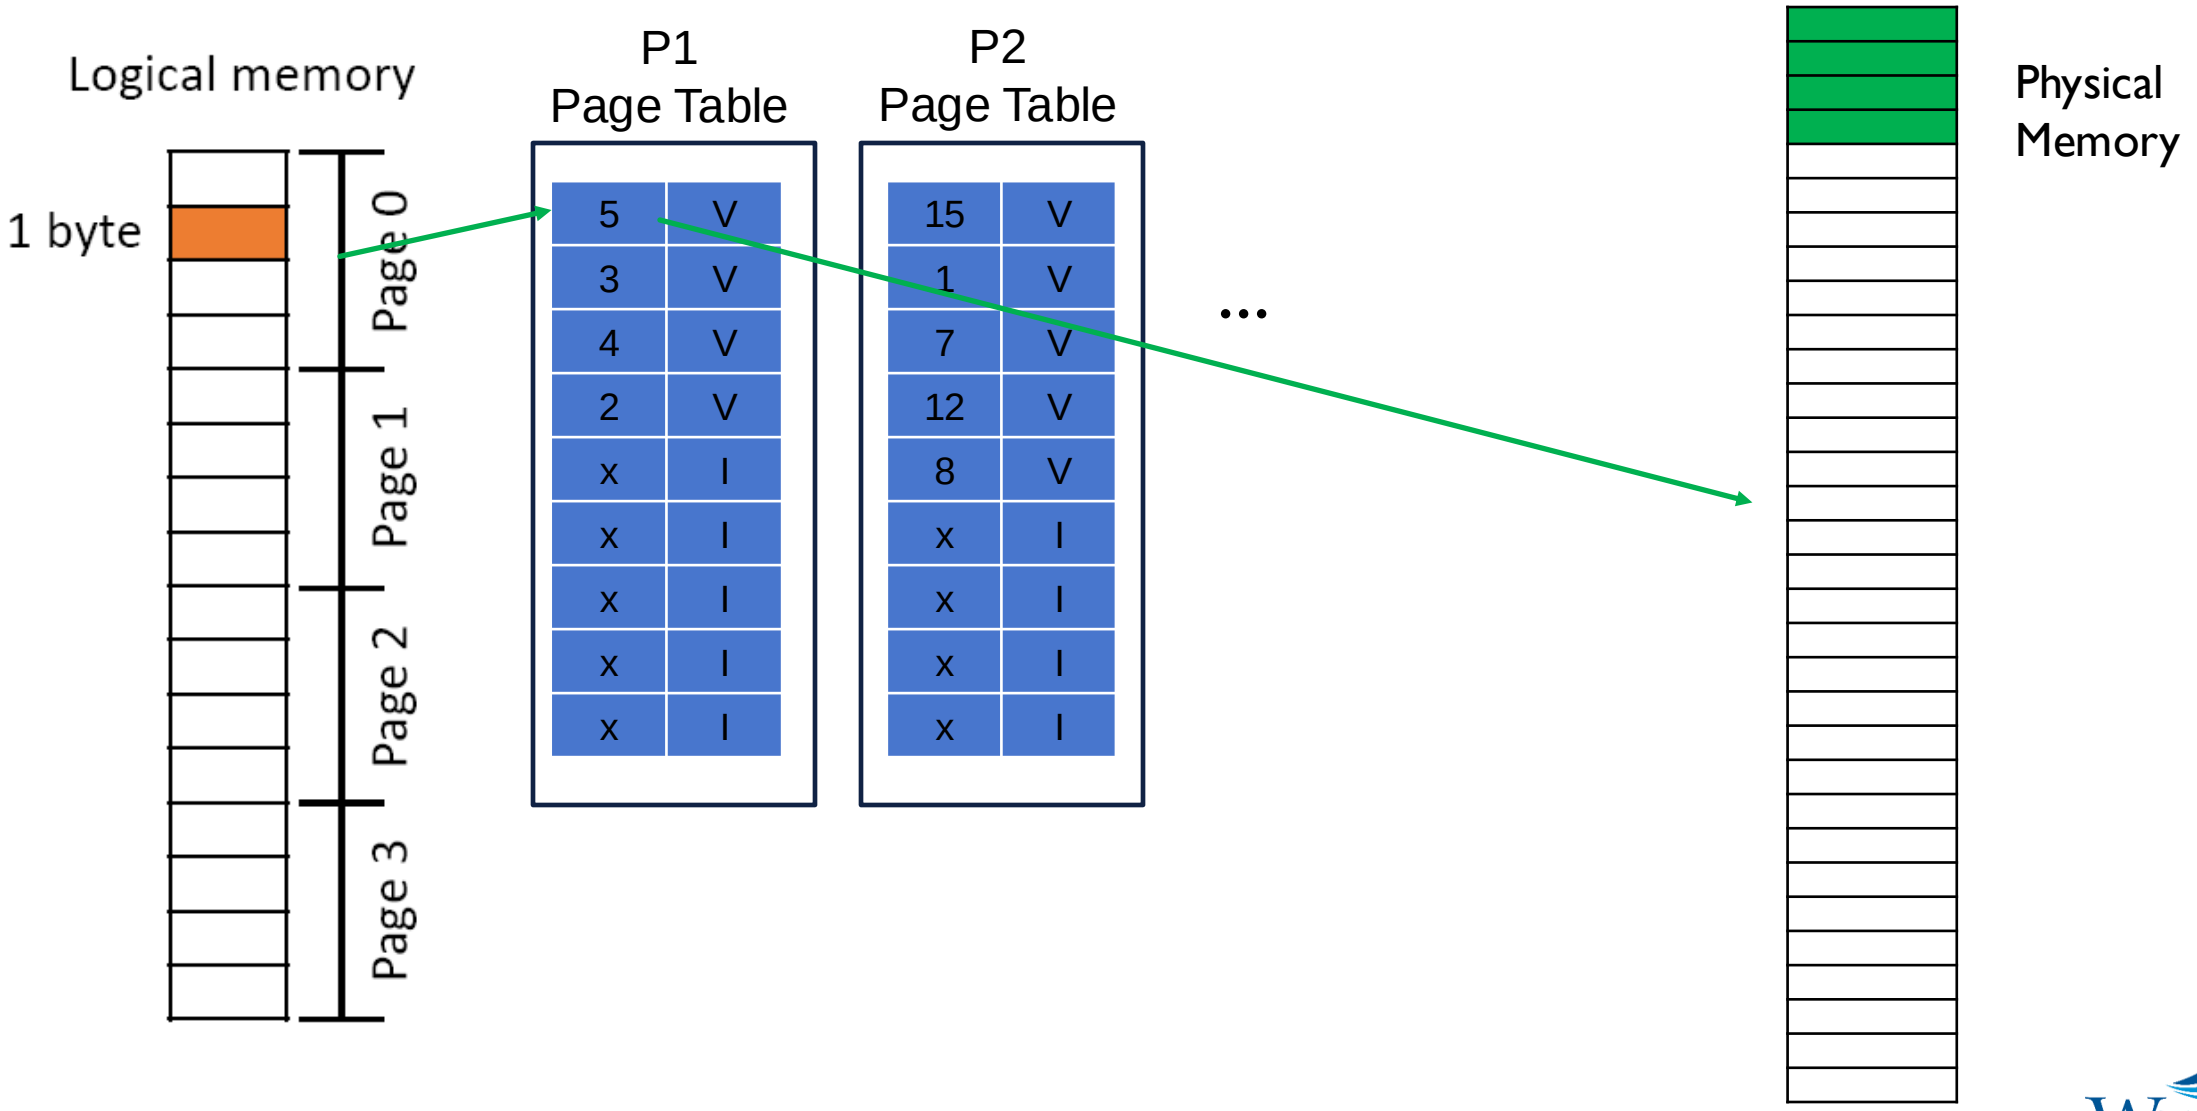

SINGLE PAGE TABLE

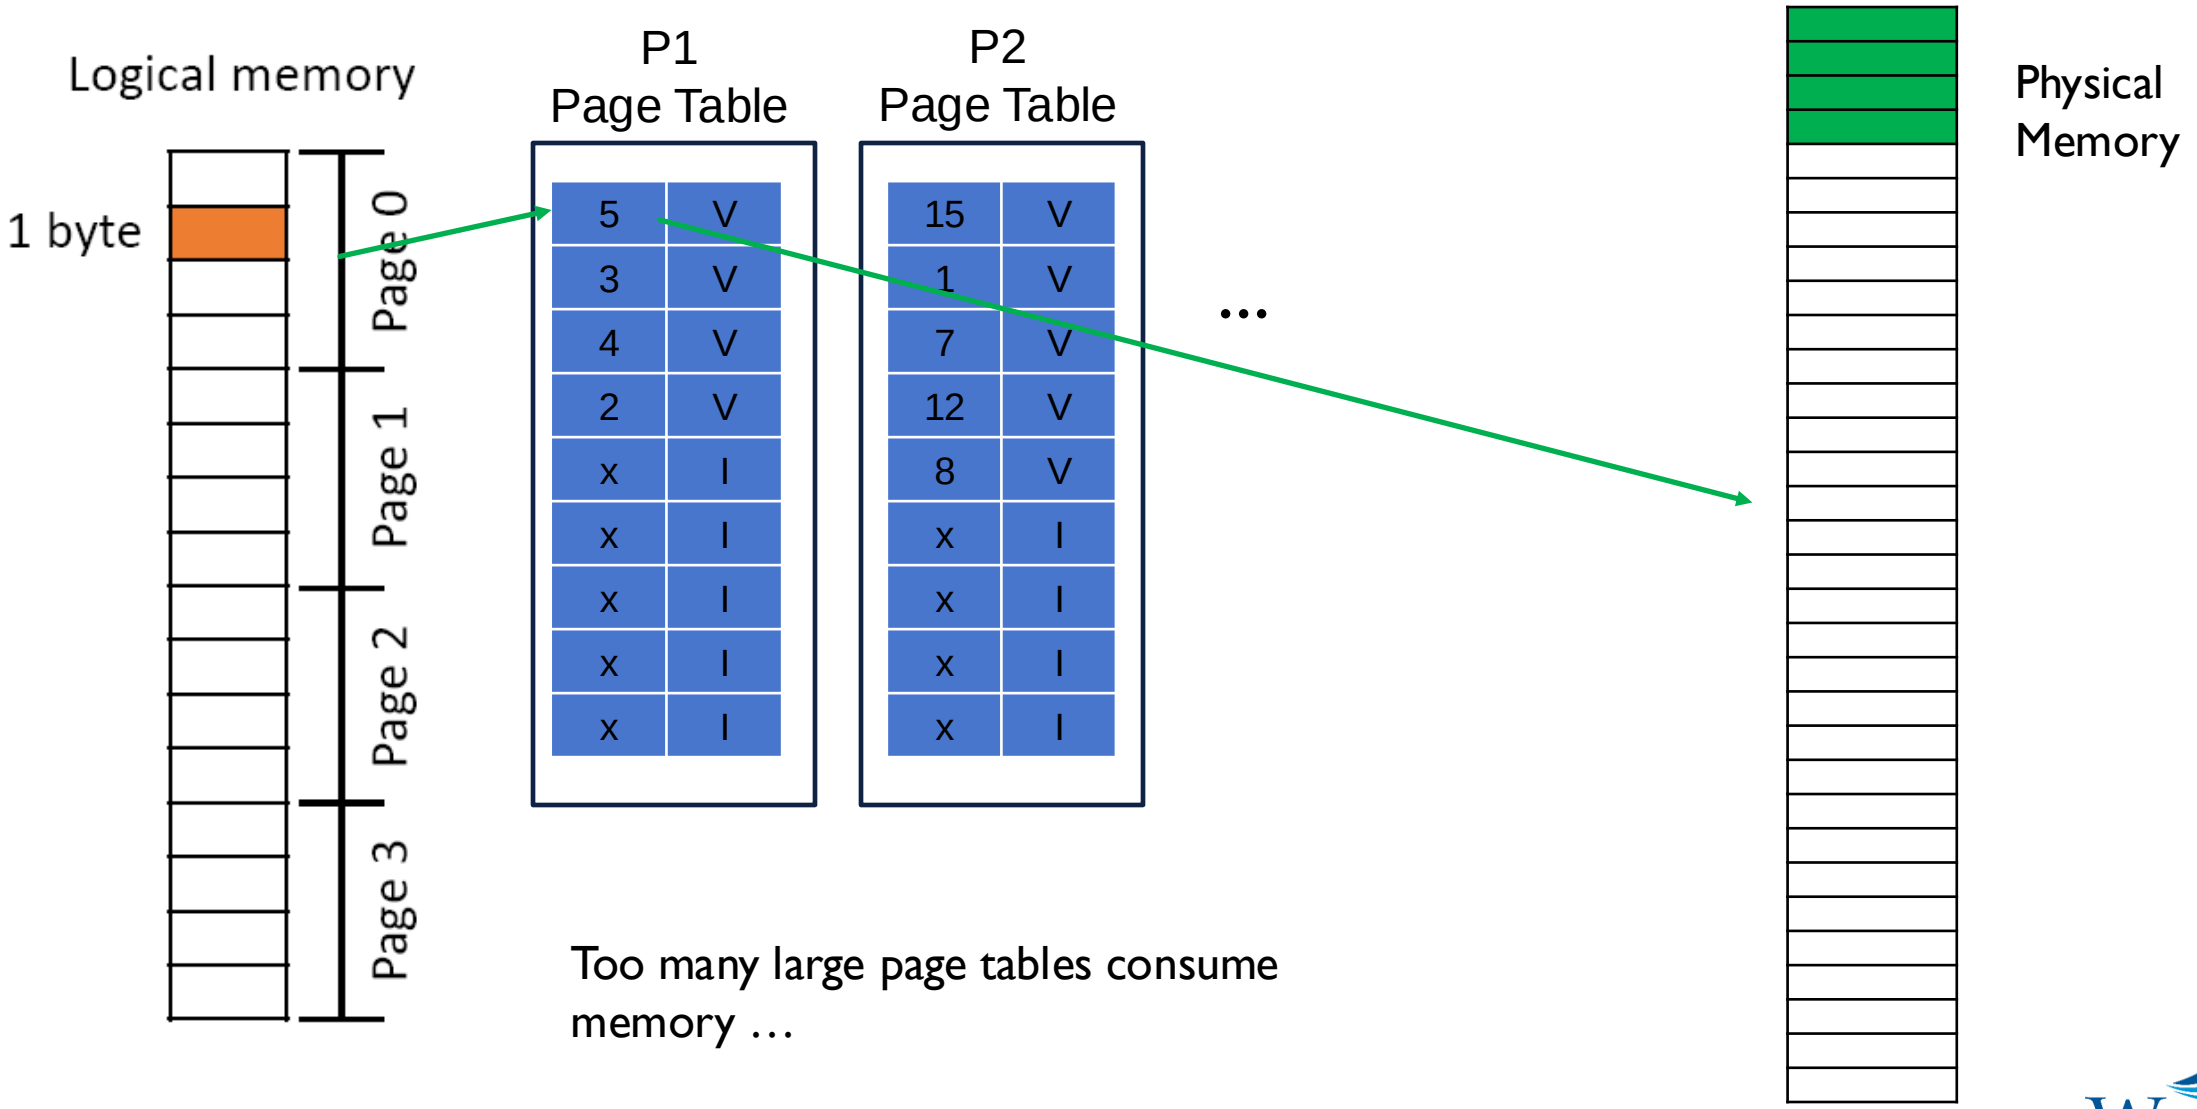

Too many large page tables consume memory …

SINGLE PAGE TABLE: INVERTED PAGE TABLE

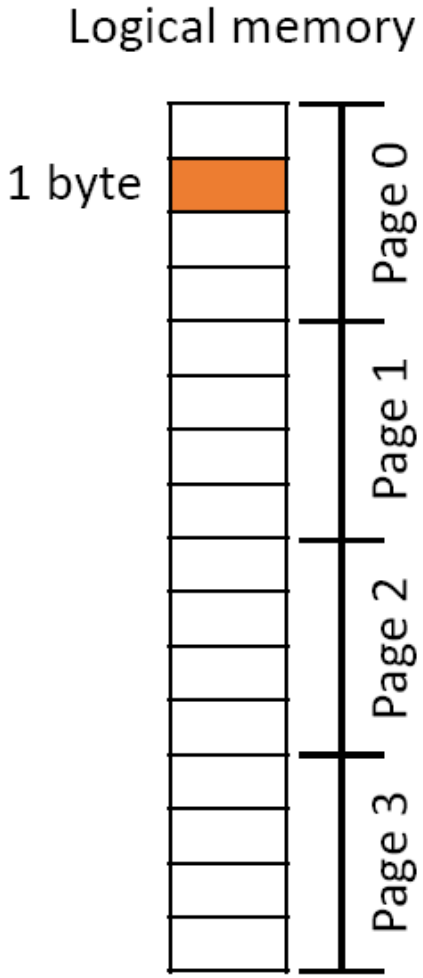

Instead of a page table, a 'frame' table with all the valid frames in physical memory with process ID and page Number assigned to.

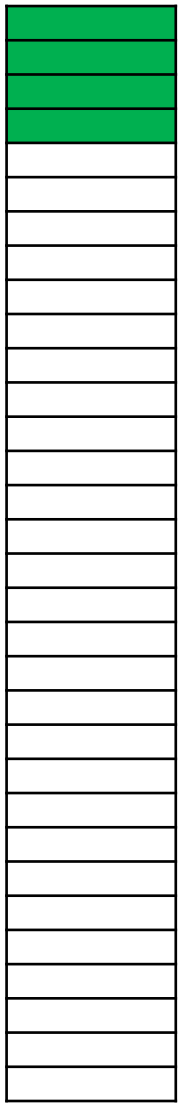

Physical Memory

SINGLE PAGE TABLE: INVERTED PAGE TABLE

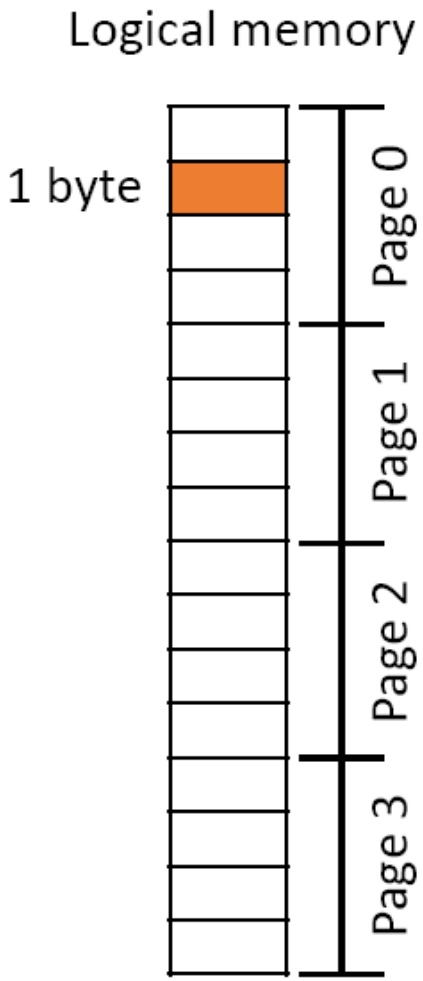

Instead of a page table, a 'frame' table with all the valid frames in physical memory with process ID and page Number assigned to.

Inverted Page Table

Indexed By Frame!

	Page No.	Process ID
Frame 0		
Frame 1		
Frame 2		
Frame 3		
Frame 4		
Frame 5		
⋮	⋮	

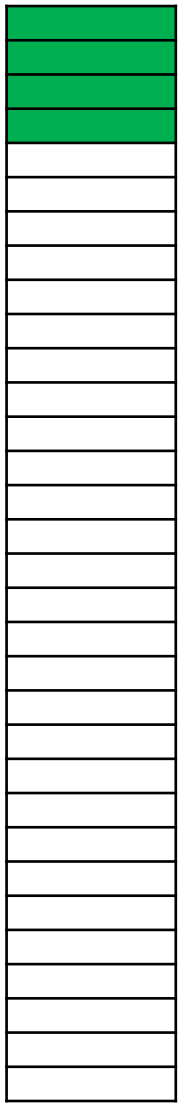

Physical Memory

SINGLE PAGE TABLE: INVERTED PAGE TABLE

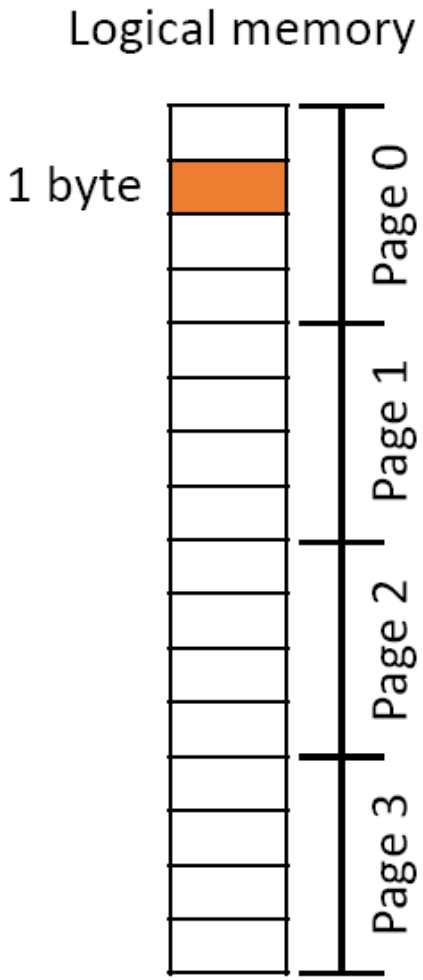

Instead of a page table, a 'frame' table with all the valid frames in physical memory with process ID and page Number assigned to.

Inverted Page Table

	Page No.	Process ID
Frame 0	Page 1	Process 3
Frame 1	Page 14	Process 1
Frame 2	Page 511	Process 3
Frame 3	Page 13	Process 2
Frame 4	None	None
Frame 5	Page 16	Process 1
⋮	⋮	⋮

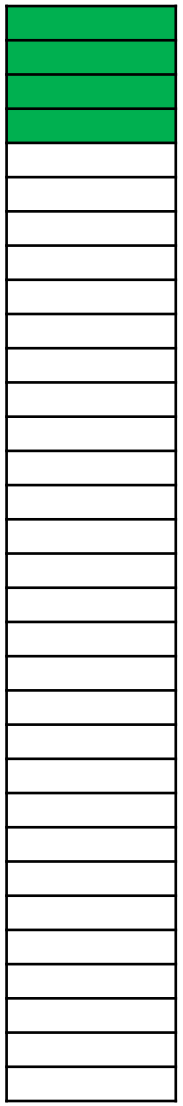

Physical Memory

INVERTED PAGE TABLE

Now the address should include the process id, since the table is shared by all processes.

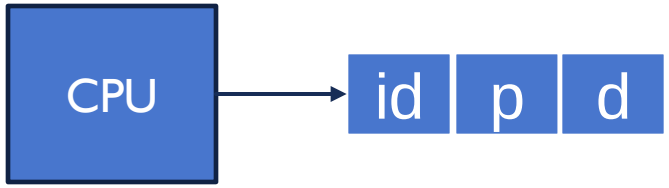

Inverted Page Table

	Page No.	Process ID
Frame 0	Page 1	Process 3
Frame 1	Page 14	Process 1
Frame 2	Page 511	Process 3
Frame 3	Page 13	Process 2
Frame 4	None	None
Frame 5	Page 16	Process 1
⋮	⋮	⋮

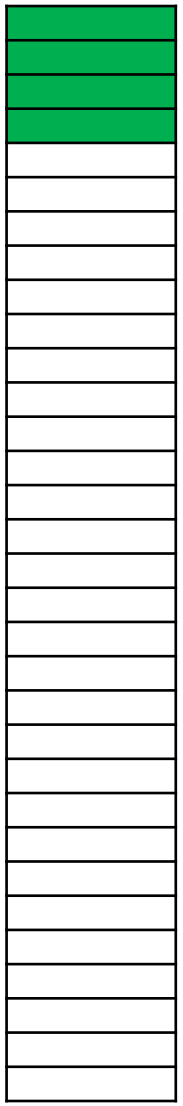

Physical Memory

PAGE TABLE STRUCTURES

- Multi-Level Paging

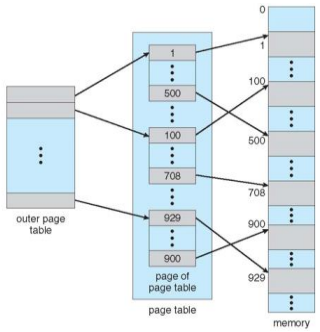

- Hashed Page Table

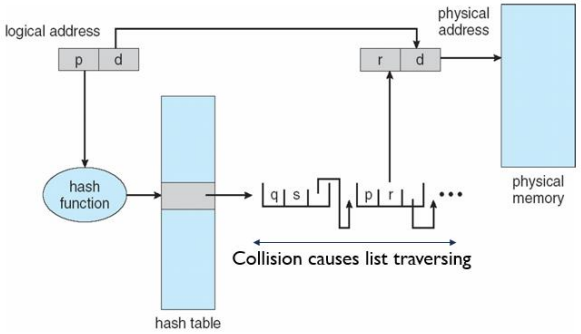

- Inverted Page Table

Page No.	Process ID
Page 1	Process 3
Page 14	Process 1
Page 511	Process 3
Page 13	Process 2
None	None
Page 16	Process 1

IA-32: INTEL ARCHITECTURE 32-BIT

- Two paging levels.
- Outer page table, called page directory.
- Two page sizes: 4-KB and 4-MB

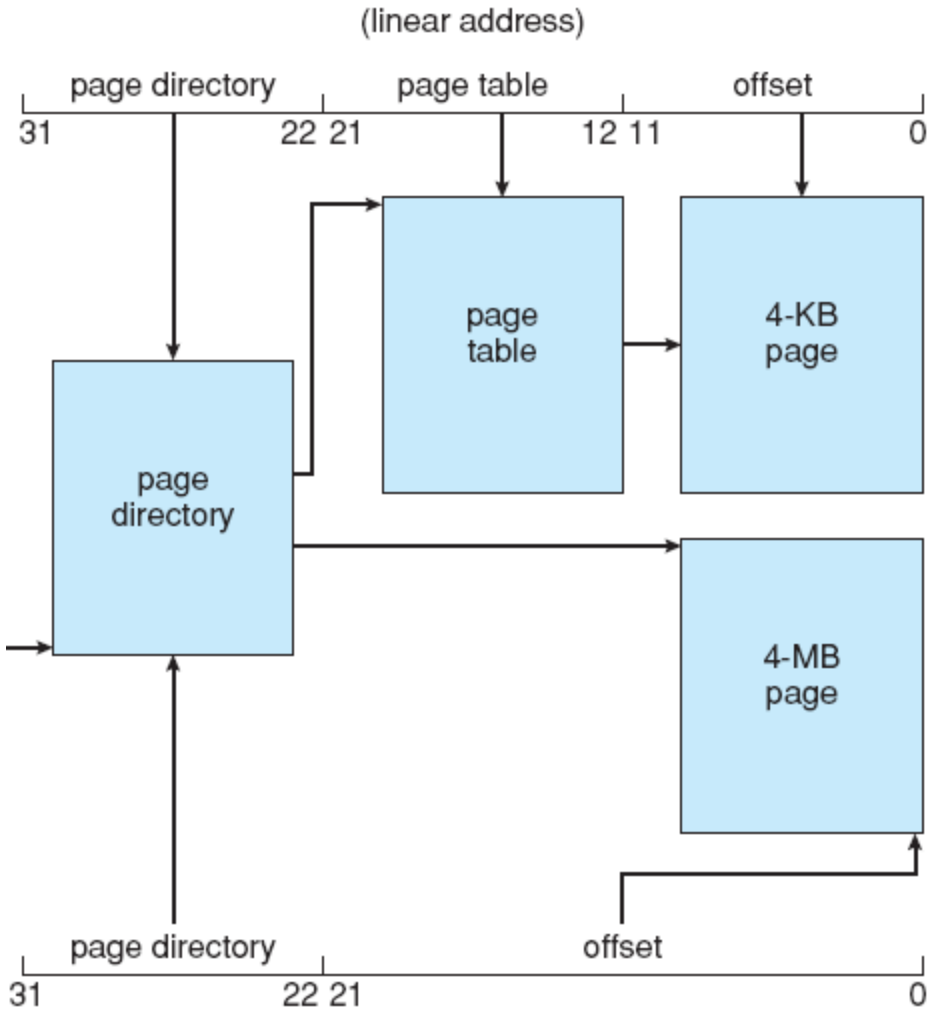

IA-32 PAE

- IA-32 Physical Address Extension.
- With 32-bit of address, only 4GB of ram can be supported.
- To solve this issue, intel added PAE: Physical Address Extension.

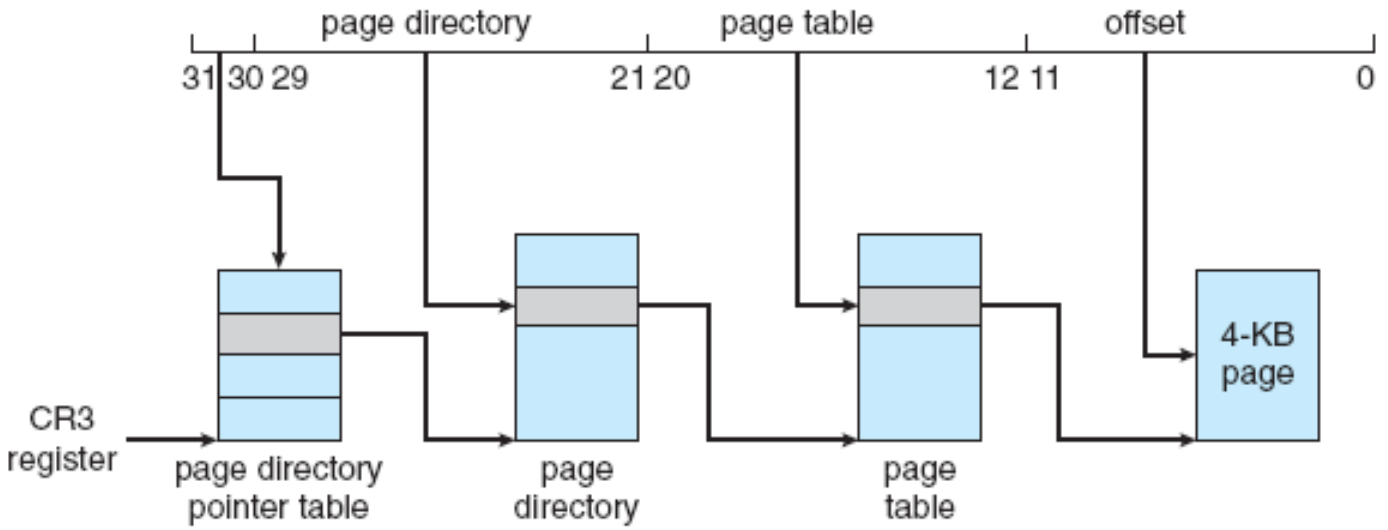

IA-32 PAE

- IA-32 Physical Address Extension.
- With 32-bit of address, only 4GB of ram can be supported.
- To solve this issue, intel added PAE: Physical Address Extension.
- Can be enabled/disabled on *supported* chips.
- The 32-bit address adopts a 3-level paging instead of 2 levels.

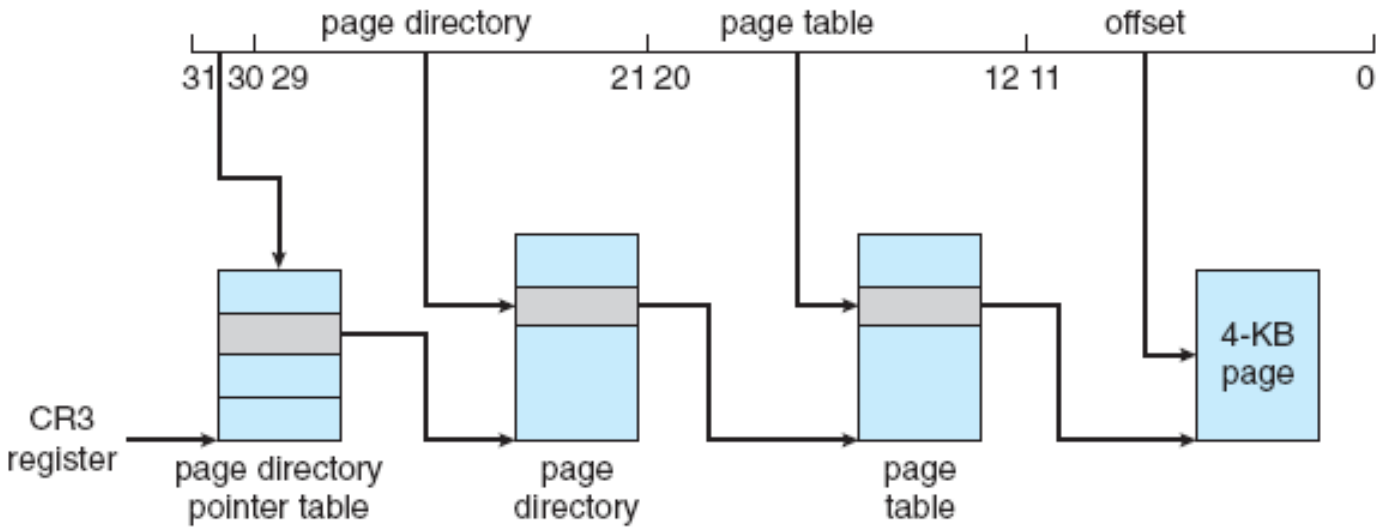

IA-32 PAE

- IA-32 Physical Address Extension.
- With 32-bit of address, only 4GB of ram can be supported.
- To solve this issue, intel added PAE: Physical Address Extension.
- Can be enabled/disabled on *supported* chips.
- The 32-bit address adopts a 3-level paging instead of 2 levels.

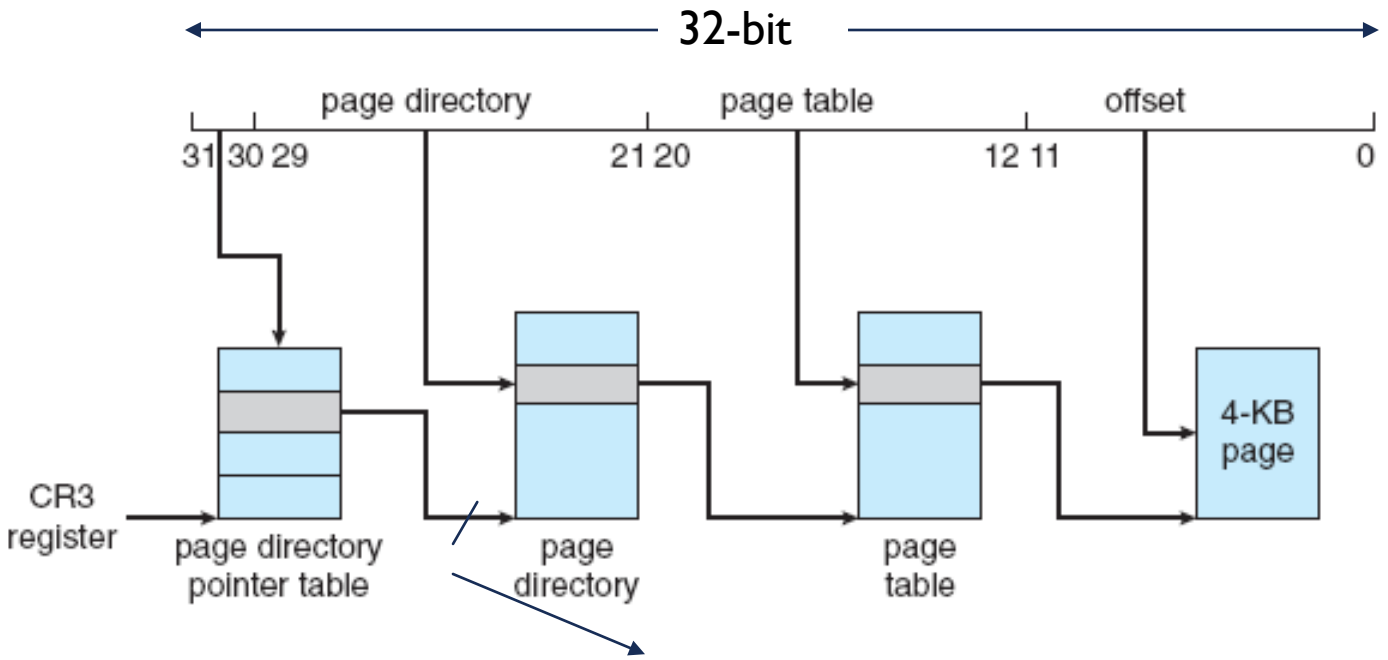

- Virtual address is still 32-bits .. But physical address is now 36-bits.

IA-32 PAE

- IA-32 Physical Address Extension.
- With 32-bit of address, only 4GB of ram can be supported.
- To solve this issue, intel added PAE: Physical Address Extension.
- Can be enabled/disabled on *supported* chips.
- The 32-bit address adopts a 3-level paging instead of 2 levels.

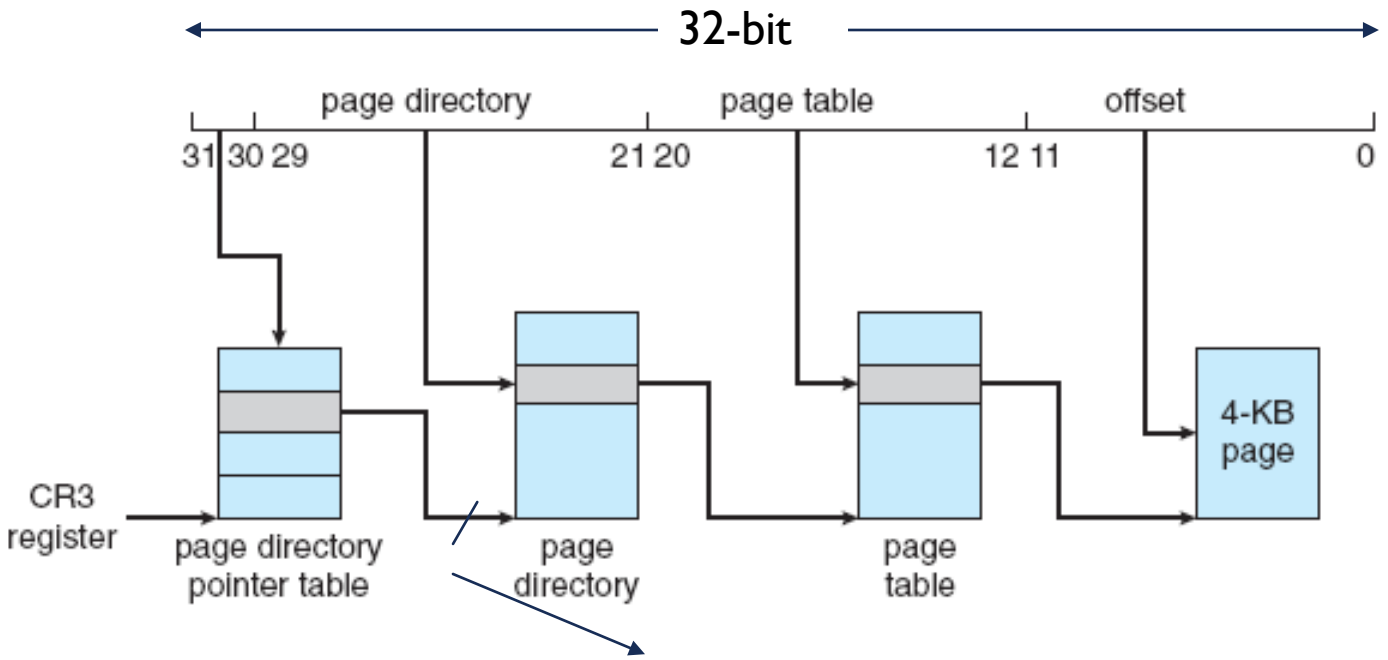

What does that imply?

- Virtual address is still 32-bits .. But physical address is now 36-bits.

IA-32 PAE

- IA-32 Physical Address Extension.
- With 32-bit of address, only 4GB of ram can be supported.
- To solve this issue, intel added PAE: Physical Address Extension.
- Can be enabled/disabled on *supported* chips.
- The 32-bit address adopts a 3-level paging instead of 2 levels.

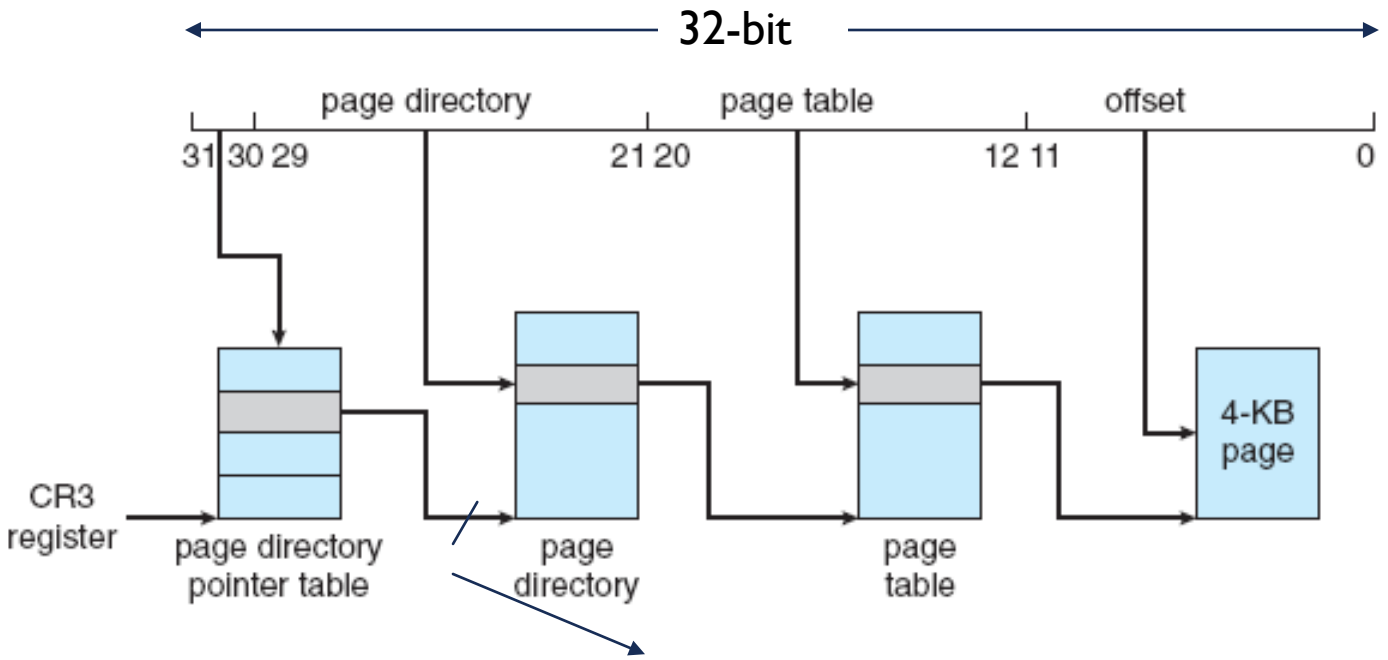

- Virtual address is still 32-bits .. But physical address is now 36-bits.
- Process is still limited to 4-GB, but system is not.

PAGING IMPLEMENTATION 64-X86

- Introduced by AMD in Opteron (2003)
- 64-bit virtual address (only 48-bits are used)
- 4-level paging
- TLB misses: very expensive, each level require a memory access.
- Memory Management Unit has its own cache dedicated for paging.

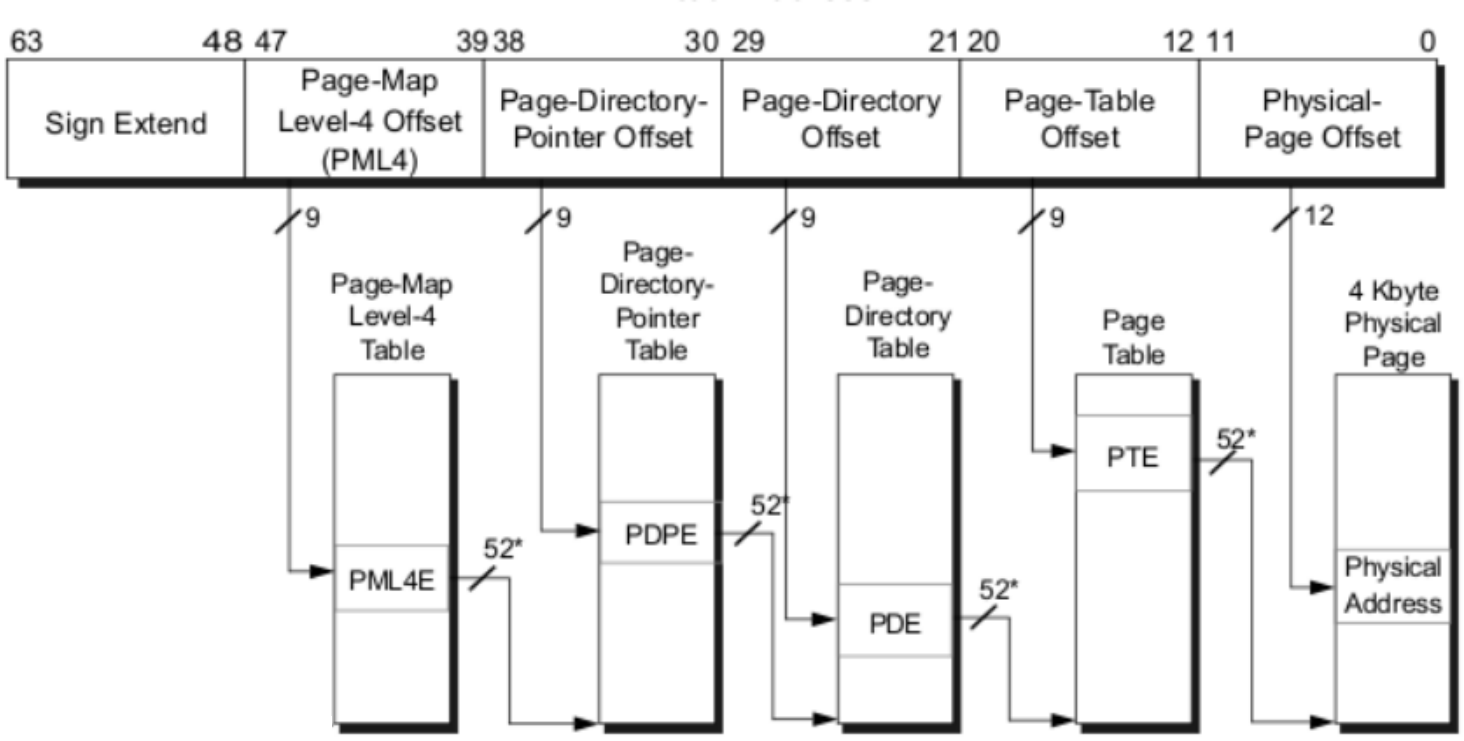

ARM V8 ARCHITECTURE

- ARM Architecture dominates mobile devices.
- Uses 64-bit (48-bit used) 4 level paging.
- Can offer contiguous memory in form of 'regions'.
- Two level TLBs:
  - Data TLB
  - Instruction TLB
  - Outer TLB for both Data and Instruction in case of TLB Miss

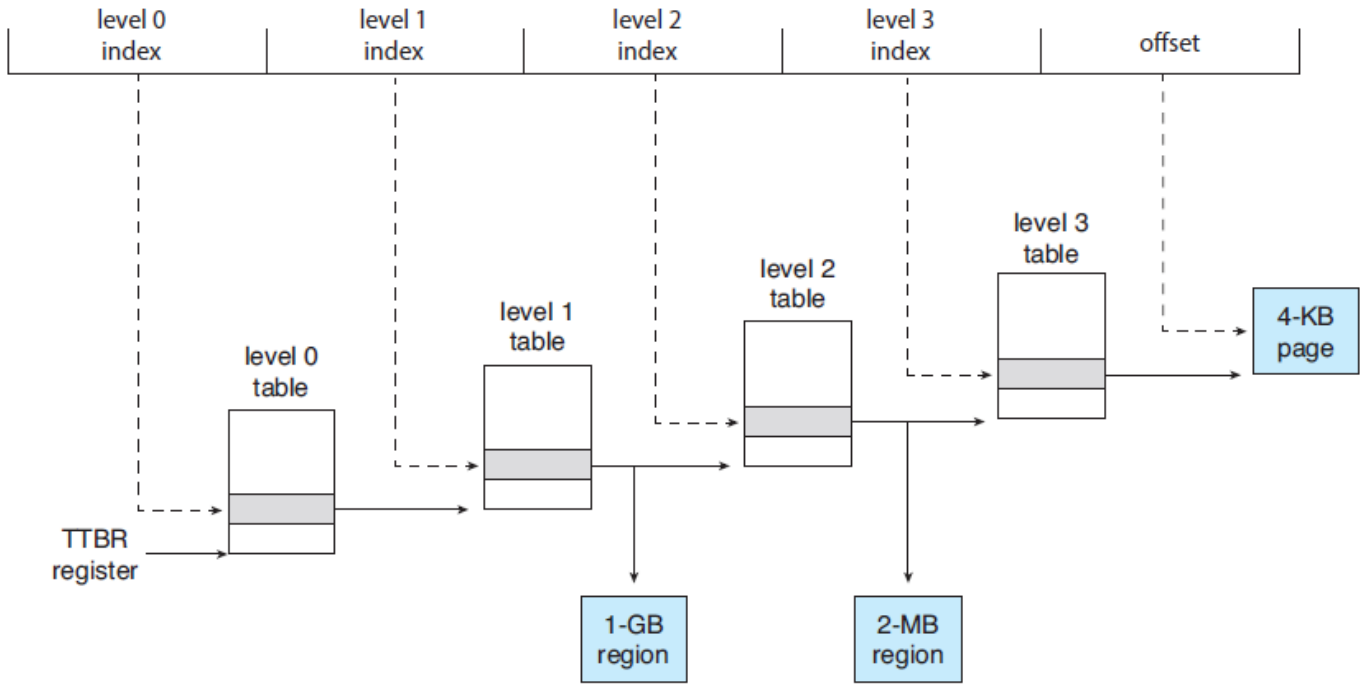

ARM V8 TWO-LEVEL TLB

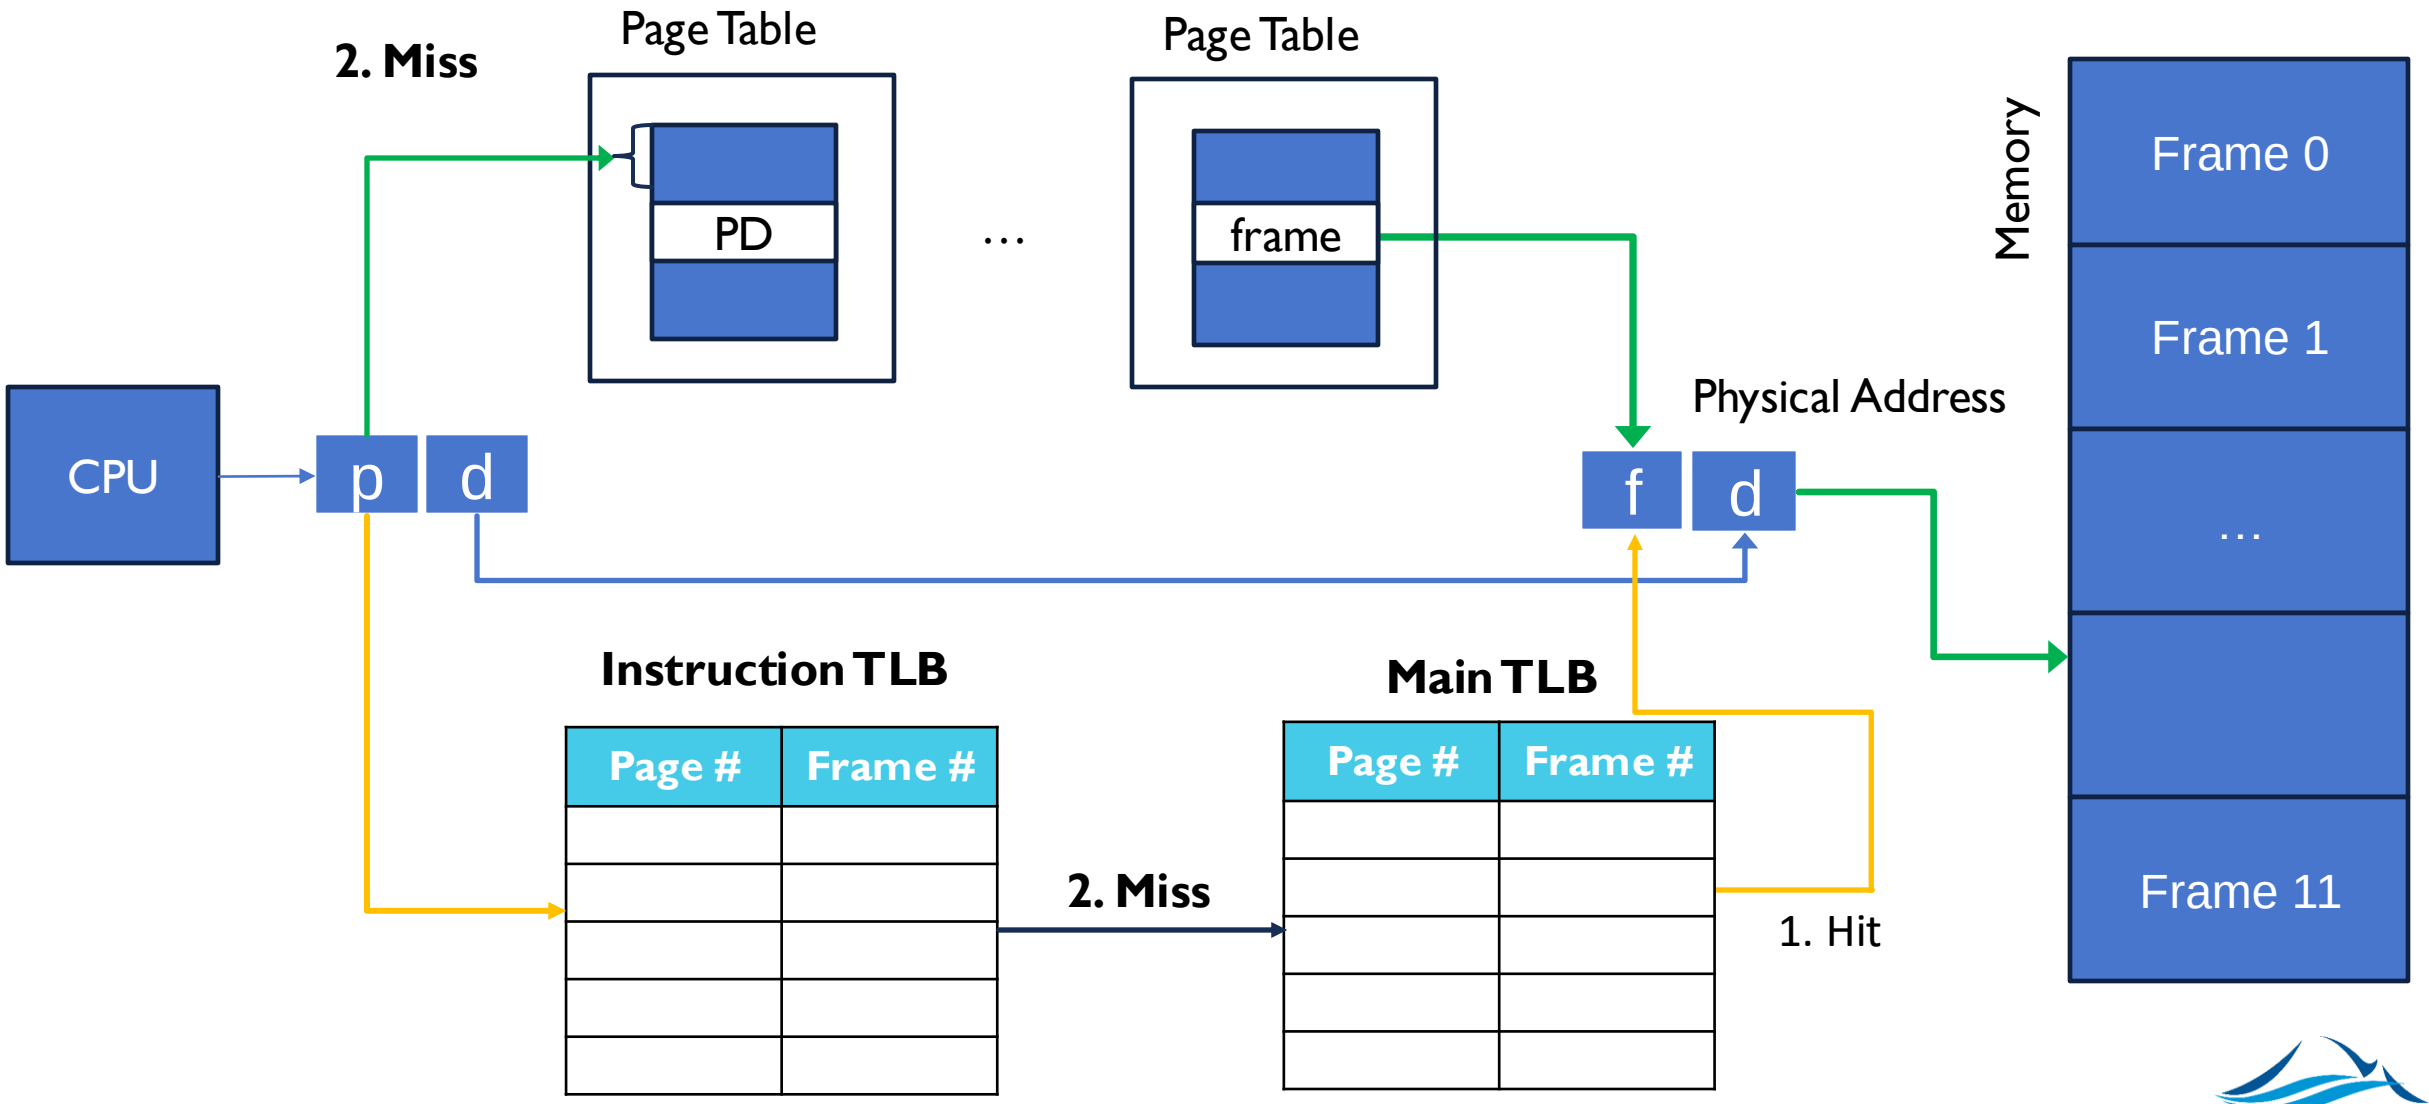

ARM V8 TWO-LEVEL TLB

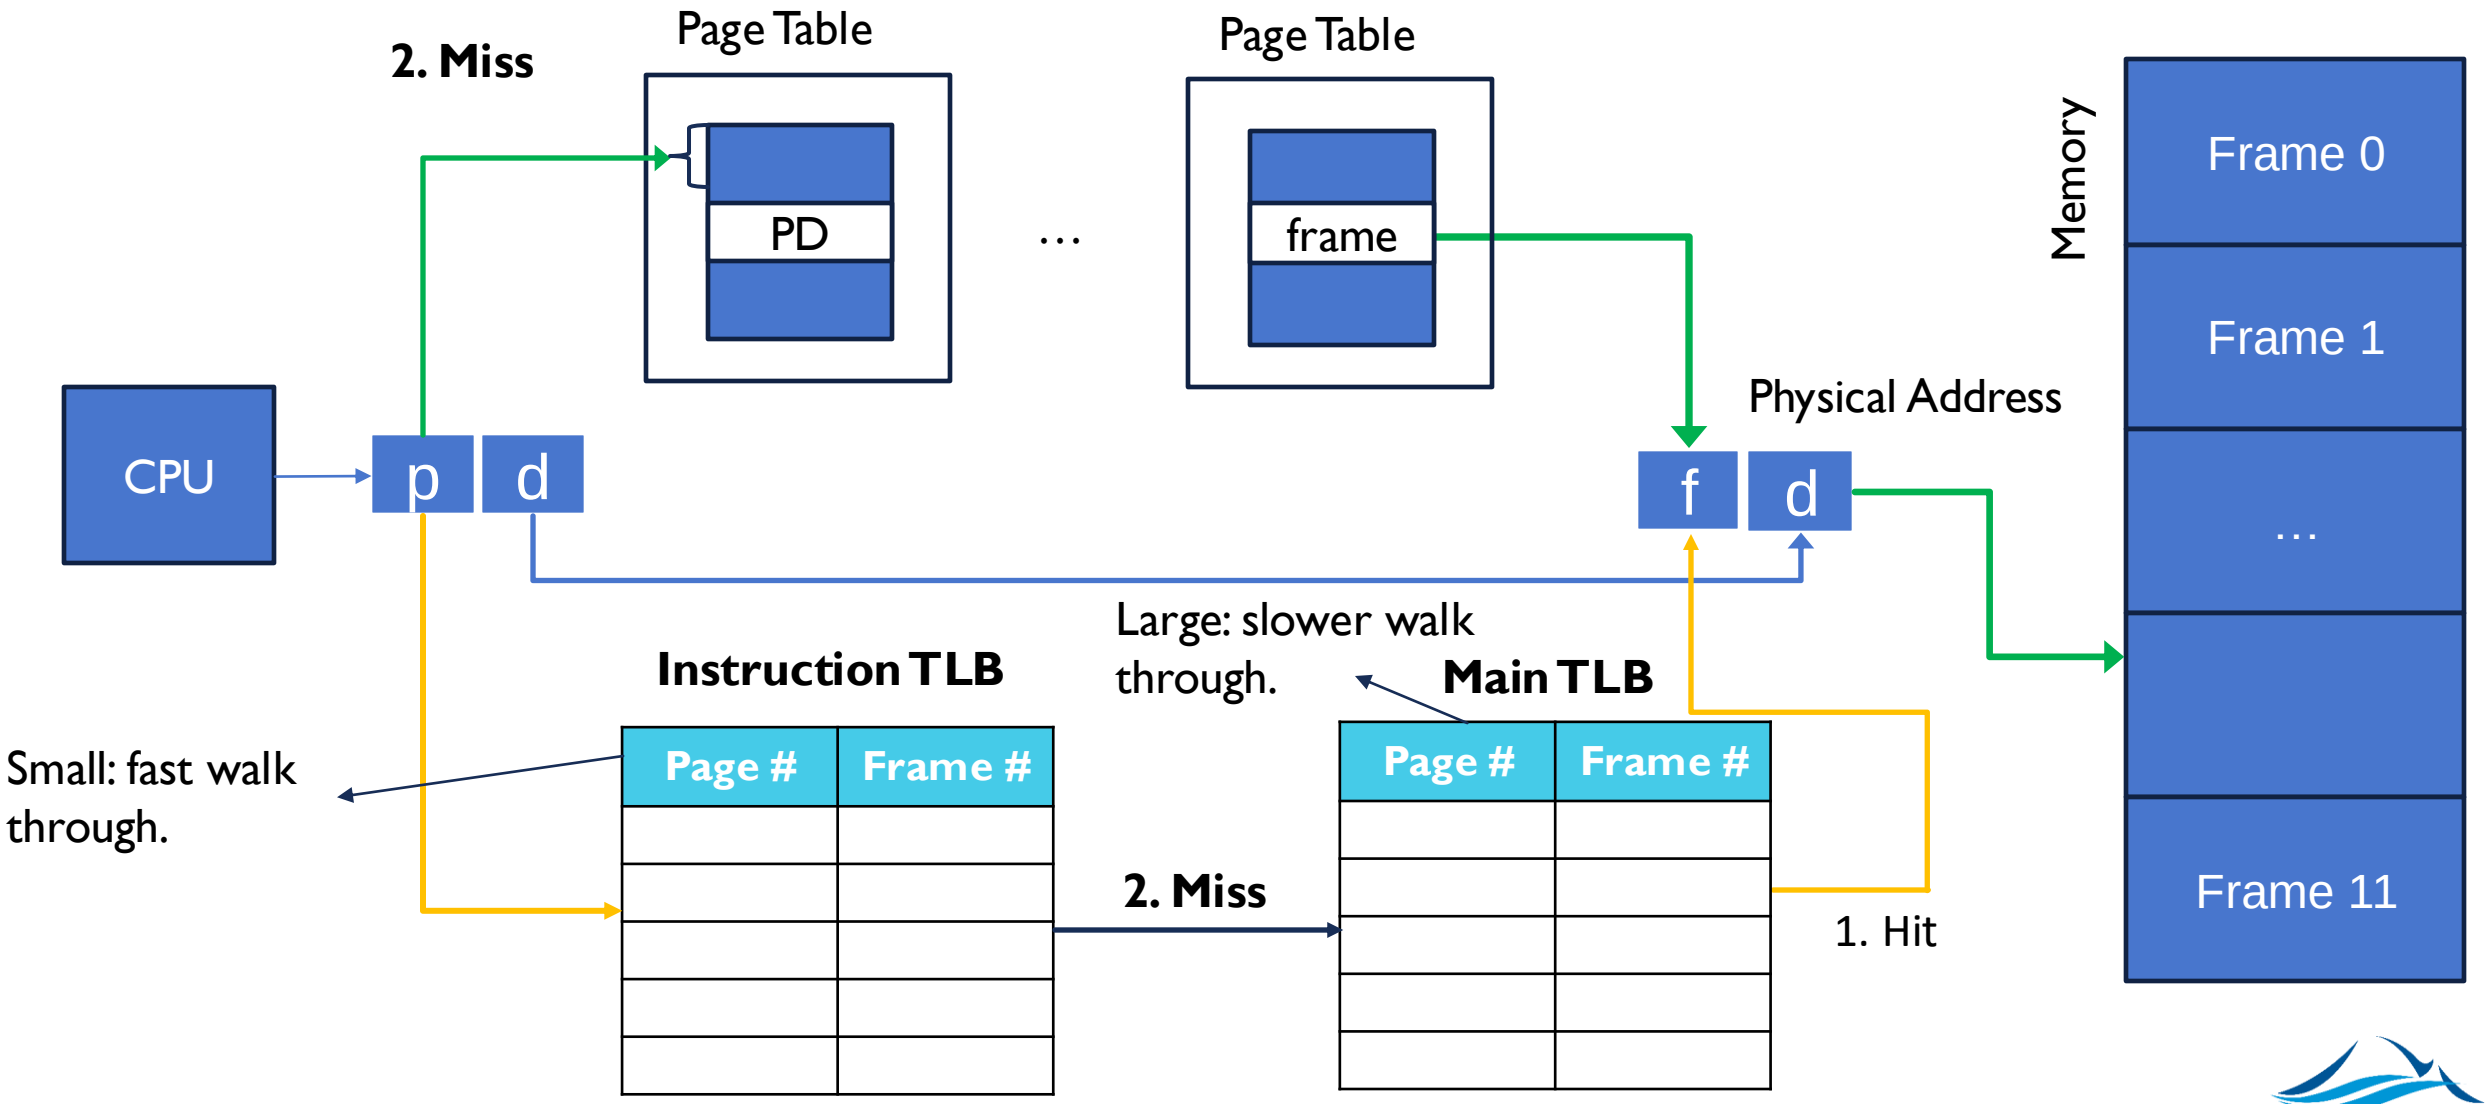

VIRTUAL MEMORY

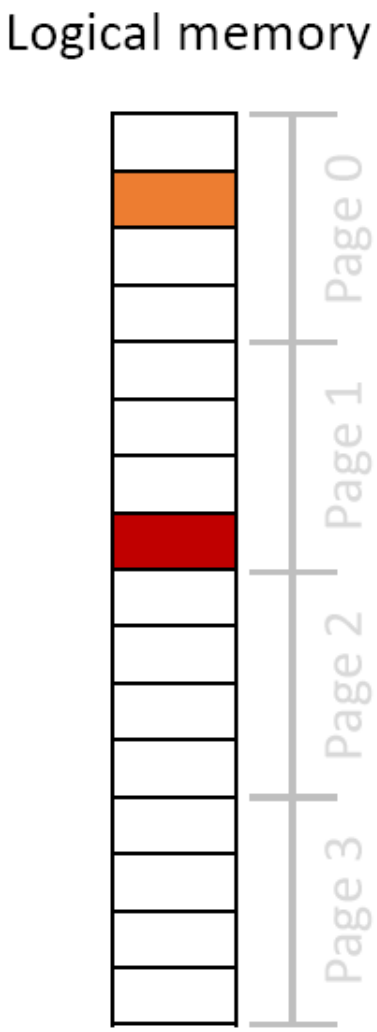

- Q: What are the two methods that can map logical memory to physical ones?

VIRTUAL MEMORY

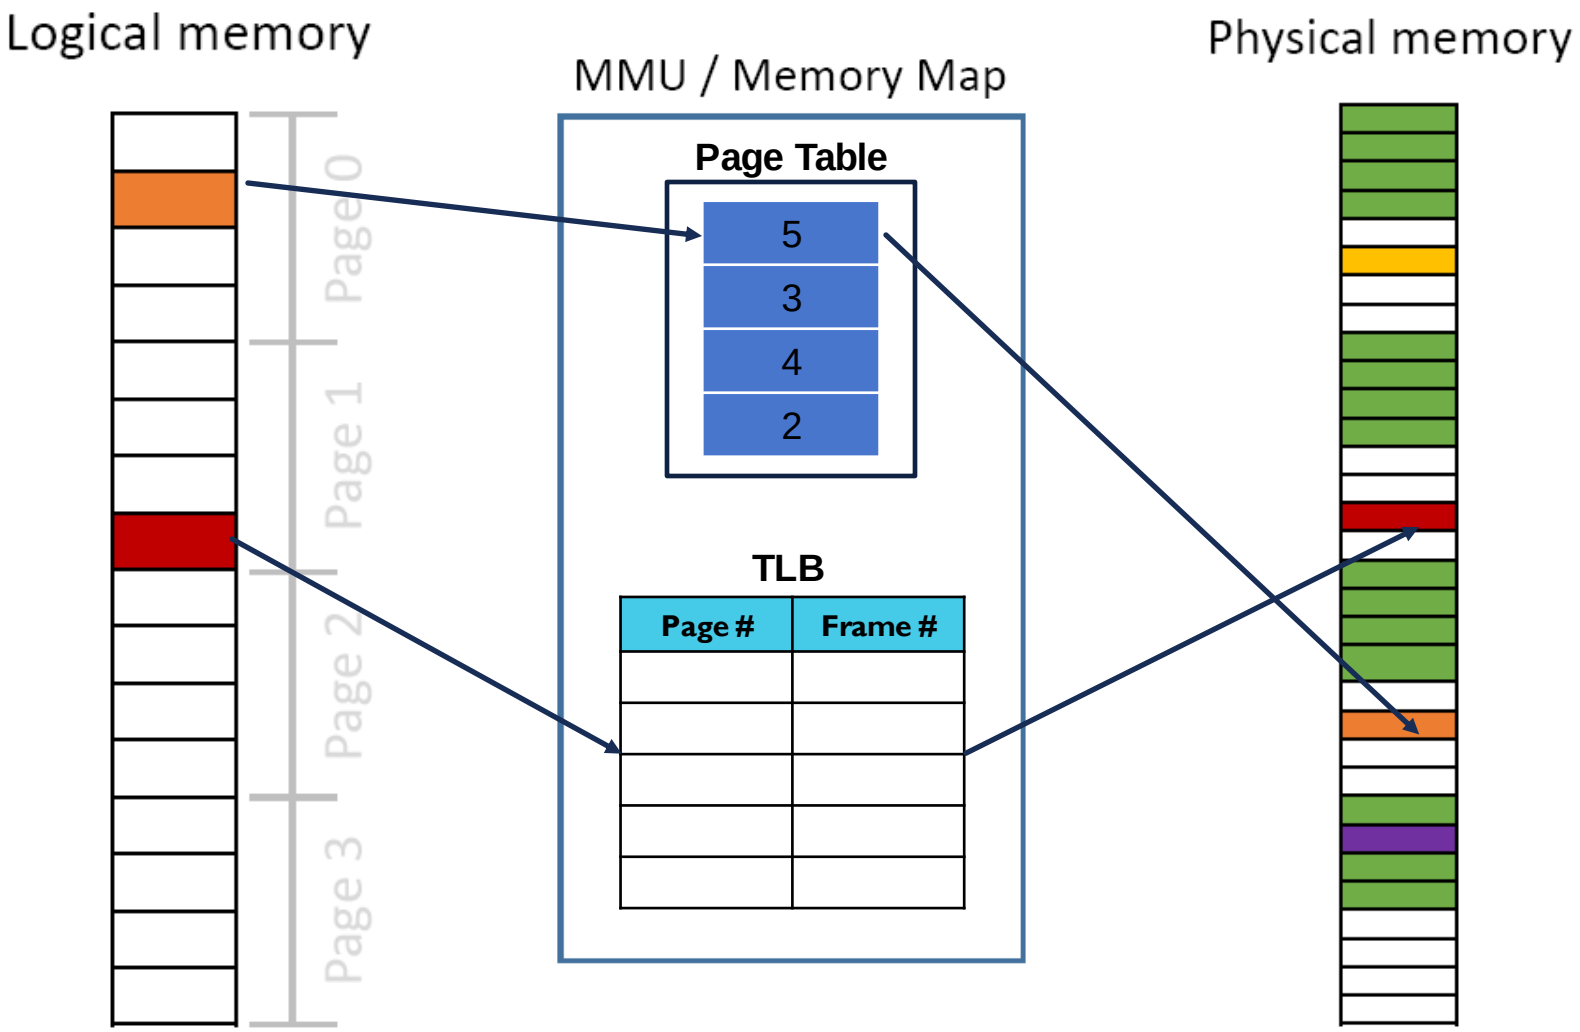

- Q:What are the two methods that can map logical memory to physical ones?
- Page Table and Translation Look Aside Buffers.

VIRTUAL MEMORY

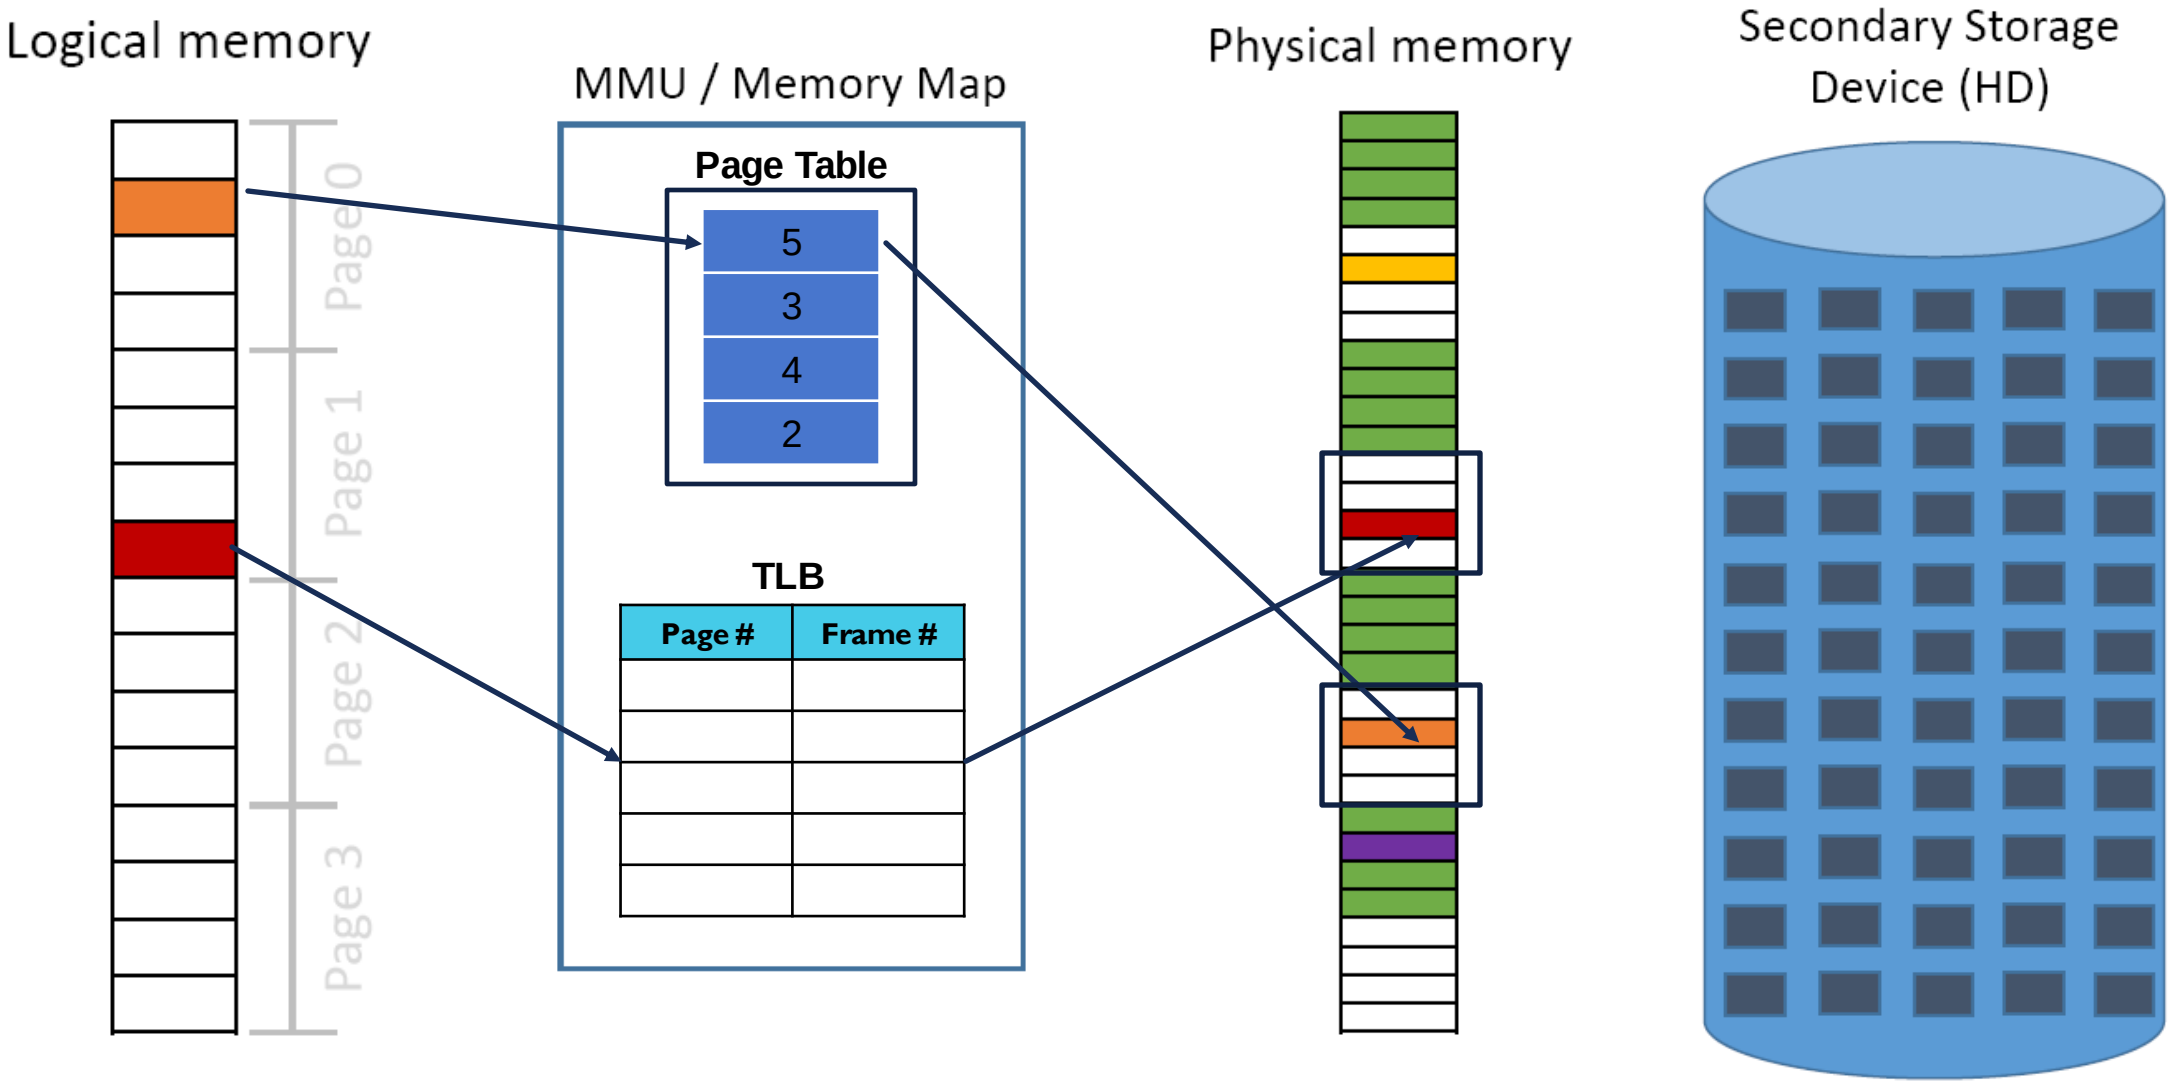

VIRTUAL MEMORY

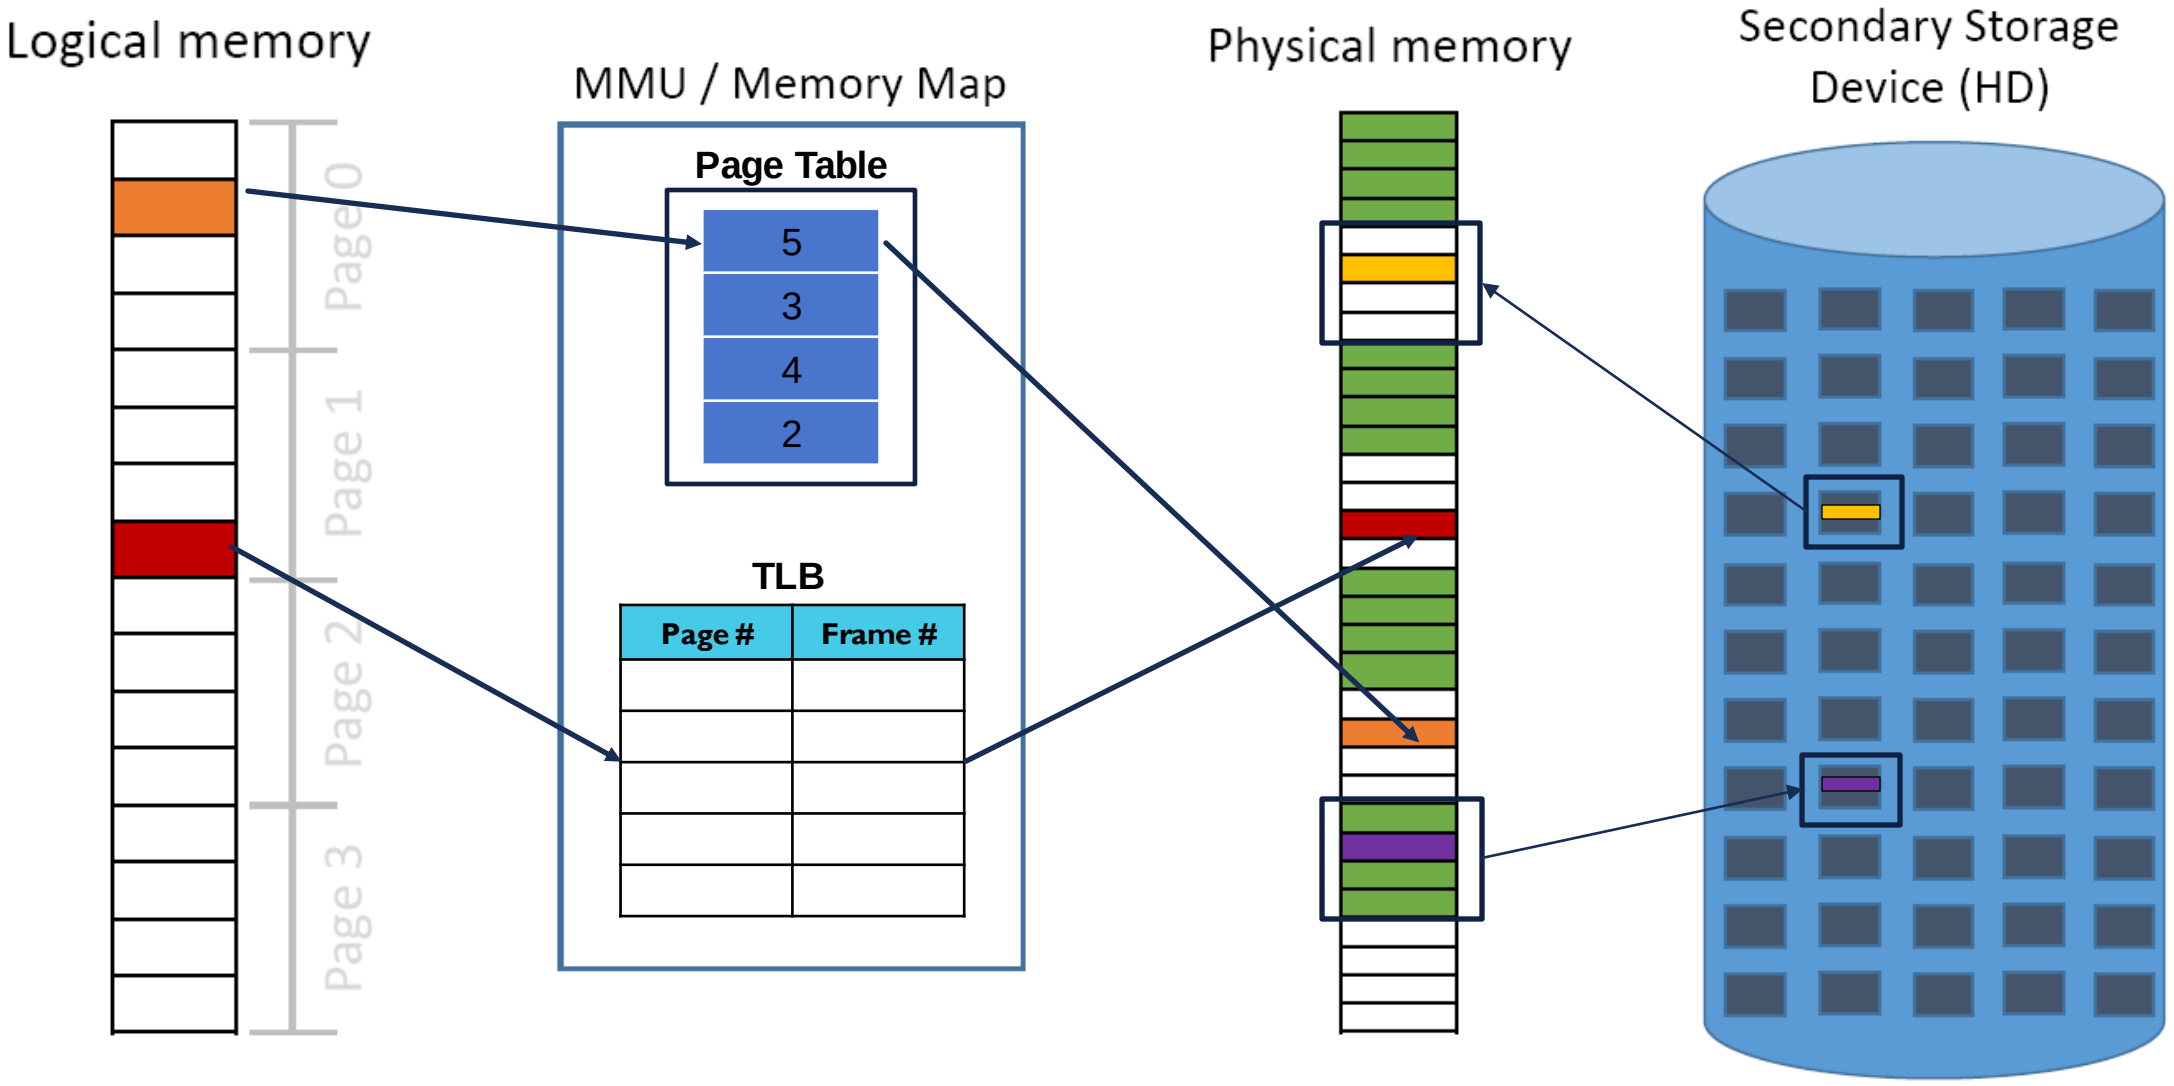

VIRTUAL MEMORY

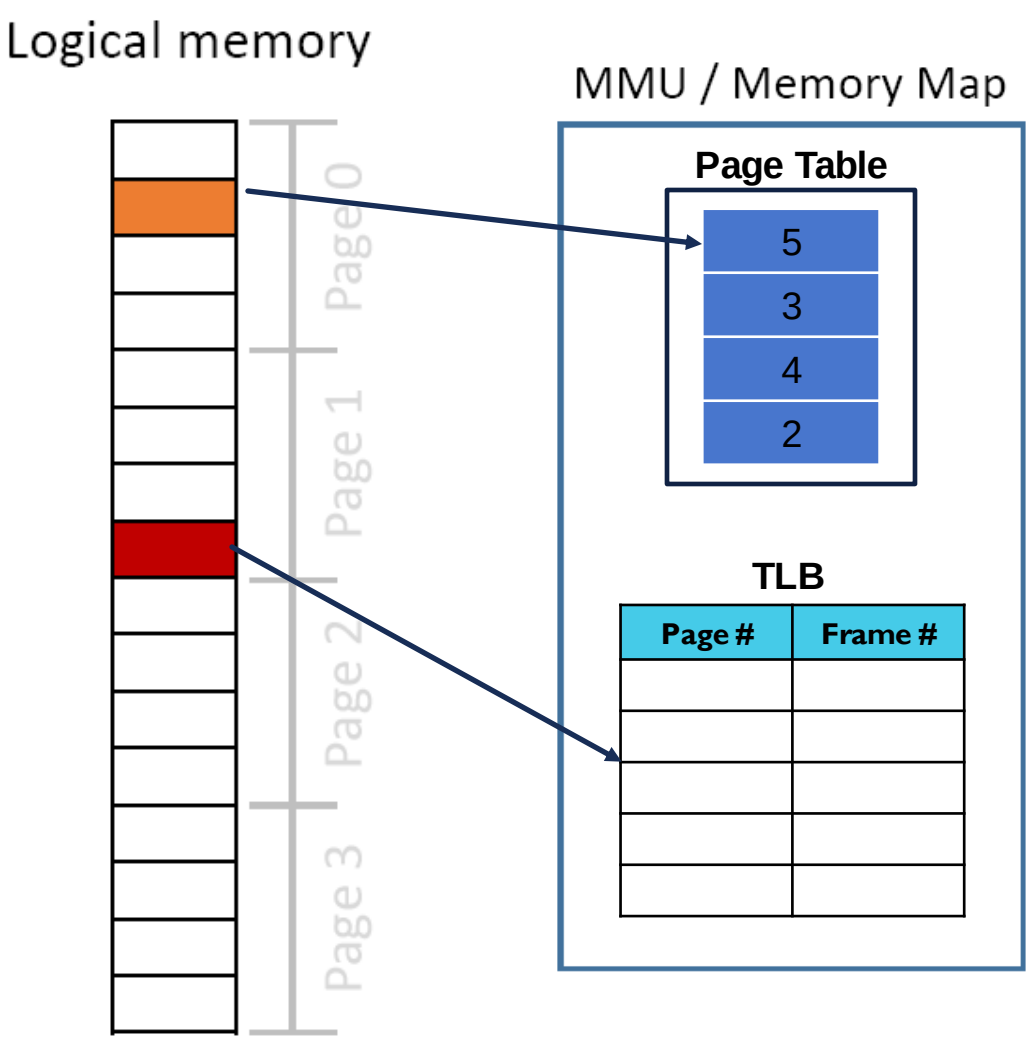

Q: How to allocate frames in physical memory?

Q: What to do if the physical memory is full?

Q: What frames need to be written out?

Q: How to decide what page to bring in memory?

Q: What's the impact on performance?

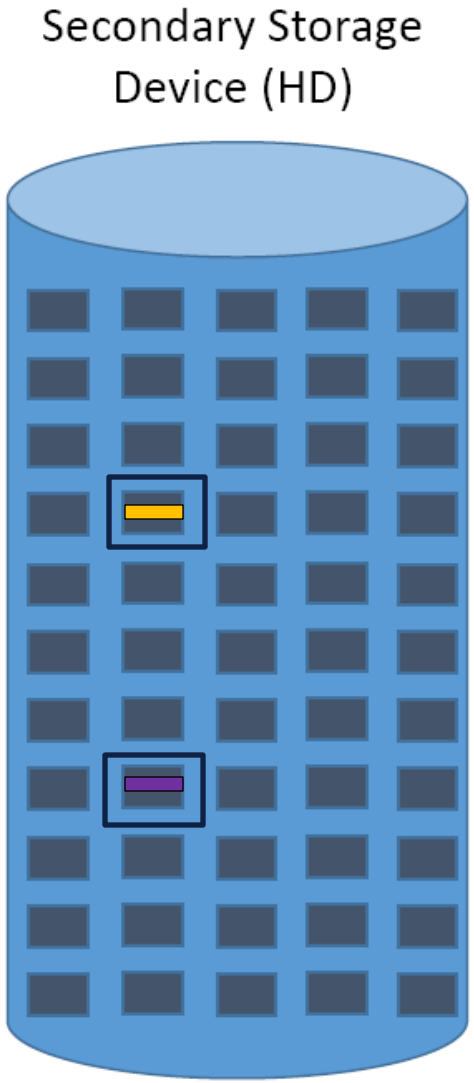

VIRTUAL MEMORY

Q: How to allocate frames in physical memory?

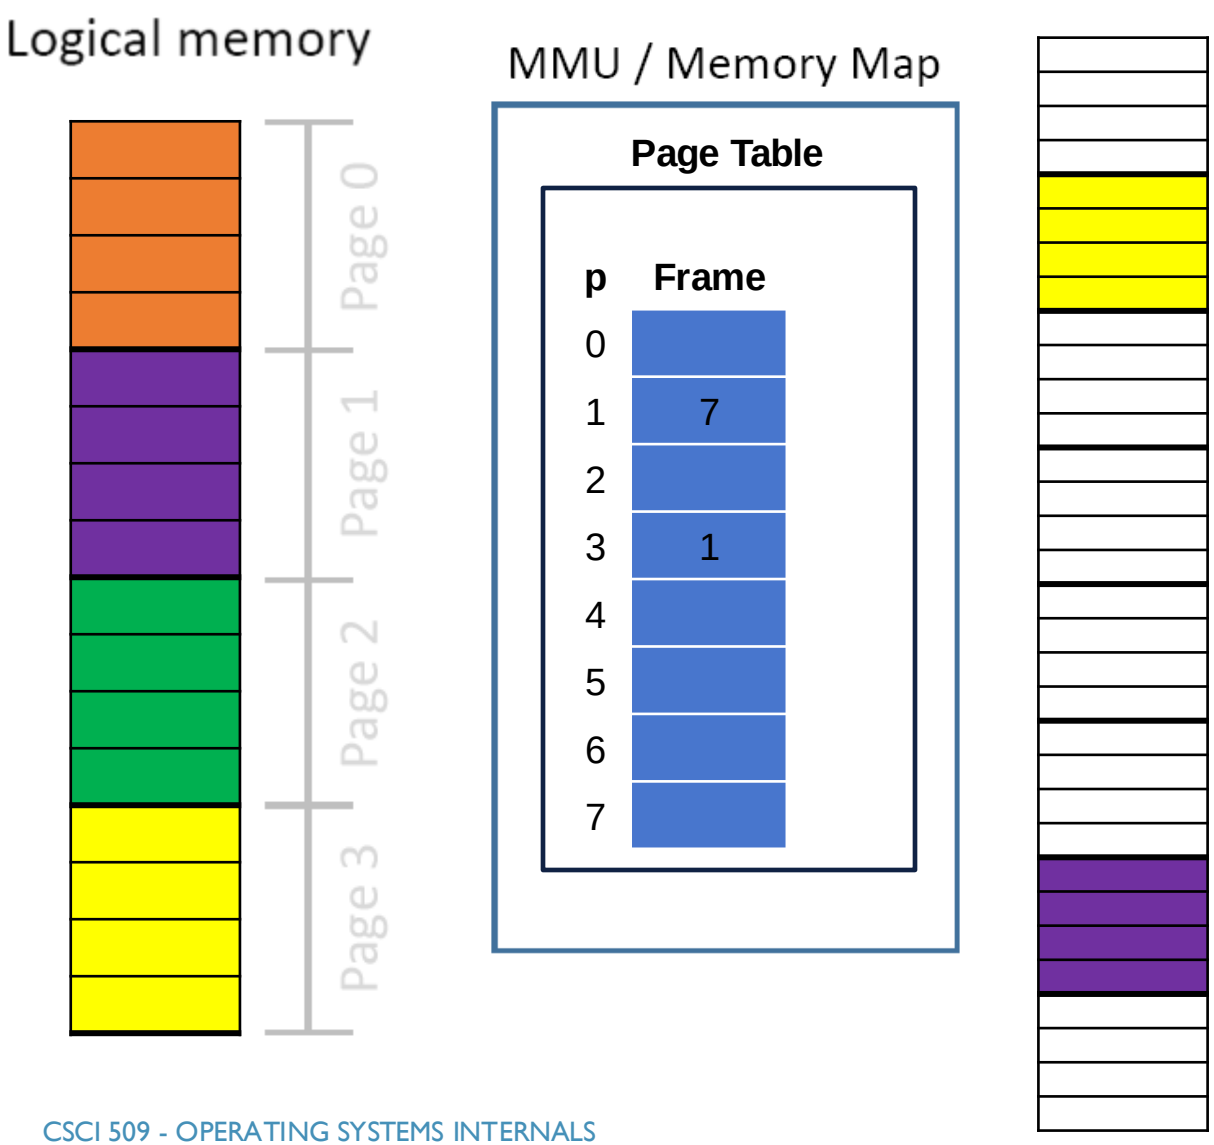

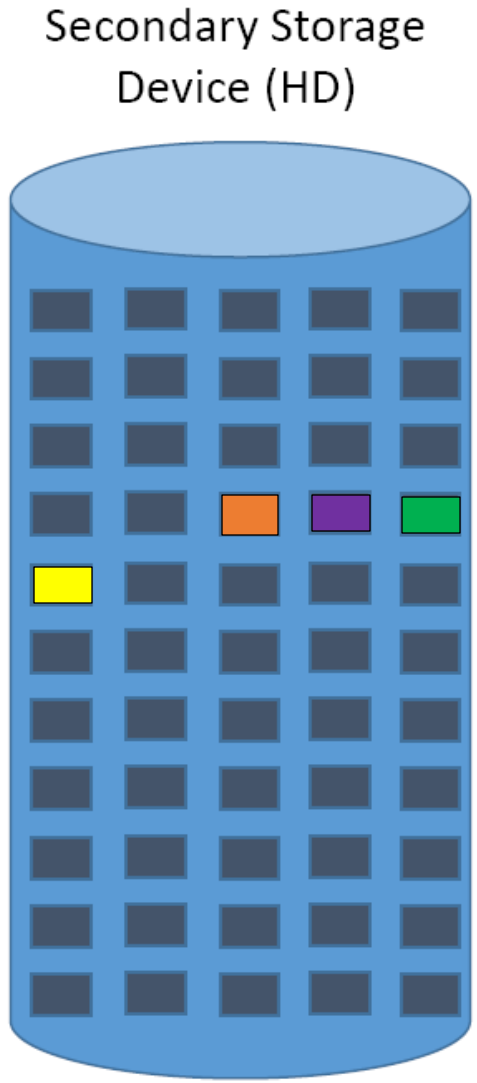

FRAME ALLOCATION

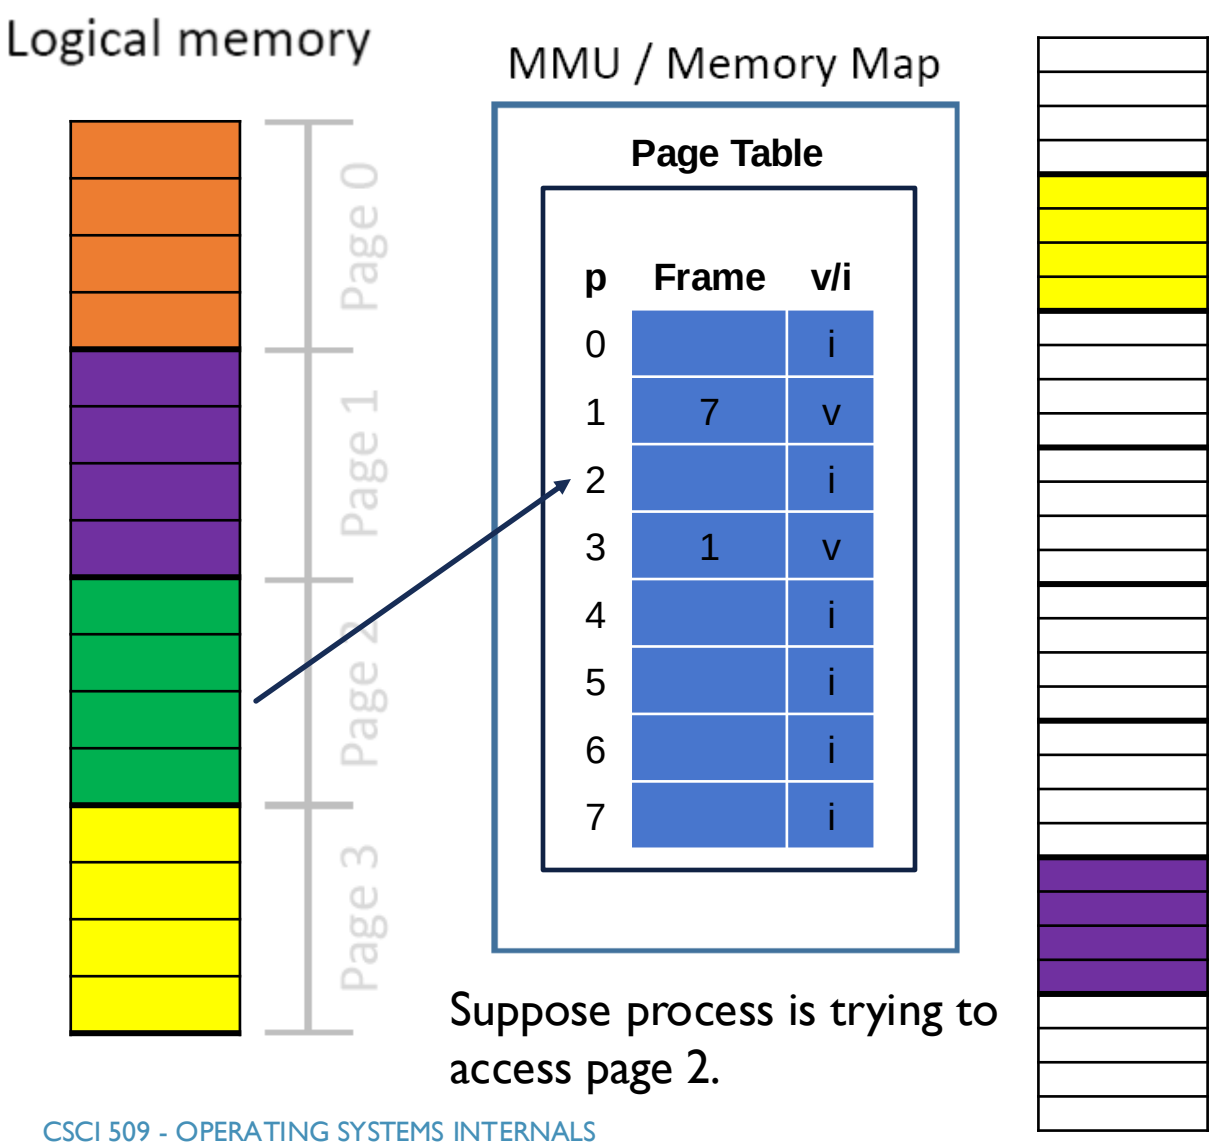

Suppose process is trying to access page 2.

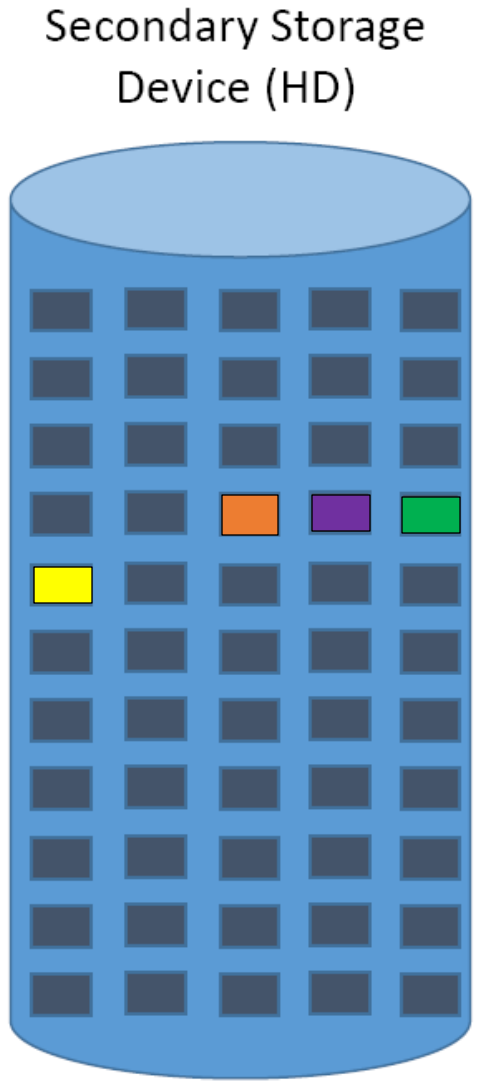

FRAME ALLOCATION

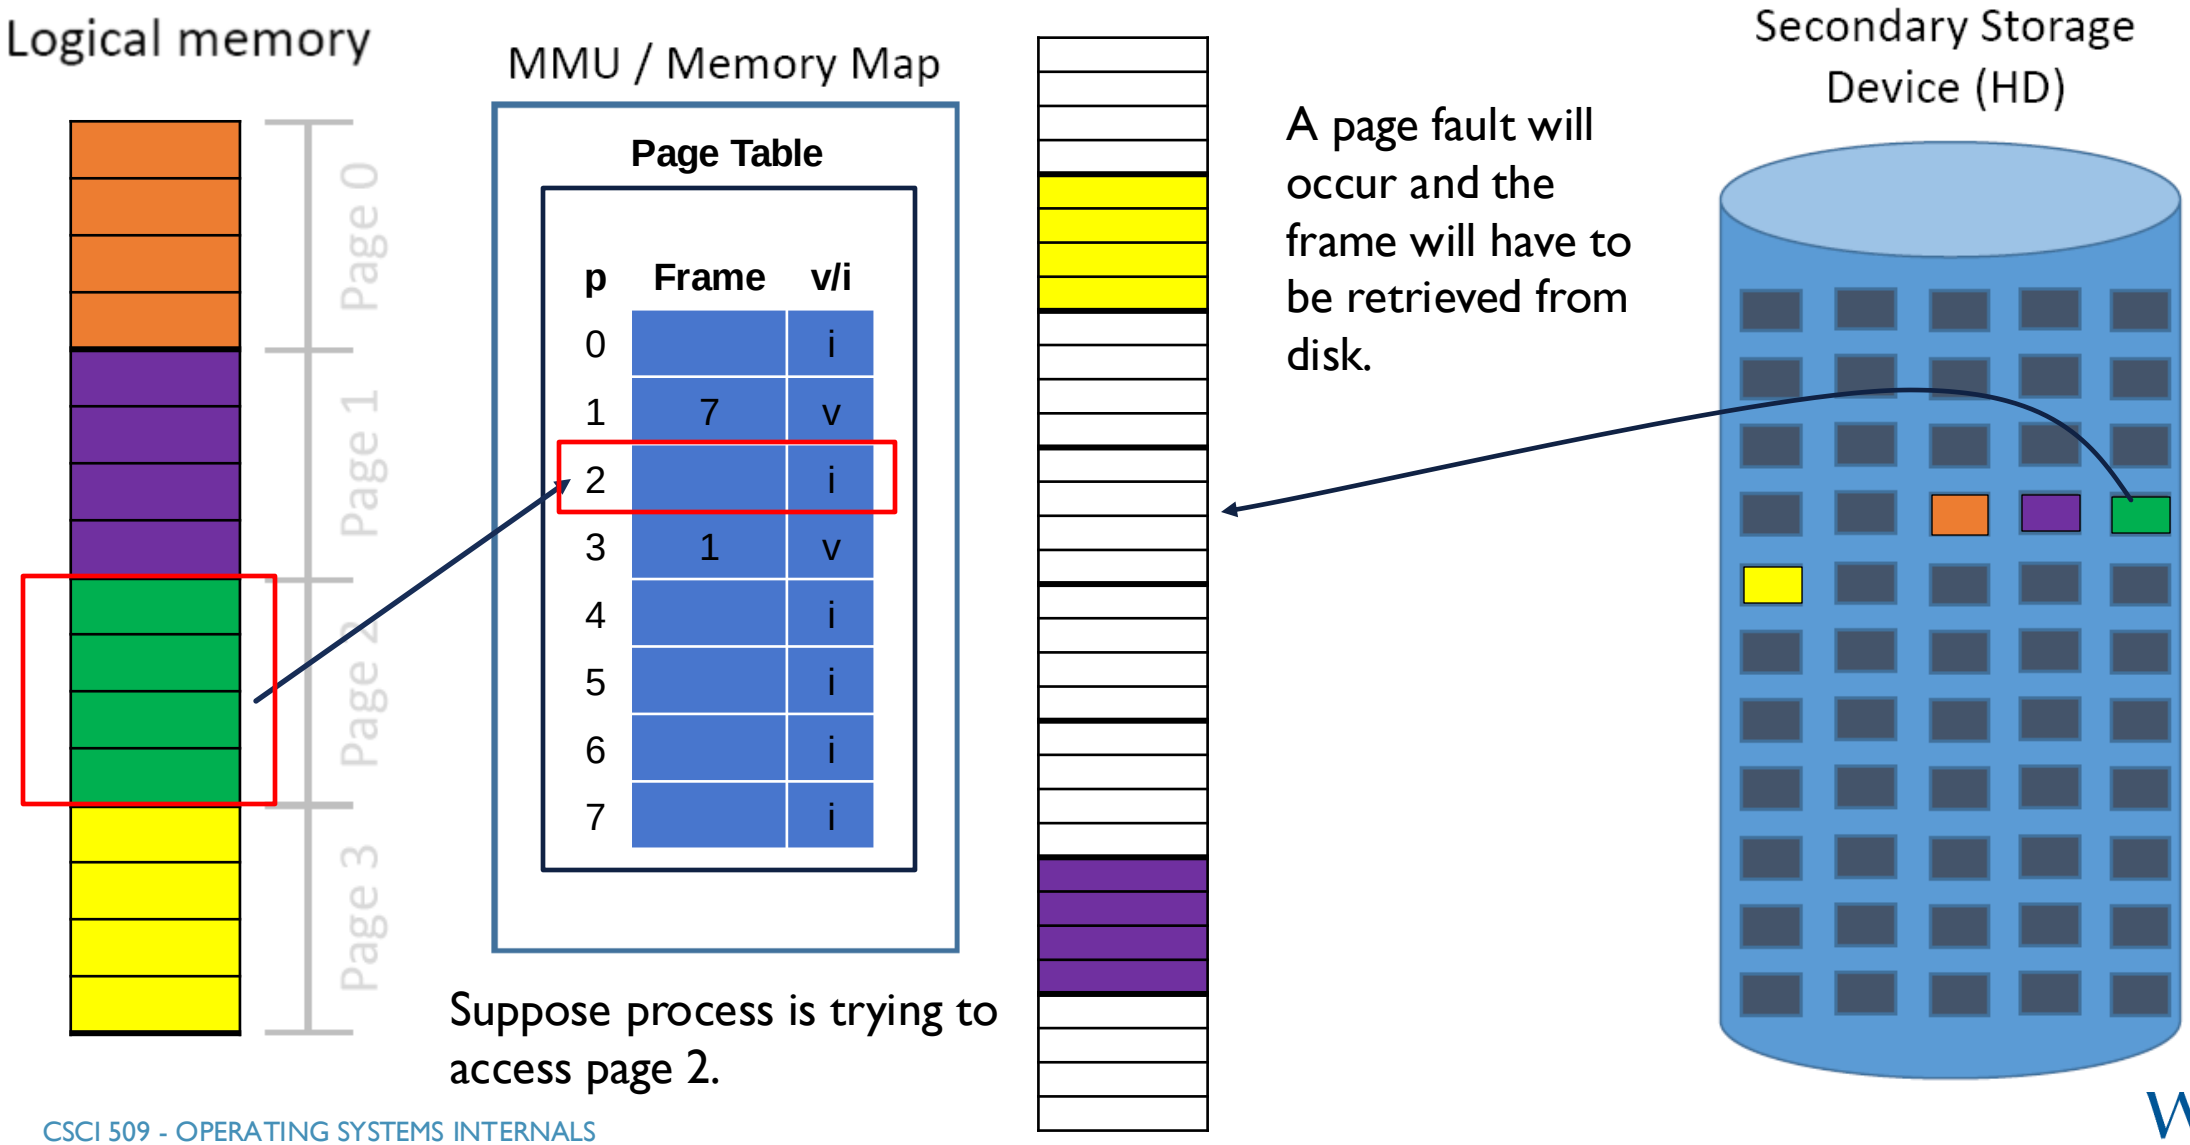

FRAME ALLOCATION

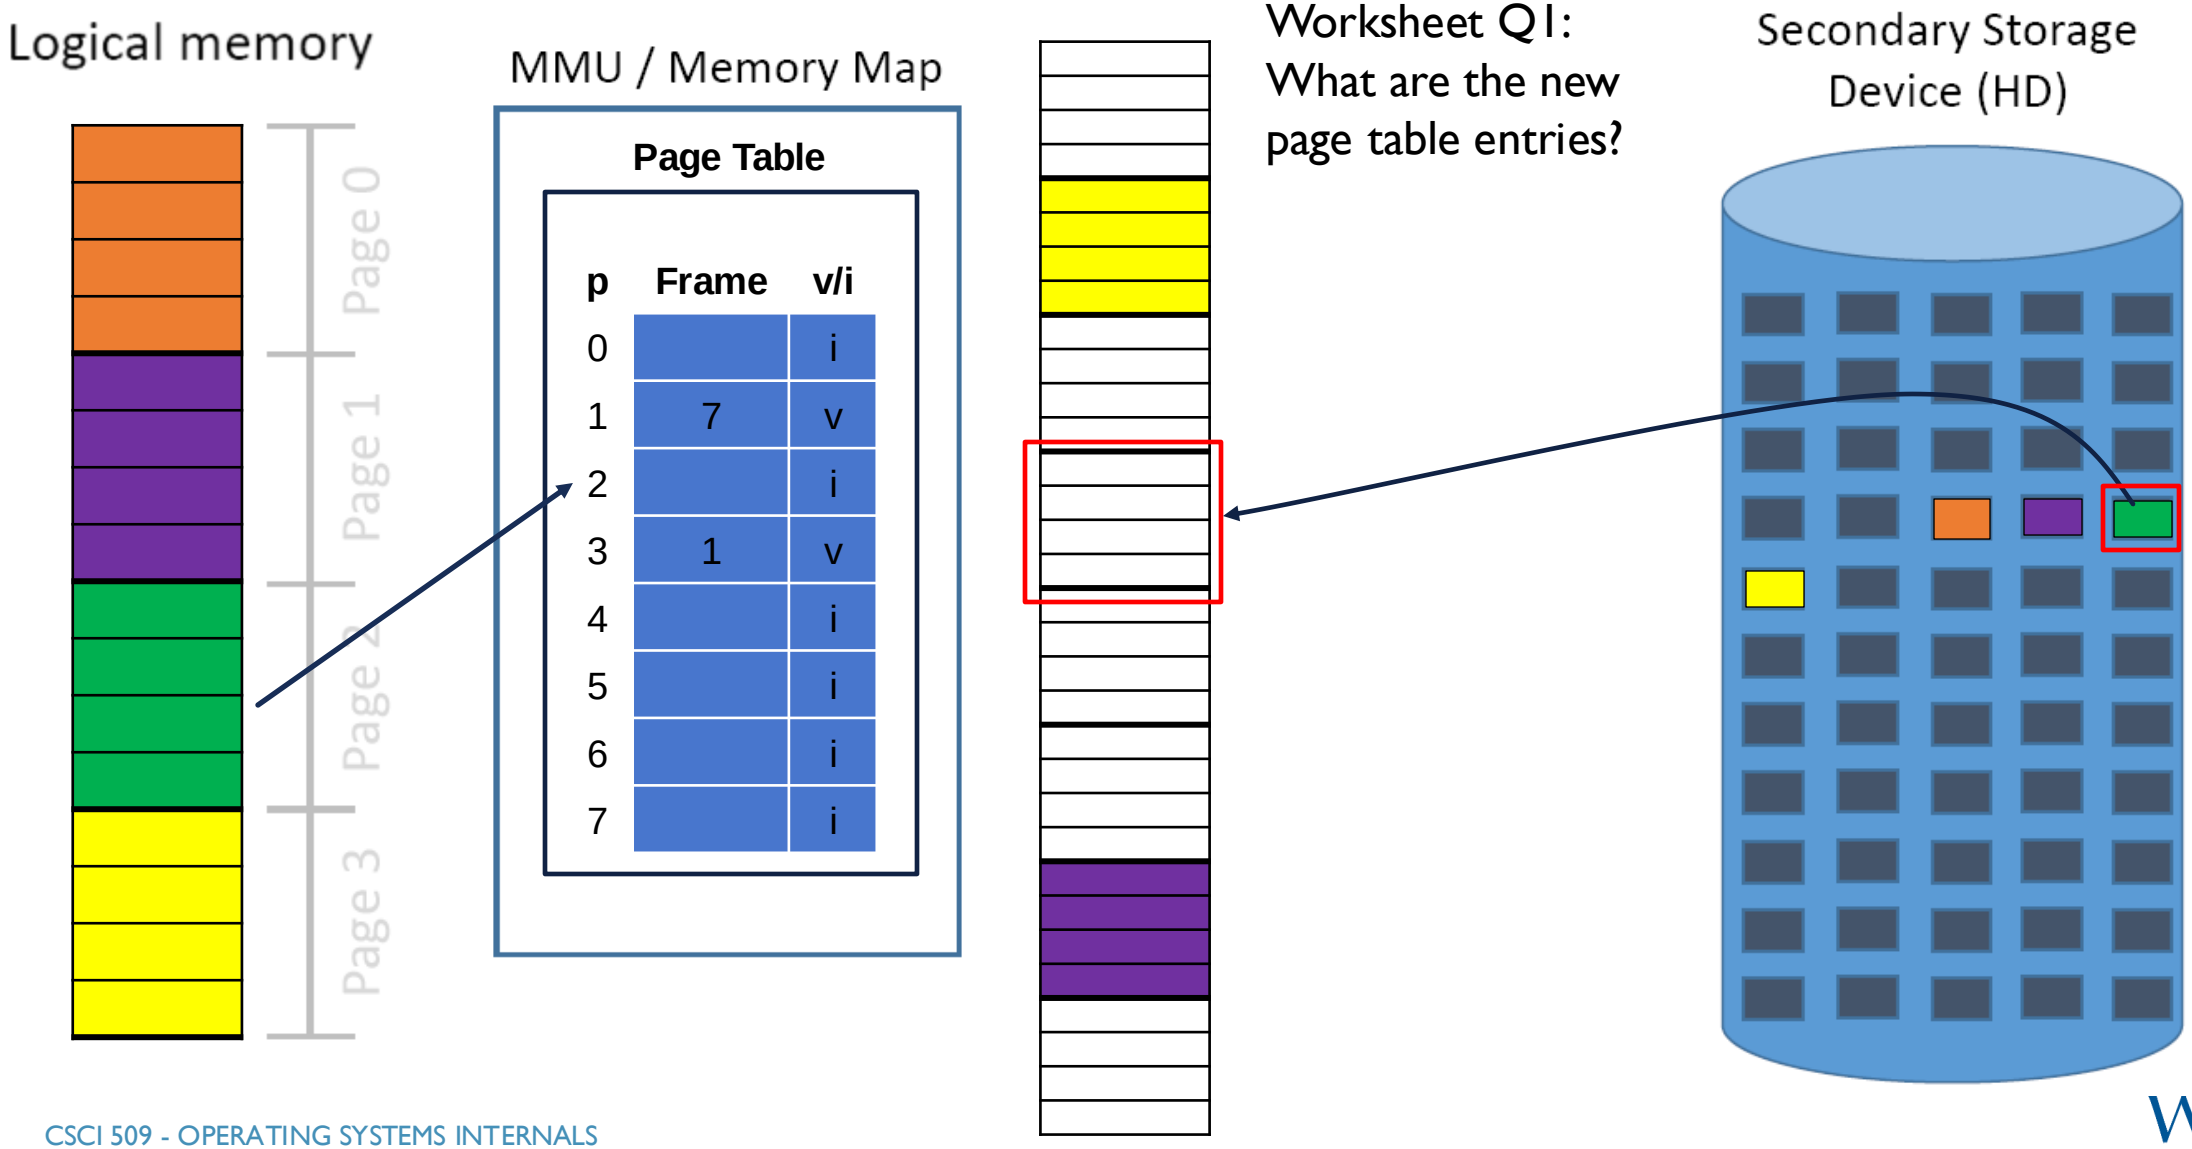

FRAME ALLOCATION

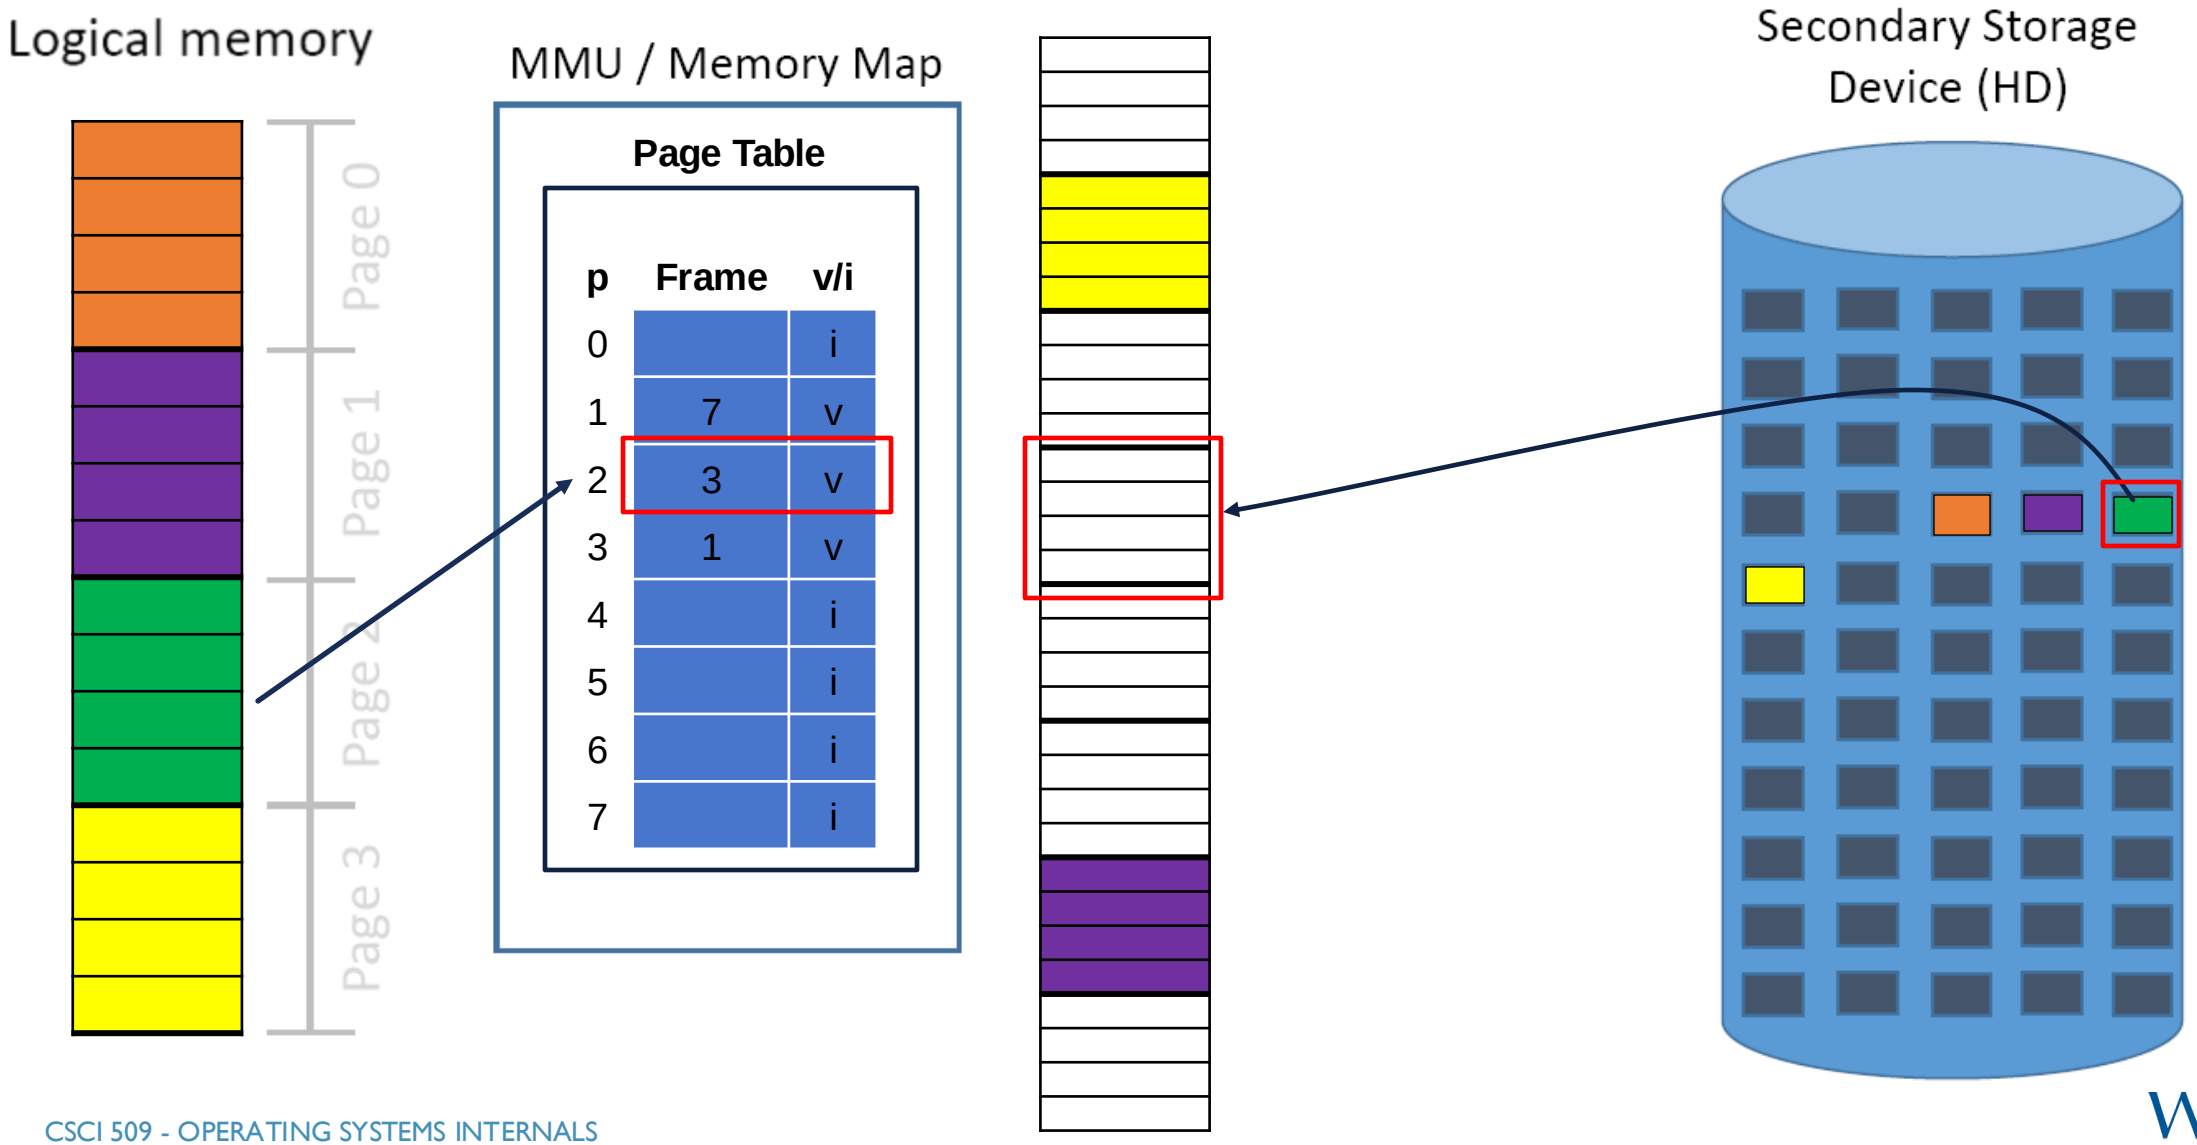

FRAME ALLOCATION

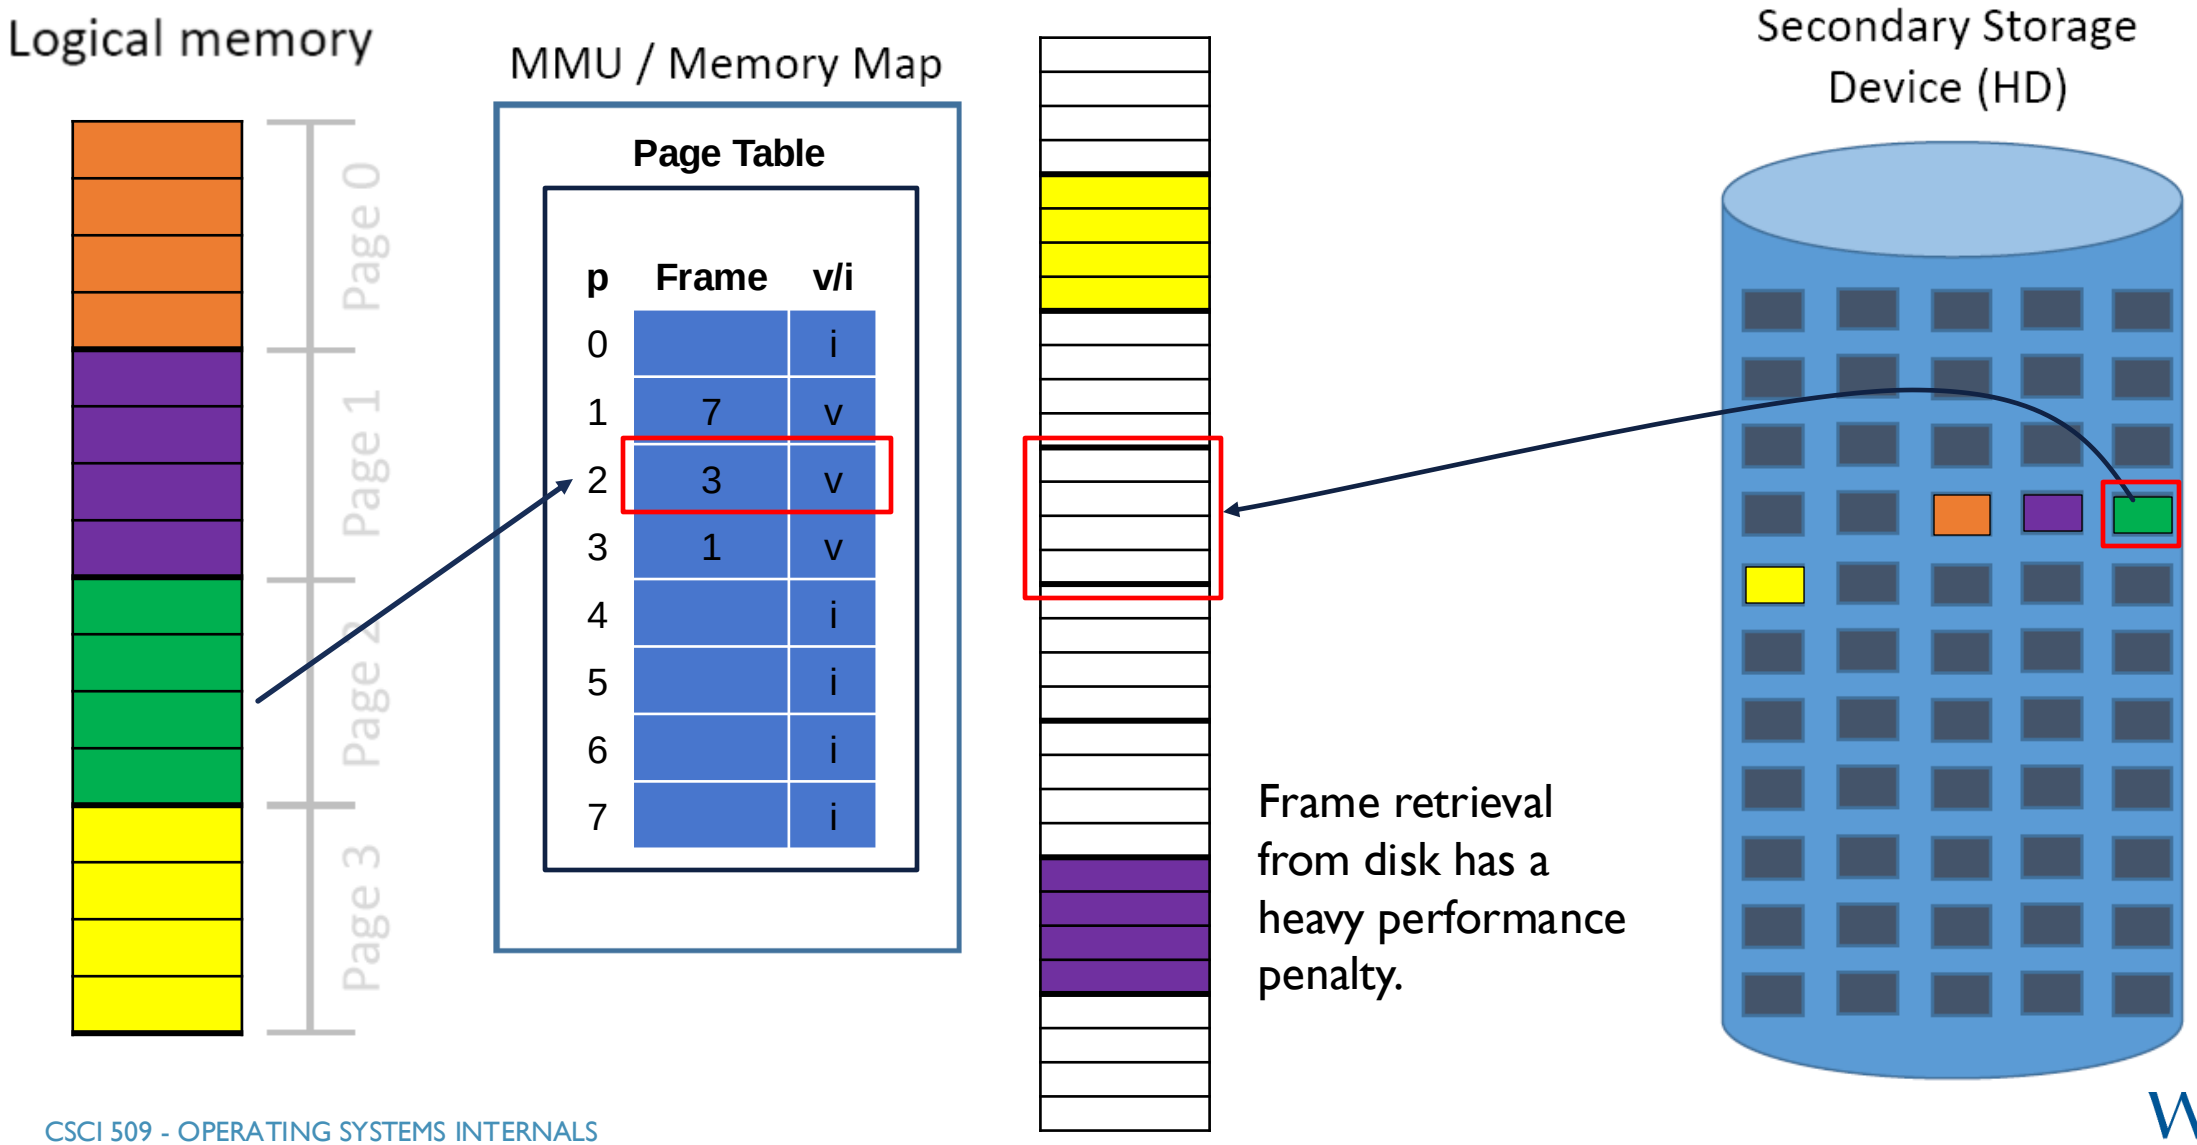

Frame retrieval from disk has a heavy performance penalty.

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

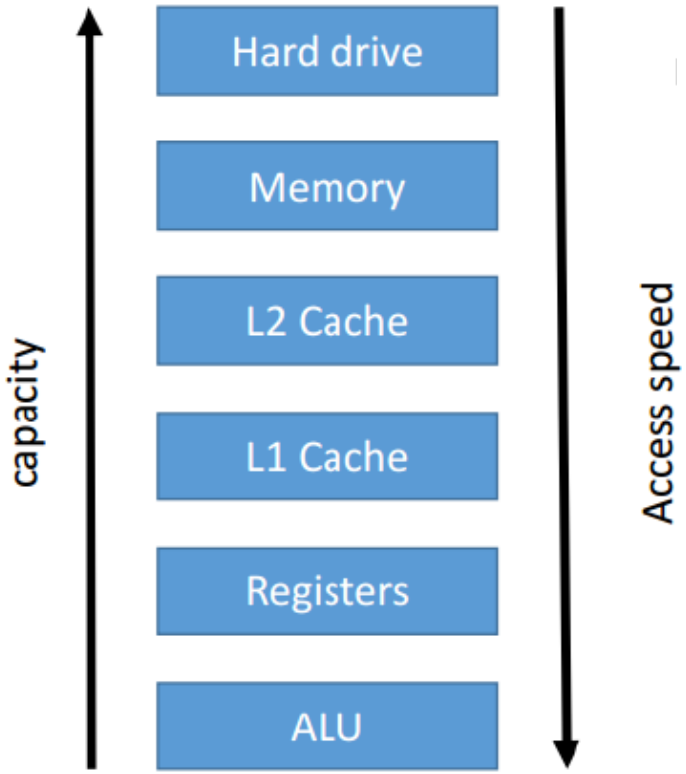

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

A : 1 millisecond
B : 3 nanoseconds
C : 2 microseconds
D : 1 picosecond

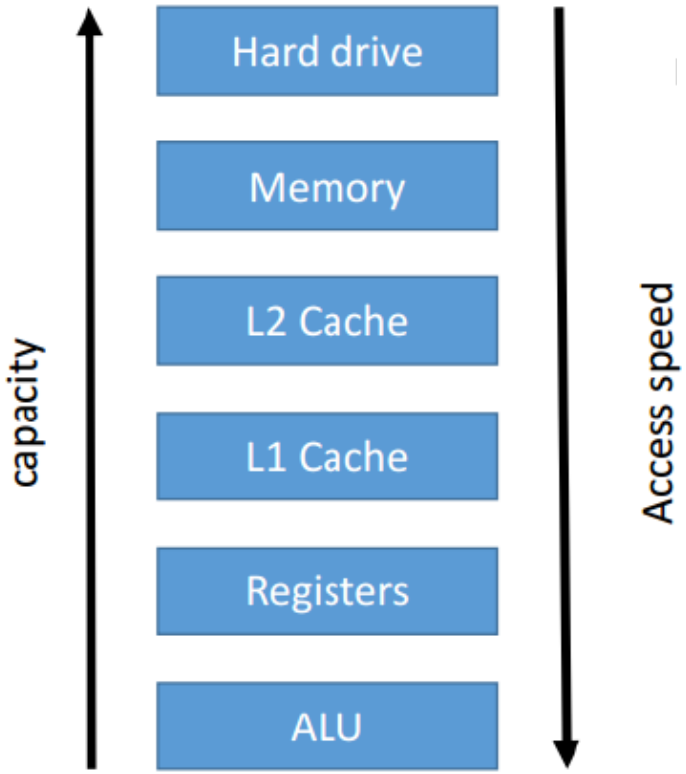

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

A : 1 millisecond
B : 3 nanoseconds
C : 2 microseconds
D : 1 picosecond

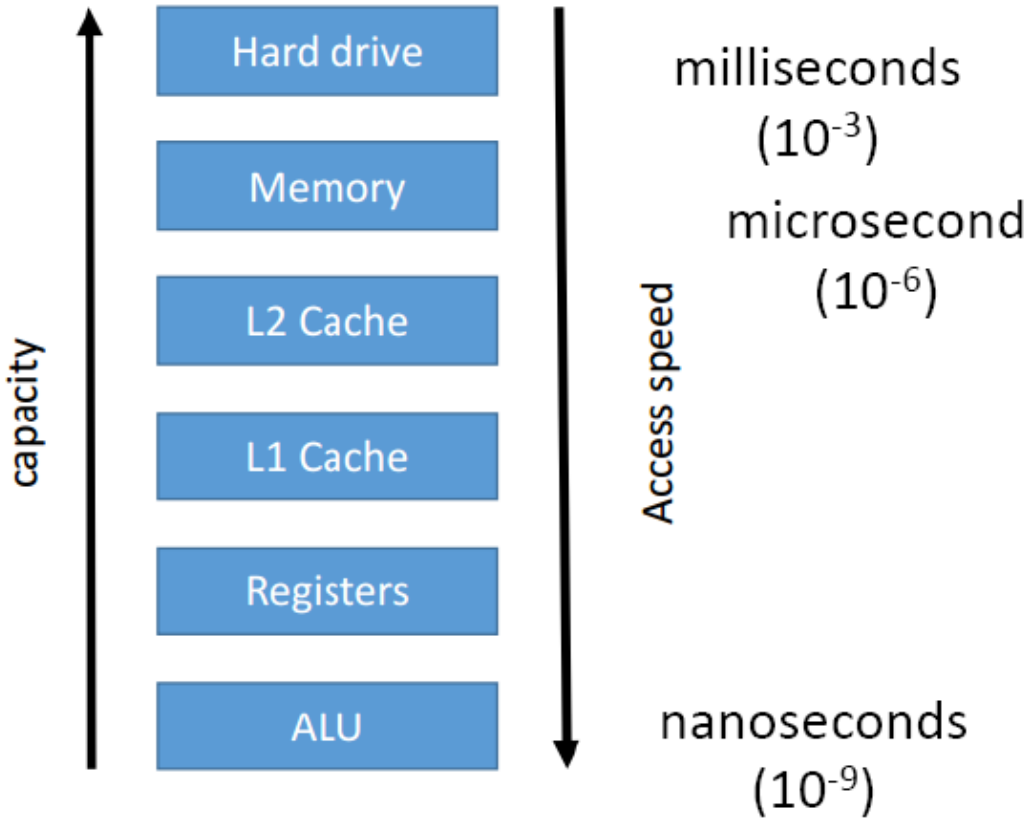

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

The approximate steps taken

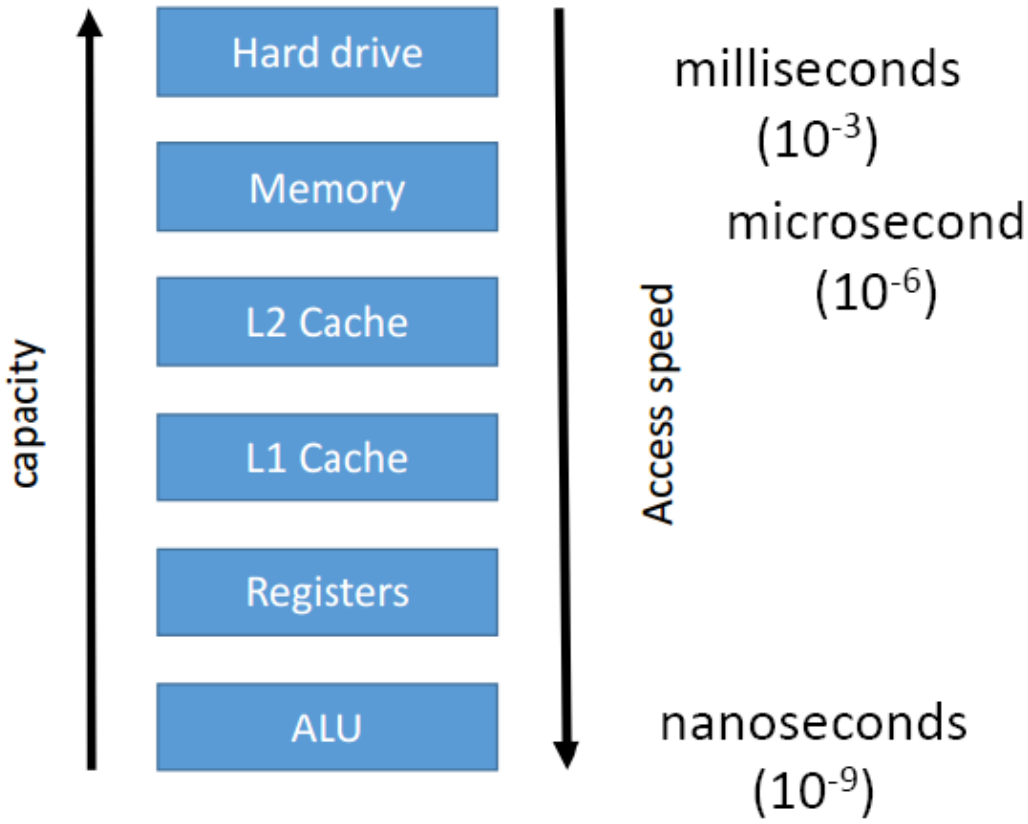

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

Servicing the page fault means interrupting the process, and saving the process's state (PCB)

Context Switch

~50 microseconds

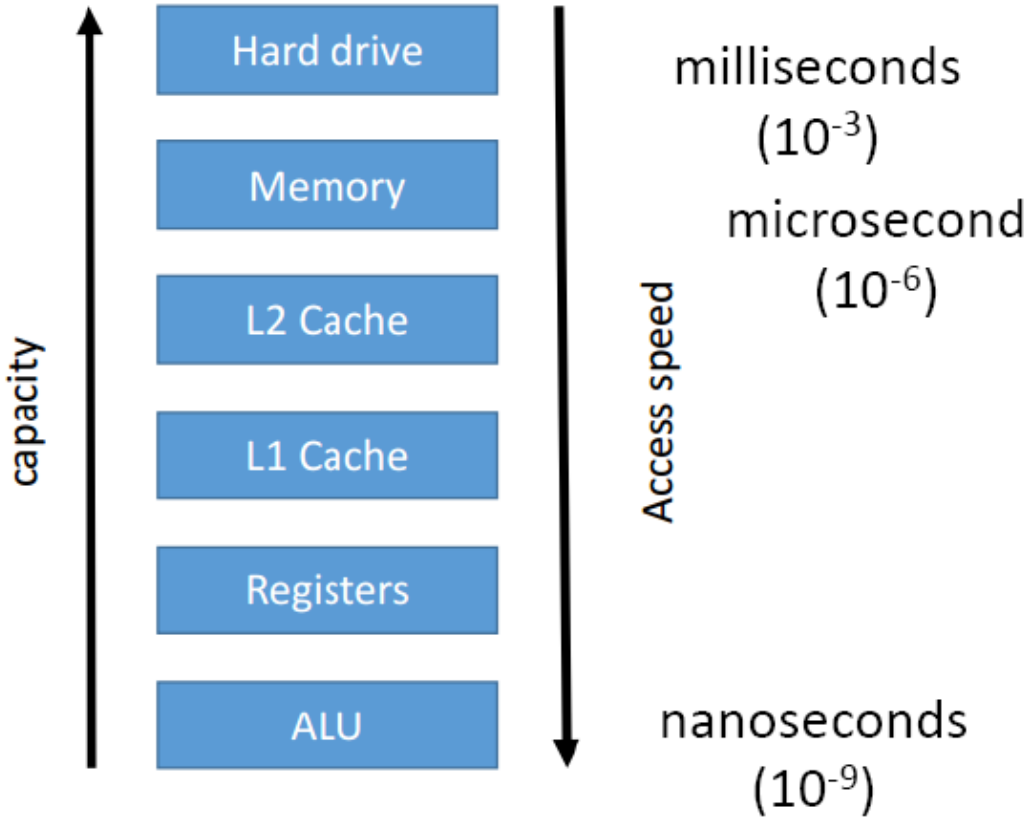

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

~50 microseconds

Issue read (I/O), and schedule other processes to run

Bring in needed frame corresponding to needed page, update page table, wait for OS to schedule the process again

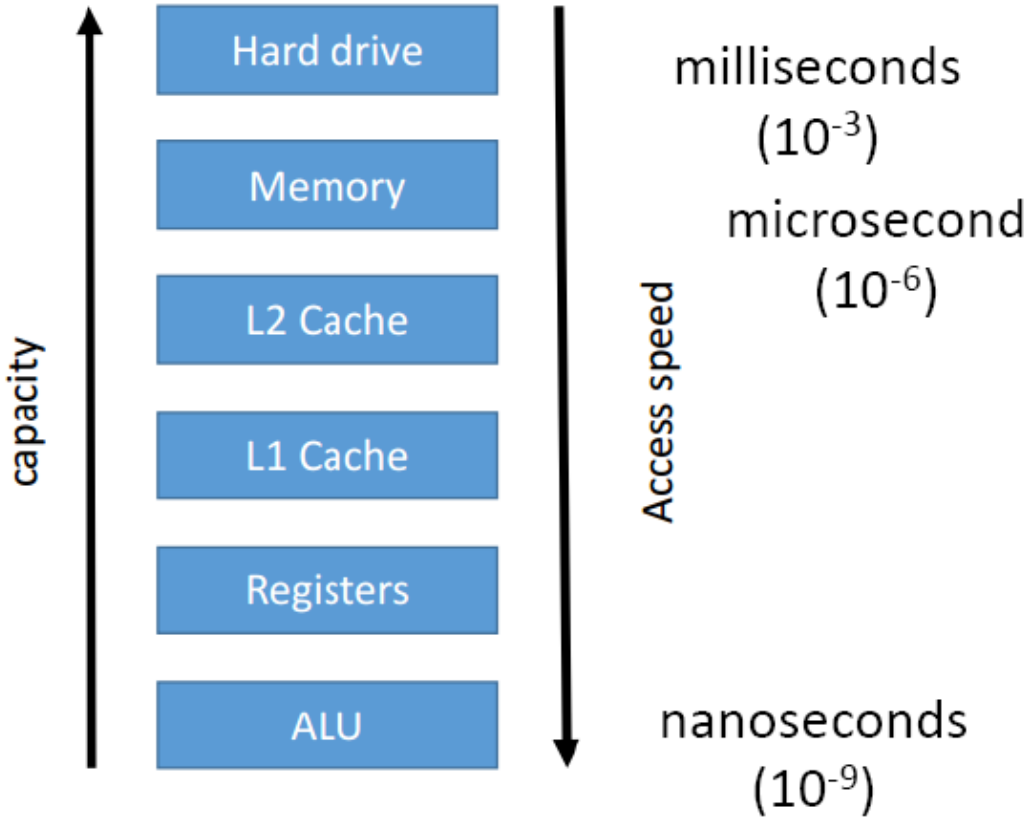

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

Issue read (I/O), and schedule other processes to run

Bring in needed frame corresponding to needed page, update page table, wait for OS to schedule the process again

~50 microseconds + ~5 milliseconds

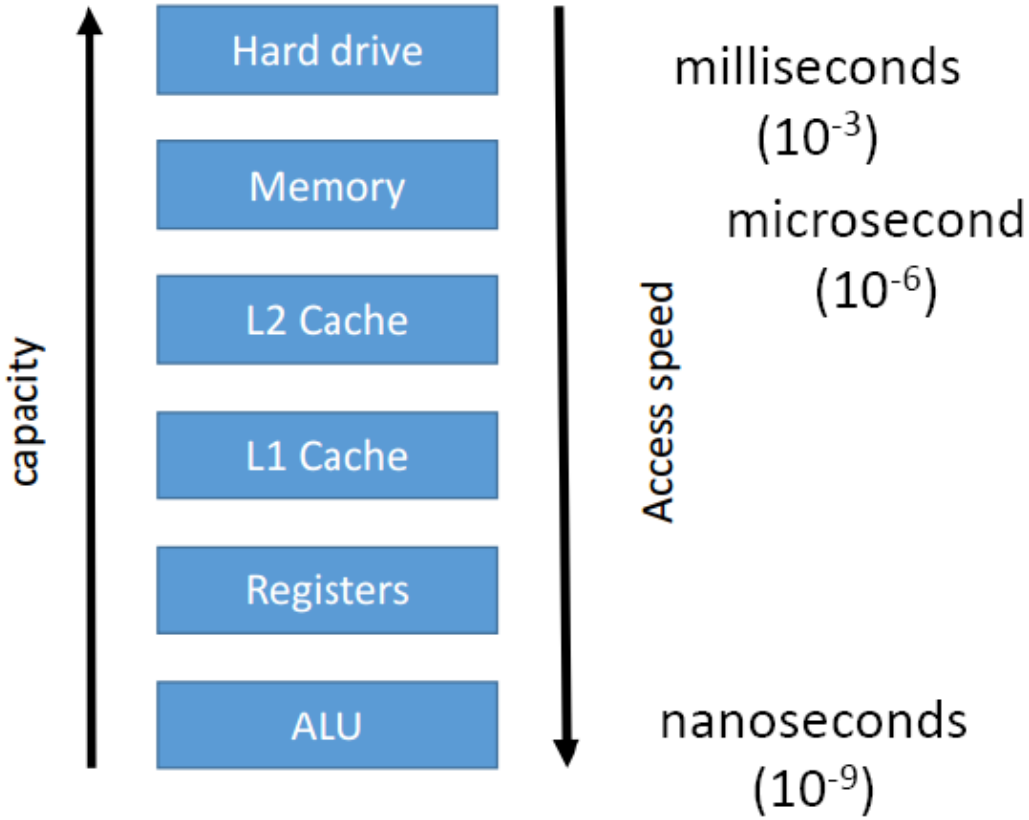

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

Assuming no other processes have a higher priority in the ready queue …

restart the process as soon as "it" is dispatched

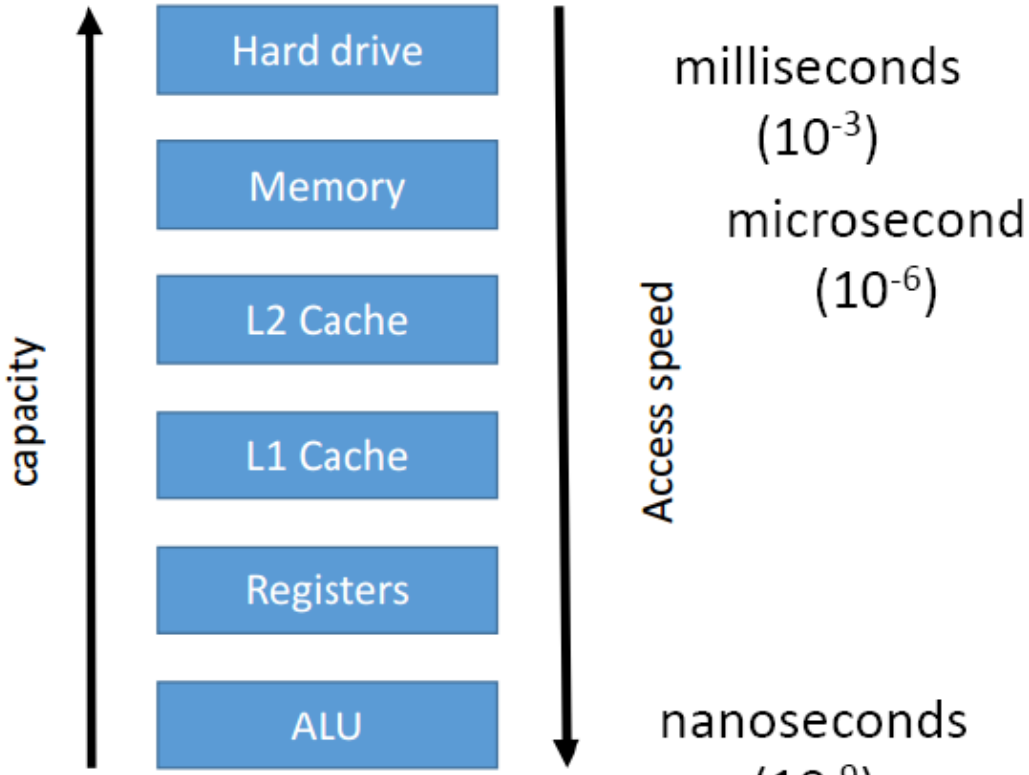

~50 microseconds + ~5 milliseconds + ~50 microseconds

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

Total time for all three, which does not include waiting for dispatch in the ready queue …

About 5 milliseconds

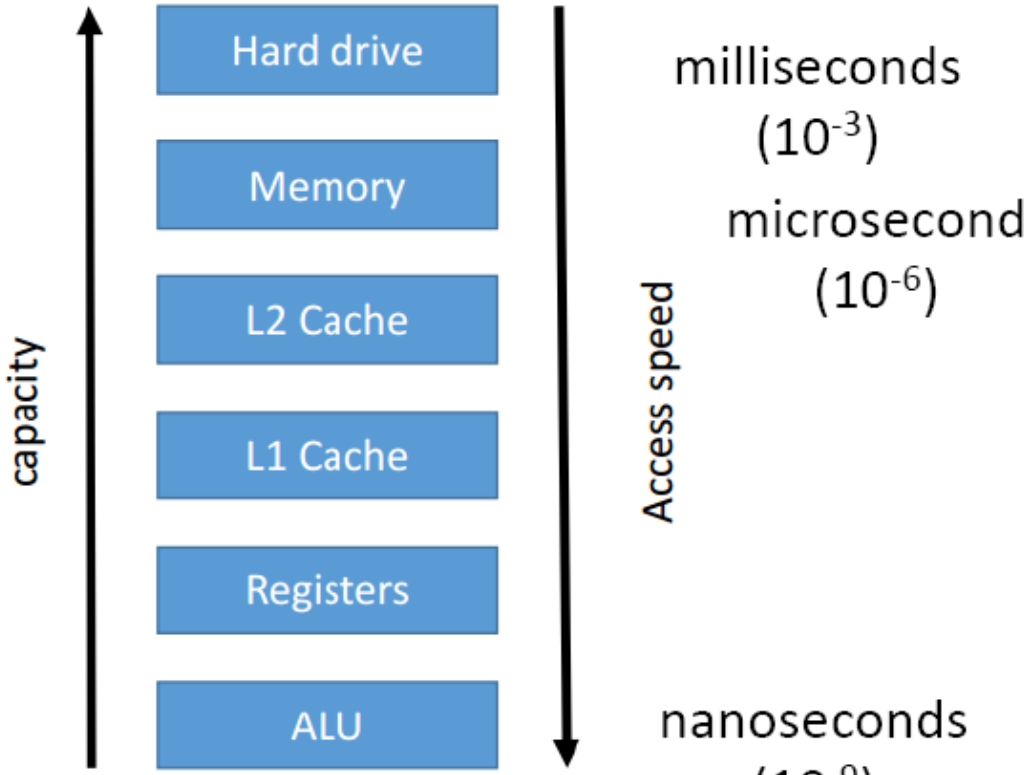

~50 microseconds + ~5 milliseconds + ~50 microseconds = 5.1 milliseconds

PAGE FAULT PENALTY

Ideal Case : all of the pages needed by a program are in physical memory

Real-world Case : Swapping in (and out) must occur

Q: What time delay (penalty) does a program incur due to demand paging in response to a single page fault?

- Service page fault
- swap in the needed frame
- Restart the process

Total time for all three, which does not include waiting for dispatch in the ready queue …

About 5 milliseconds

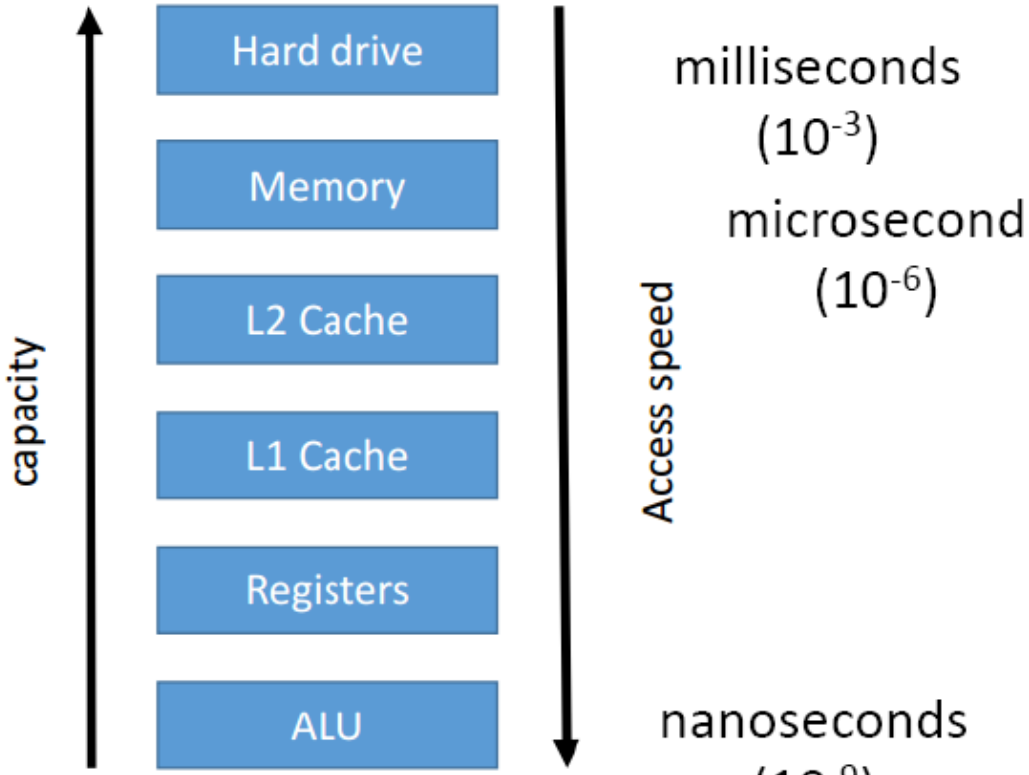

~50 microseconds + ~5 milliseconds + ~50 microseconds = 5.1 milliseconds

PAGE REPLACEMENT

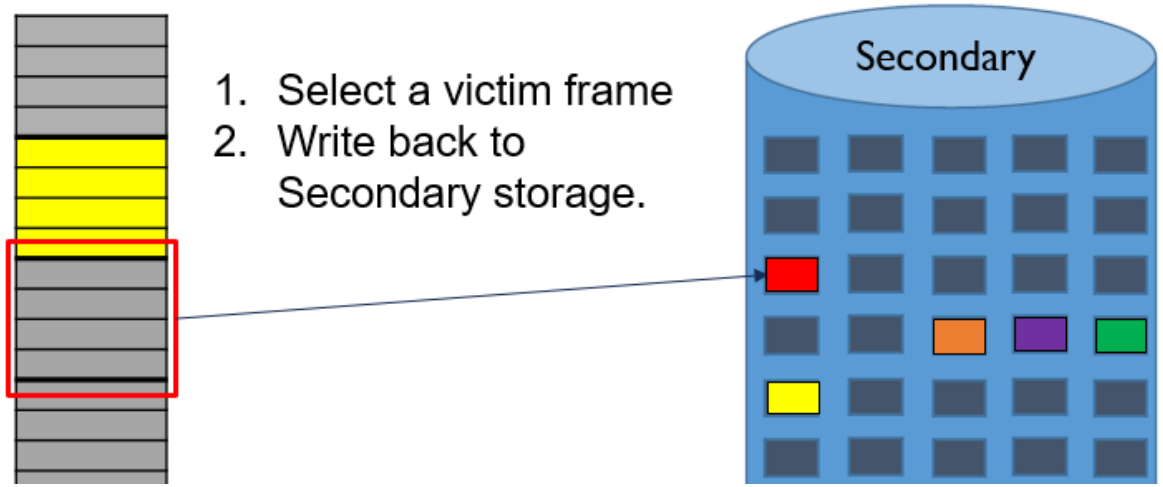

PAGE REPLACEMENT

- How is the victim selected?

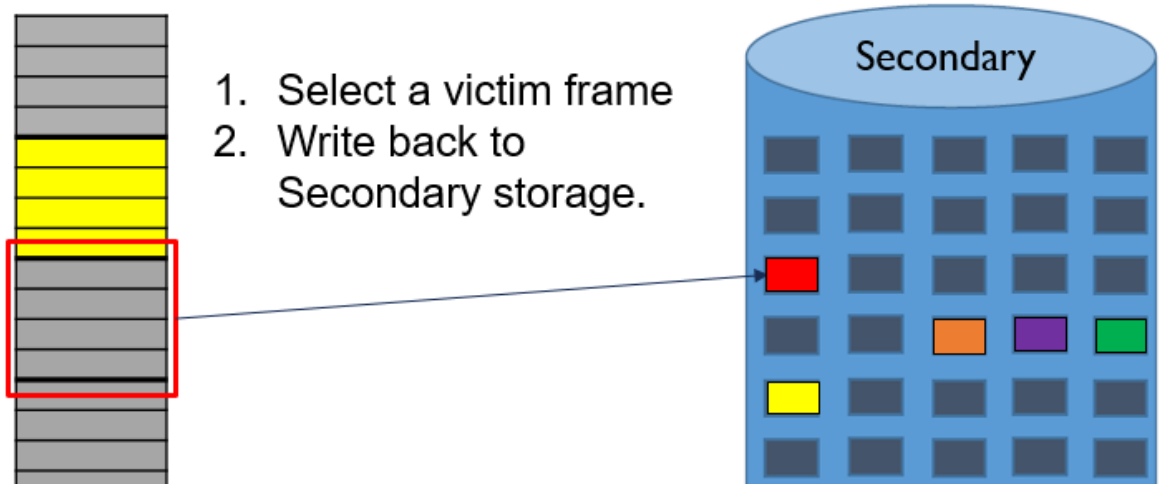

PAGE REPLACEMENT

- The choice of the "victim" page to select is important.
- How is the victim selected?

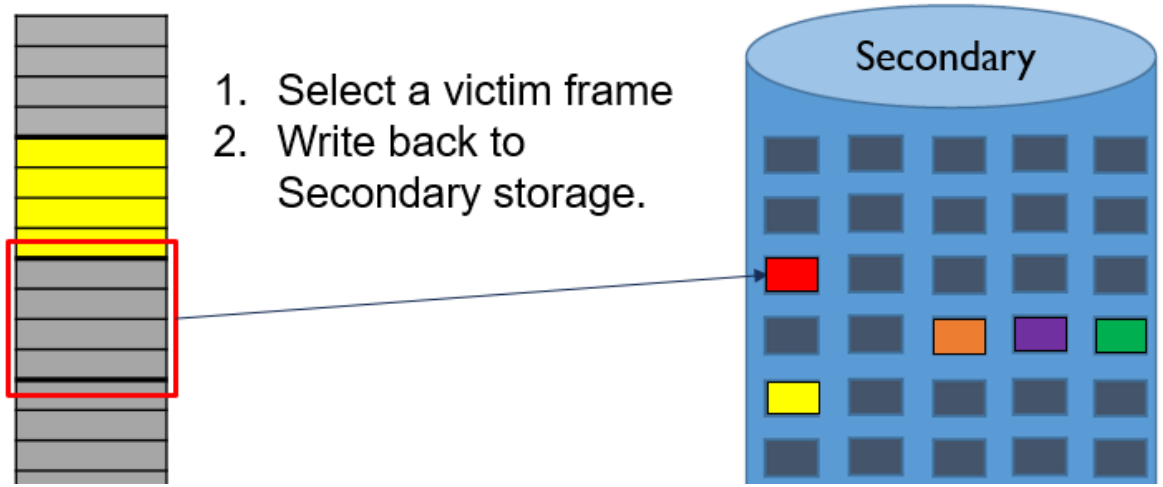

PAGE REPLACEMENT

- The choice of the "victim" page to select is important.
- How is the victim selected?
- Worksheet Q2: what are possible criteria or factors for selecting the victim page?

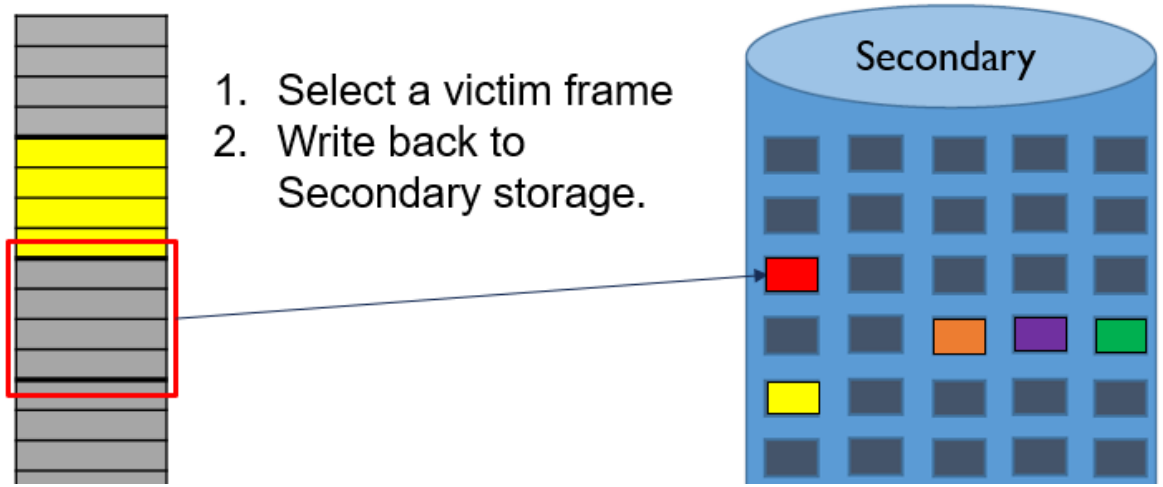

PAGE REPLACEMENT

- The choice of the "victim" page to select is important.
- How is the victim selected?
- Worksheet Q2: what are possible criteria or factors for selecting the victim page?

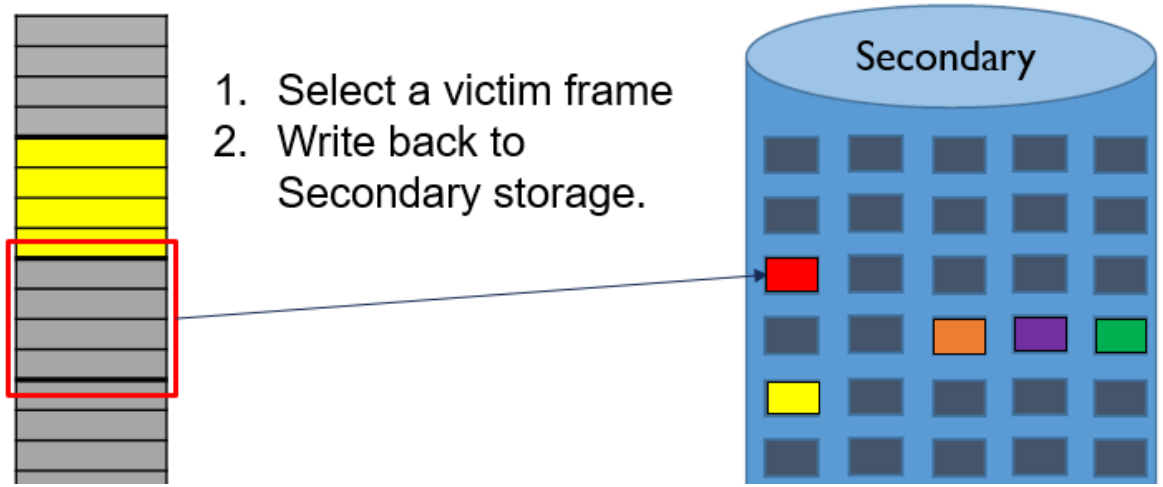

objective: reduce page faults

PAGE REPLACEMENT

- How is the victim selected?
- Victim Selection: FIFO
- Select the oldest page in the frame and remove it.
- How does it work? How does it perform?

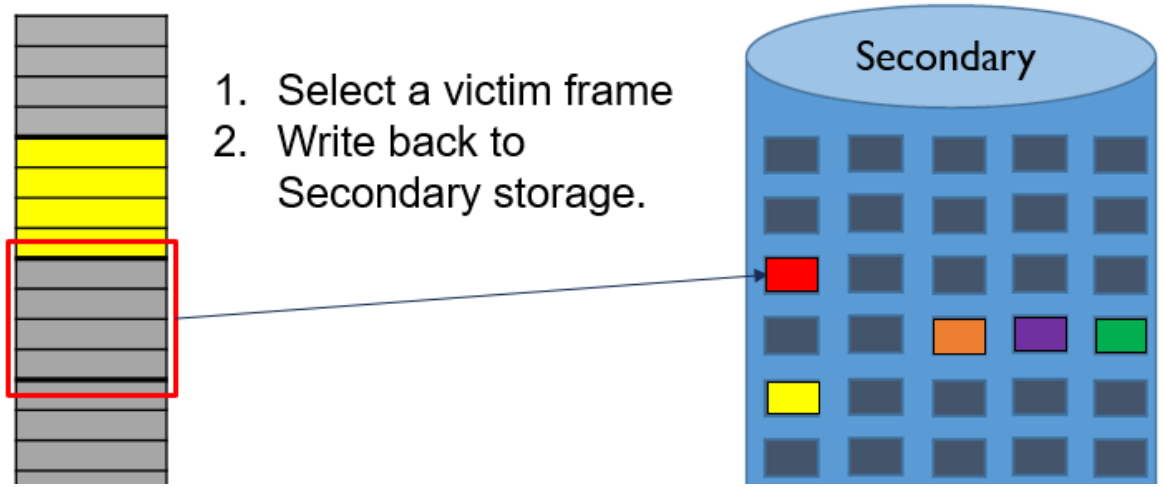

FIFO PAGE REPLACEMENT

- Victim Selection: FIFO

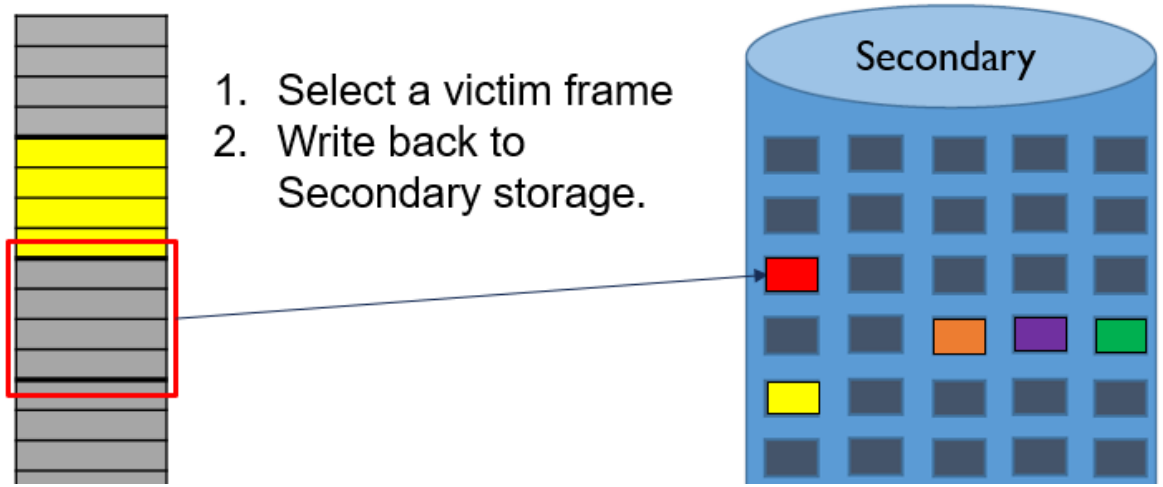

FIFO PAGE REPLACEMENT

- Victim Selection: FIFO
- Select the oldest page in the frame and remove it.
- How does it work? How does it perform?

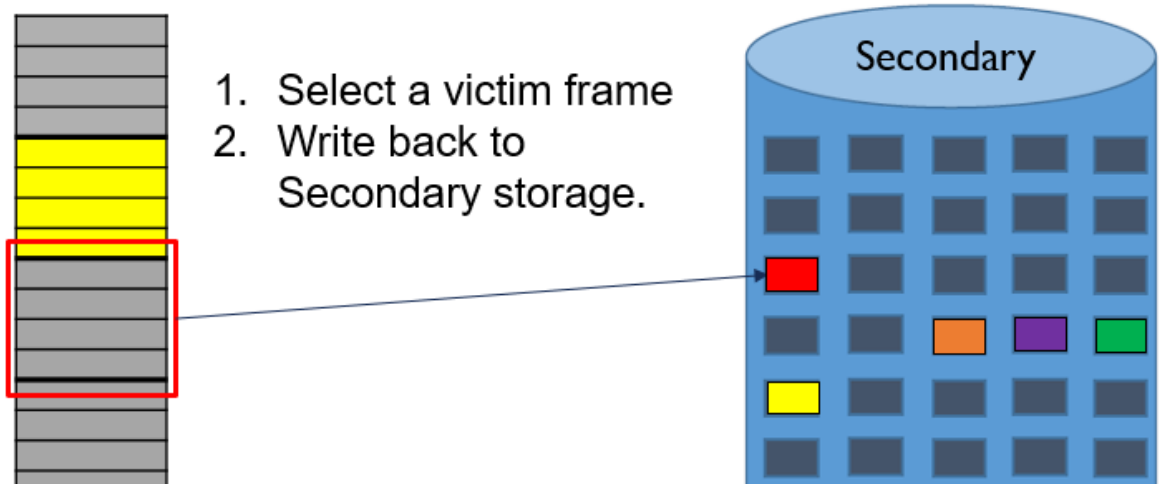

FIFO REPLACEMENT

- Assume a page/frame architecture where each page has **100 bytes.** Assume the following byte (address) requests, left to right (written in base 10) generated by a process.

1011 0656 0692 1466 0605 1141 1222

FIFO REPLACEMENT

- Assume a page/frame architecture where each page has **100 bytes.** Assume the following byte (address) requests, left to right (written in base 10) generated by a process.

1011 0656 0692 1466 0605 1141 1222

Starting with an empty 3-frame memory:

- How will the pages be places?
- What is the number of page faults?

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

10	06	06	14	06	11	12

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

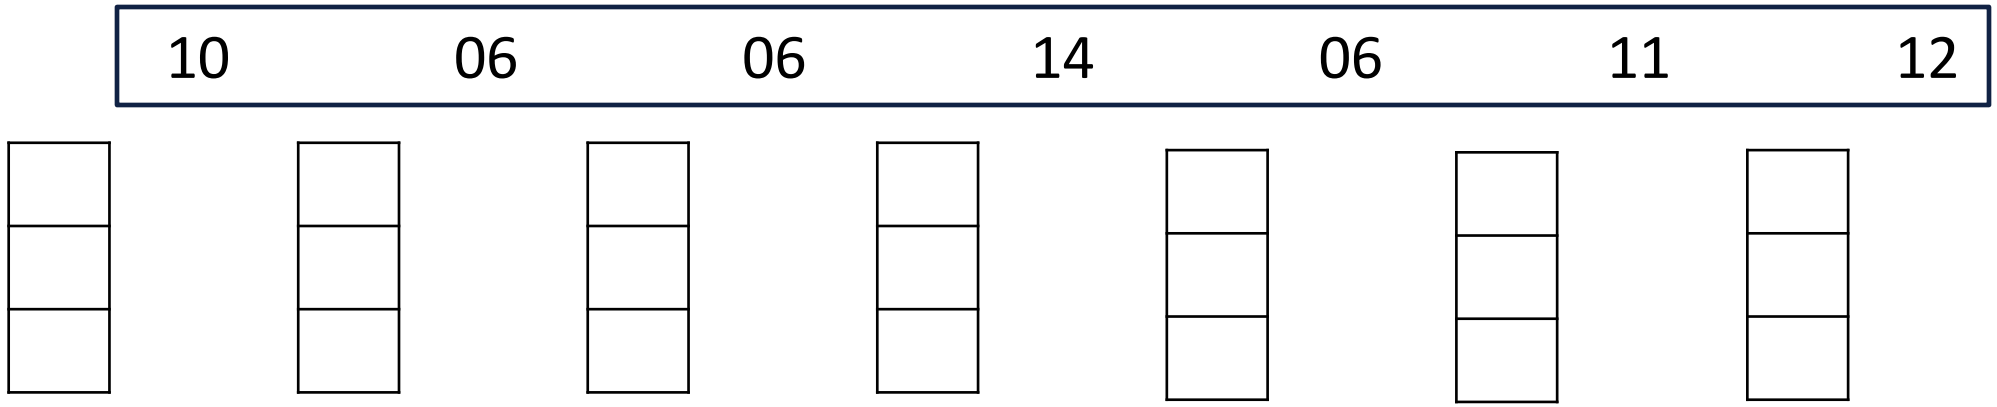

Page Fault Count: 0

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

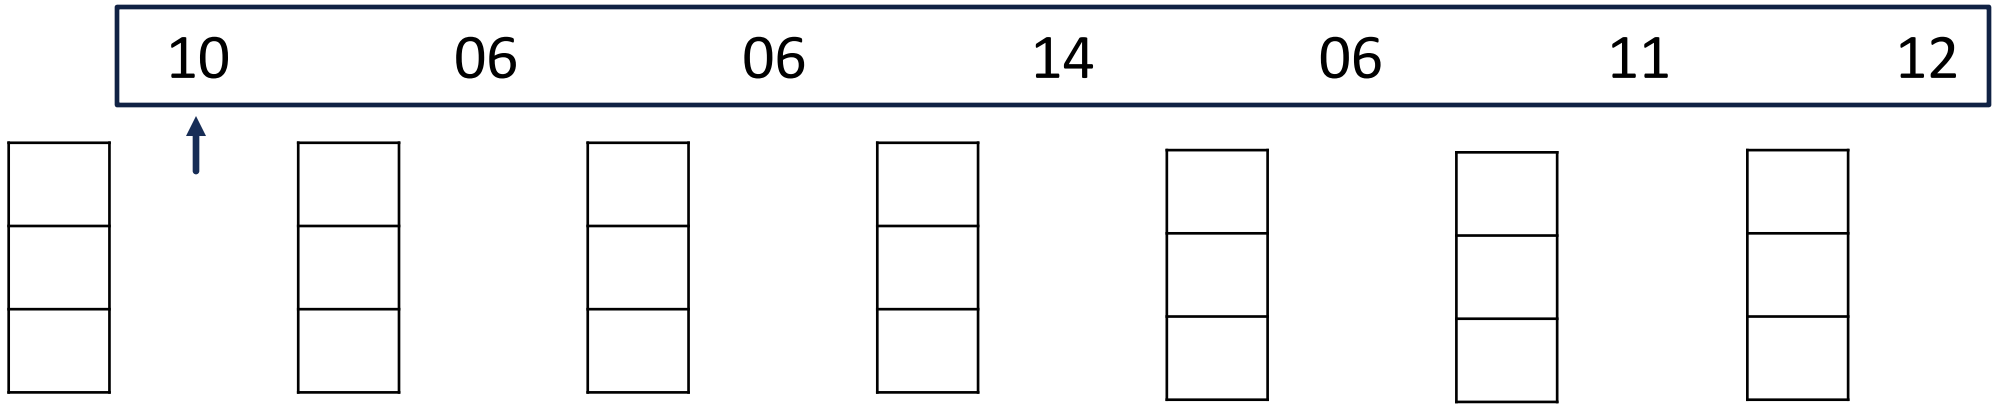

Page Fault Count: 0

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

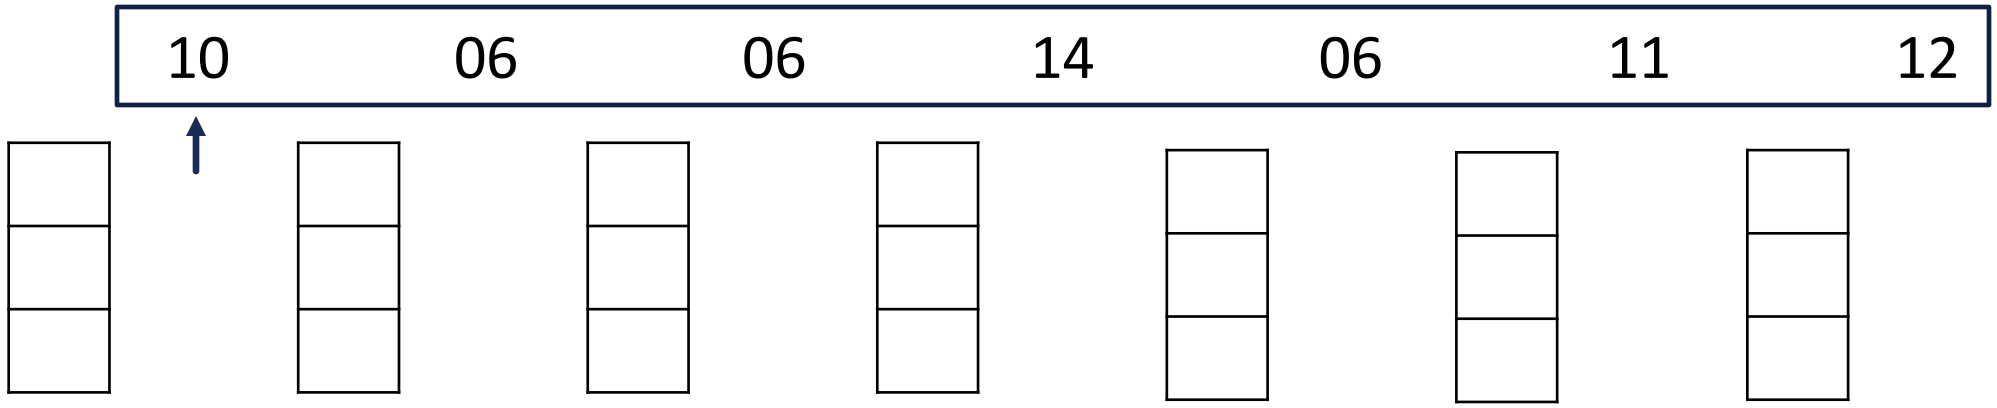

Page Fault Count: 0

Q: Do we get a page fault?

A: Yes
B: No

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

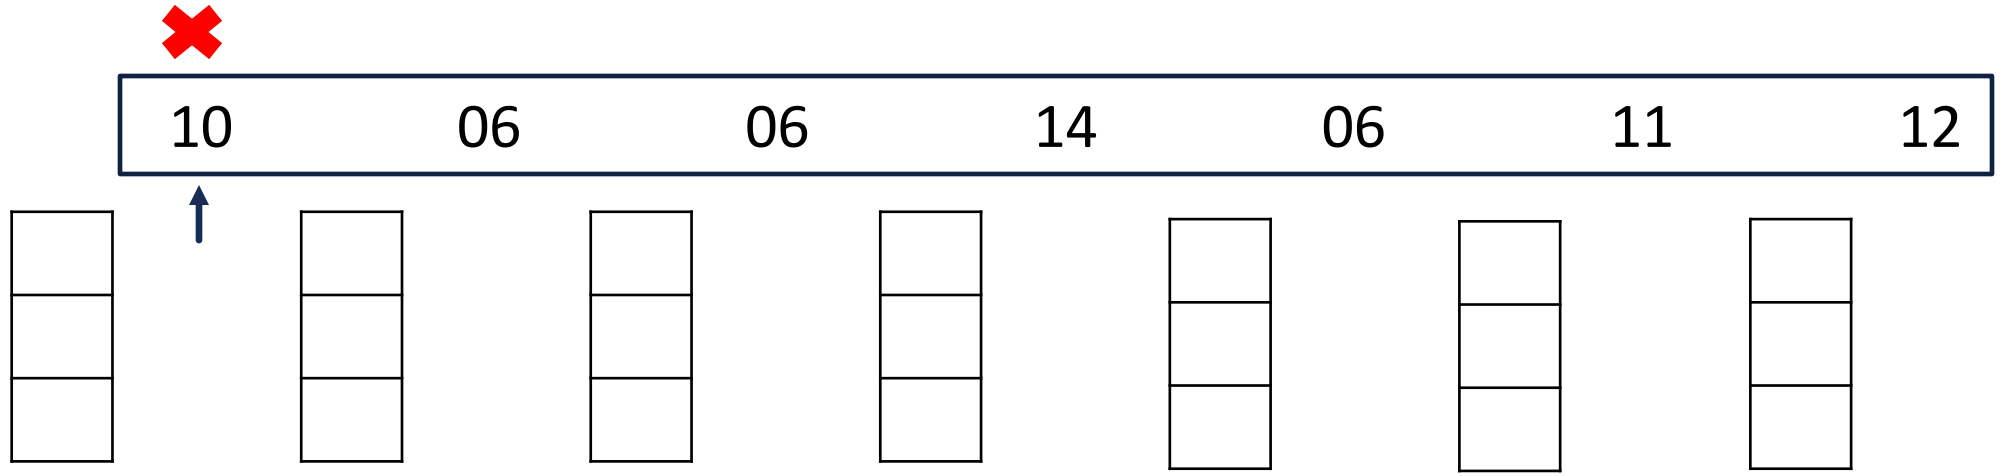

Page Fault Count: 1

Q: Do we get a page fault?

A: Yes
B: No

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

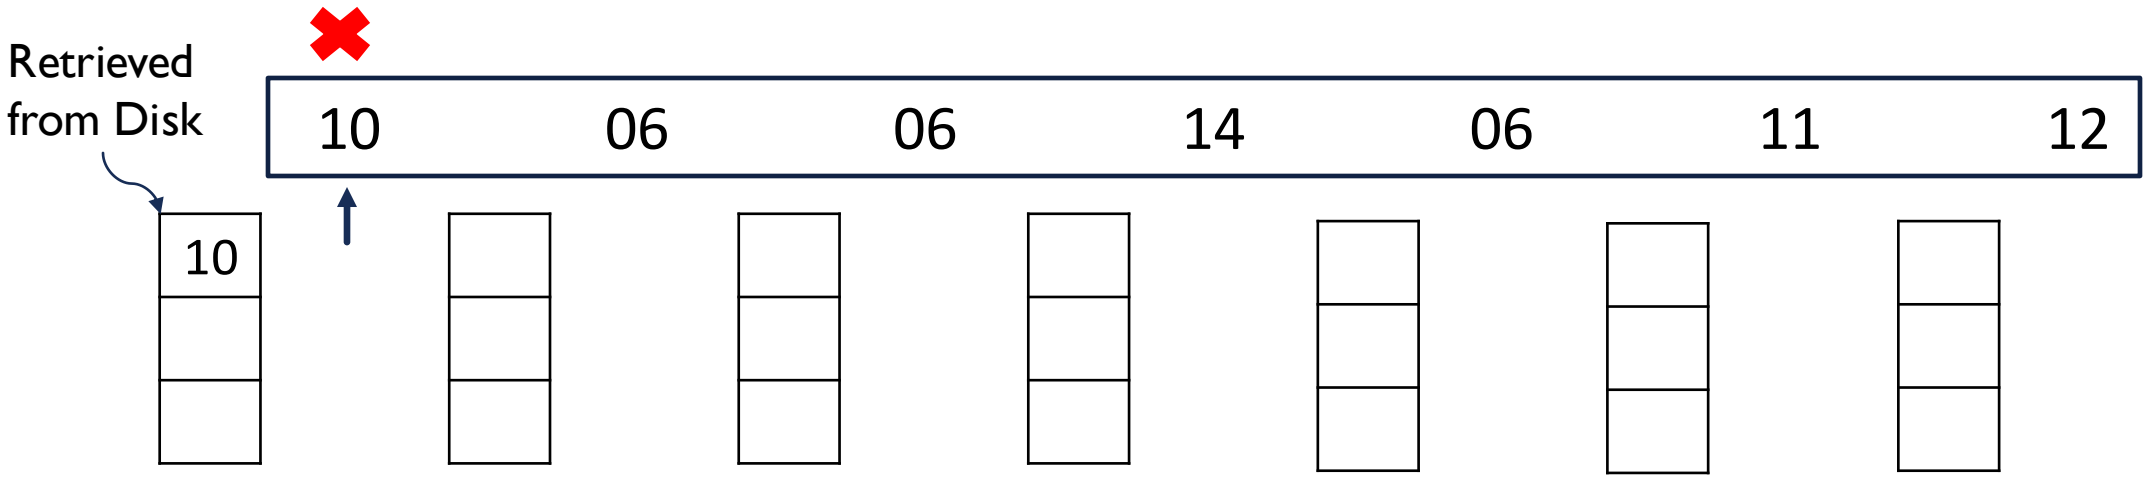

Page Fault Count: 0

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

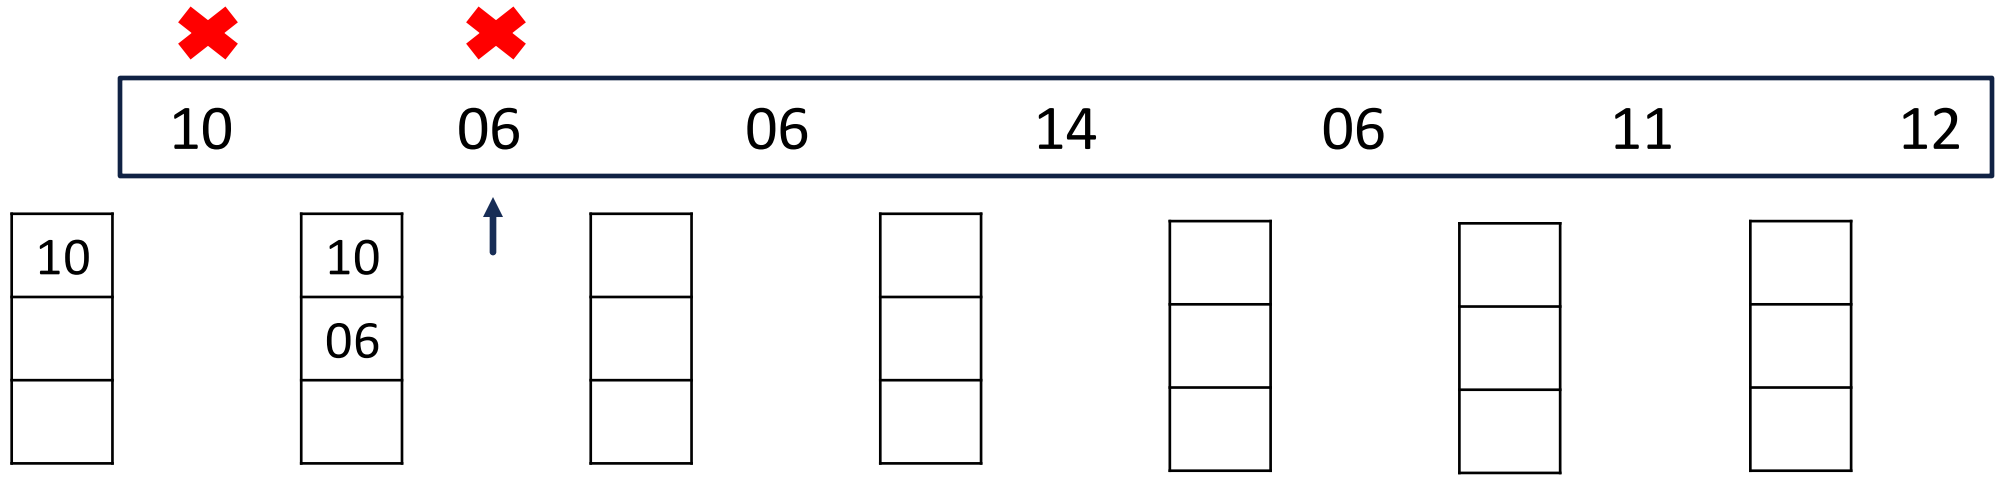

Page Fault Count: 2

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

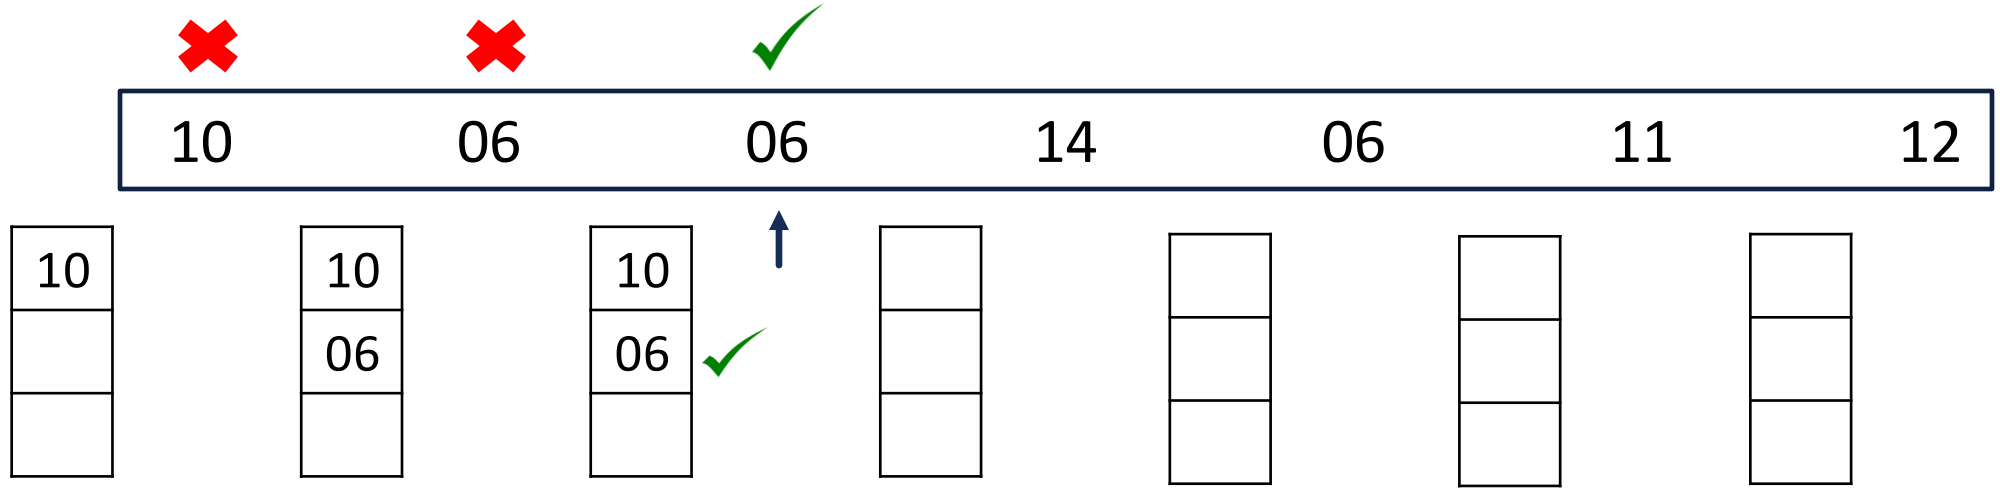

Page Fault Count: 2

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

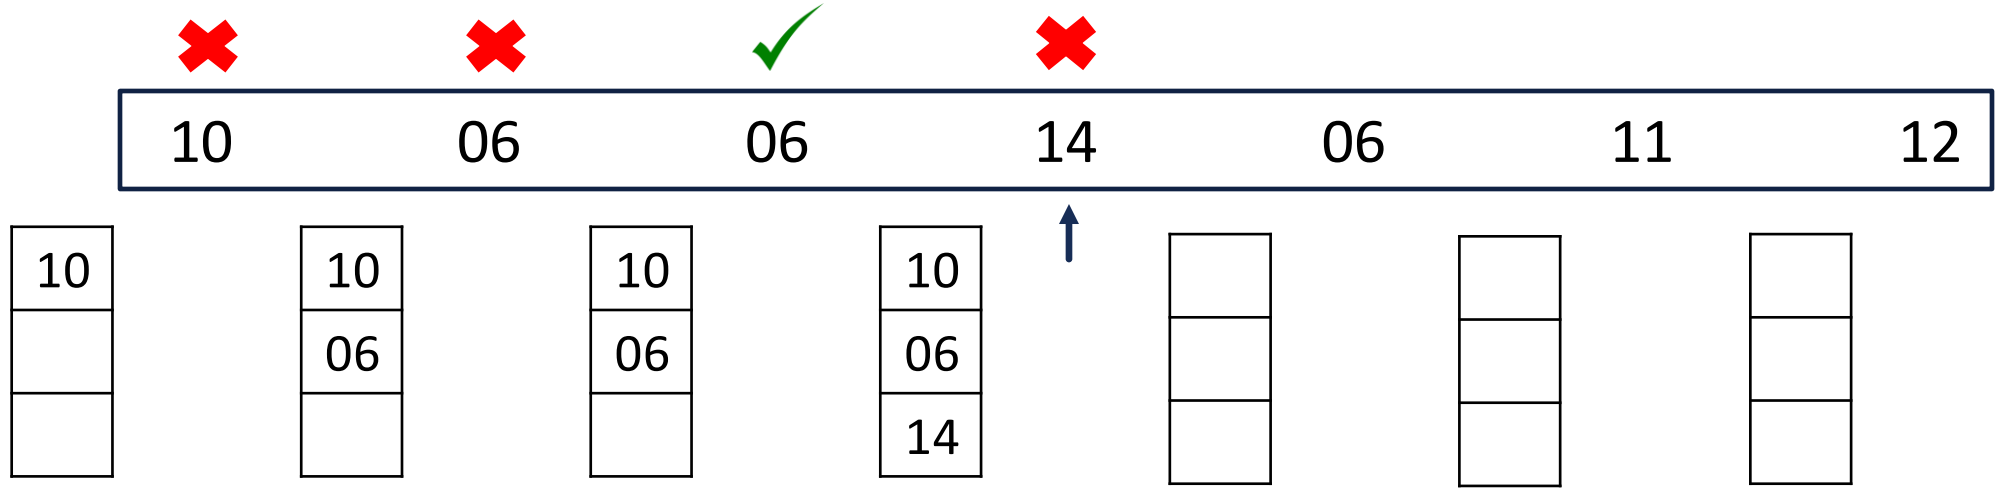

Page Fault Count: 3

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

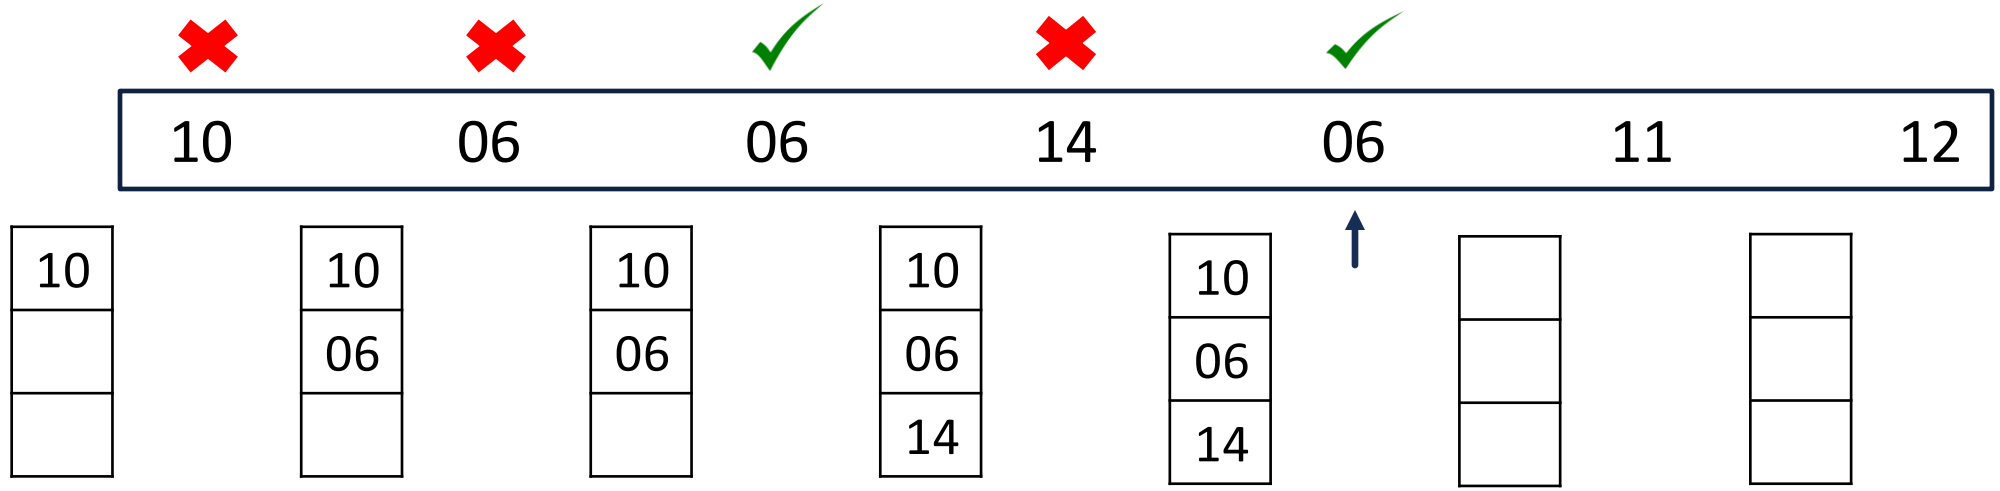

Page Fault Count: 3

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

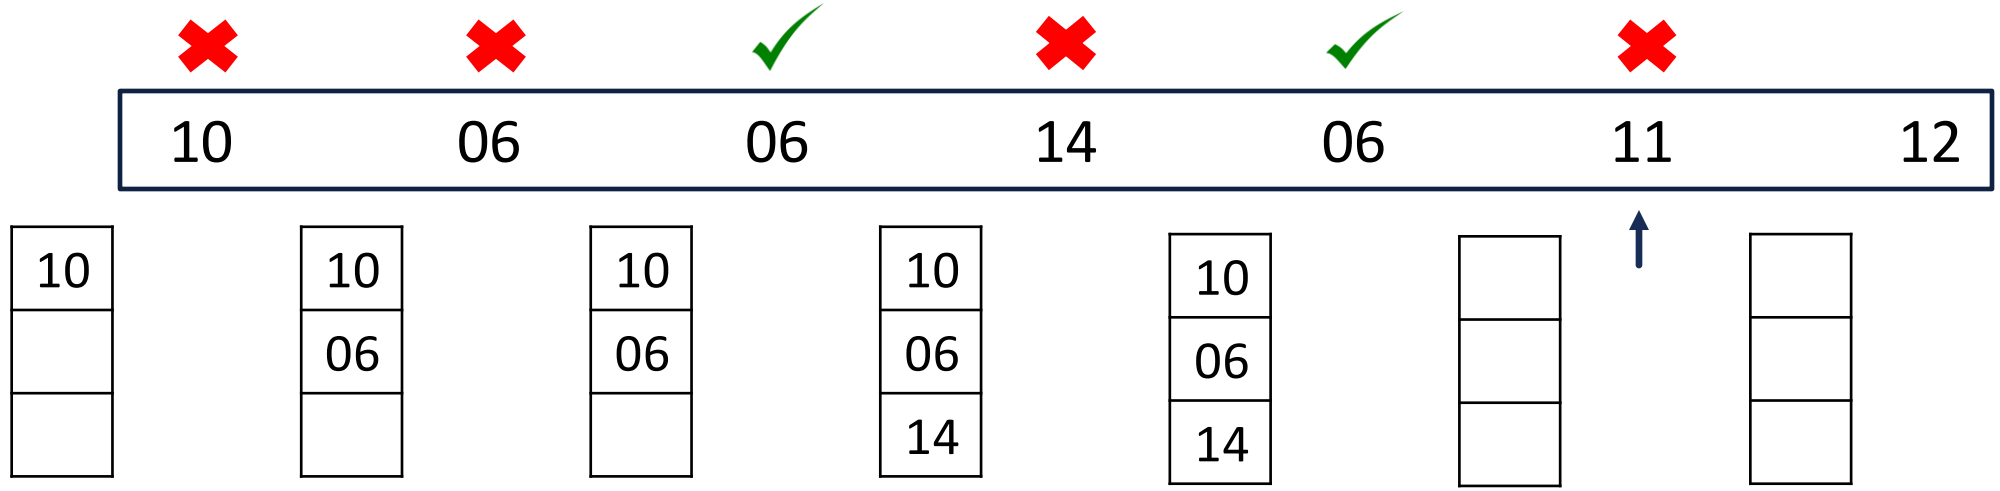

Page Fault Count: 4

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

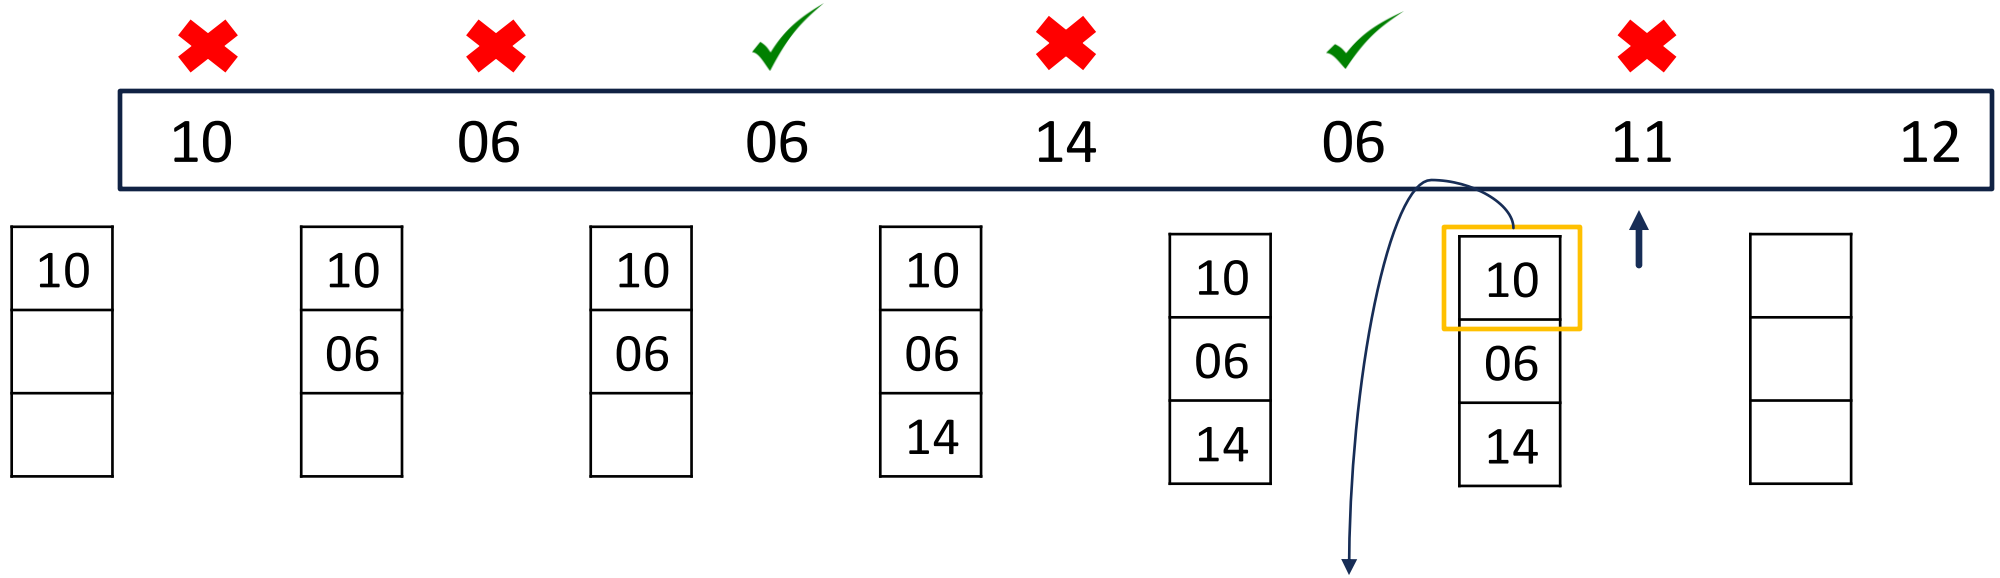

First Frame In
This is our victim.

Page Fault Count: 4

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

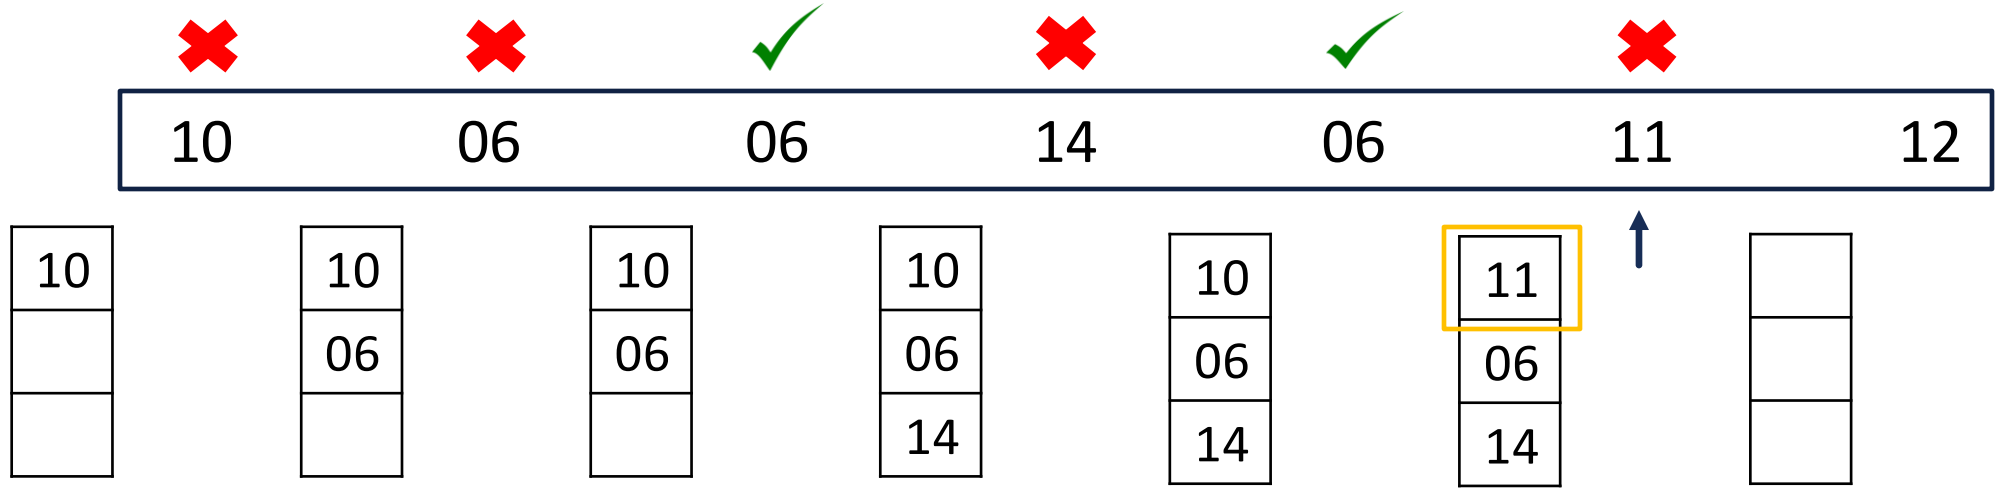

Page Fault Count: 4

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

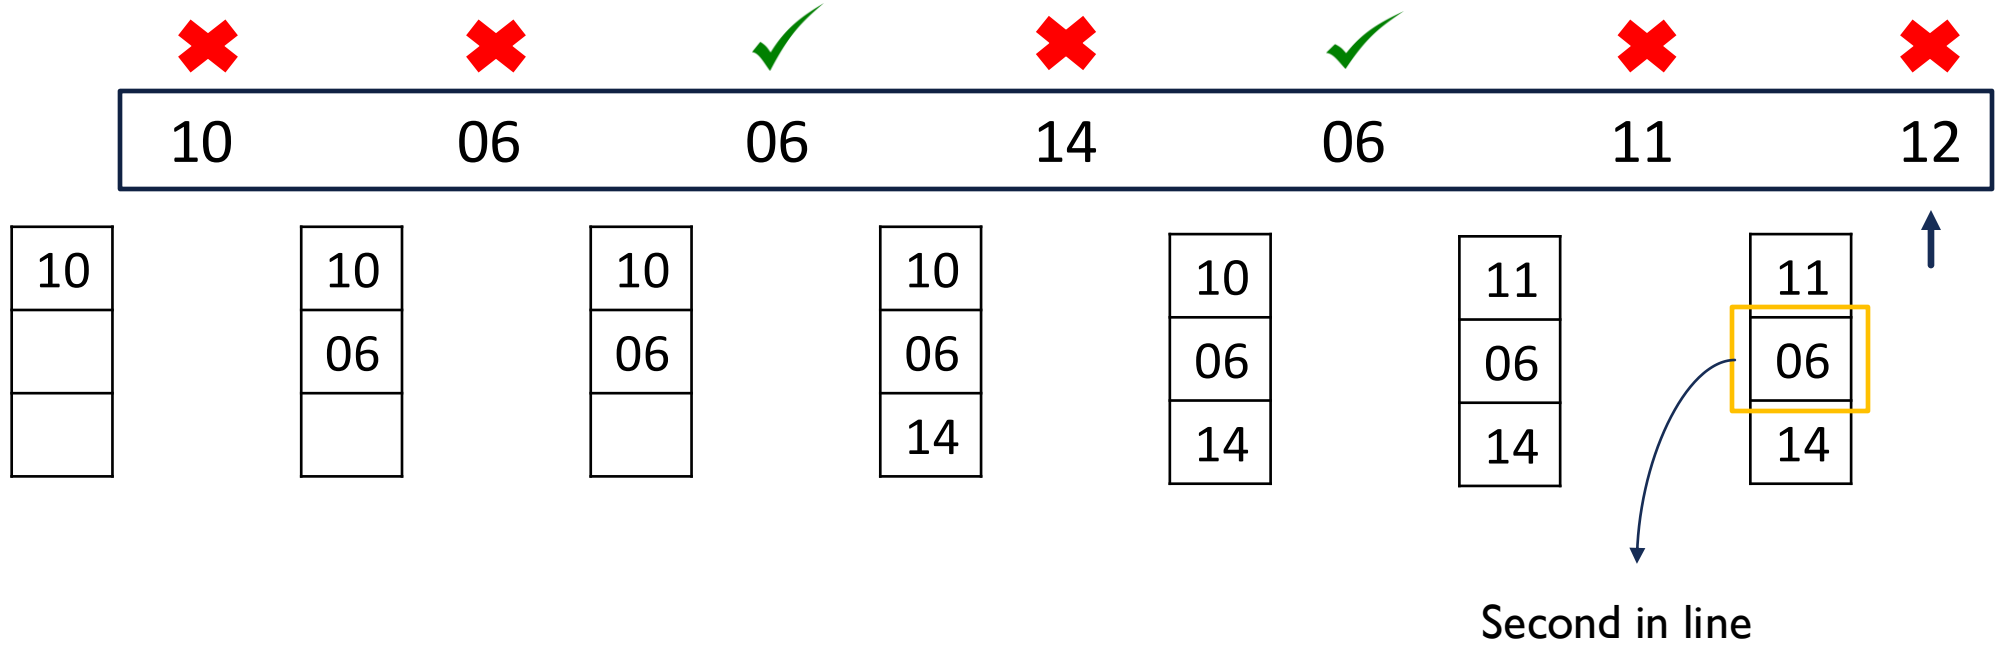

Page Fault Count: 5

FIFO EXAMPLE

Consider this sequence of pages
Assume a physical memory size of 3 frames (for the process).

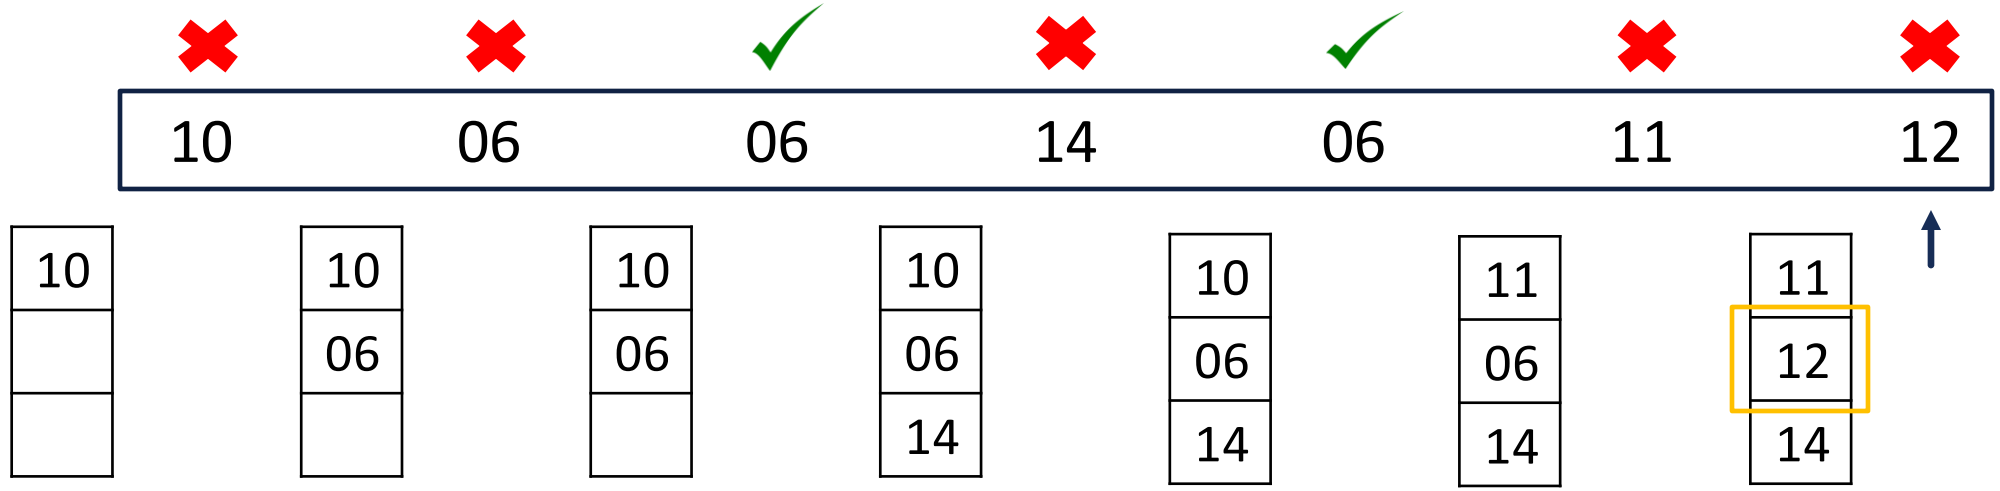

Page Fault Count: 5

WORKSHEET

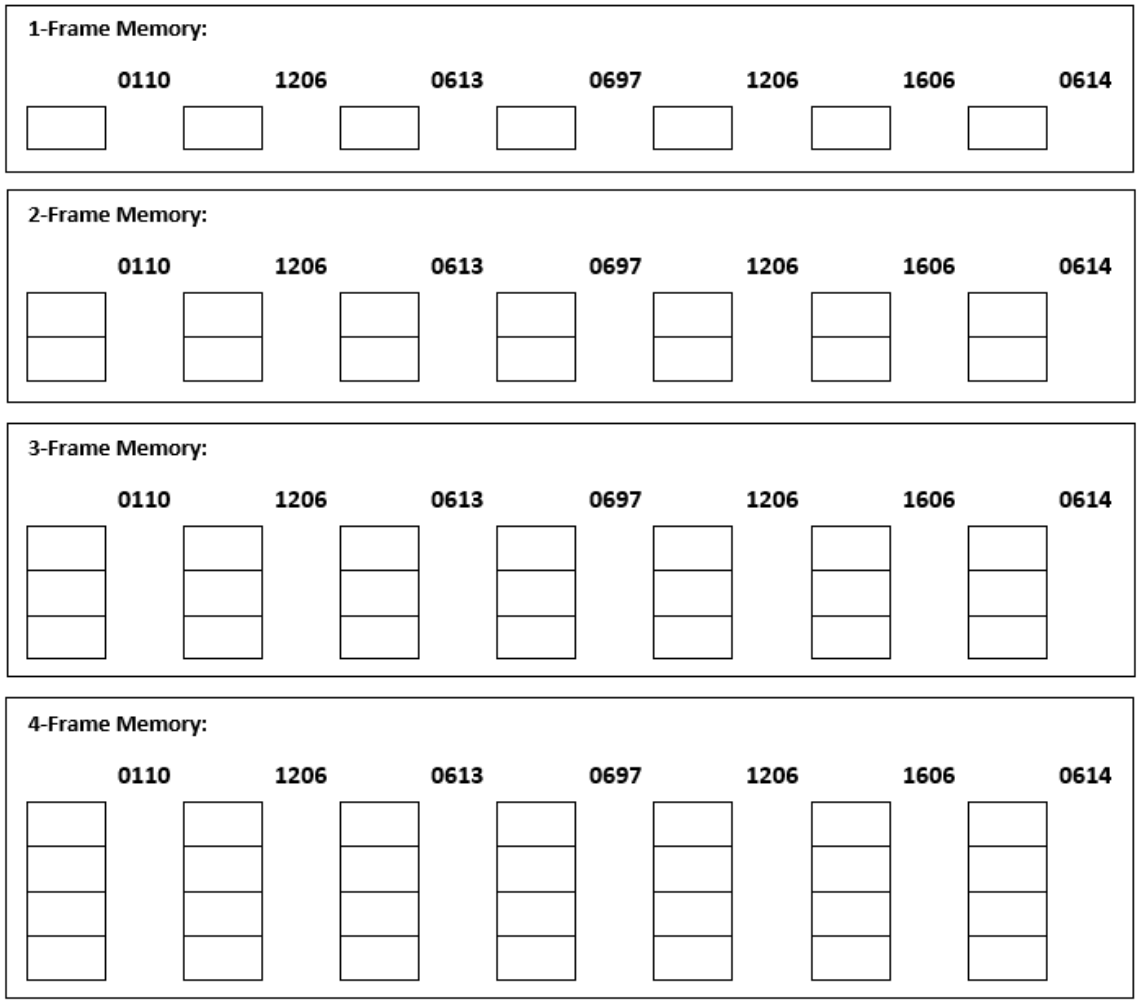

1-Frame Fault Count:	
2-Frame Fault Count:	
3-Frame Fault Count:	
4-Frame Fault Count:	

WORKSHEET SOLUTION

1-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	✖ 1206	✖ 1606	✖ 0614
01	12	06	06	12	16	06

2-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	06	06	06	06	06
	12	12	12	12	16	16

3-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	01	01	01	16	16
	12	12	12	12	12	12
		06	06	06	06	06

4-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	01	01	01	01	01
	12	12	12	12	12	12
		06	06	06	06	06
					16	16

1-Frame Fault Count:	6
2-Frame Fault Count:	4
3-Frame Fault Count:	4
4-Frame Fault Count:	4

WORKSHEET SOLUTION

- It's a page fault when any frame in memory changes content.

1-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	✖ 1206	✖ 1606	✖ 0614
01	12	06	06	12	16	06

2-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	06	06	06	06	06
	12	12	12	12	16	16

3-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	01	01	01	16	16
	12	12	12	12	12	12
		06	06	06	06	06

4-Frame Memory:

✖ 0110	✖ 1206	✖ 0613	0697	1206	✖ 1606	0614
01	01	01	01	01	01	01
	12	12	12	12	12	12
		06	06	06	06	06
					16	16

1-Frame Fault Count:	6
2-Frame Fault Count:	4
3-Frame Fault Count:	4
4-Frame Fault Count:	4

WORKSHEET SOLUTION

- It's a page fault when any frame in memory changes content.
- Page faults can't be eliminated completely even with infinite memory … every page needs to be retrieved at least once causing a single page fault for each page.

1-Frame Memory:

	0110	1206	0613	0697	1206	1606	0614
	✖	✖	✖		✖	✖	✖
	01	12	06	06	12	16	06

2-Frame Memory:

	0110	1206	0613	0697	1206	1606	0614
	✖	✖	✖			✖	
	01	01	06	06	06	06	06
		12	12	12	12	16	16

3-Frame Memory:

	0110	1206	0613	0697	1206	1606	0614
	✖	✖	✖			✖	
	01	01	01	01	01	16	16
		12	12	12	12	12	12
			06	06	06	06	06

4-Frame Memory:

	0110	1206	0613	0697	1206	1606	0614
	✖	✖	✖			✖	
	01	01	01	01	01	01	01
		12	12	12	12	12	12
			06	06	06	06	06
						16	16

1-Frame Fault Count:	6
2-Frame Fault Count:	4
3-Frame Fault Count:	4
4-Frame Fault Count:	4

INCREASING THE FRAME SIZE

Can increasing the number of memory frames result in the increase of Page Faults in a FIFO Replacement Strategy?

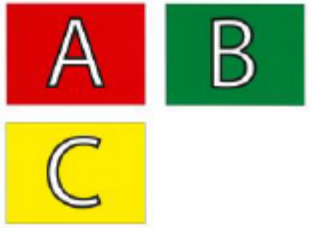

A : Yes
B : No
C : Unsure

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

1 2 3 4 1 2 5 1 2 3 4 5

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

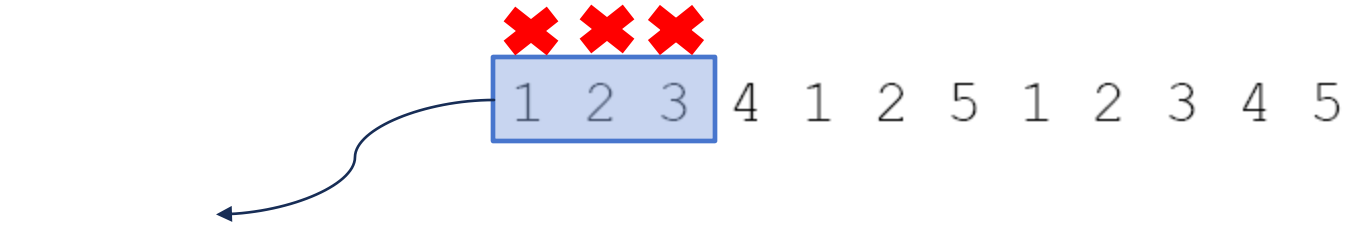

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

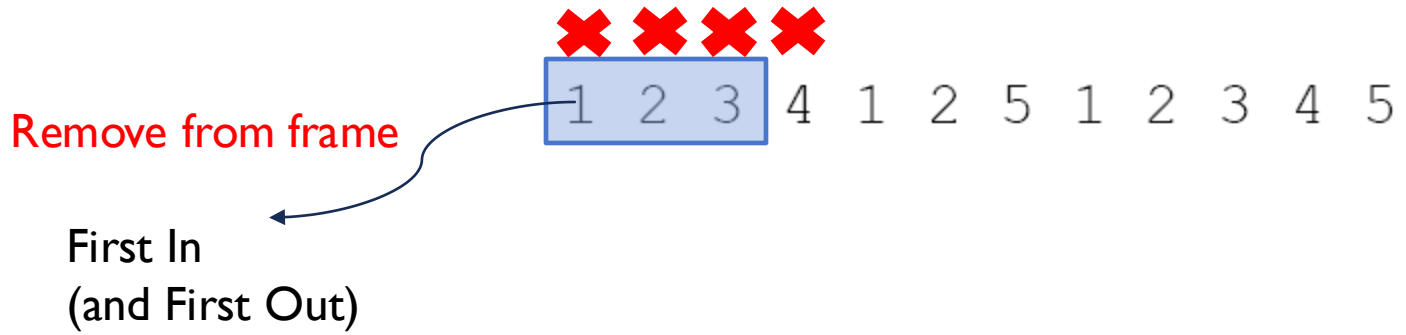

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

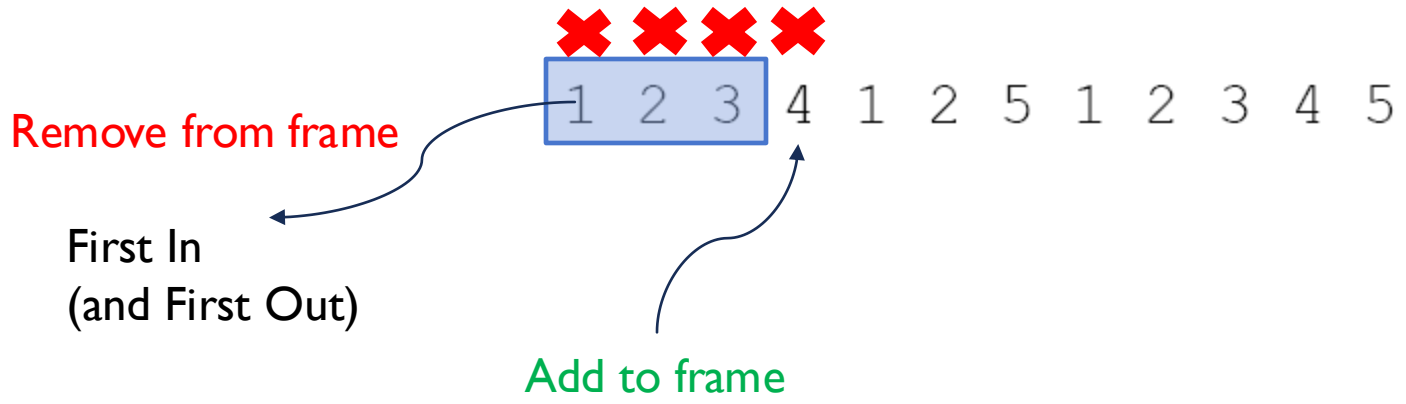

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

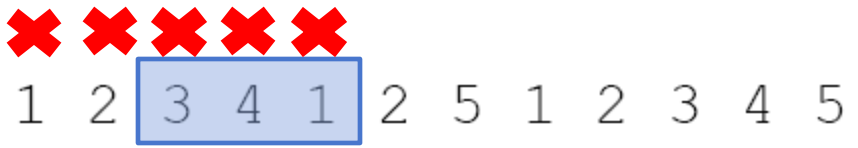

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

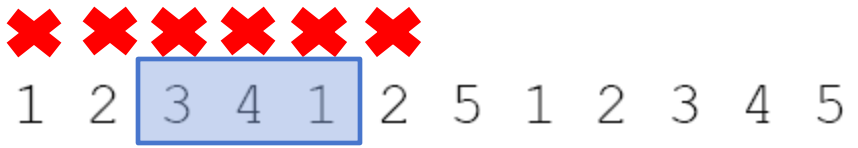

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

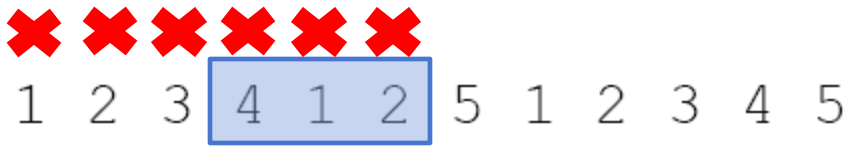

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

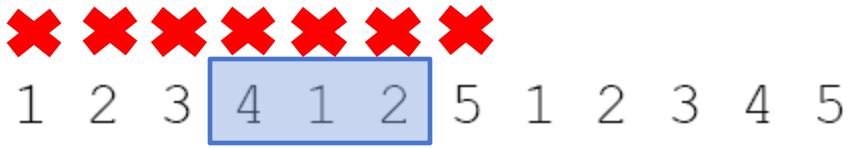

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

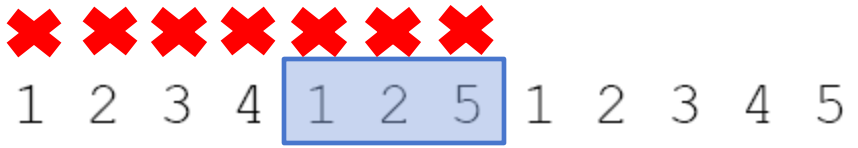

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

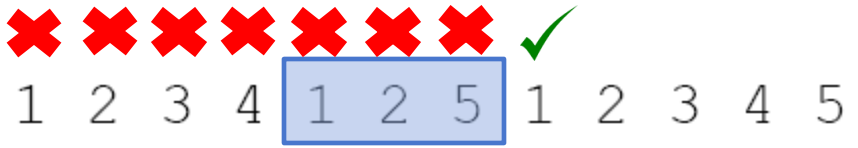

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

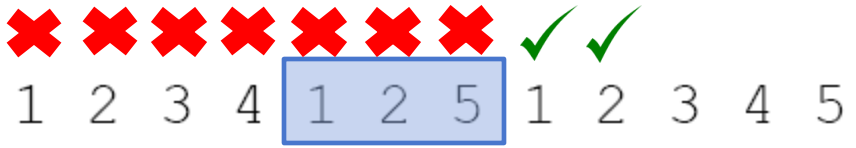

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

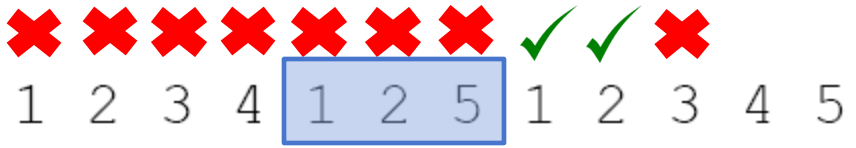

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

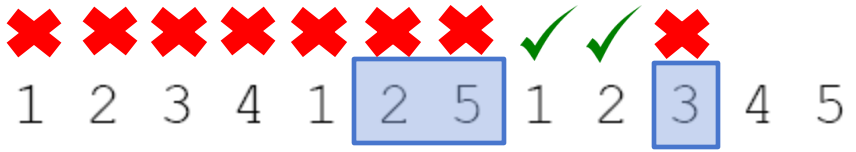

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

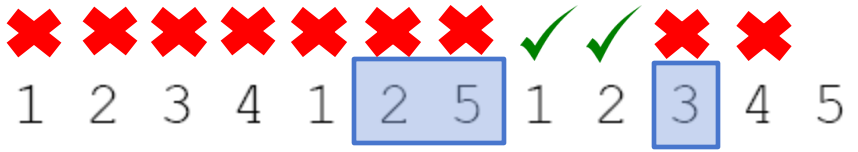

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

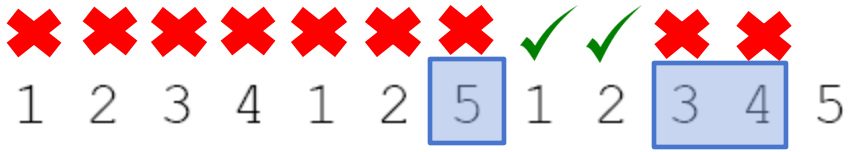

3-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

Page Fault Counter: 9

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

4-FRAME FIFO

FIFO page replacement

Assume the following are the page components of addresses (left to right, written in base 10) generated by a process

Page Fault Counter: 10

- **Increasing the frame size increased the page fault rate!**
- **Belady's Anomaly**

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?
- Well … optimally, you want to remove the page that won't be used for the longest period.

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?
- Well … optimally, you want to remove the page that won't be used for the longest period.
- So if we have P1 and P2 .. P1 will be used in 10 clock cycles and P2 in 20, we remove P2.

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?
- Well … optimally, you want to remove the page that won't be used for the longest period.
- So if we have P1 and P2 .. P1 will be used in 10 clock cycles and P2 in 20, we remove P2.
- Any particular challenge that you foresee with this algorithm?

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?
- Well … optimally, you want to remove the page that won't be used for the longest period.
- So if we have P1 and P2 .. P1 will be used in 10 clock cycles and P2 in 20, we remove P2.
- Any particular challenge that you foresee with this algorithm?
- You need to see the future to know which page will be used when.
- Why is it important then?

OPTIMAL REPLACEMENT

- What would be the optimal replacement strategy?
- Well … optimally, you want to remove the page that won't be used for the longest period.
- So if we have P1 and P2 .. P1 will be used in 10 clock cycles and P2 in 20, we remove P2.
- Any particular challenge that you foresee with this algorithm?
- You need to see the future to know which page will be used when.
- Why is it important then? It provides a benchmark/compass for new strategies.

3-FRAME OPTIMAL REPLACEMENT

Page requests : 1 2 3 4 1 2 5 1 2 3 4 5

3-FRAME OPTIMAL REPLACEMENT

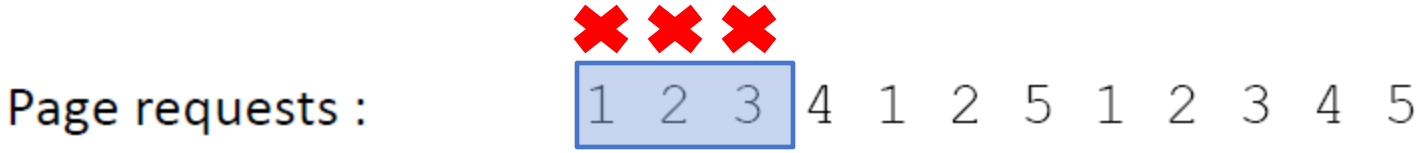

3-FRAME OPTIMAL REPLACEMENT

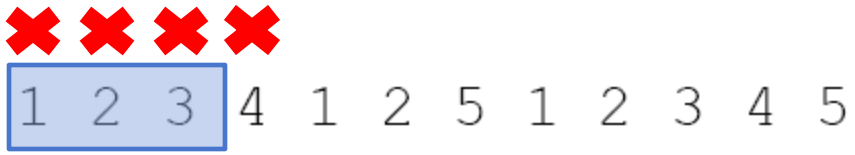

3-FRAME OPTIMAL REPLACEMENT

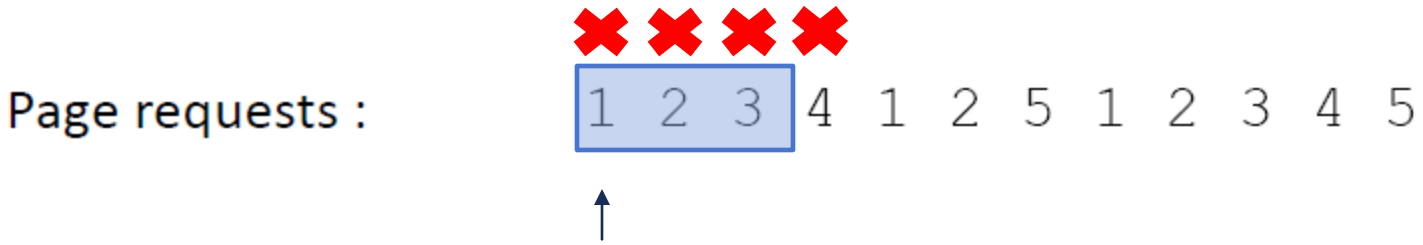

3-FRAME OPTIMAL REPLACEMENT

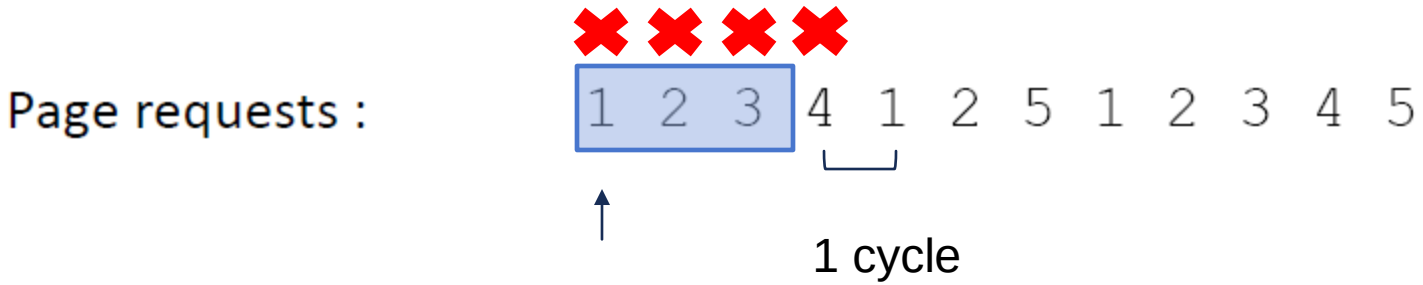

3-FRAME OPTIMAL REPLACEMENT

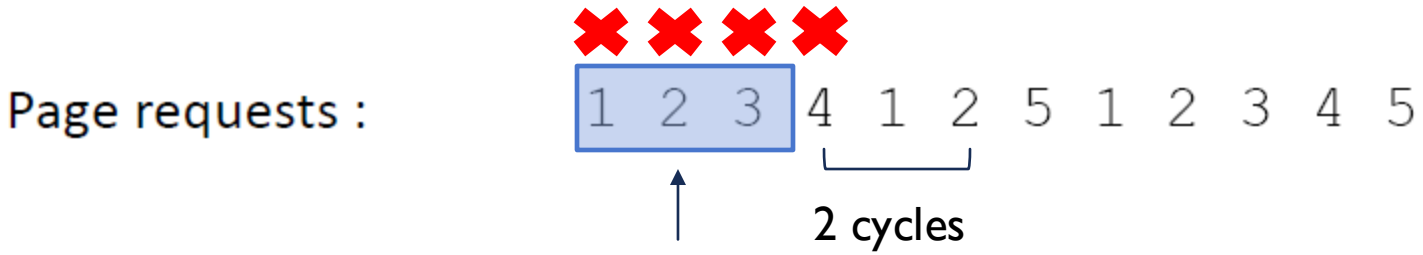

3-FRAME OPTIMAL REPLACEMENT

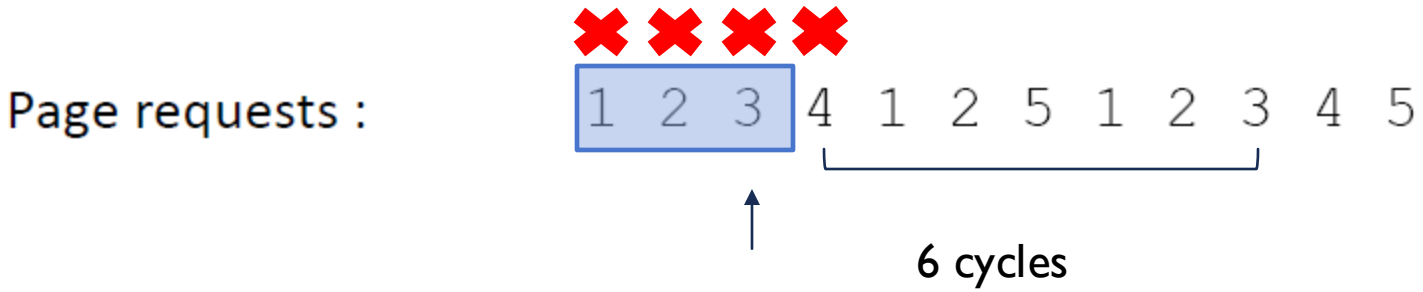

3-FRAME OPTIMAL REPLACEMENT

Page requests :

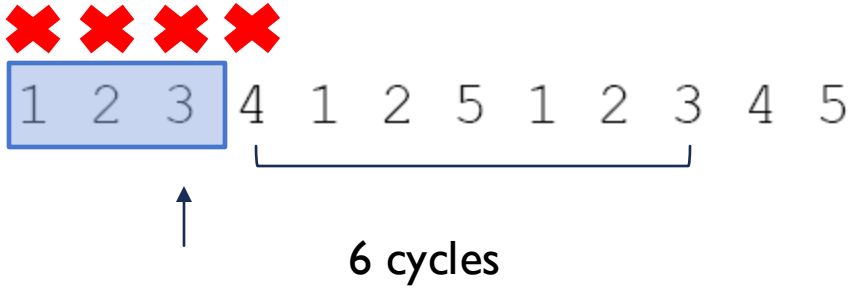

Which page should we remove from memory?

A : 1
B : 2
C : 3

3-FRAME OPTIMAL REPLACEMENT

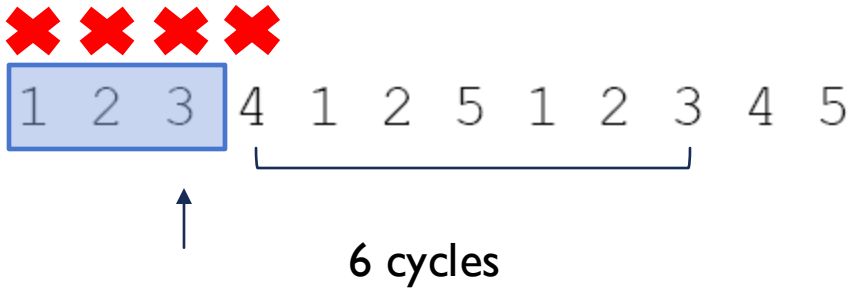

Which page should we remove from memory?

A : 1
B : 2
C : 3

Longest duration until used again.

3-FRAME OPTIMAL REPLACEMENT

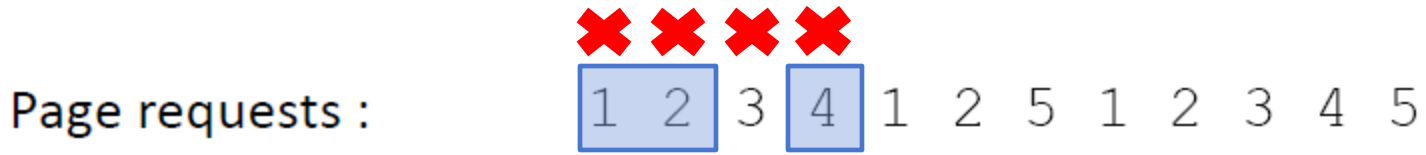

3-FRAME OPTIMAL REPLACEMENT

3-FRAME OPTIMAL REPLACEMENT

Worksheet Q1

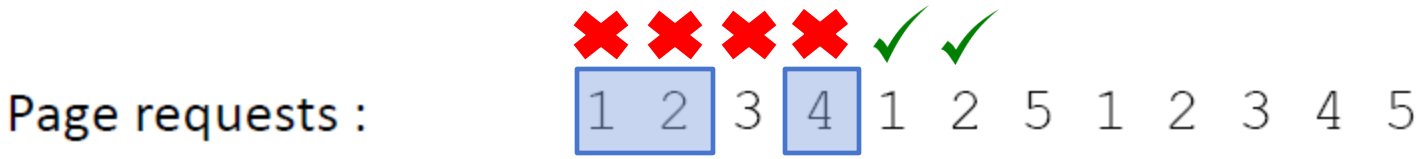

3-FRAME OPTIMAL REPLACEMENT

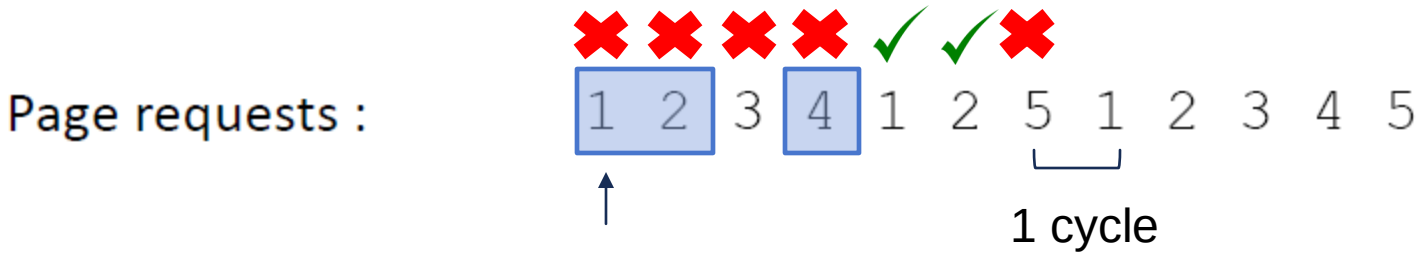

3-FRAME OPTIMAL REPLACEMENT

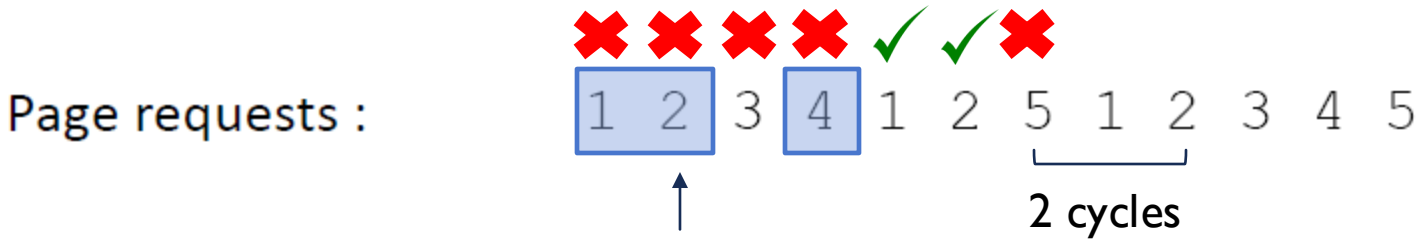

3-FRAME OPTIMAL REPLACEMENT

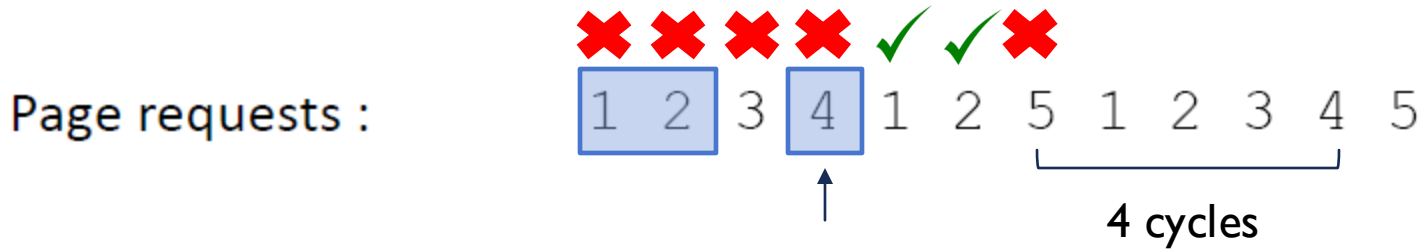

3-FRAME OPTIMAL REPLACEMENT

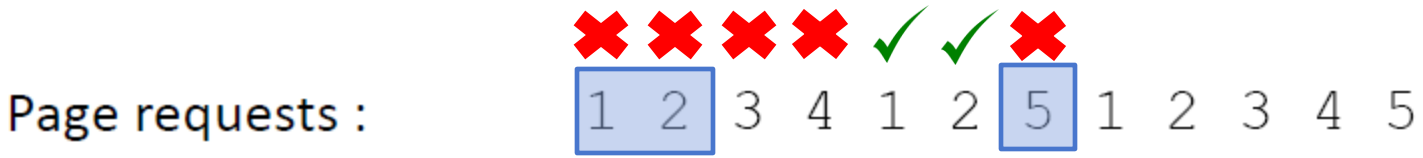

3-FRAME OPTIMAL REPLACEMENT

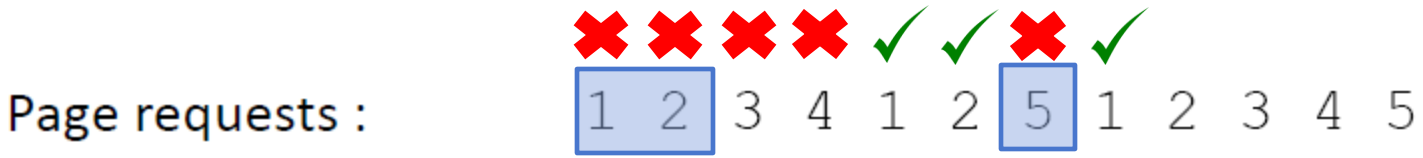

3-FRAME OPTIMAL REPLACEMENT

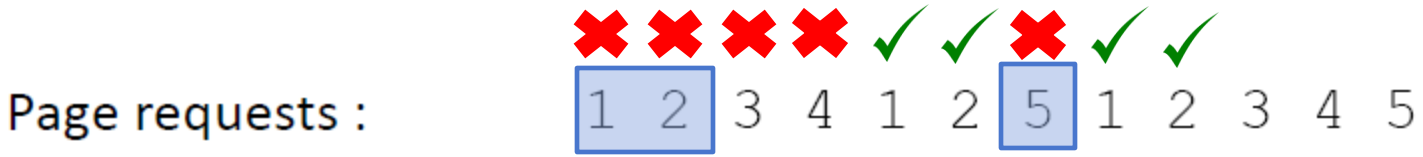

3-FRAME OPTIMAL REPLACEMENT

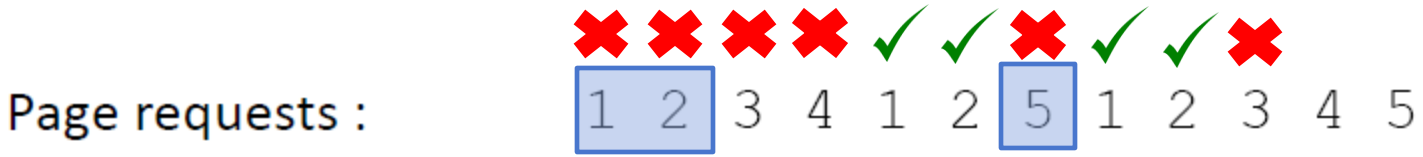

3-FRAME OPTIMAL REPLACEMENT

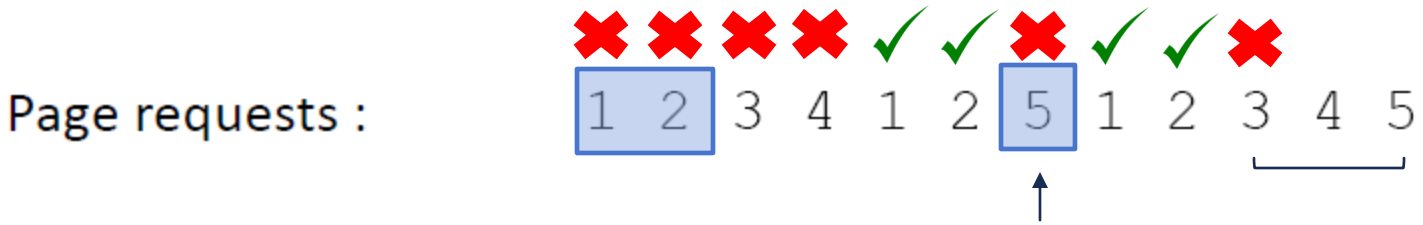

3-FRAME OPTIMAL REPLACEMENT

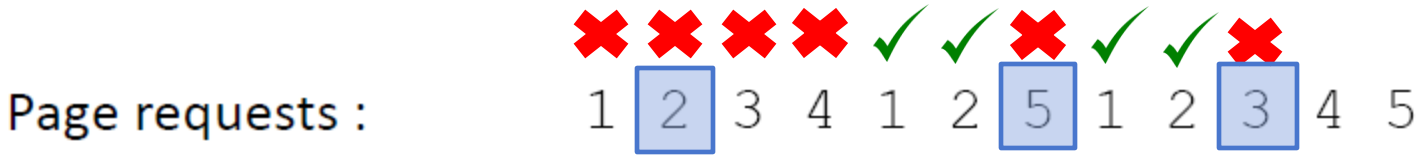

3-FRAME OPTIMAL REPLACEMENT

Page requests :

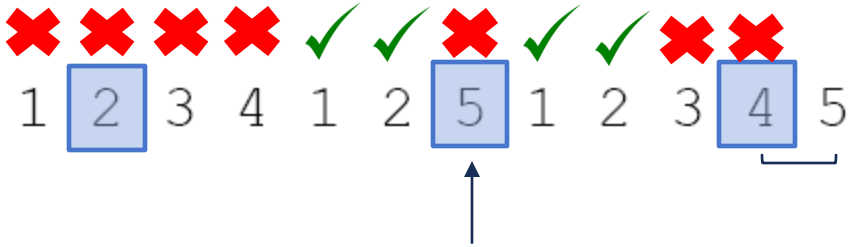

3-FRAME OPTIMAL REPLACEMENT

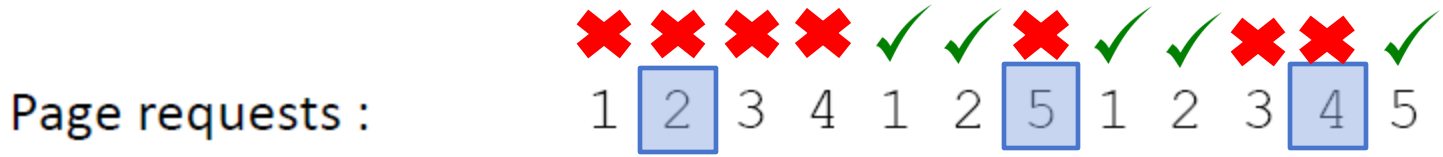

- Total Page Faults: 7
- 3-Frame FIFO: 9
- 4-Frame FIFO: 10

LEAST RECENTLY USED

- We obviously can't use Optimal Replacement …
- What's a good approximation for Optimal Replacement?

LEAST RECENTLY USED

- We obviously can't use Optimal Replacement …
- What's a good approximation for Optimal Replacement? Least Recently Used.

LEAST RECENTLY USED

- Least Recently Used attempts to estimate which of the pages in memory would be not needed for the longest period.
- The idea is that a page that we have not accessed for a long while, won't be used any time soon .. We're probably done with it for now.
- How is that different from FIFO replacement?
- FIFO looks at the time the page was ***brought into memory*** … LRU looks at the time the page was ***last accessed.***

FIFO VS LRU

Worksheet Q2

LRU

FIFO VS LRU

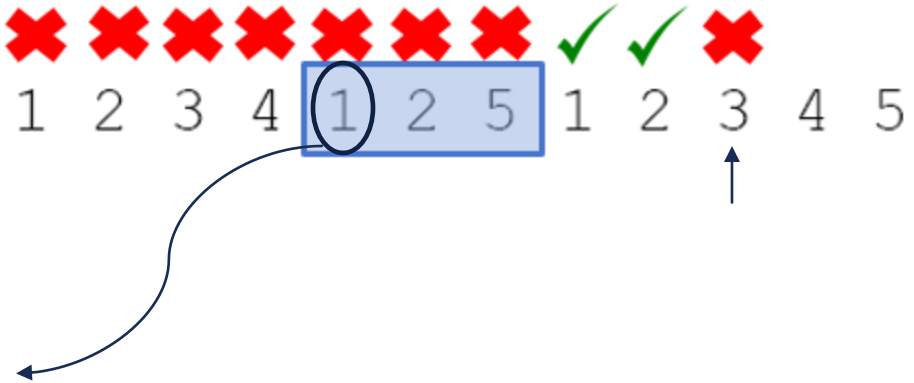

LRU: Page 1 has been recently used (2 cycles ago)

FIFO VS LRU

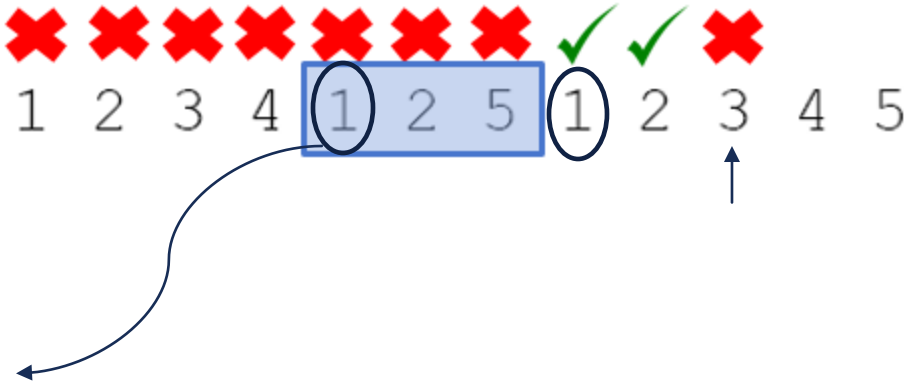

LRU: Page 1 has been recently used (2 cycles ago)

FIFO VS LRU

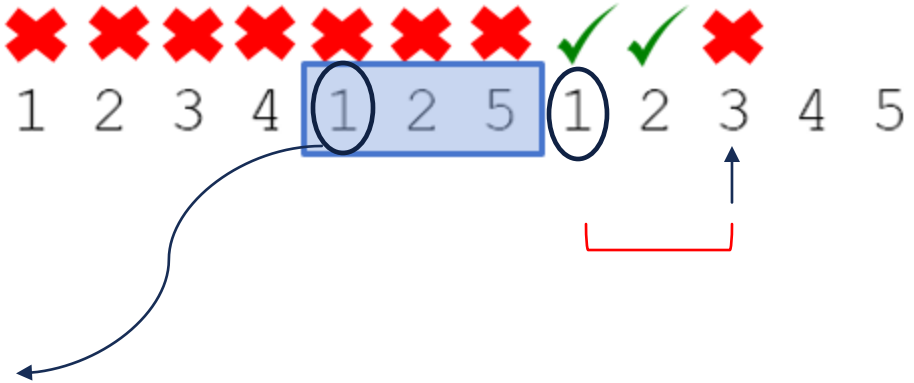

LRU: Page 1 has been recently used (2 cycles ago)

FIFO VS LRU

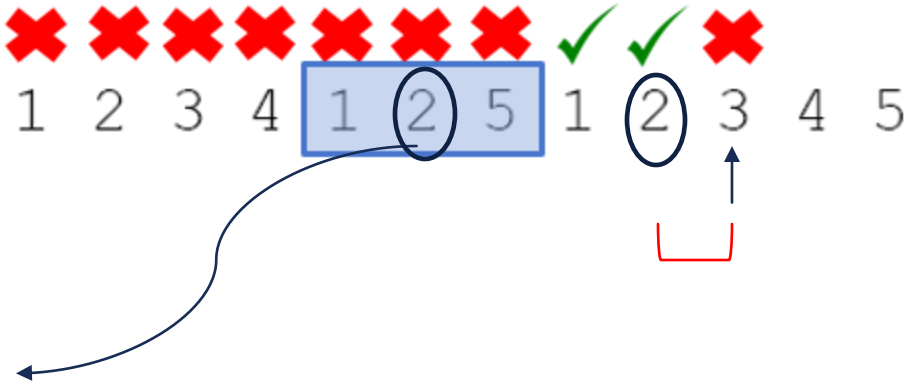

LRU: Page 2 has been recently used (1 cycle ago)

FIFO VS LRU

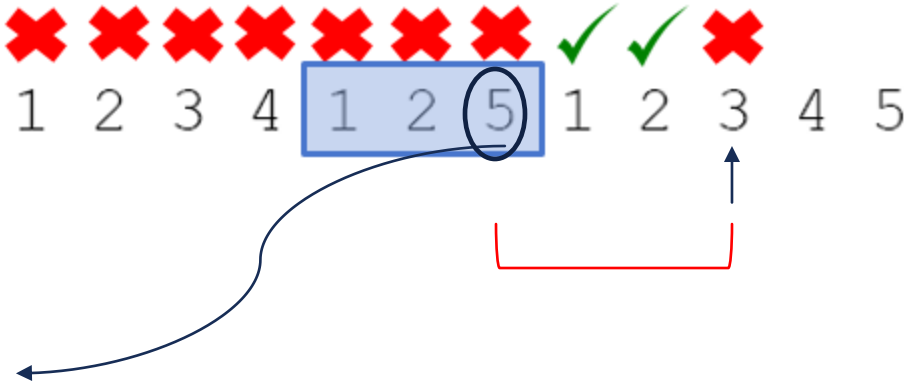

LRU: Page 5 has been less recently used (3 cycles ago)

FIFO VS LRU

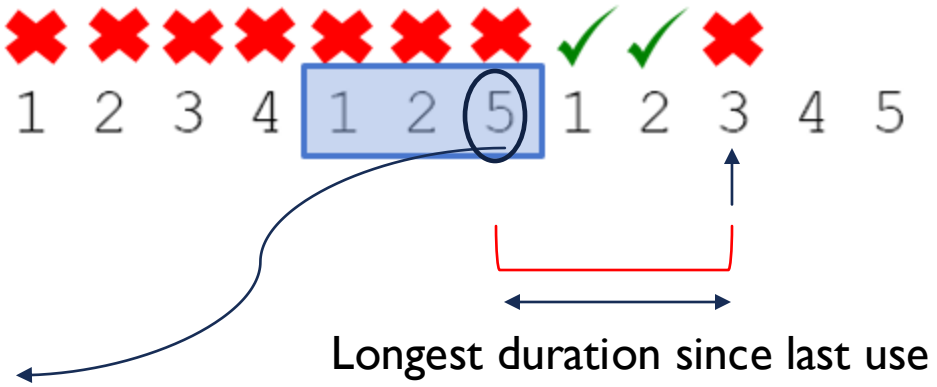

LRU: Page 5 has been less recently used (3 cycles ago)

FIFO VS LRU

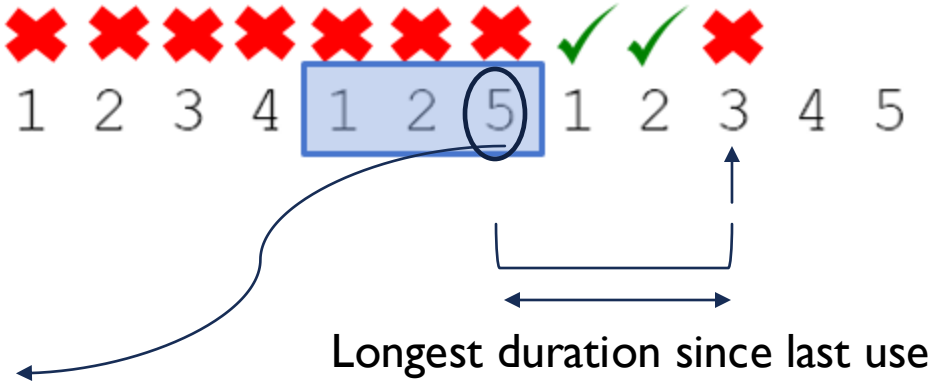

LRU: Page 5 has been less recently used (3 cycles ago)

LRU: Remove 5 (the last page inserted)
FIFO: Remove 1 (the first/oldest page inserted).

LRU DATA STRUCTURE

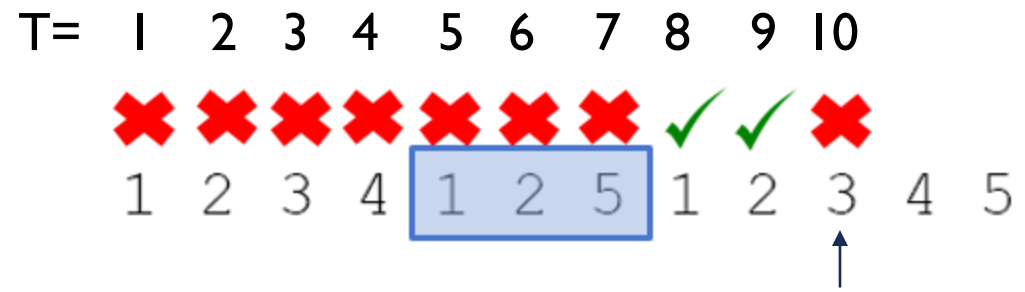

Q: How to keep track of last access/reference?

LRU DATA STRUCTURE

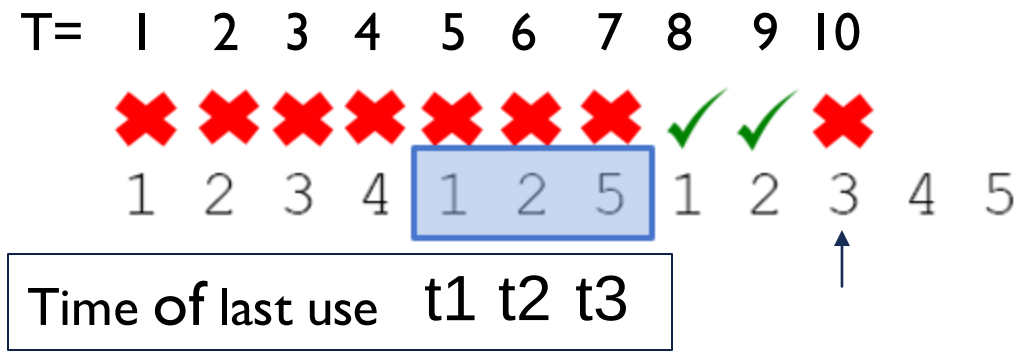

Q: How to keep track of last access/reference?

- Need a data structure that keeps track of last time use of each frame in memory.

LRU DATA STRUCTURE

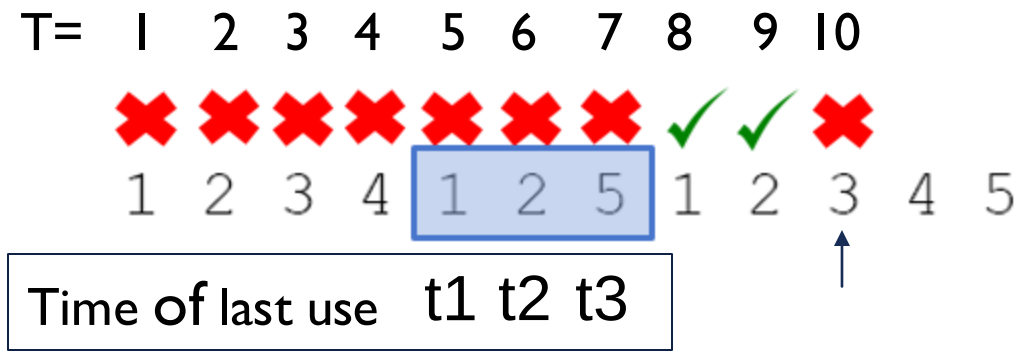

Q: How to keep track of last access/reference?

- Need a data structure that keeps track of last time use of each frame in memory.

Worksheet Q3: What would be the complications?

LRU DATA STRUCTURE

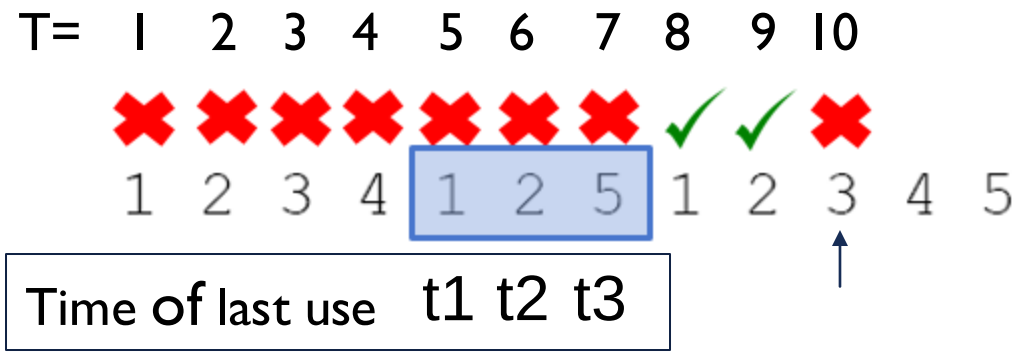

Q: How to keep track of last access/reference?

- Need a data structure that keeps track of last time use of each frame in memory.
- How many frames in memory? For an 8GB ram, 4KB page size … that's around 2 million entries.

LRU DATA STRUCTURE

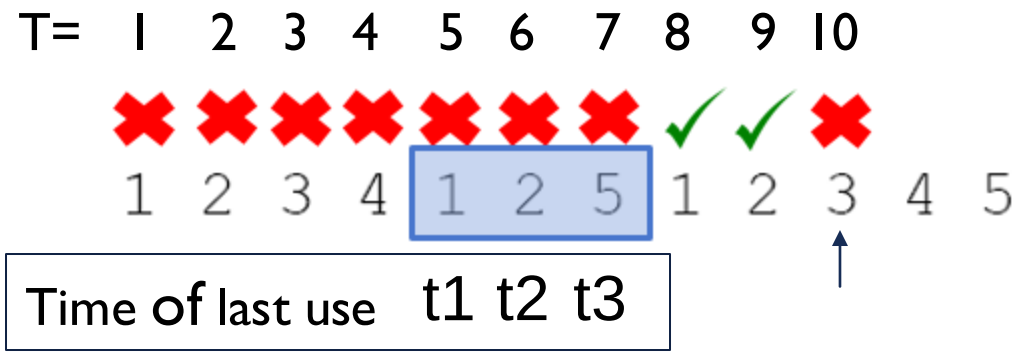

Q: How to keep track of last access/reference?

- Need a data structure that keeps track of last time use of each frame in memory.
- How many frames in memory? For an 8GB ram, 4KB page size … that's around 2 million entries.

- Impractical:
  - Too large
  - Need to sort 2 million entries every access!

LRU SUPPORT

LRU support

Common architecture support provides a single reference bit

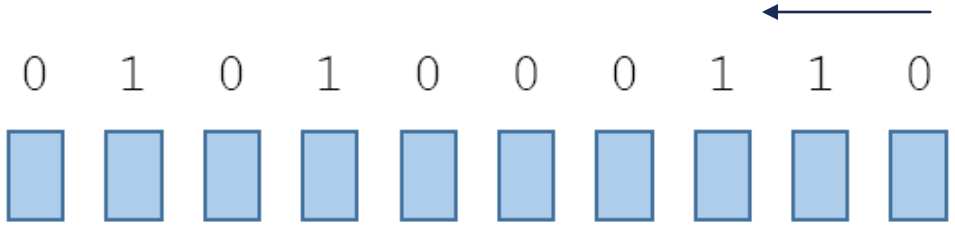

- When a page is swapped in, the bit it set to 0
- When a page is read from (referenced) or written to, the bit is set to 1

LRU SUPPORT

LRU support

- Reduces size of data structure saved.
- No need to sort, remove the first '0': less time traversing.

Common architecture support provides a single reference bit

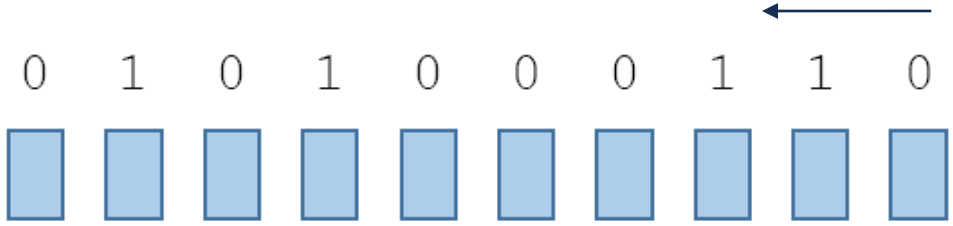

- When a page is swapped in, the bit it set to 0
- When a page is read from (referenced) or written to, the bit is set to 1

LRU SUPPORT

LRU support

Common architecture support provides a single reference bit

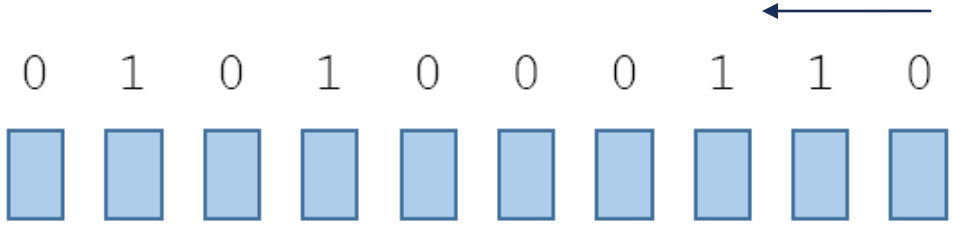

- Reduces size of data structure saved.
- No need to sort, remove the first '0': less time traversing.

- When a page is swapped in, the bit it set to 0
- When a page is read from (referenced) or written to, the bit is set to 1

- Can only track whether a page has been read or not …
- Not very useful

RECENT HISTORY

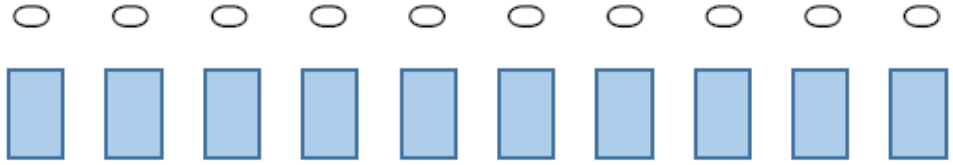

- Improvement: Last n-cycles history.

RECENT HISTORY

Periodically reset
all reference bits

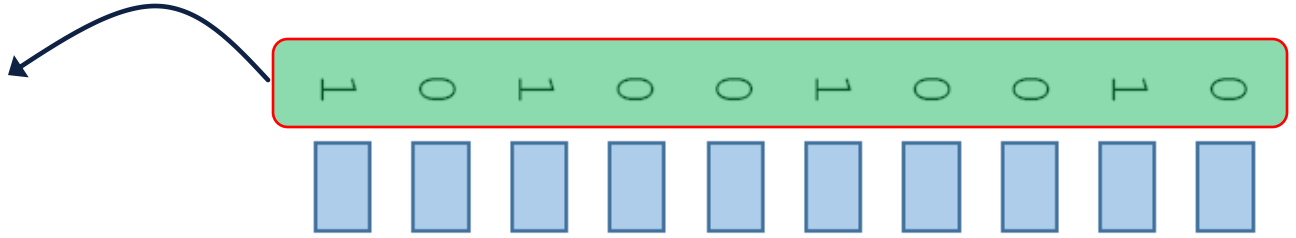

- Improvement: Last n-cycles history.

RECENT HISTORY

Periodically reset
all reference bits

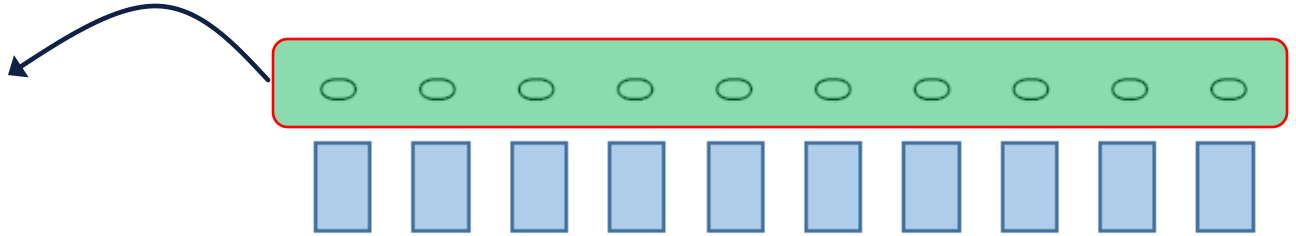

- Now the reference bit stores whether a frame have been referenced since last period started … rather than since it was brought in.

TRACKING RECENT HISTORY

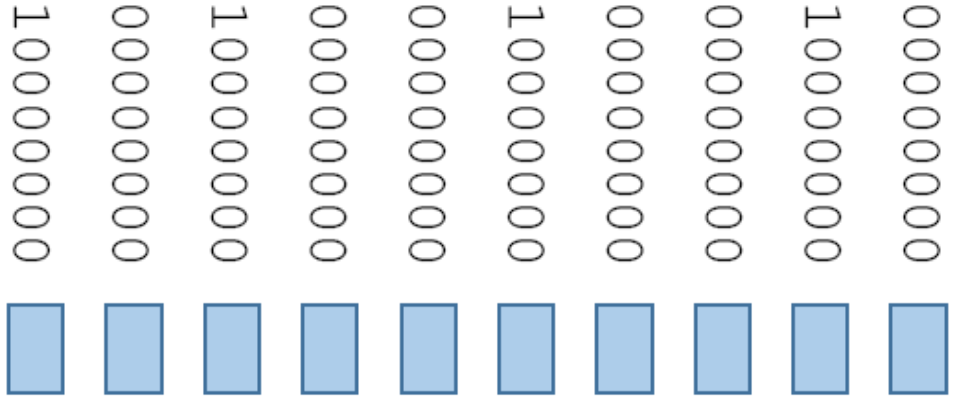

- Tracks a longer history
- Still smaller than data structure.
- Remove first page that has all ‘zeros’ in its history: no need to sort.

TRACKING RECENT HISTORY

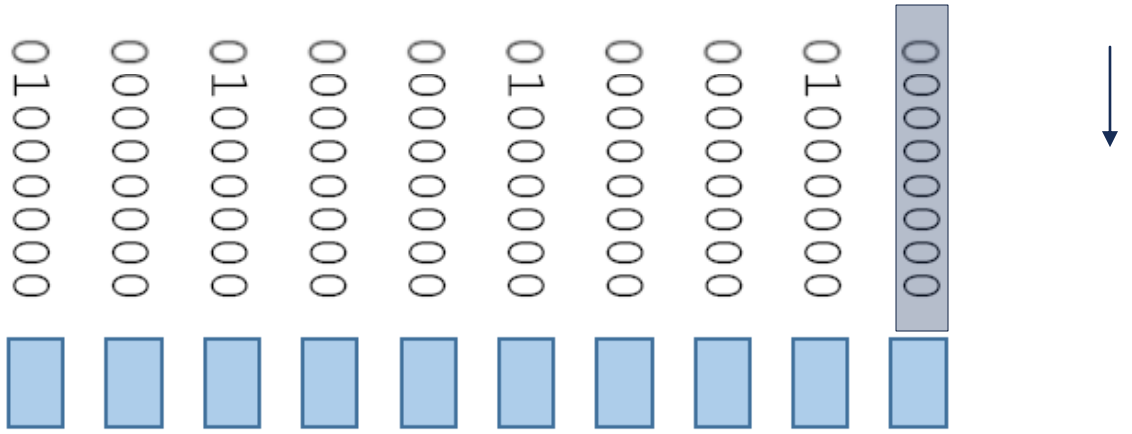

- Tracks a longer history
- Still smaller than data structure.
- Remove first page that has all ‘zeros’ in its history:
  - no need to sort.
  - No need to go through all memory

DIRTY BIT FOR PAGE REPLACEMENT

Q: Can we use dirty bit in page replacement algorithm?

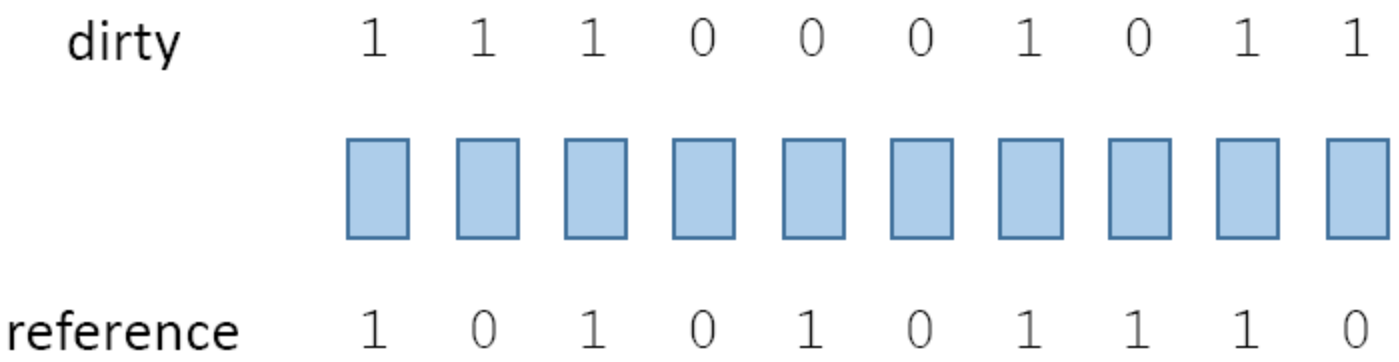

DIRTY BIT FOR PAGE REPLACEMENT

Q: Can we use dirty bit in page replacement algorithm?

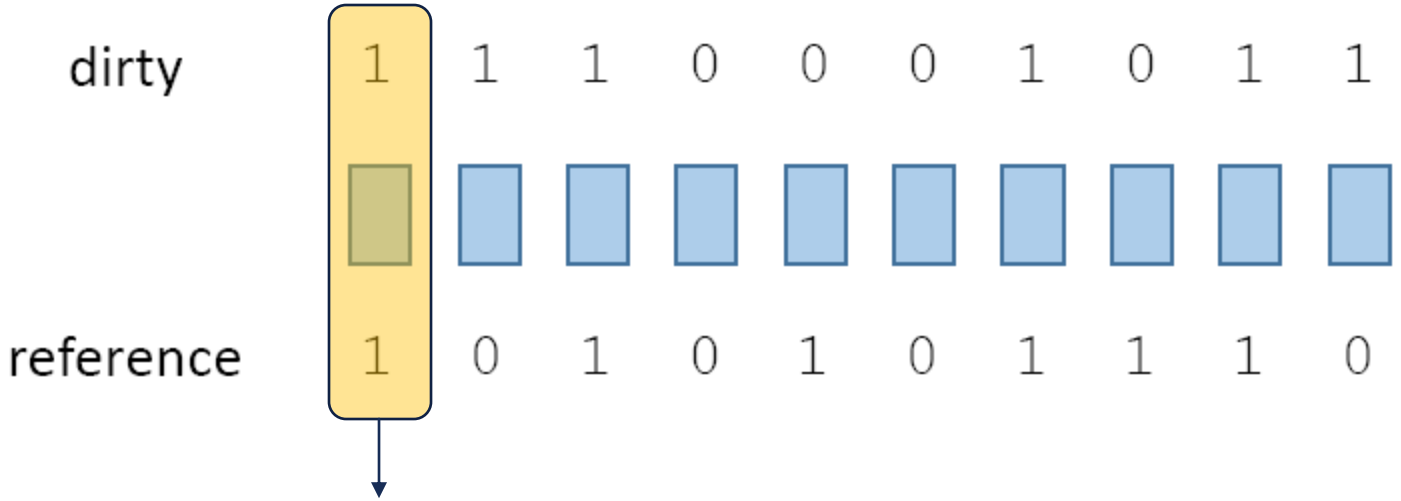

DIRTY BIT FOR PAGE REPLACEMENT

Q: Can we use dirty bit in page replacement algorithm?

Not Modified: No need to write back!

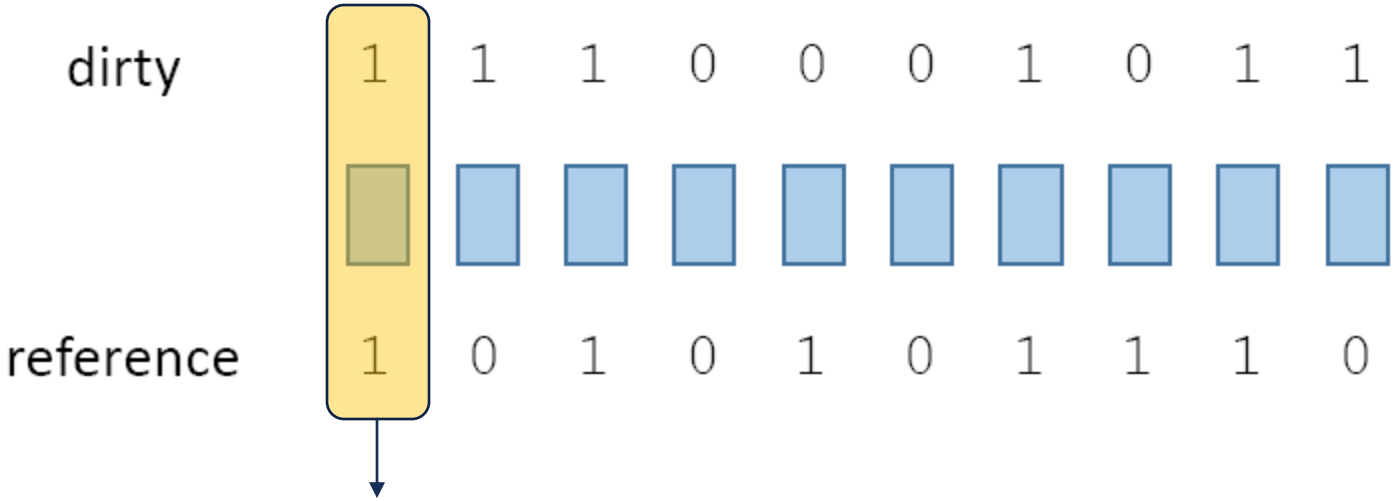

DIRTY BIT FOR PAGE REPLACEMENT

Q: Can we use dirty bit in page replacement algorithm?

Not Modified: No need to write back!

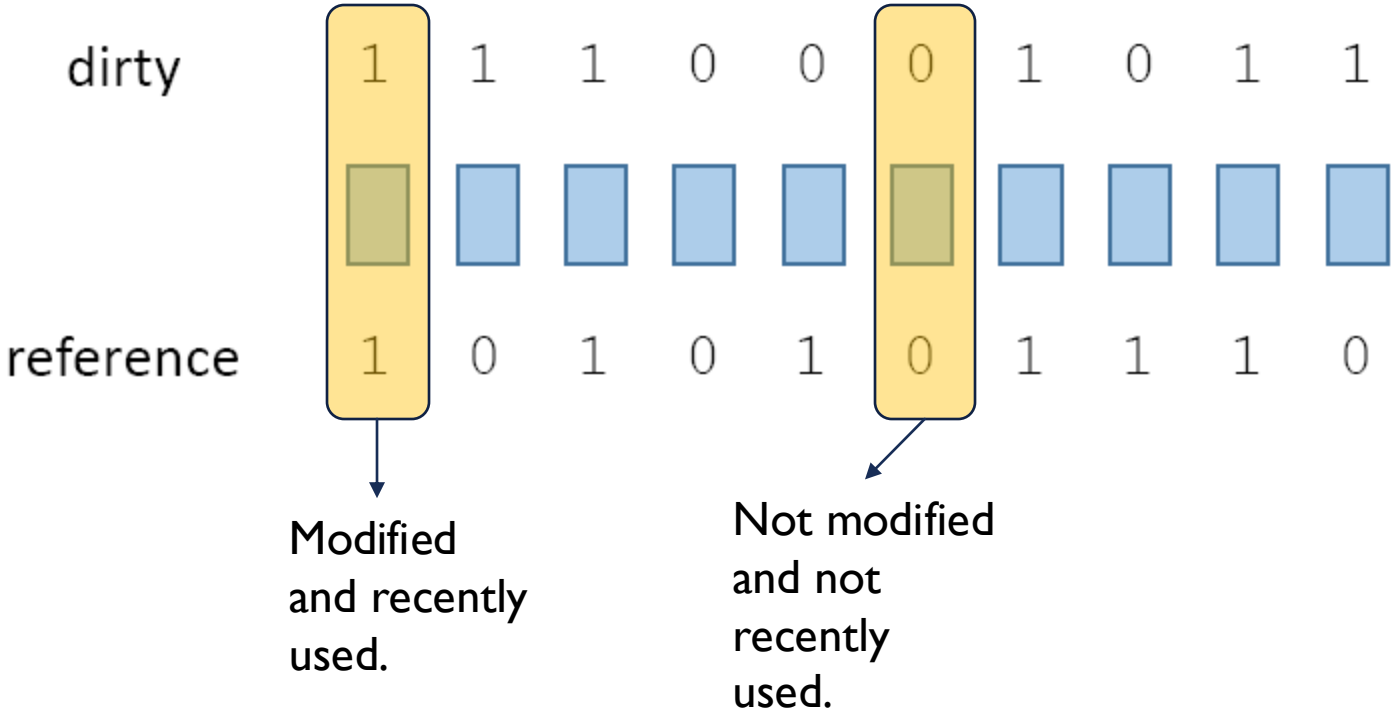

DIRTY BIT FOR PAGE REPLACEMENT

Q: Can we use dirty bit in page replacement algorithm?

Not Modified: No need to write back!

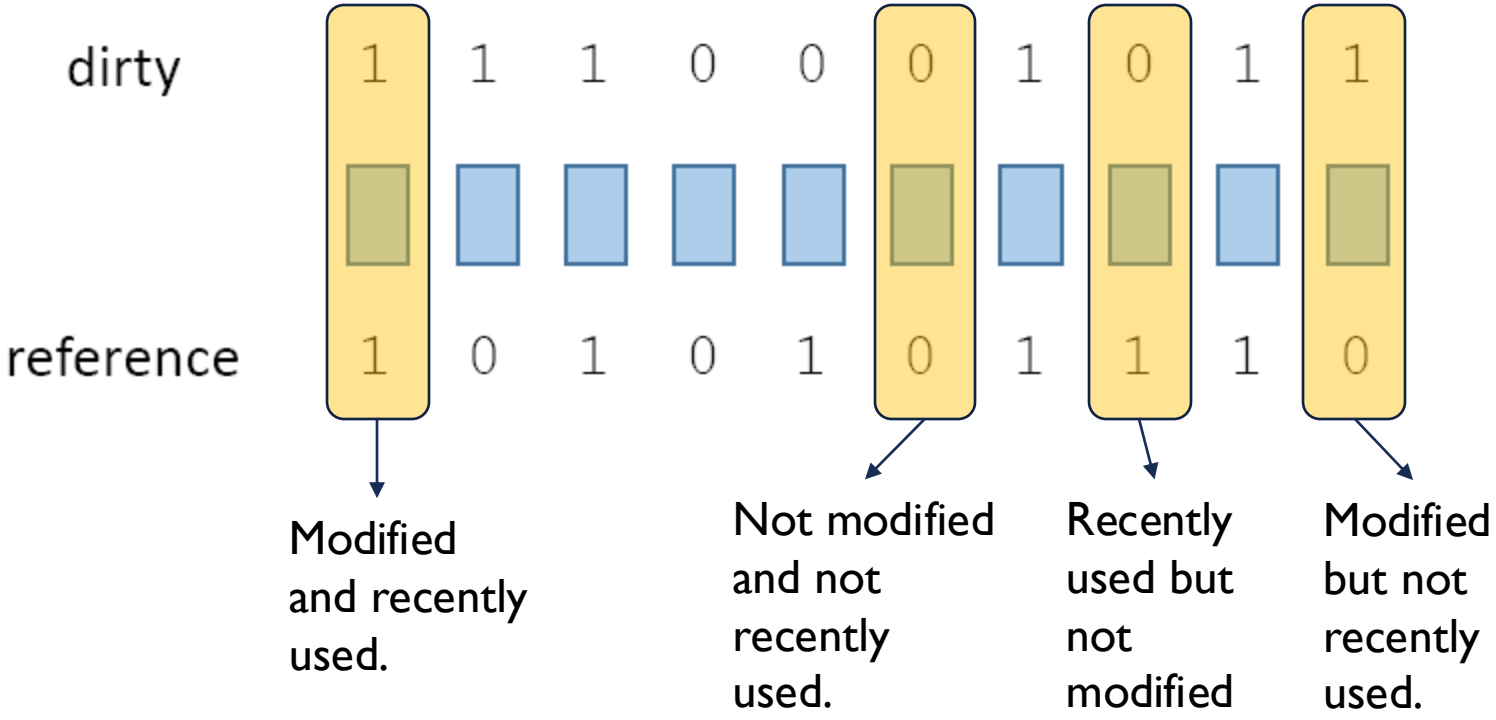

DIRTY BIT FOR PAGE REPLACEMENT

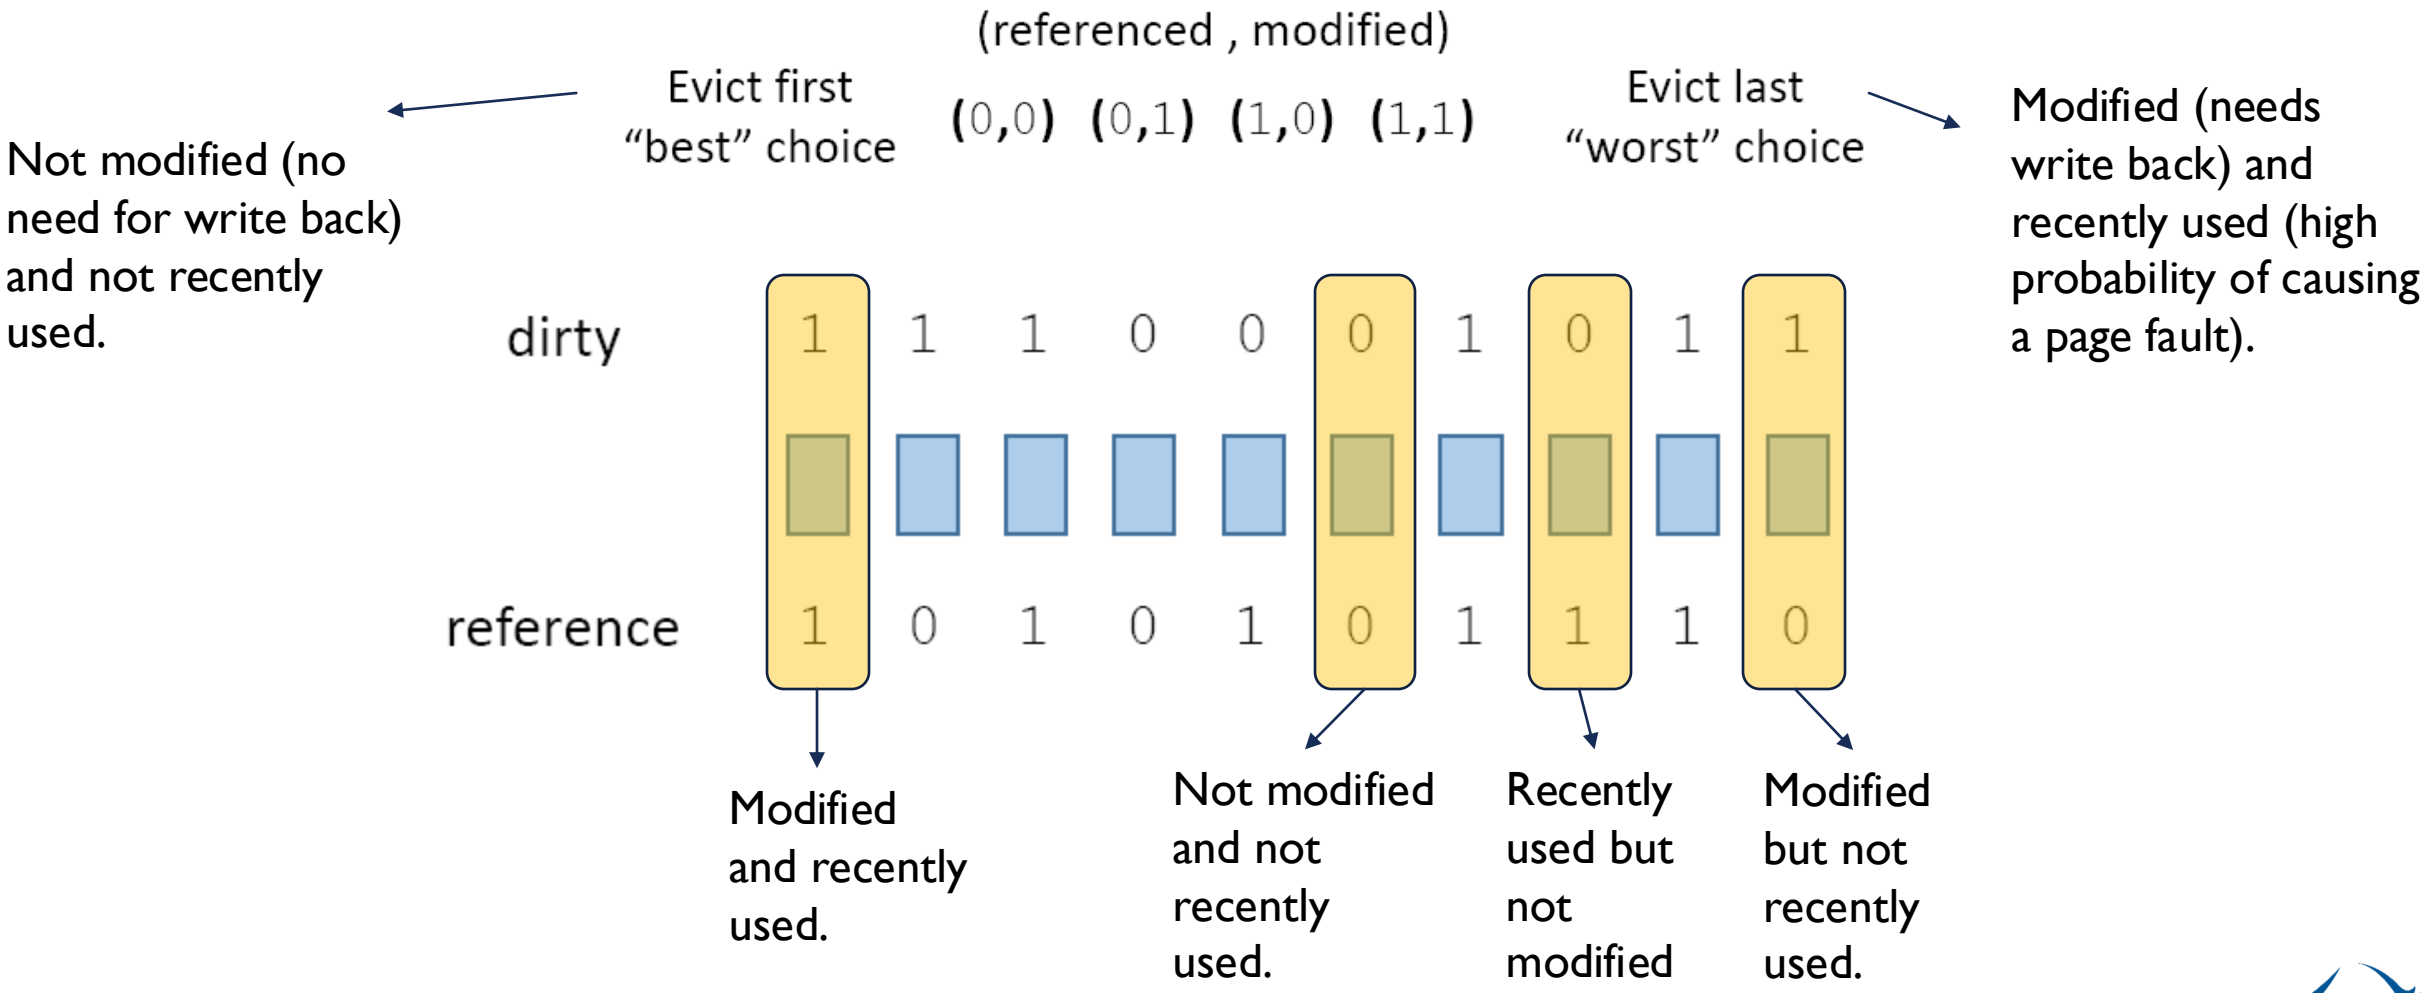

DIRTY BIT FOR PAGE REPLACEMENT

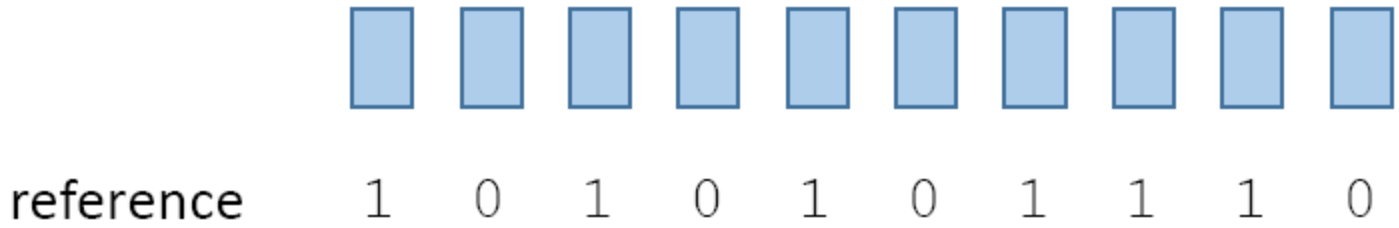

- Dirty bits indicate that a page needs to be written back.
- This slows down page replacement …
- Look at both, reference bit (history) and dirty bit.
- Pick a page that doesn't need replacement!

DIRTY BIT

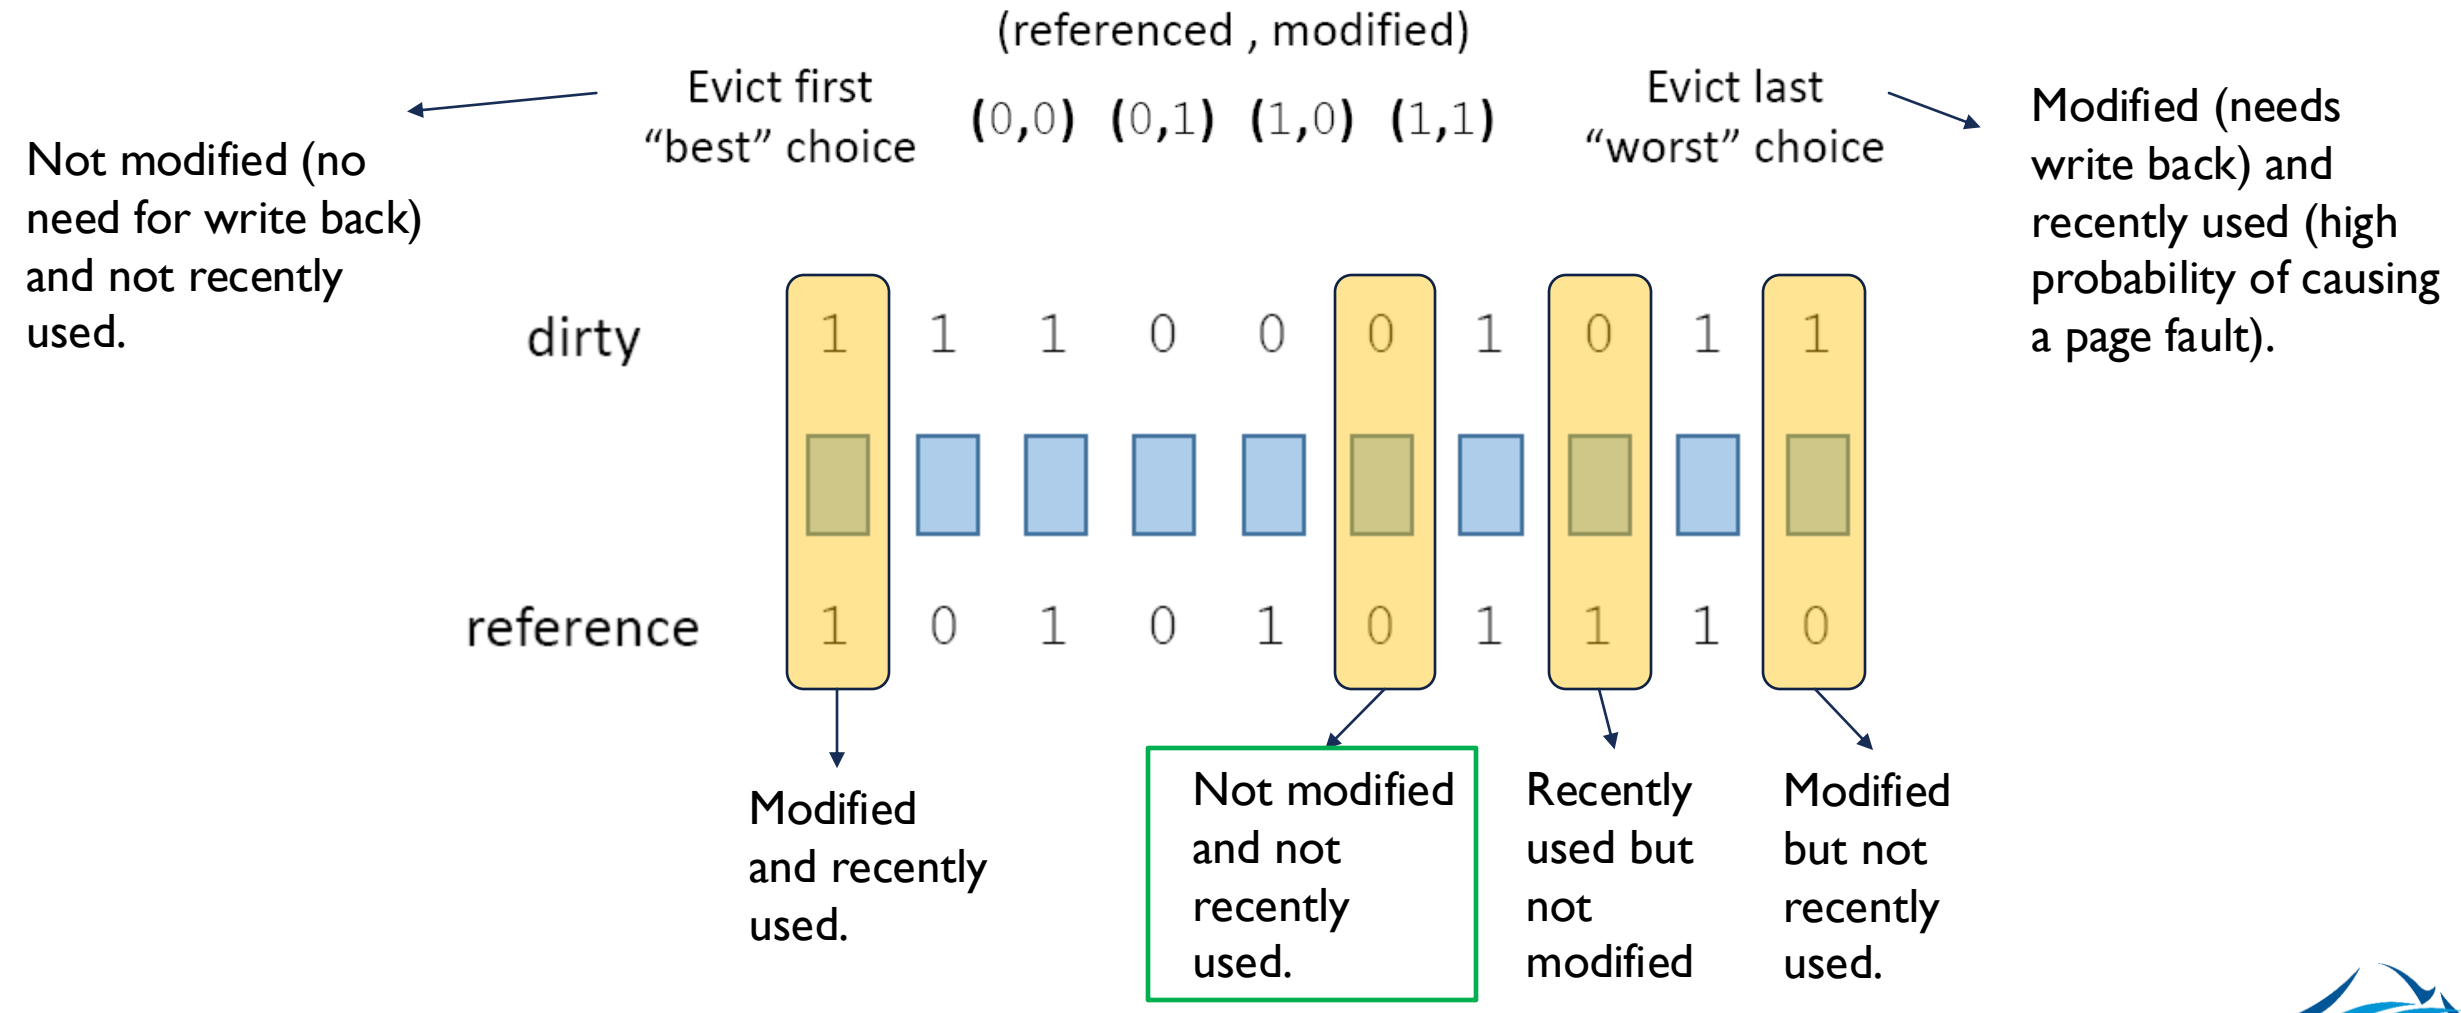

SECOND CHANCE

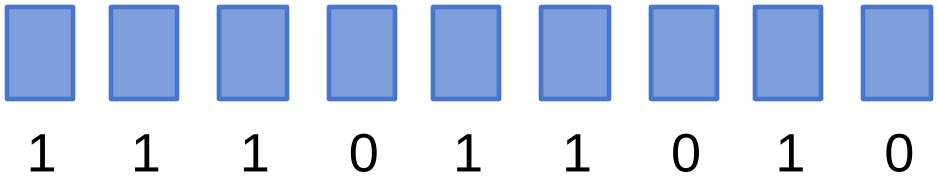

SECOND CHANCE

Similar to periodic reset …
with a simpler implementation.

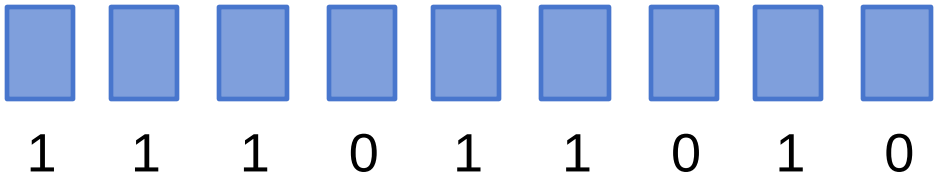

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

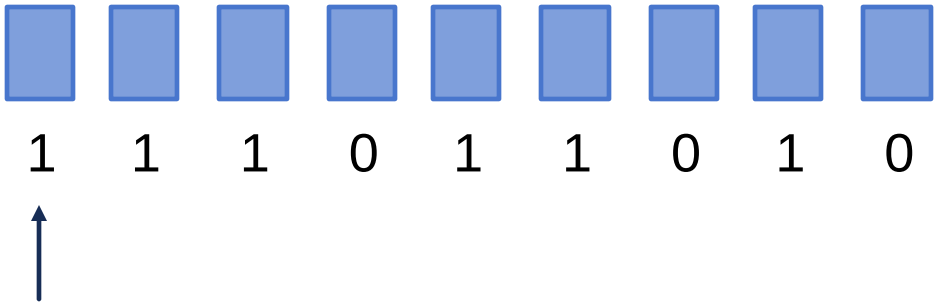

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

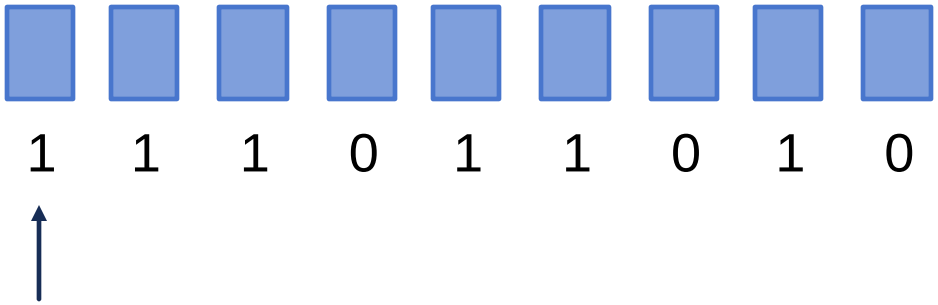

Page fault just occurred, need to replace frame … walking through memory

SECOND CHANCE

As OS walks through memory, it resets the ‘1’s to ‘0’s.

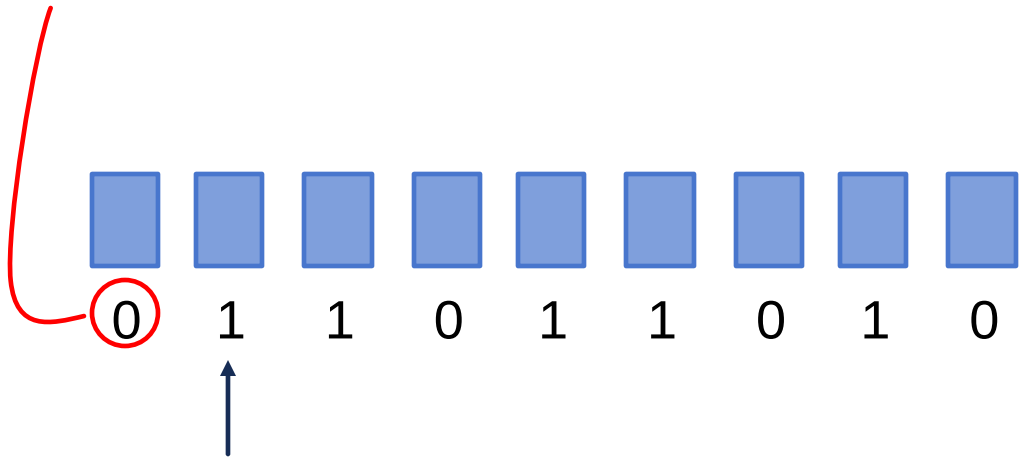

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

Page fault just occurred, need to replace frame … walking through memory

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

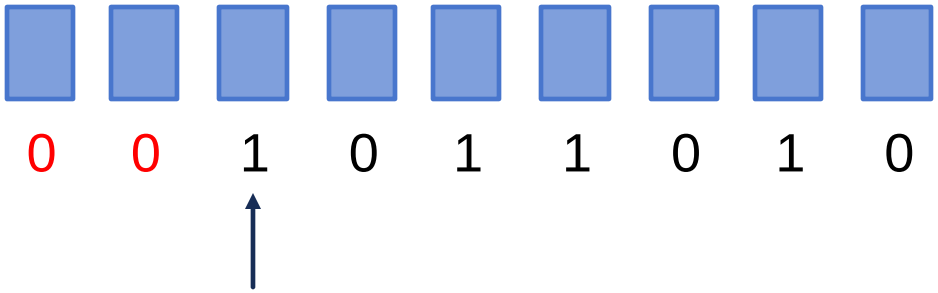

Page fault just occurred, need to replace frame … walking through memory

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

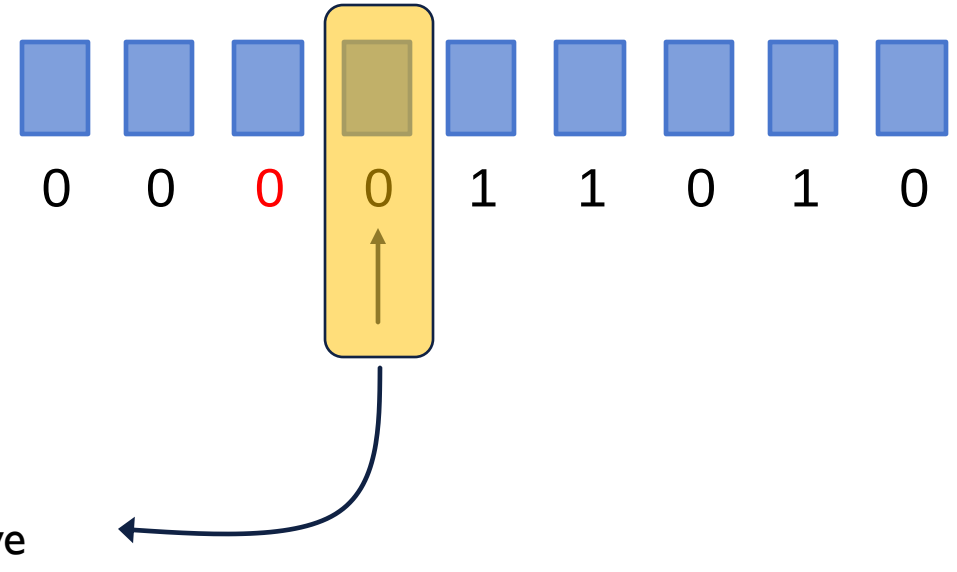

When we find a zero, we remove the frame.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

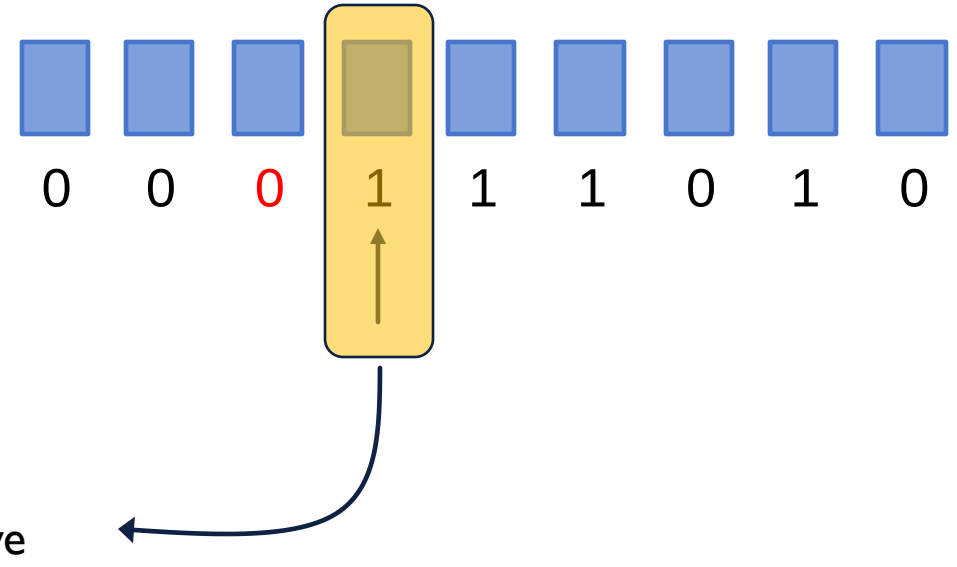

When we find a zero, we remove the frame.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

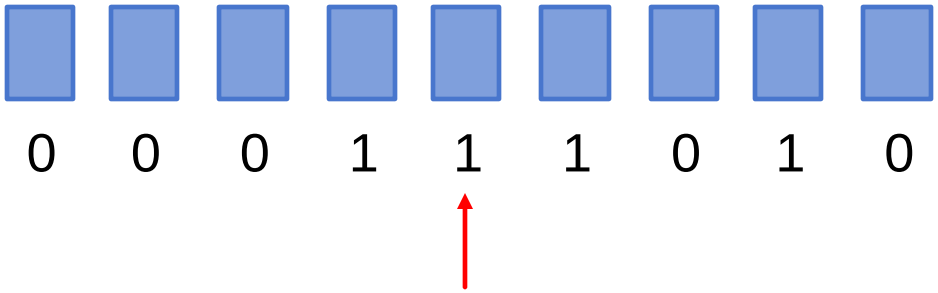

Worksheet Q1

Next replacement search continue from last pointer position.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

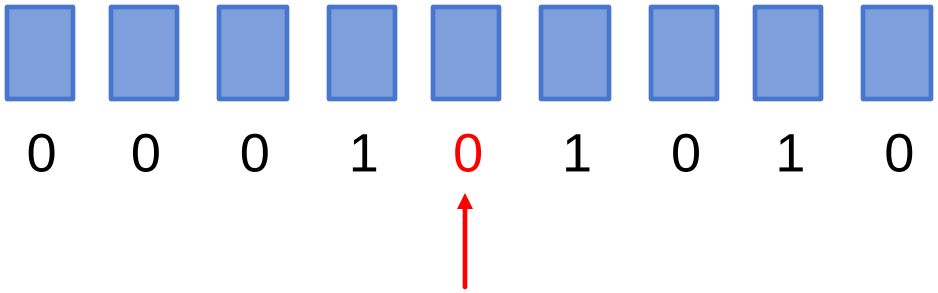

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

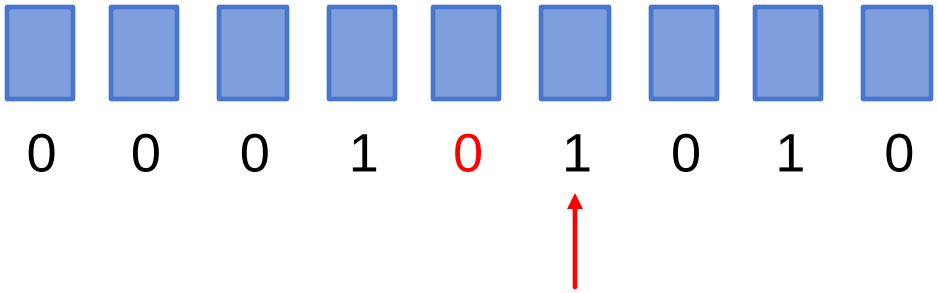

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

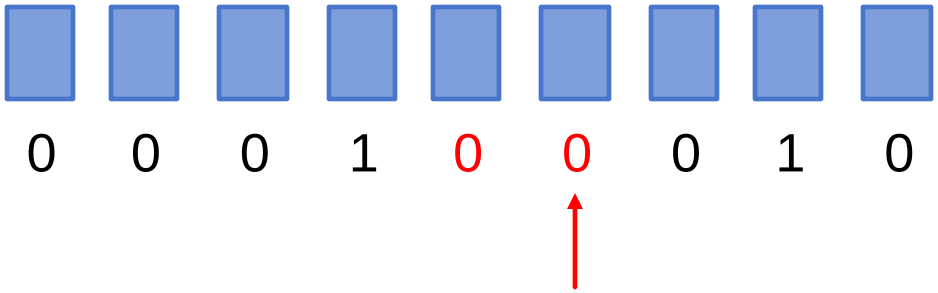

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

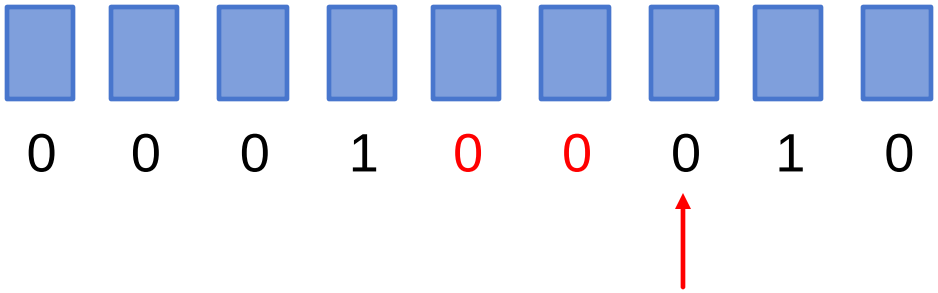

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

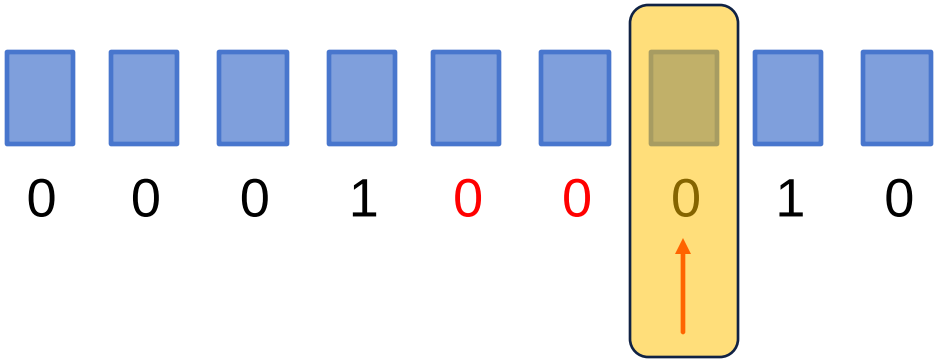

Next replacement search continue from last pointer position.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

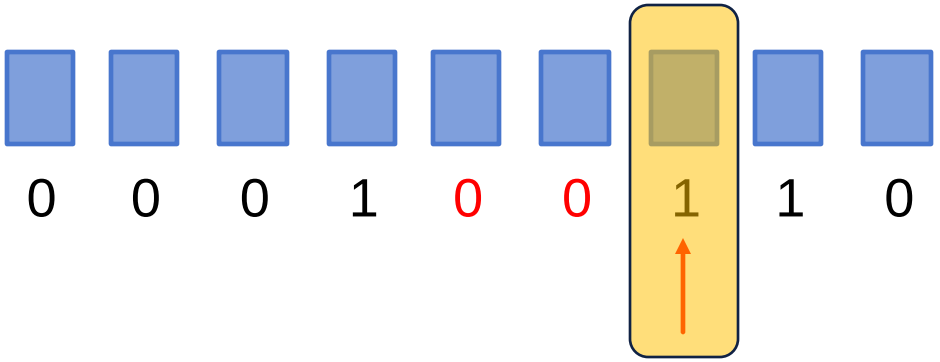

Next replacement search continue from last pointer position.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

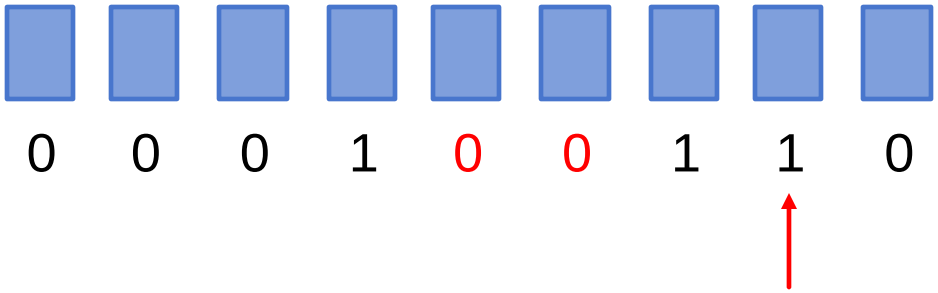

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

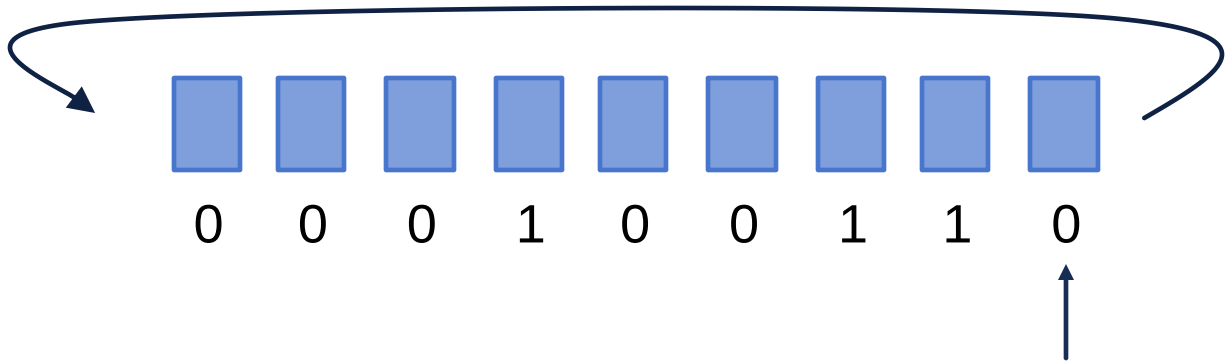

Circle back once you go through all memory frame …

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

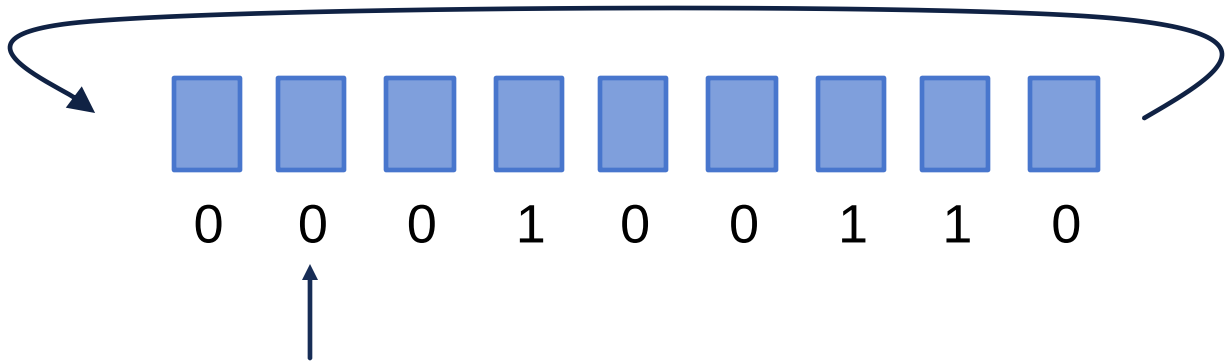

- Why is it called second chance?
- We remove frames that has not been referenced since our last "visit".
- This frame has stayed unreferenced since we reset it to '0'.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

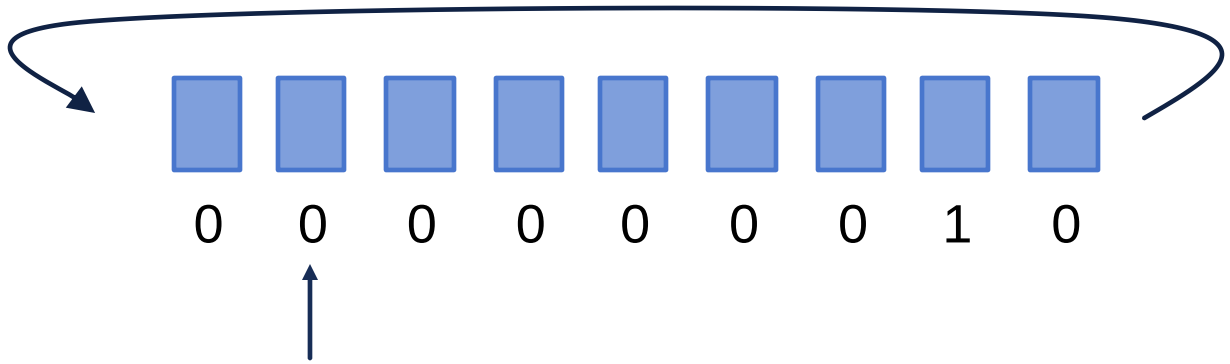

- Advantages?

Worksheet Q2

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

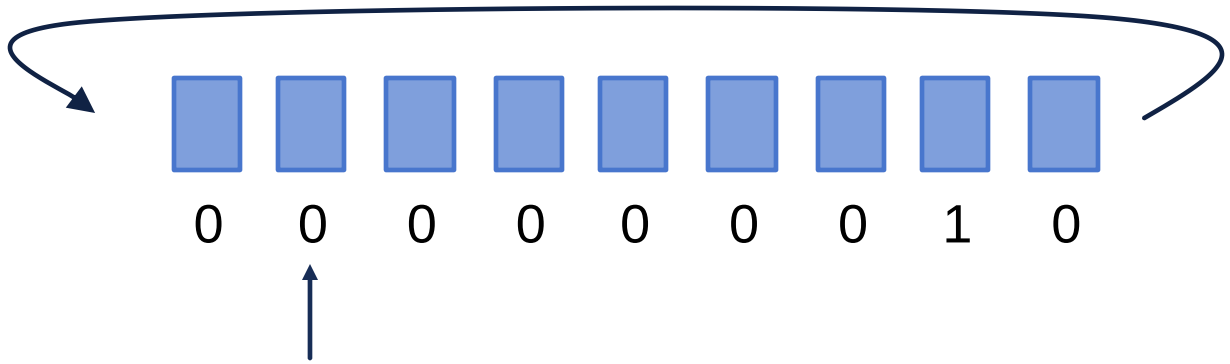

- Advantages?
- No need to pause programs to reset bits, we reset them during search.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

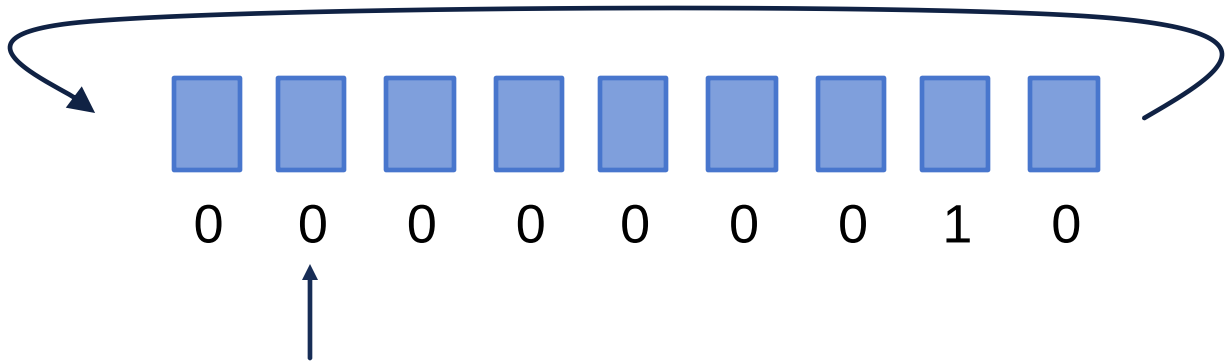

- Advantages?
- No need to pause programs to reset bits, we reset them during search.
- No need to reset all bits, we rest the ones that are relevant.

SECOND CHANCE

Instead of resetting periodically, we reset reference bits *while* searching for a not-recently-used page.

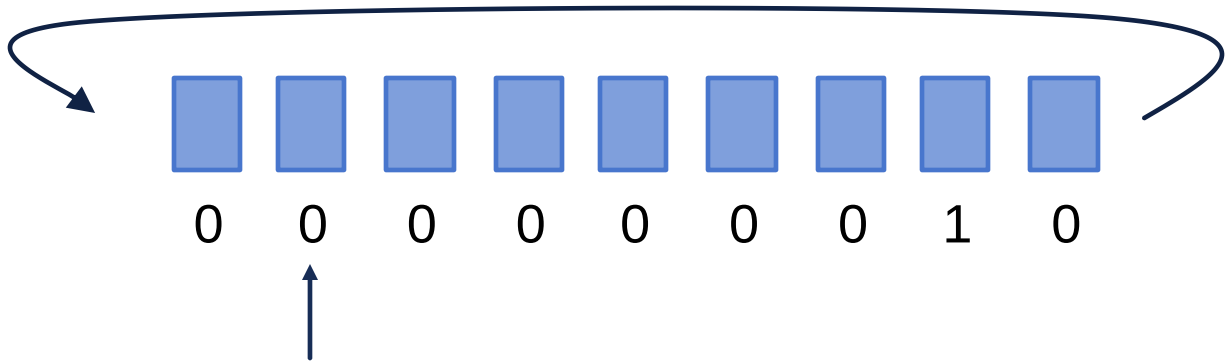

- Advantages?
- No need to pause programs to reset bits, we reset them during search.
- No need to reset all bits, we rest the ones that are relevant.
- Less stalls and shorter penalty.

FRAME ALLOCATION

- So far we talked about how to replace frames.

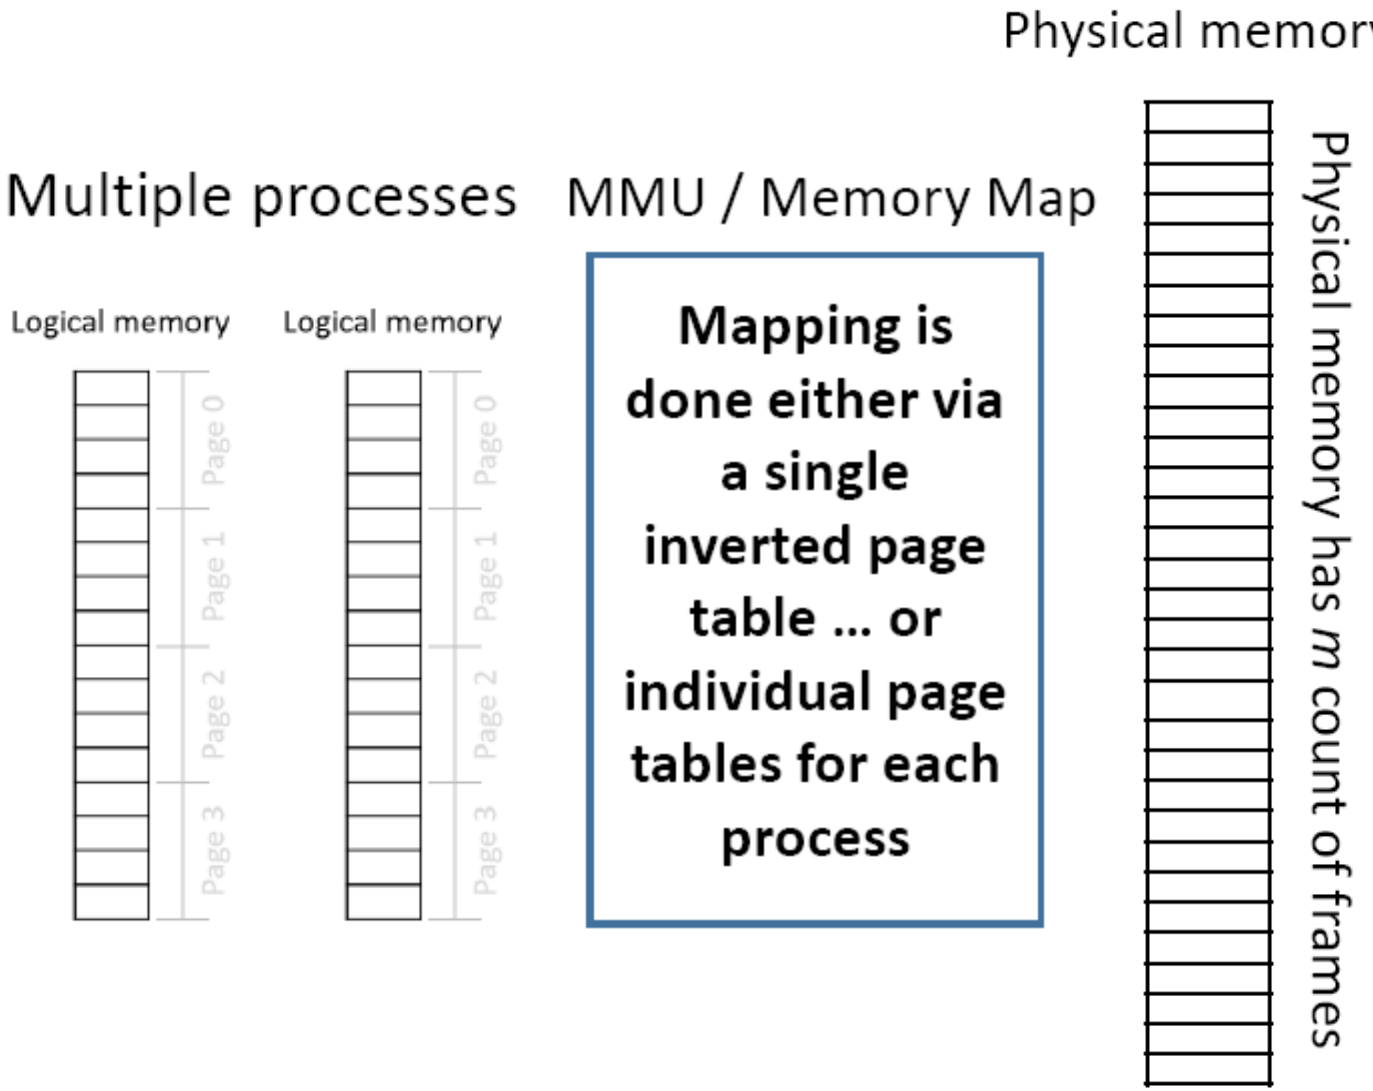

FRAME ALLOCATION

- How to allocate frames?
- Divide all frames among process equally?
- Allocate on demand?

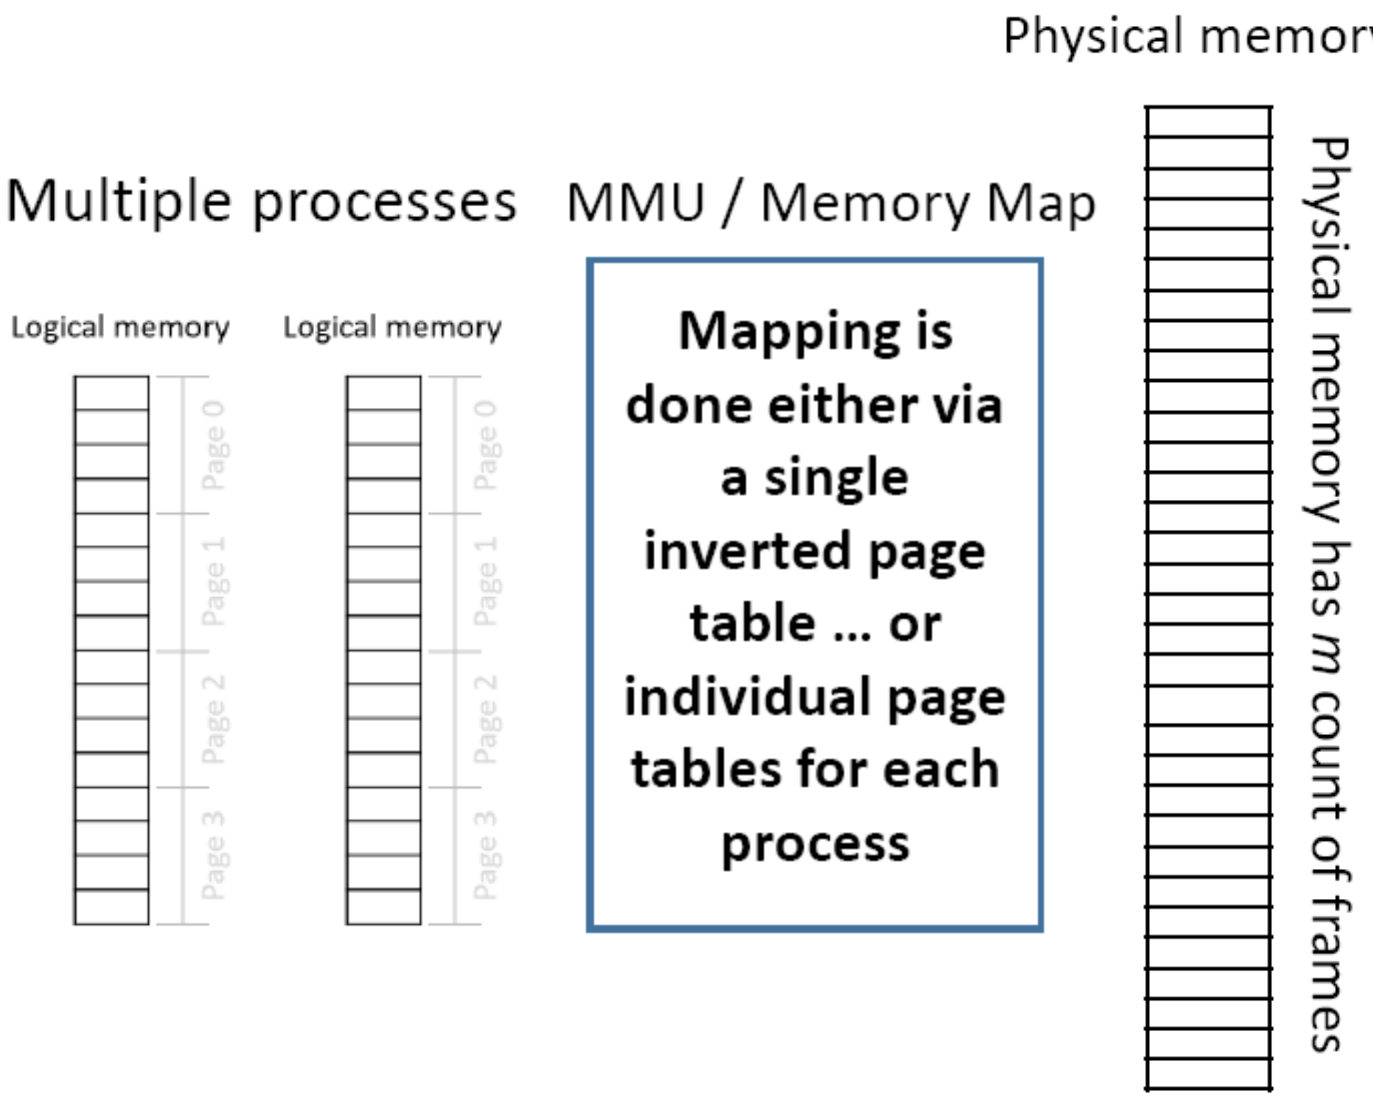

FRAME ALLOCATION

- As the number of frames allocated for each process decreases, page fault rate goes up.

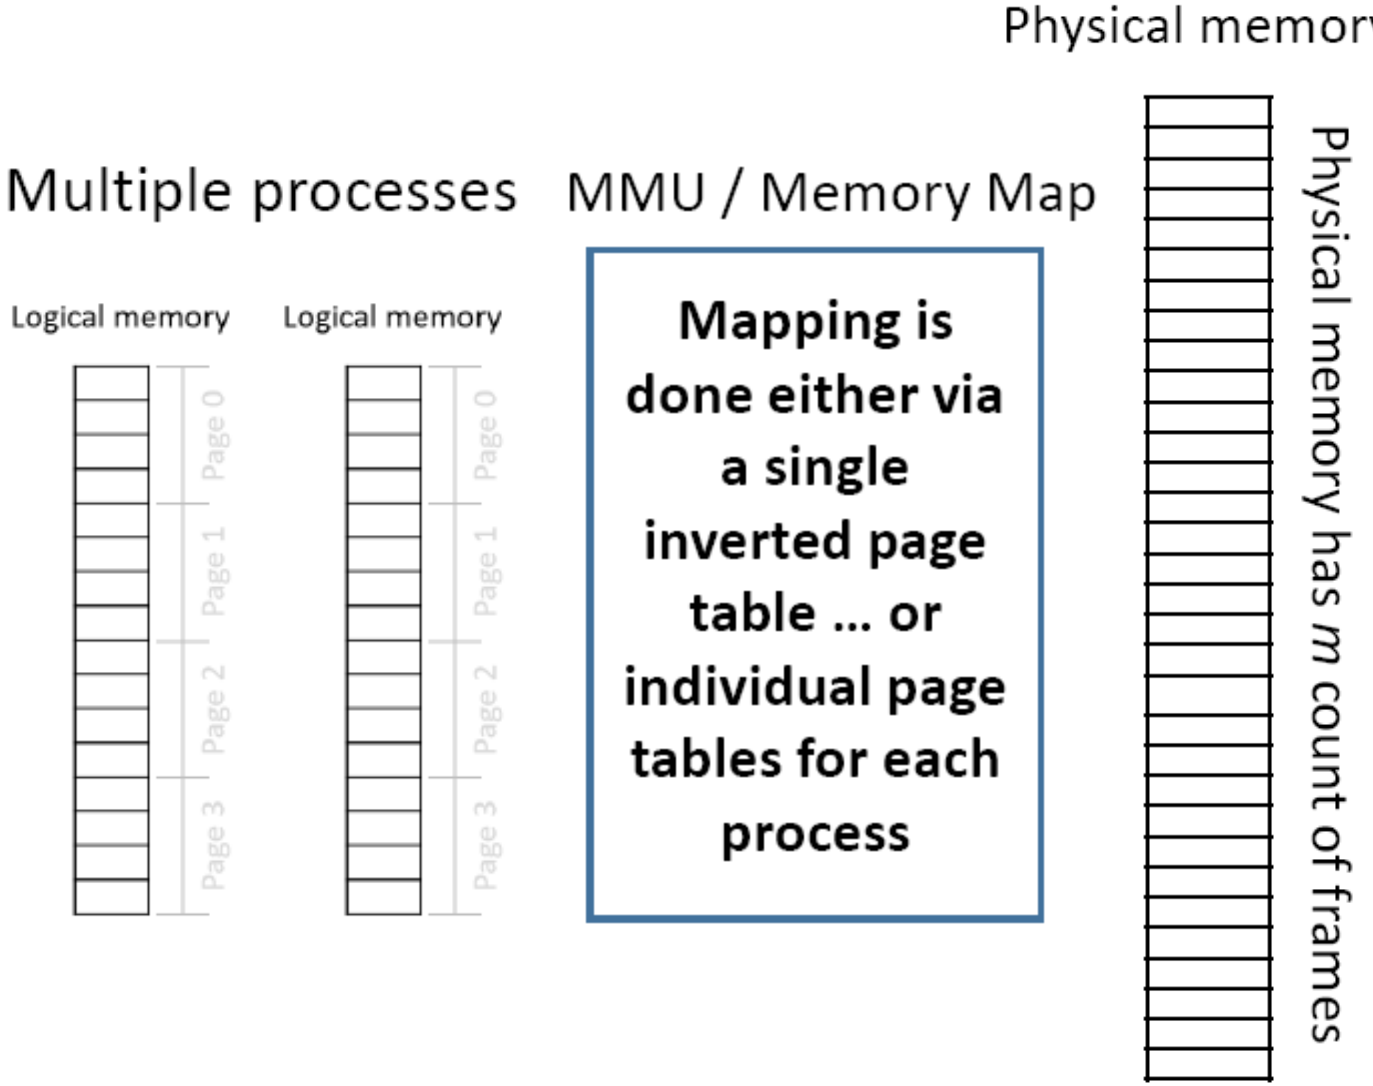

FRAME ALLOCATION

- Q: What is the minimum number of frames to allocate for each process?

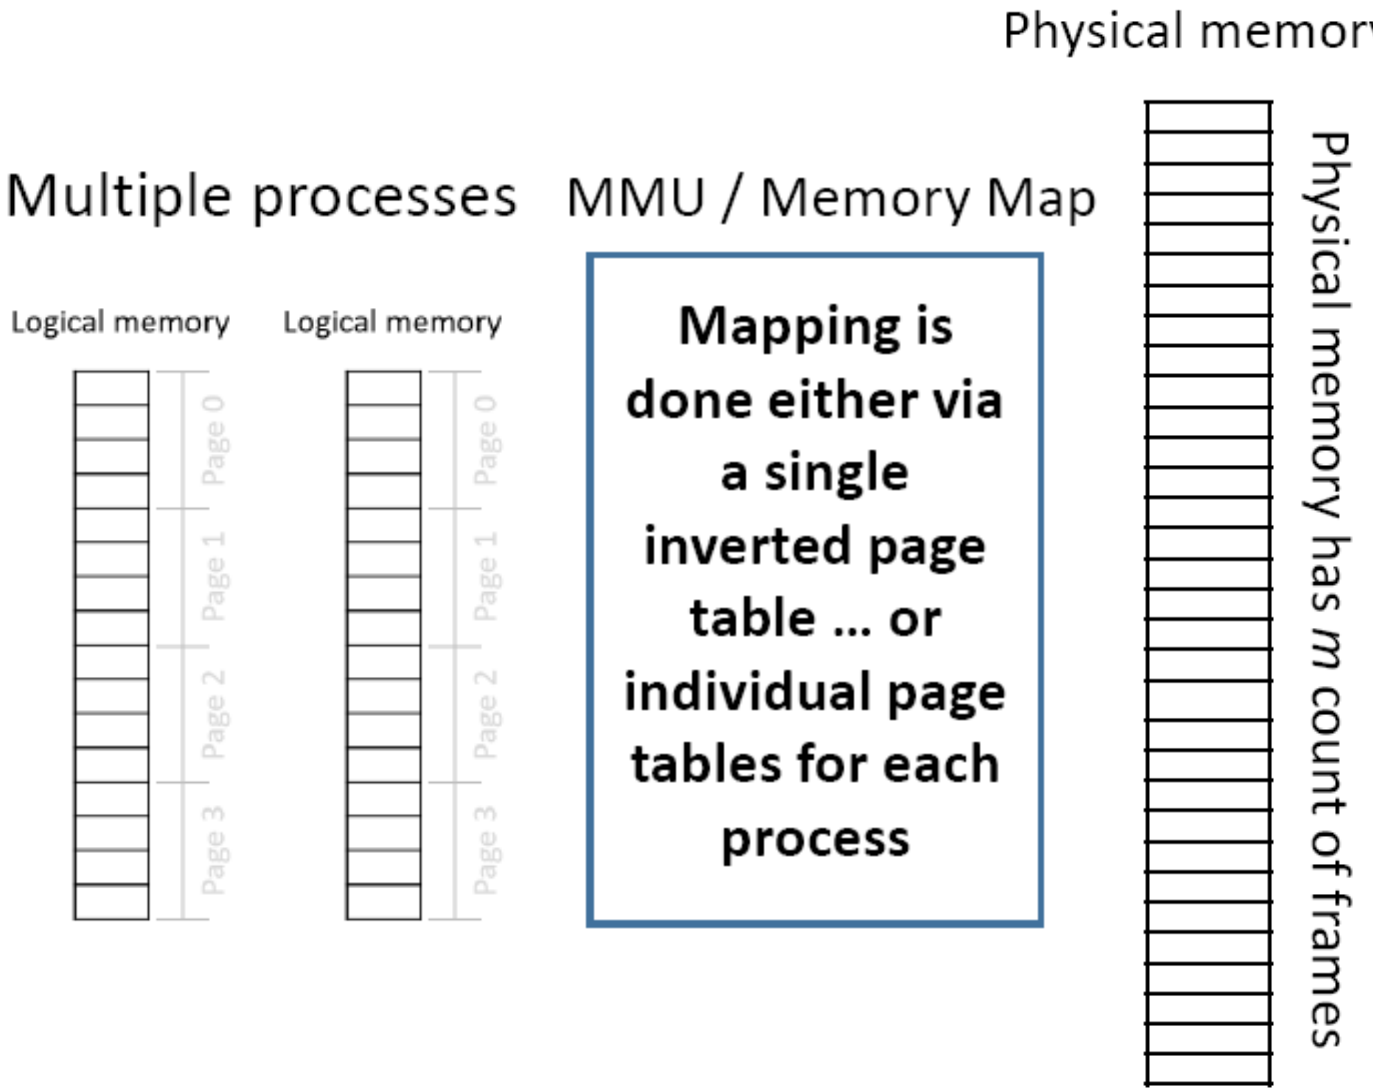

FRAME ALLOCATION

- Q: What is the minimum number of frames to allocate for each process?

Decided by architecture …

Q: How many memory references can occur in a single instruction?

- This can be a lot, 8 or even 16 in some architectures … why?
- Indirect addressing, where address can point to another address ..

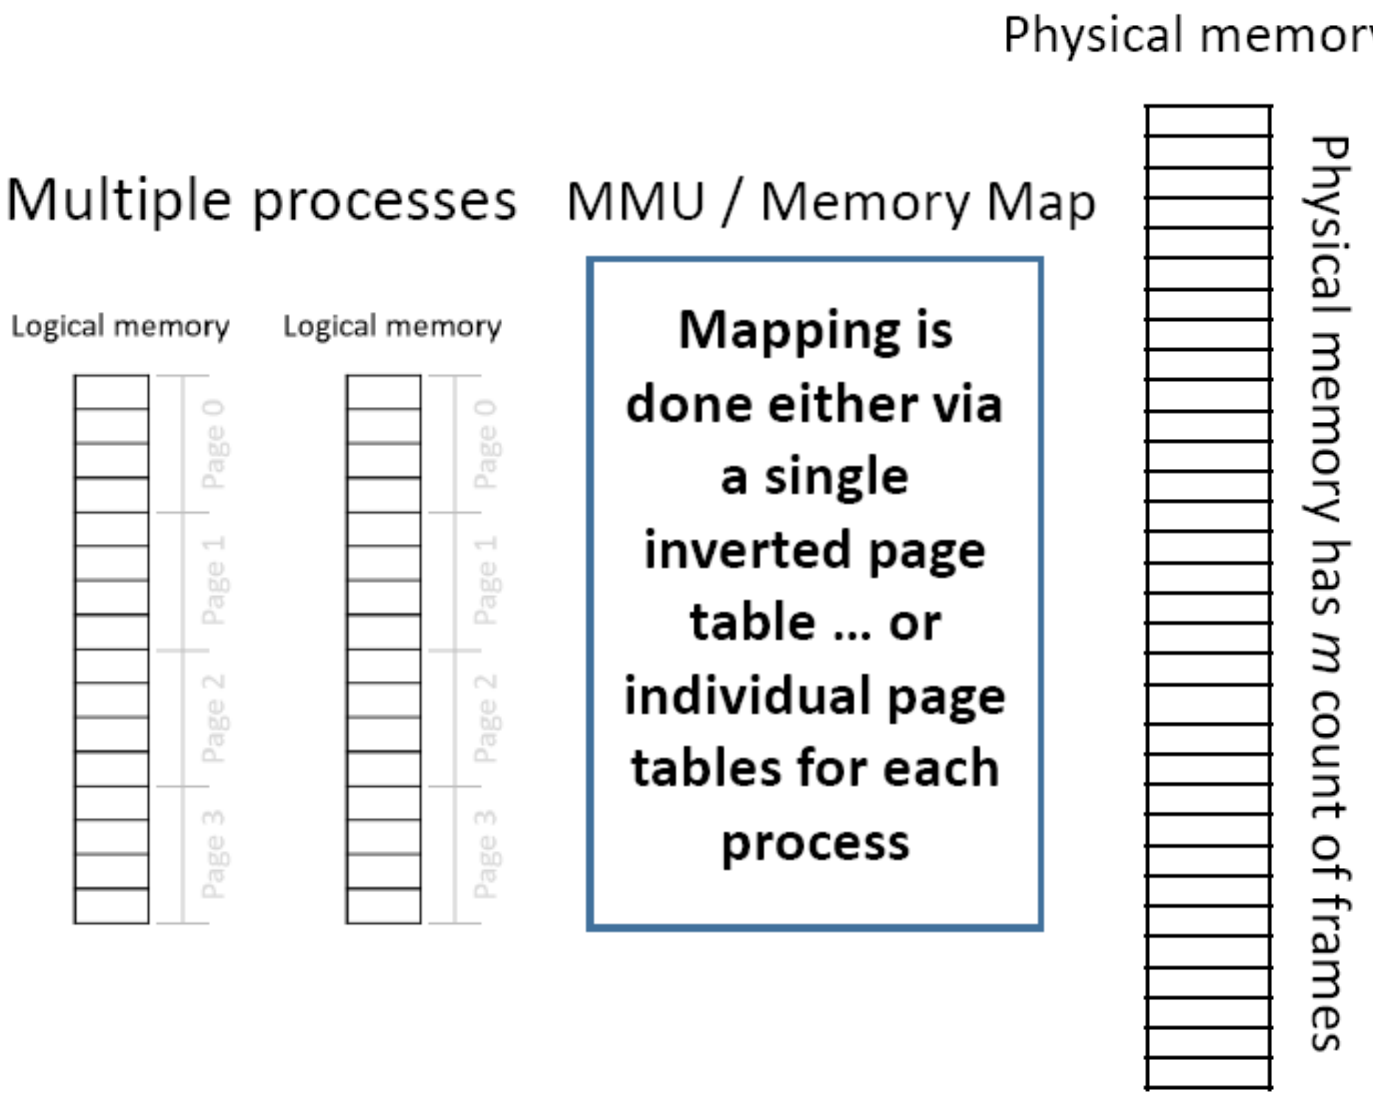

FRAME ALLOCATION

- Q: What is the minimum number of frames to allocate for each process?

Decided by architecture …

Q: How many memory references can occur in a single instruction?

- This can be a lot, 8 or even 16 in some architectures … why?
- Indirect addressing, where address can point to another address ..

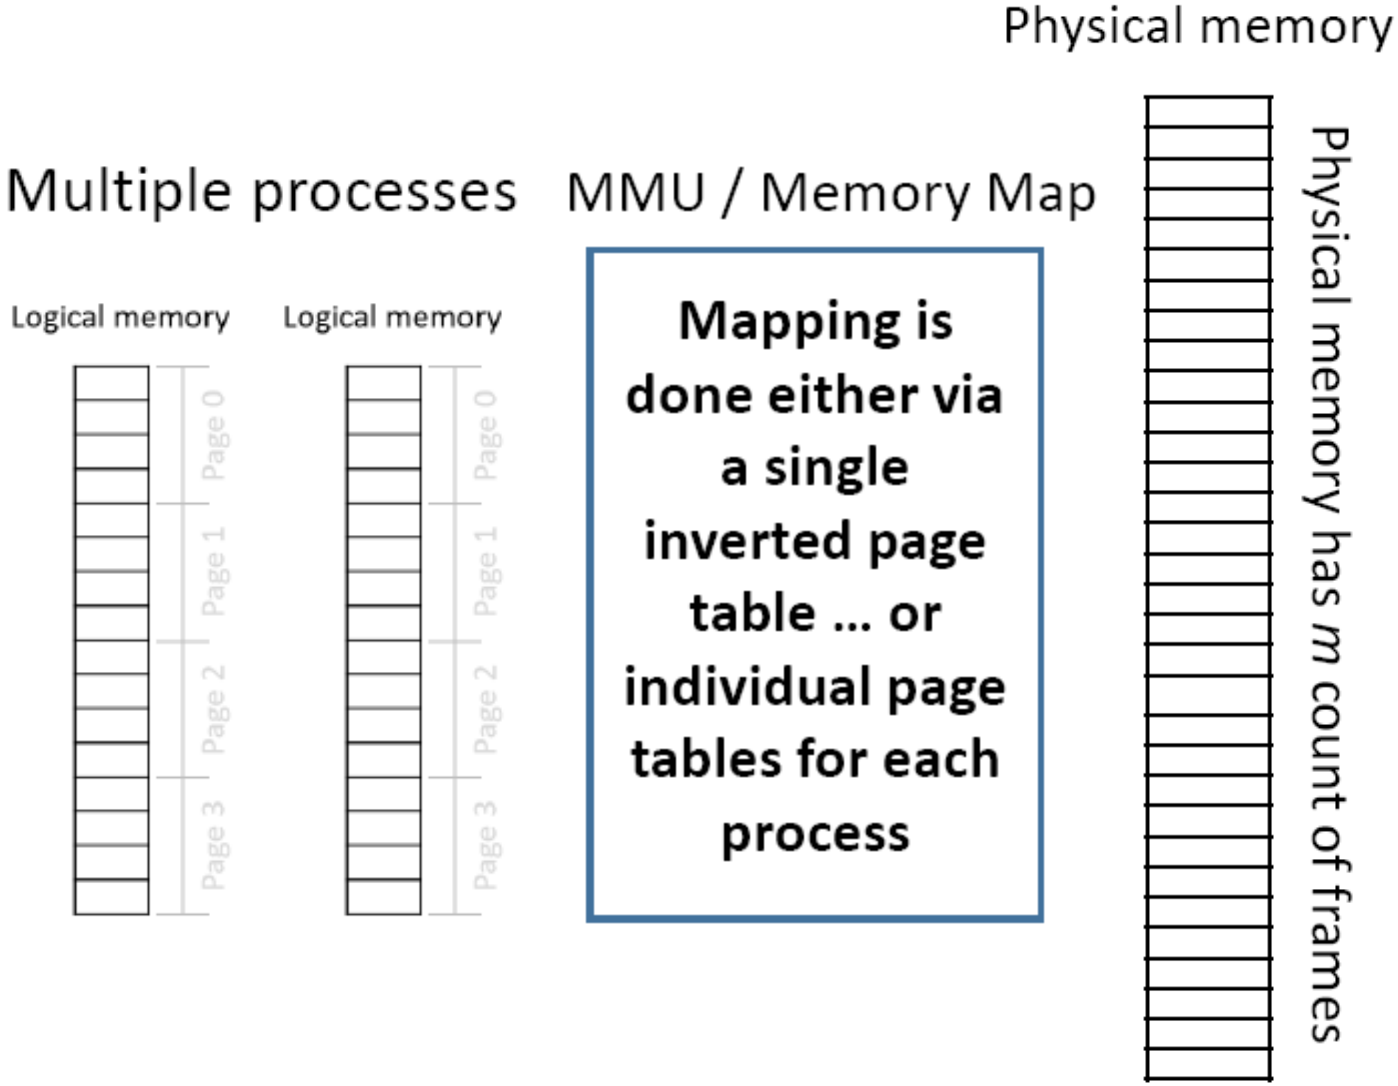

It's simply the number of different memory locations that can be accessed in a *single* instruction.

EQUAL ALLOCATION

- Split the memory equally among all processes.
- Suppose we have 100 frames and 5 process … each would have 20 frames.
- Pretty simple …
- **Q: What are the drawbacks of this approach?**

EQUAL ALLOCATION

- Split the memory equally among all processes.
- Suppose we have 100 frames and 5 process … each would have 20 frames.
- Pretty simple …
- **Q: What are the drawbacks of this approach?**
  - Not all processes need the same amount of frames!
  - Some might need many others might need fewer pages … not very efficient use of memory

EQUAL ALLOCATION

- Split the memory equally among all processes.
- Suppose we have 100 frames and 5 process … each would have 20 frames.
- Pretty simple …
- **Q: What are the drawbacks of this approach?**
  - Not all processes need the same amount of frames!
  - Some might need many others might need fewer pages … not very efficient use of memory
  - Alternative: proportional memory allocation.

PROPORTIONAL MEMORY ALLOCATION

- Allocate frames proportional to process memory size.
- Process twice the size will get twice the number of frames.

Worksheet Q3: Any inefficiencies?

PROPORTIONAL MEMORY ALLOCATION

- Allocate frames proportional to process memory size.
- Process twice the size will get twice the number of frames.

Worksheet Q3: Any inefficiencies?

- Processes do not need all their pages in memory …
- The number of frames required is not always proportional to process size …
- The number for frames required *might vary* during run time.

THRASHING

- As we add more processes for multitasking we increase CPU utilization.

THRASHING

- As we add more processes for multitasking we increase CPU utilization.

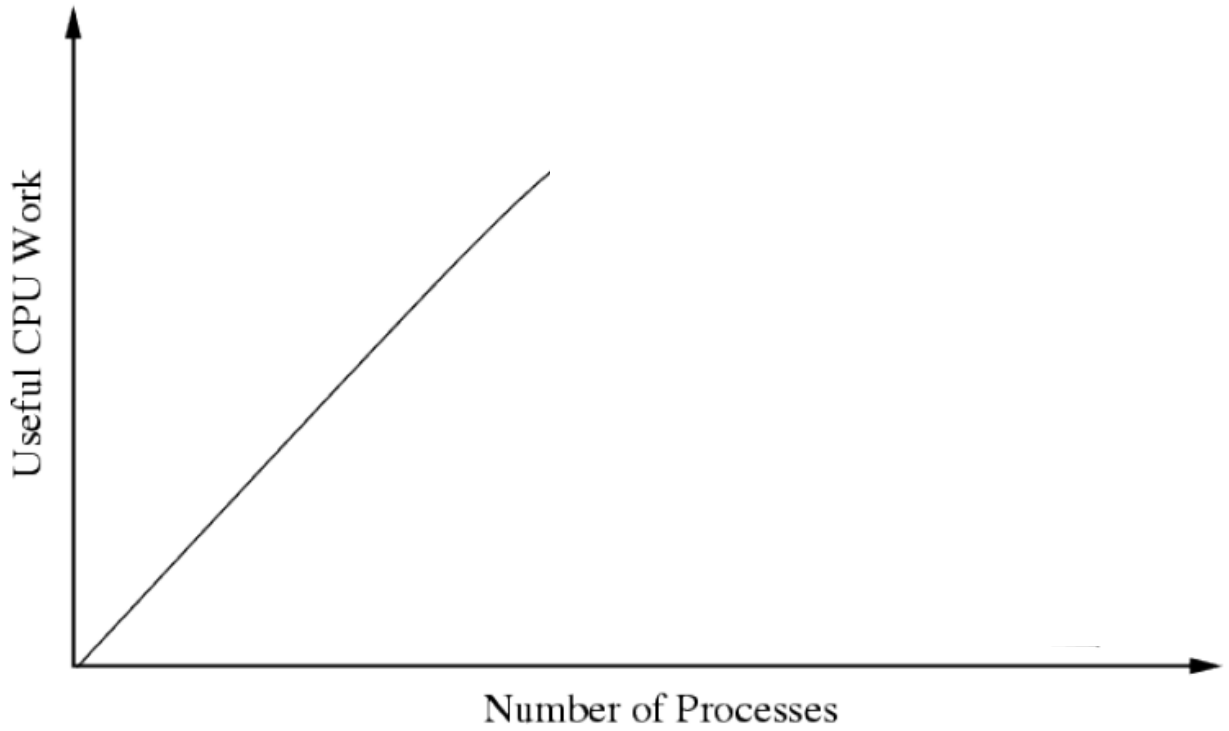

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.

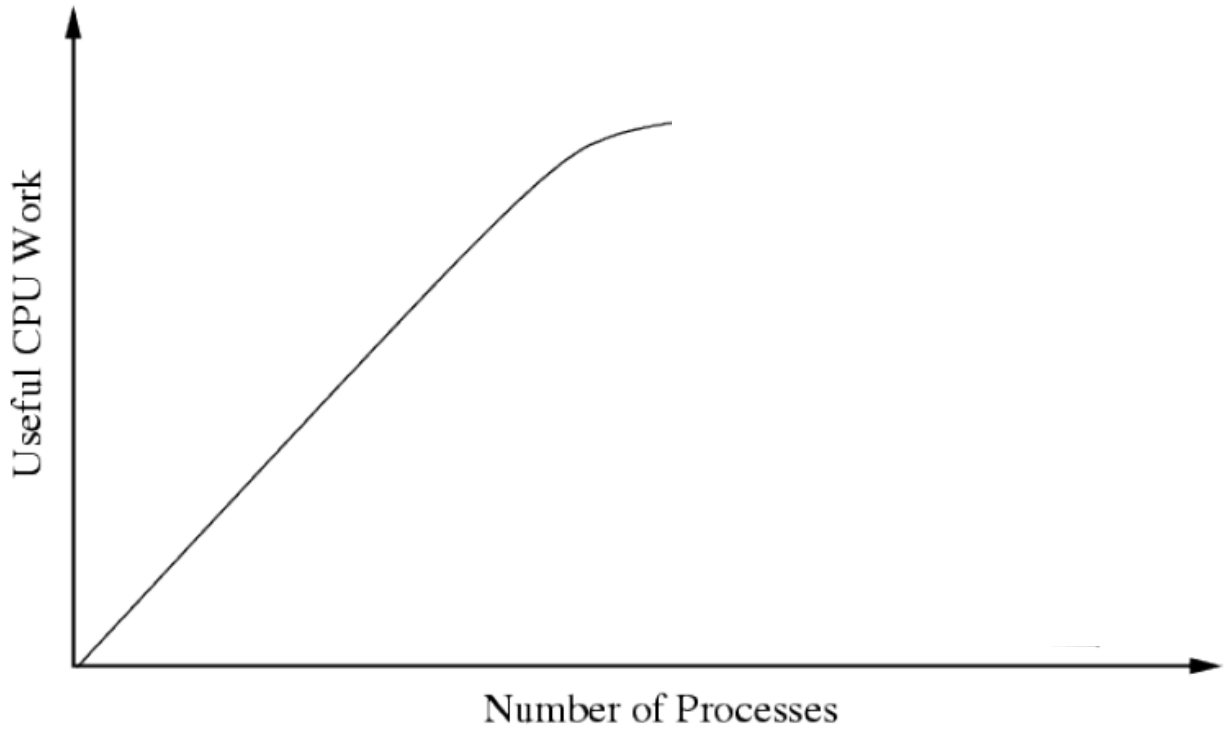

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.

Q: When page faults increase, what happens to CPU utilization

A B

A : Increases

B : Decreases

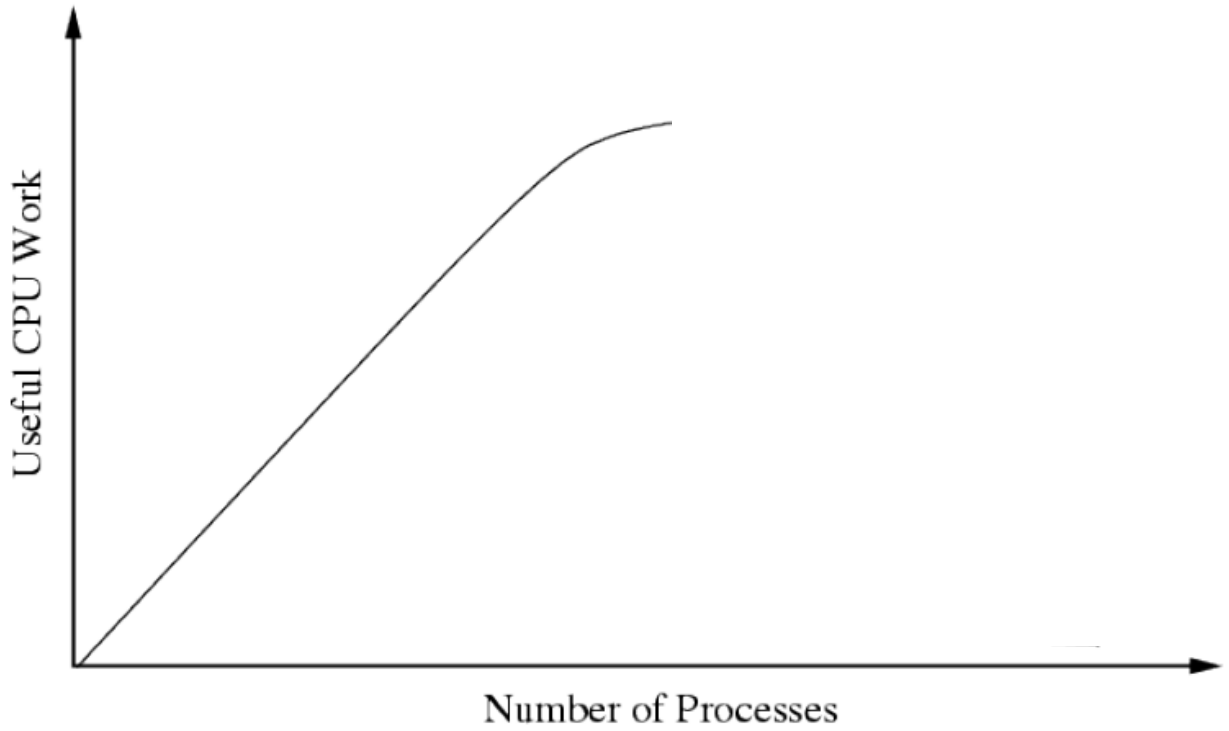

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.

Q: When page faults increase, what happens to CPU utilization

A B

A : Increases

B : Decreases

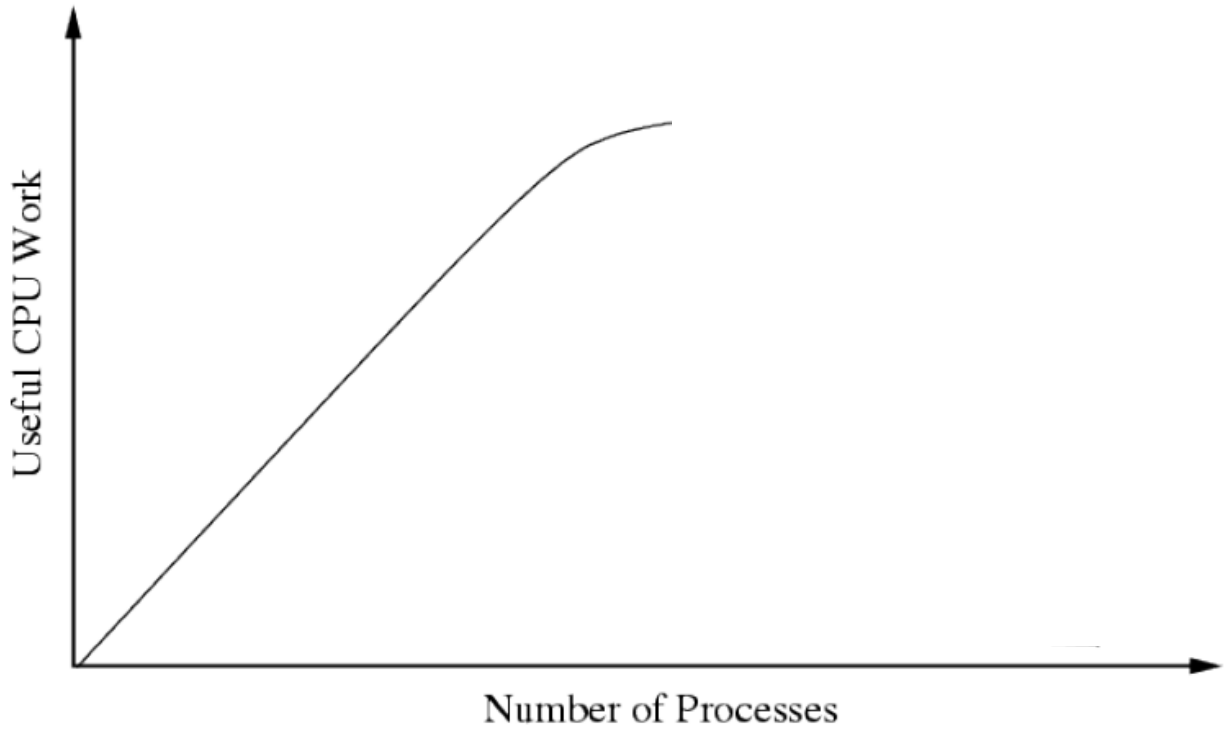

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.
- How does the Operating System react?
- The OS would think that CPU is under utilized, and as such it would retrieve more processes for concurrent execution.

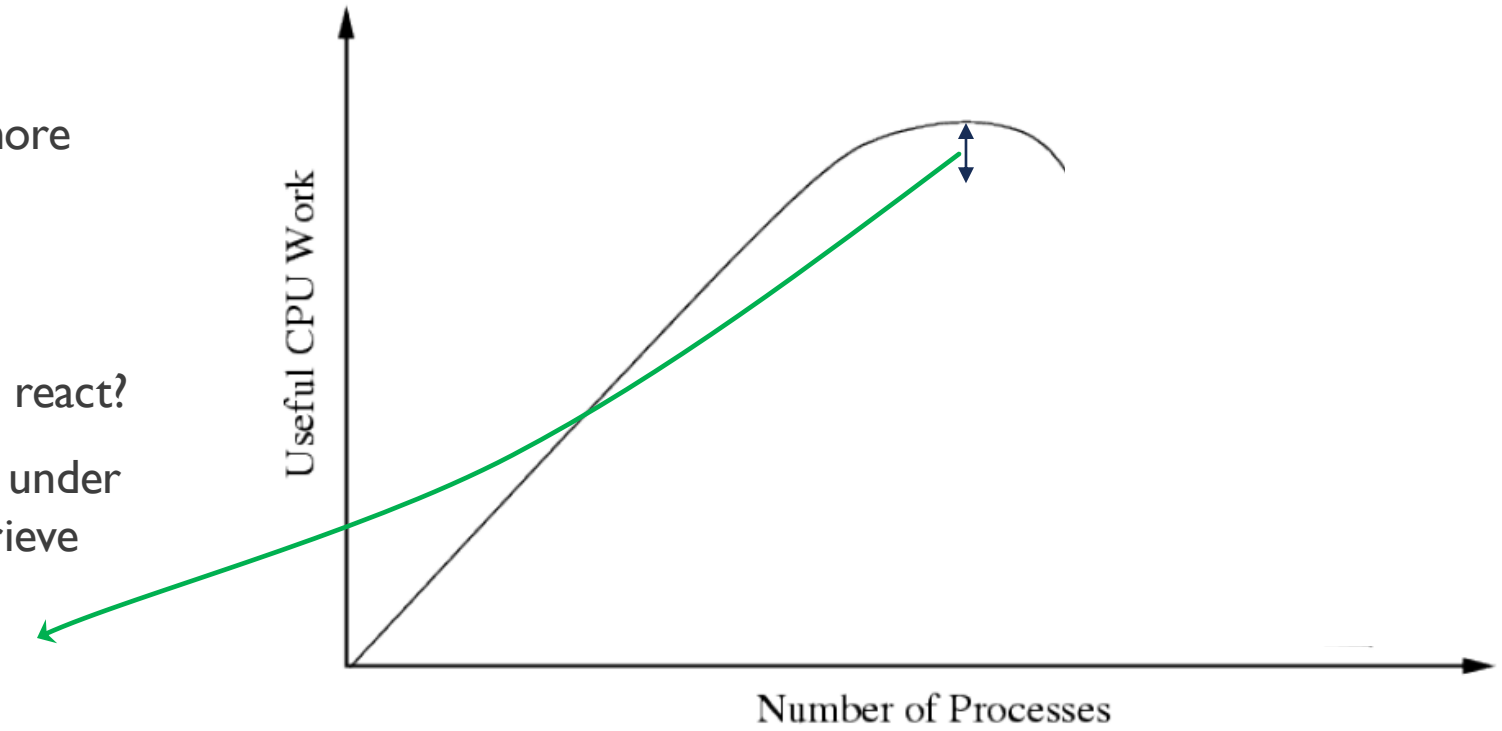

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.
- How does the Operating System react?
- The OS would think that CPU is under utilized, and as such it would retrieve more processes for concurrent execution.
- This causes even more page faults.

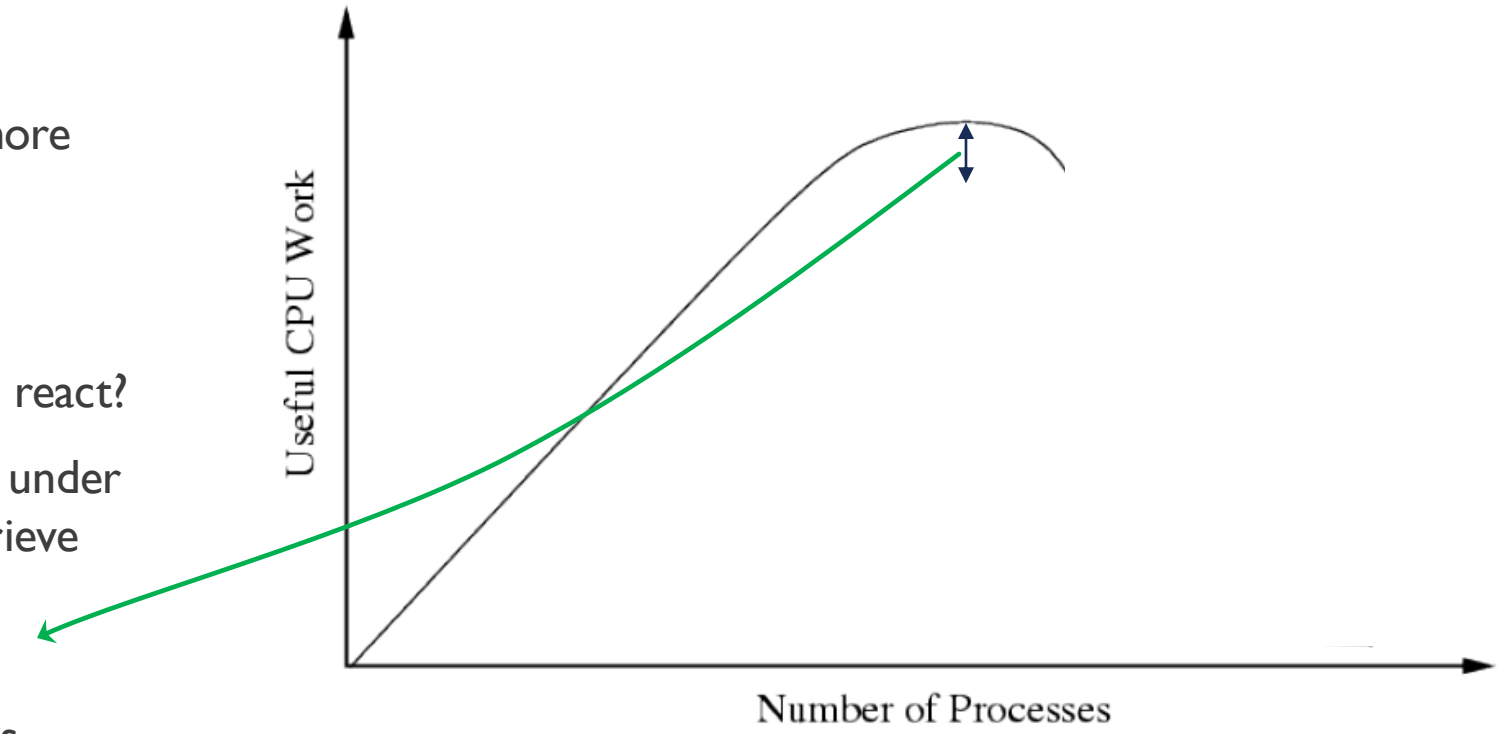

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.
- How does the Operating System react?
- The OS would think that CPU is under utilized, and as such it would retrieve more processes for concurrent execution.
- This causes even more page faults.

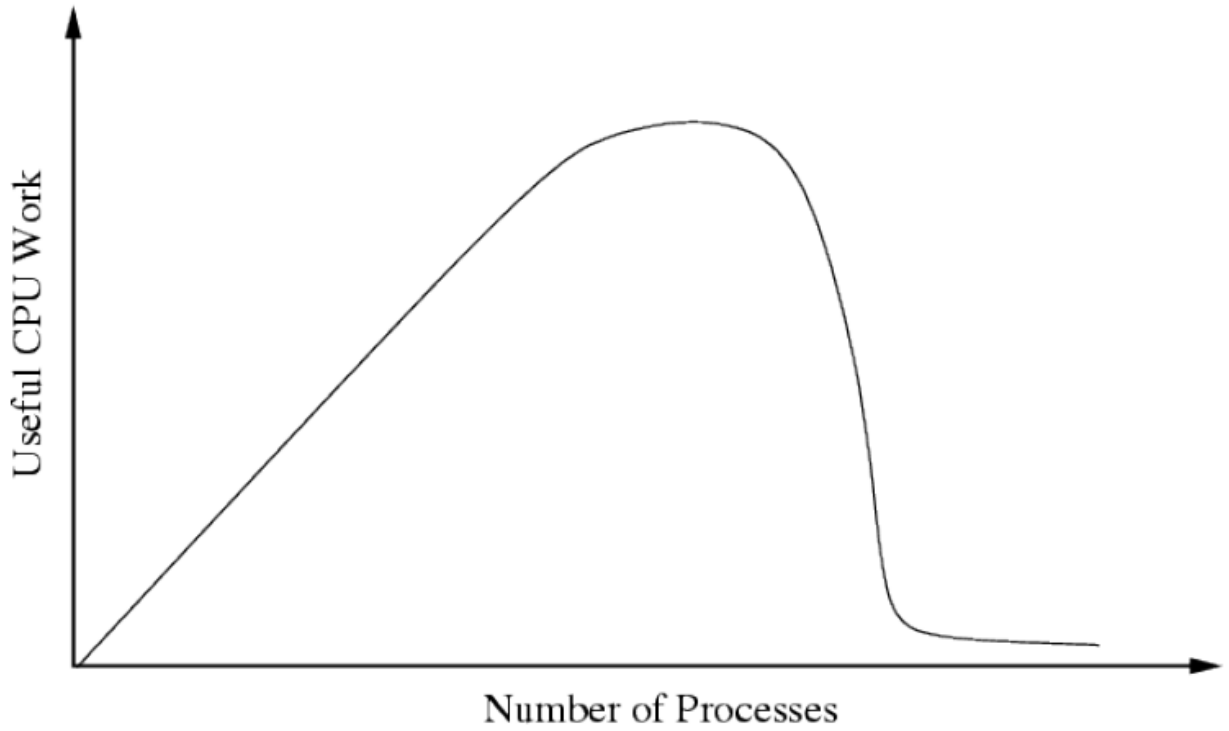

THRASHING

- As we add more processes for multitasking we increase CPU utilization.
- We reach a point when adding more processes compete for frames in memory and page faults start to increase.
- How does the Operating System react?
- The OS would think that CPU is under utilized, and as such it would retrieve more processes for concurrent execution.

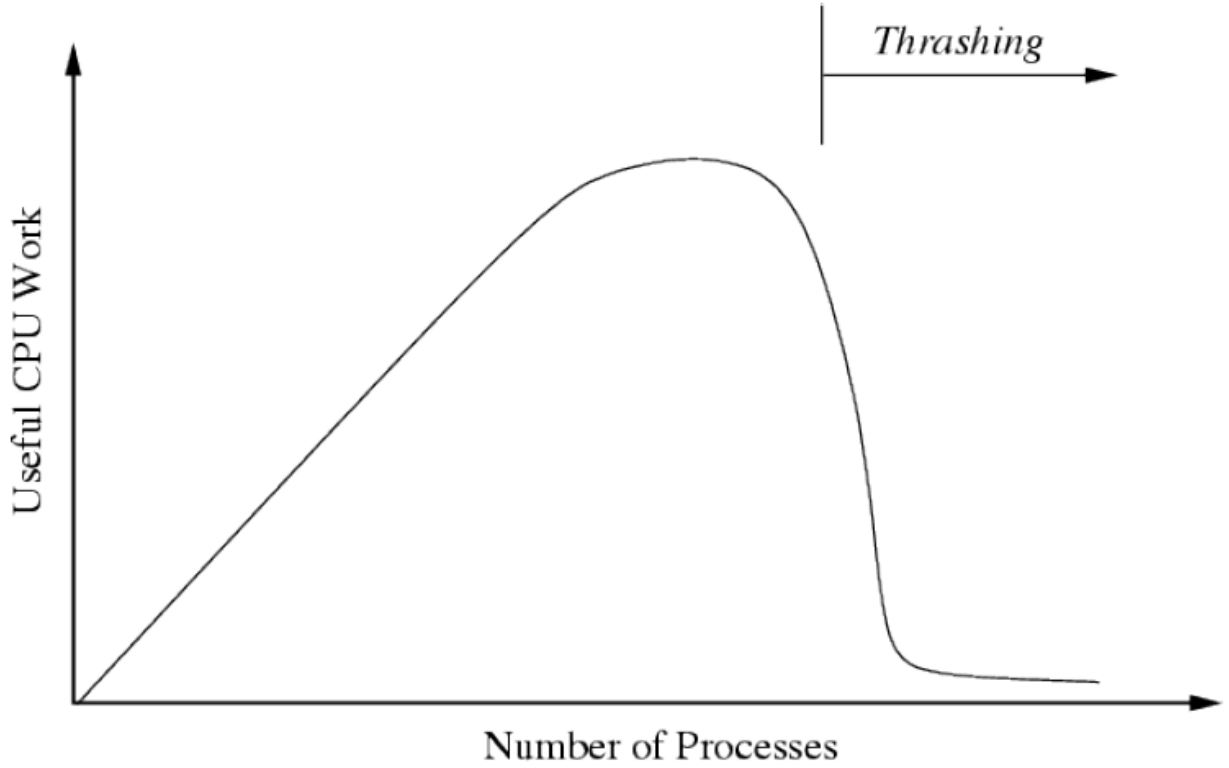

HOW TO PREVENT THRASHING

- We can allow for only local frame replacement … a process can only replace frames it owns.

HOW TO PREVENT THRASHING

- We can allow for only local frame replacement … a process can only replace frames it owns.
- To truly prevent thrashing, we need to know how much frames each process needs for "healthy" execution.
- This is different than the minimum amount of frames that is dependent on architecture.

HOW TO PREVENT THRASHING

- We can allow for only local frame replacement … a process can only replace frames it owns.
- To truly prevent thrashing, we need to know how much frames each process needs for "healthy" execution.
- This is different than the minimum amount of frames that is dependent on architecture.
- How to know how much frames do we need for each process?

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

1 2 4 5 6 4 2 5 6 5 1 2 6 8 2 4 5 6 7 2 3 3 3 2 2 3 3 3 2 3

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

1 2 4 5 6 4 2 5 6 5 1 2 6 8 2 4 5 6 7 2 3 3 3 2 2 3 3 3 2 3

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

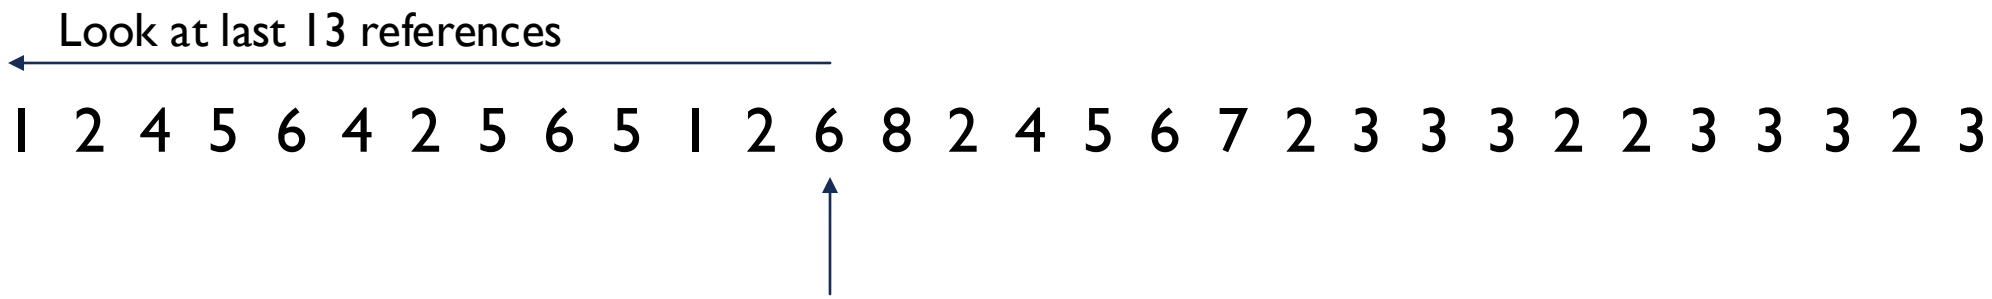

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

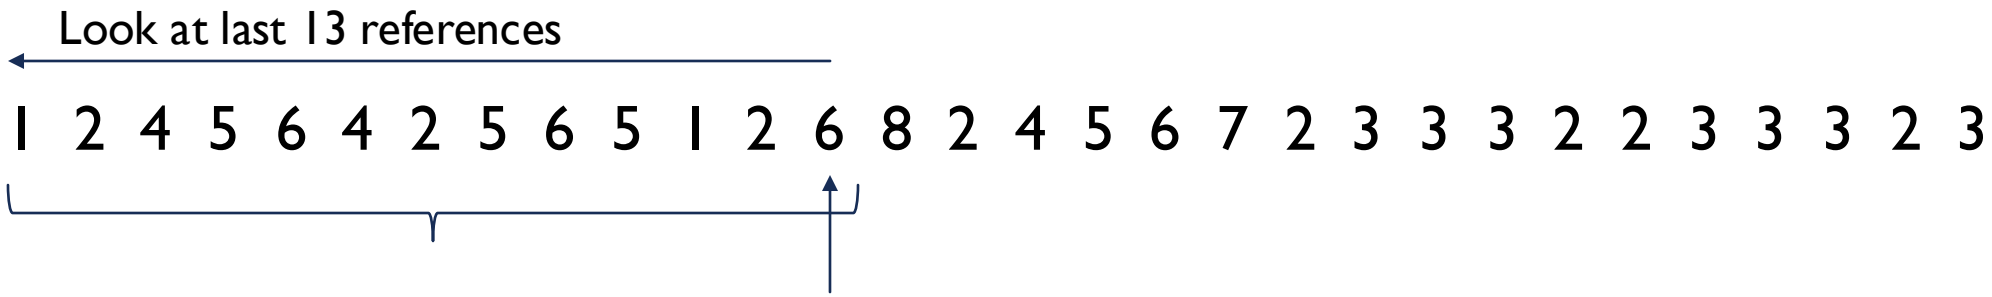

{1,2,4,5,6} = Working Set

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

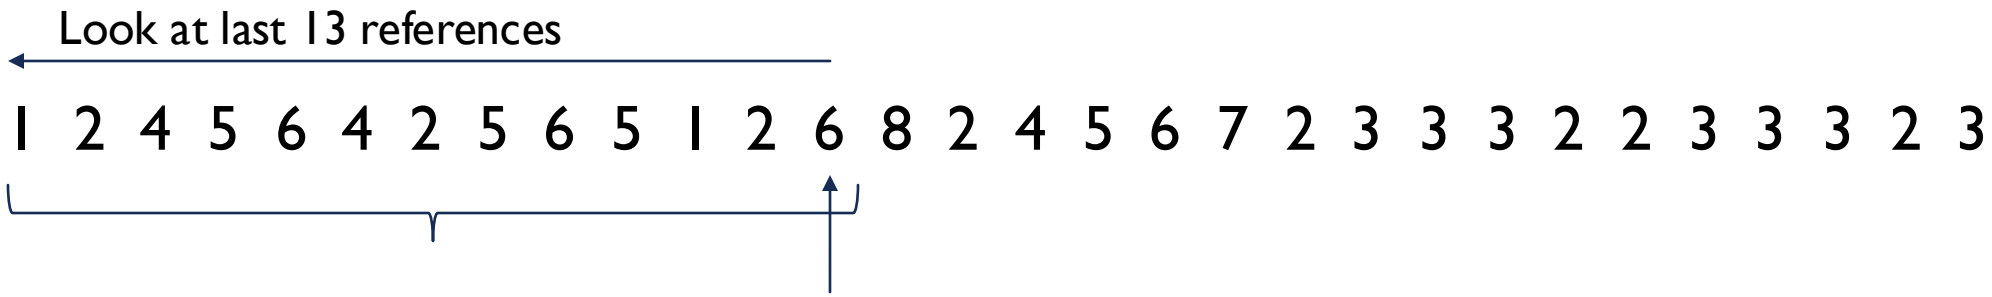

{1,2,4,5,6} = Working Set
Working Set Size = 5

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

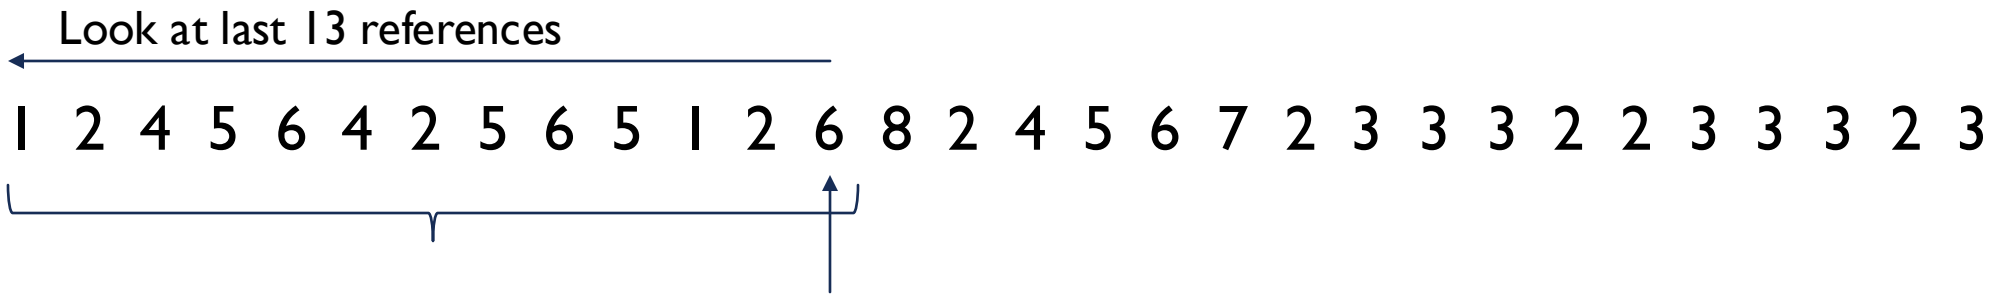

{1,2,4,5,6} = Working Set
Working Set Size = 5
OS allocates 5 frames for P1

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

Worksheet Q4

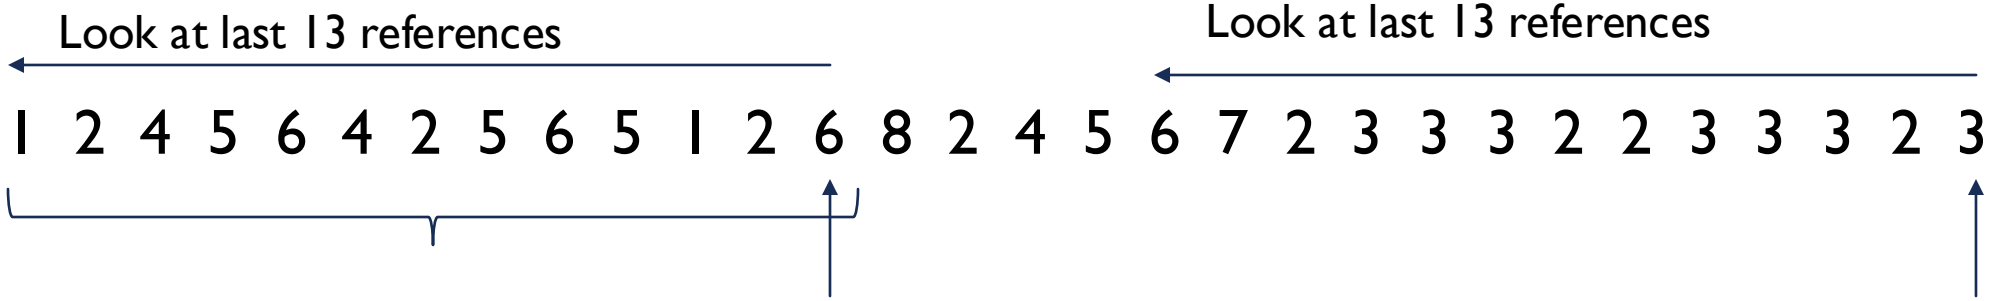

{1,2,4,5,6} = Working Set
Working Set Size = 5
OS allocates 5 frames for P1

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

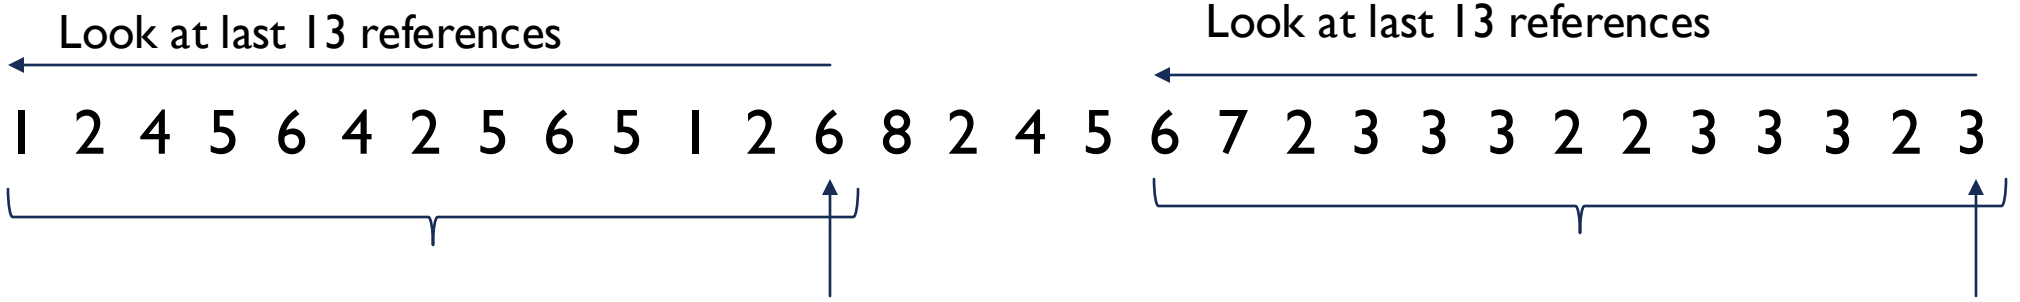

{1,2,4,5,6} = Working Set
Working Set Size = 5
OS allocates 5 frames for P1

{2,3,6,7} = Working Set

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

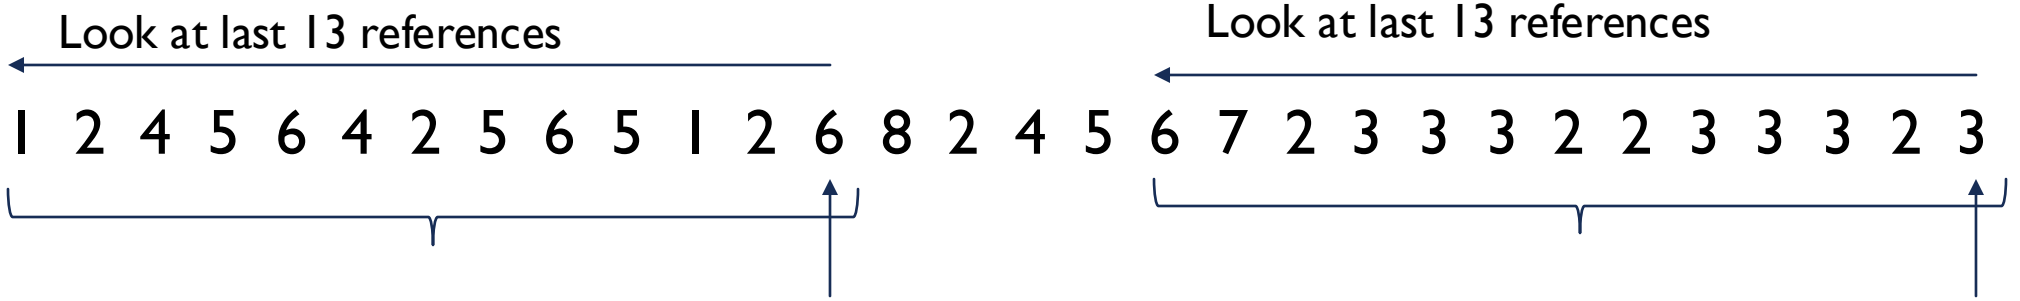

{1,2,4,5,6} = Working Set
Working Set Size = 5
OS allocates 5 frames for P1

{2,3,6,7} = Working Set
Working Set Size = 4

WORKING SET

Goal : Prevent Thrashing

Main idea : Consider what pages a process has needed in the recent past, as an indicator of the pages it will need in the future

Working Set: a set of pages that are actively used together

Recent Memory References: $\Delta = 13$

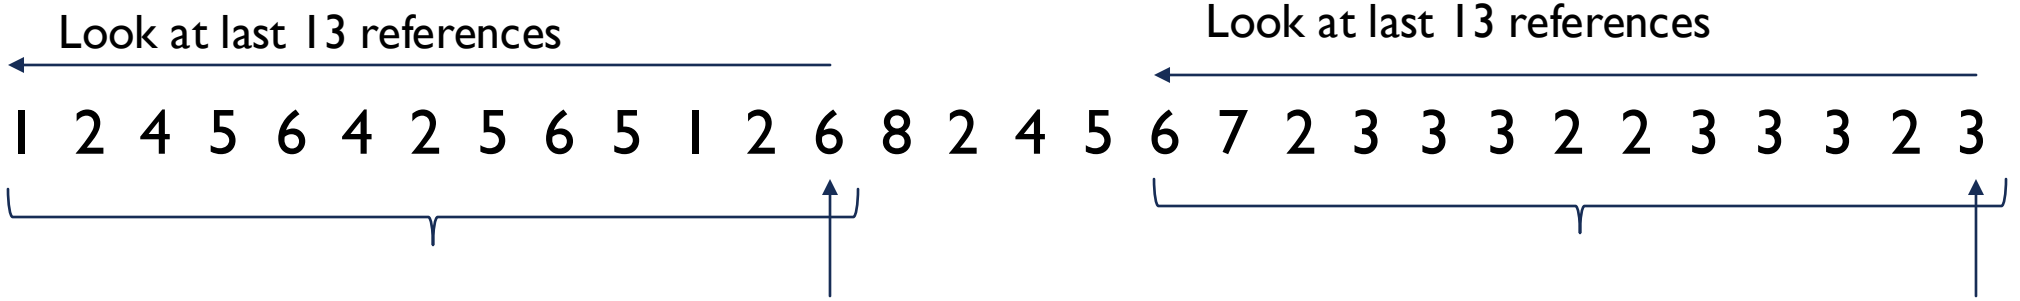

{1,2,4,5,6} = Working Set
Working Set Size = 5
OS allocates 5 frames for P1

{2,3,6,7} = Working Set
Working Set Size = 4
OS allocates 4 frames for P1

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

Q: If the WS Size (WSS) of process *i* is WSS_i, then what is the total demand of the most active frames for *n* processes?

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

Q: If the WS Size (WSS) of process *i* is WSS_i, then what is the total demand of the most active frames for *n* processes?

$$D = \sum_{i=0}^{n} WSS_i$$

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

Q: If the WS Size (WSS) of process *i* is WSS_i, then what is the total demand of the most active frames for *n* processes?

$$D = \sum_{i=0}^{n} WSS_i$$

If the total demand D is greater than the count of available frames, thrashing will occur.

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

Q: If the WS Size (WSS) of process *i* is WSS_i, then what is the total demand of the most active frames for *n* processes?

$$D = \sum_{i=0}^{n} WSS_i$$

If the total demand D is greater than the count of available frames, thrashing will occur.

Q: What can the OS do?

WORKING SIZE SET AND THRASHING

Most important property of the working set (WS) : size

Q: If the WS Size (WSS) of process *i* is WSS_i, then what is the total demand of the most active frames for *n* processes?

$$D = \sum_{i=0}^{n} WSS_i$$

If the total demand D is greater than the count of available frames, thrashing will occur.

The OS monitors WSS for each process, if D is approaching the memory limit, suspend a process.

WORKING SET

Q: What is the working set size?

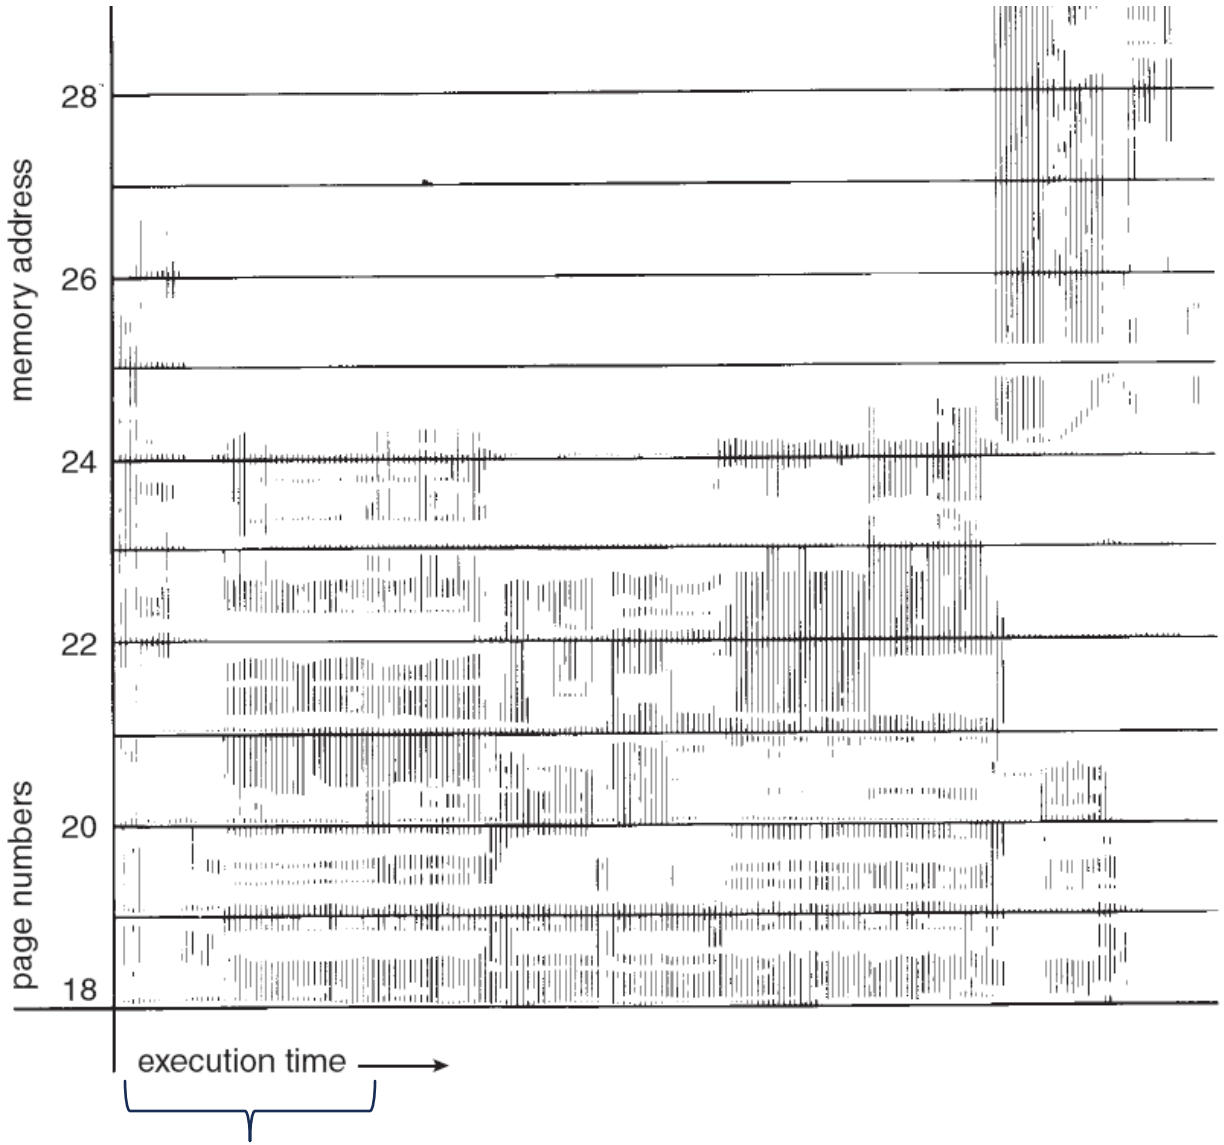

WORKING SET

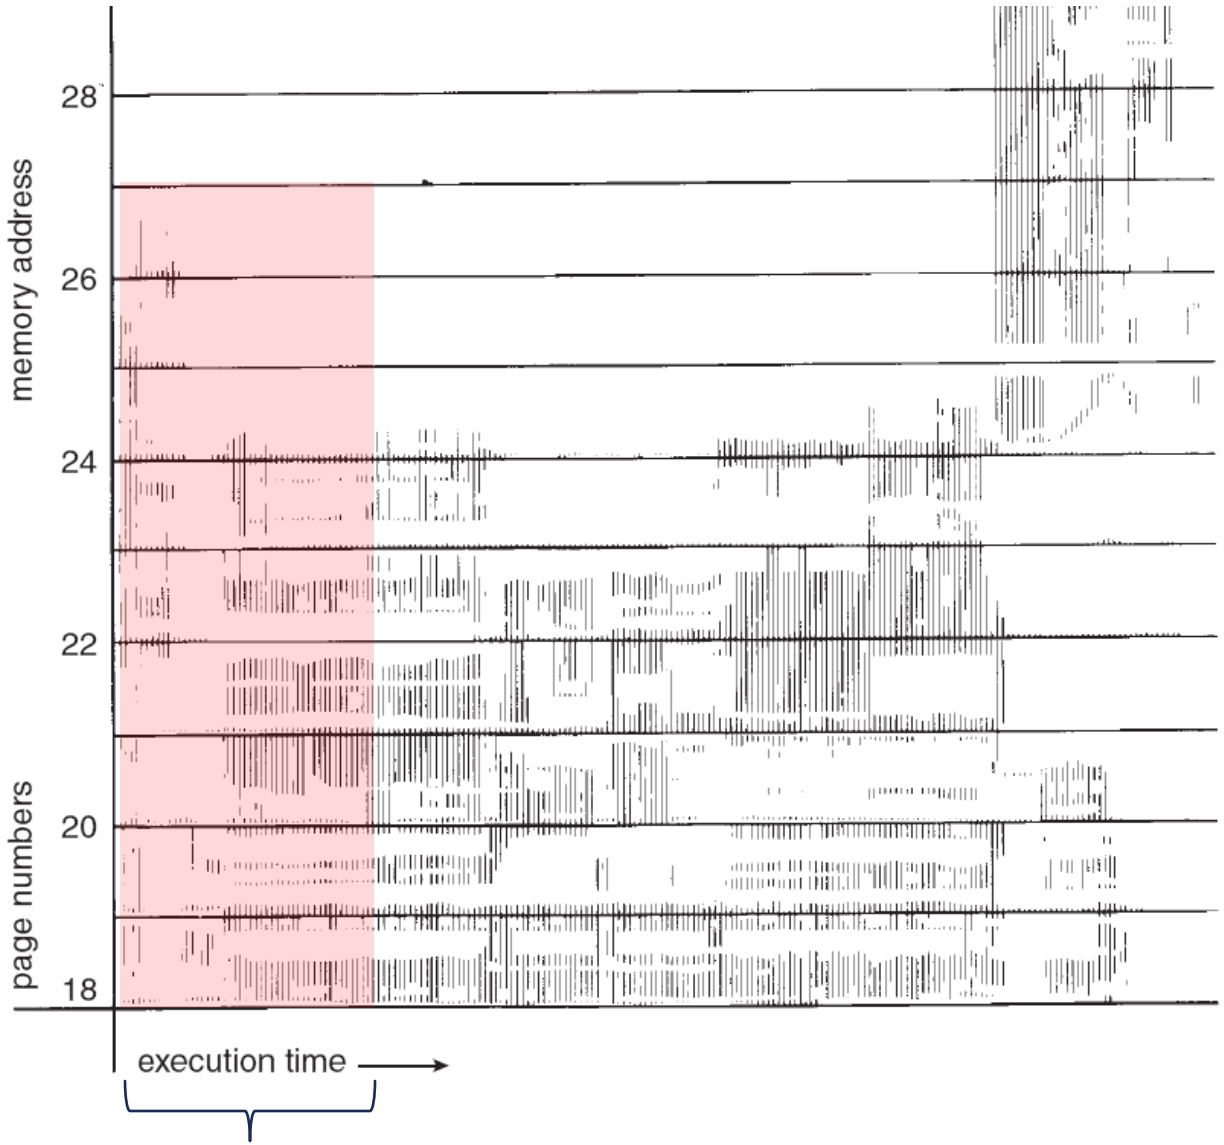

Q: What is the working set size?

Pages 18 to 26
WSS = 9

WORKING SET

Q: What is the working set size?

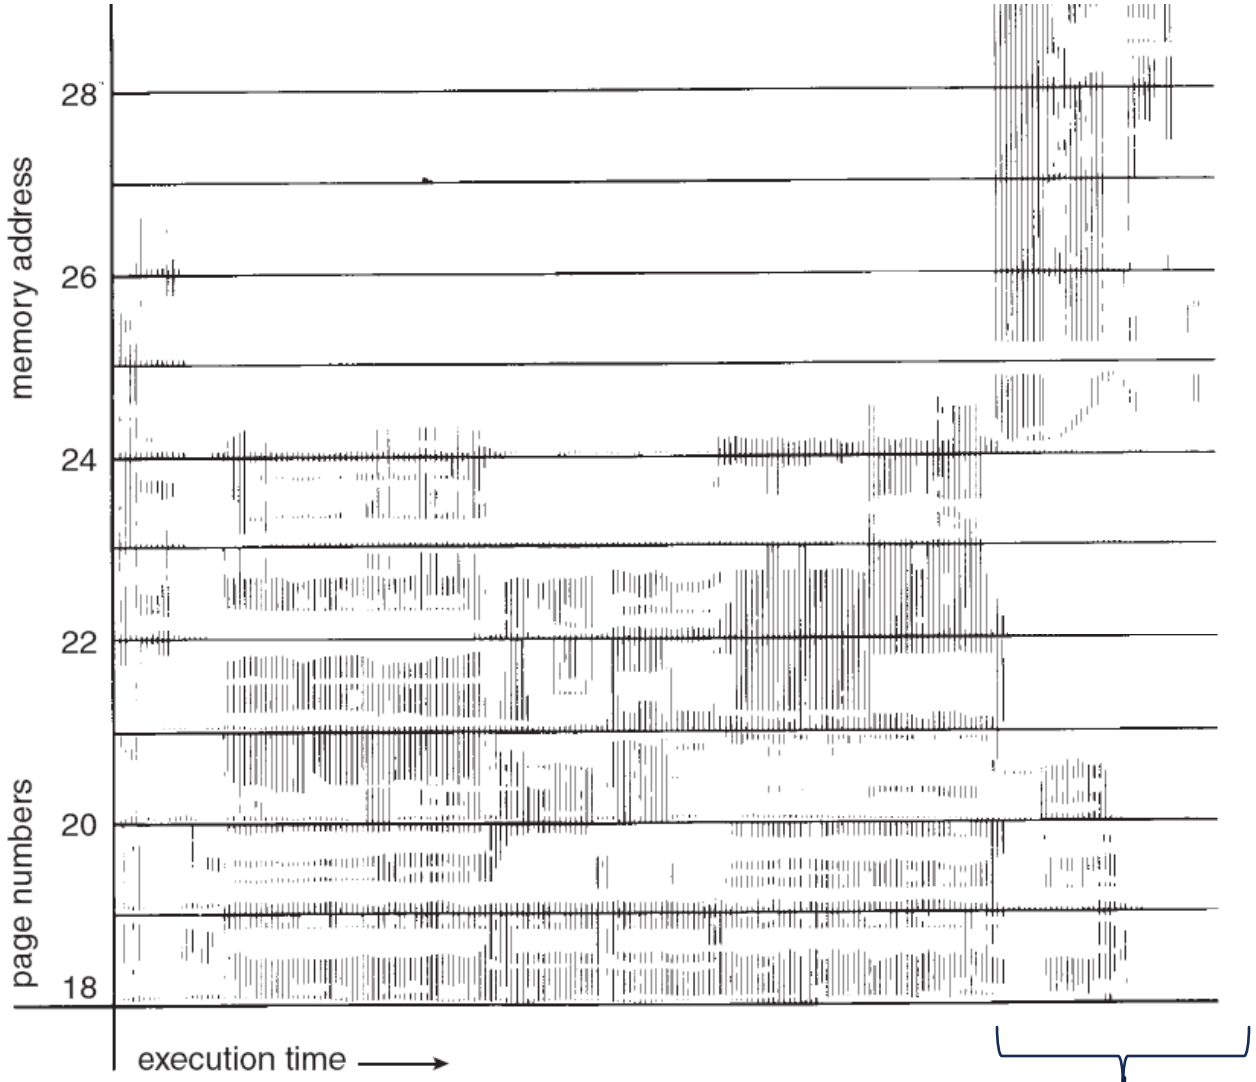

WORKING SET

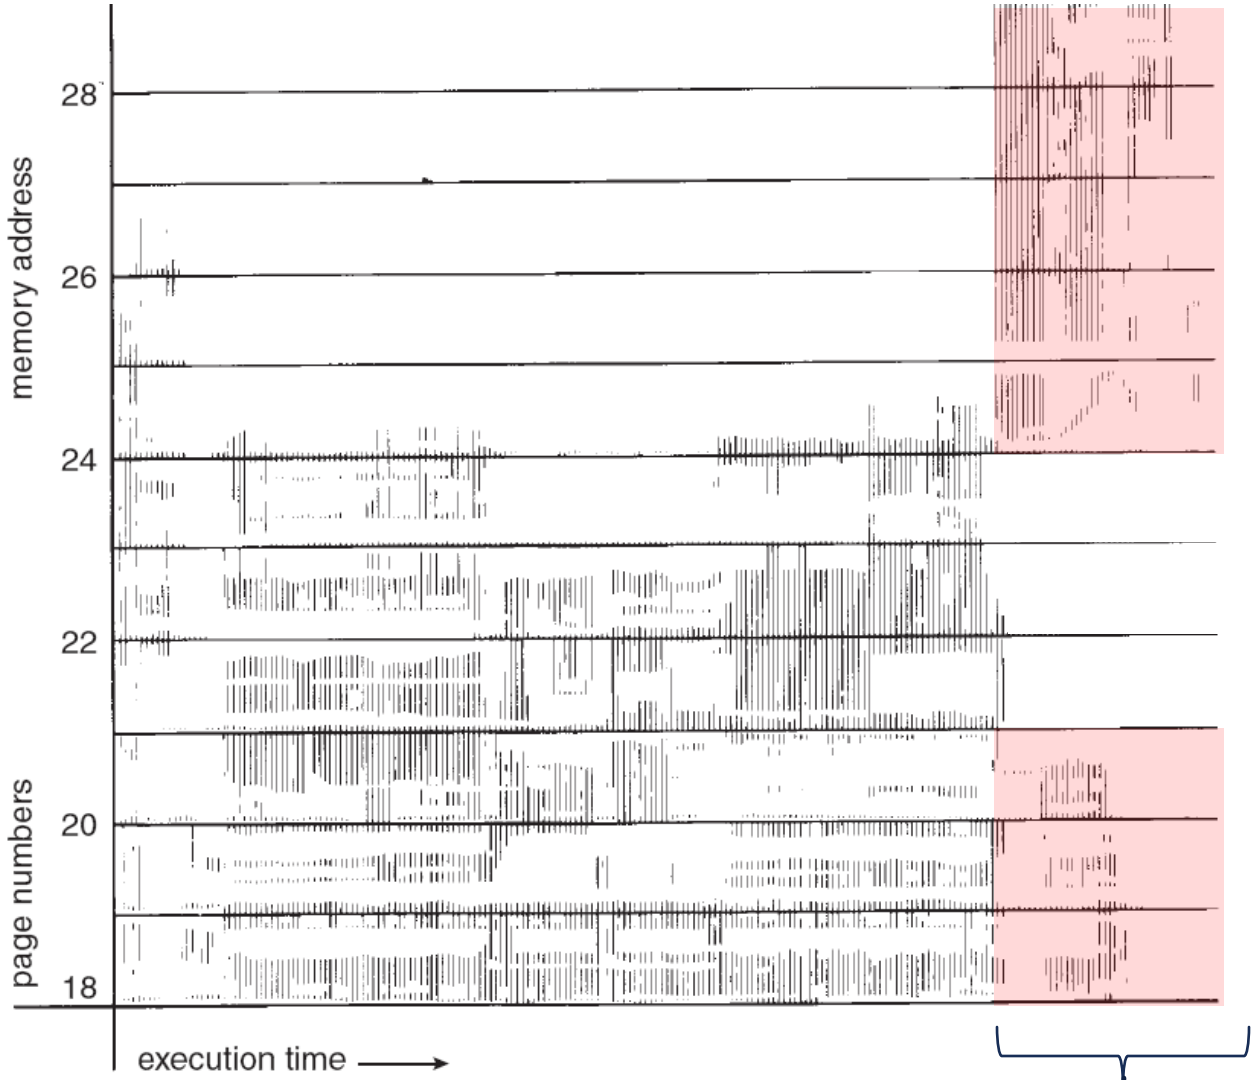

Q: What is the working set size?

Pages 18 to 20 and 24 to 28
WSS = 8