


# OPERATING SYSTEMS



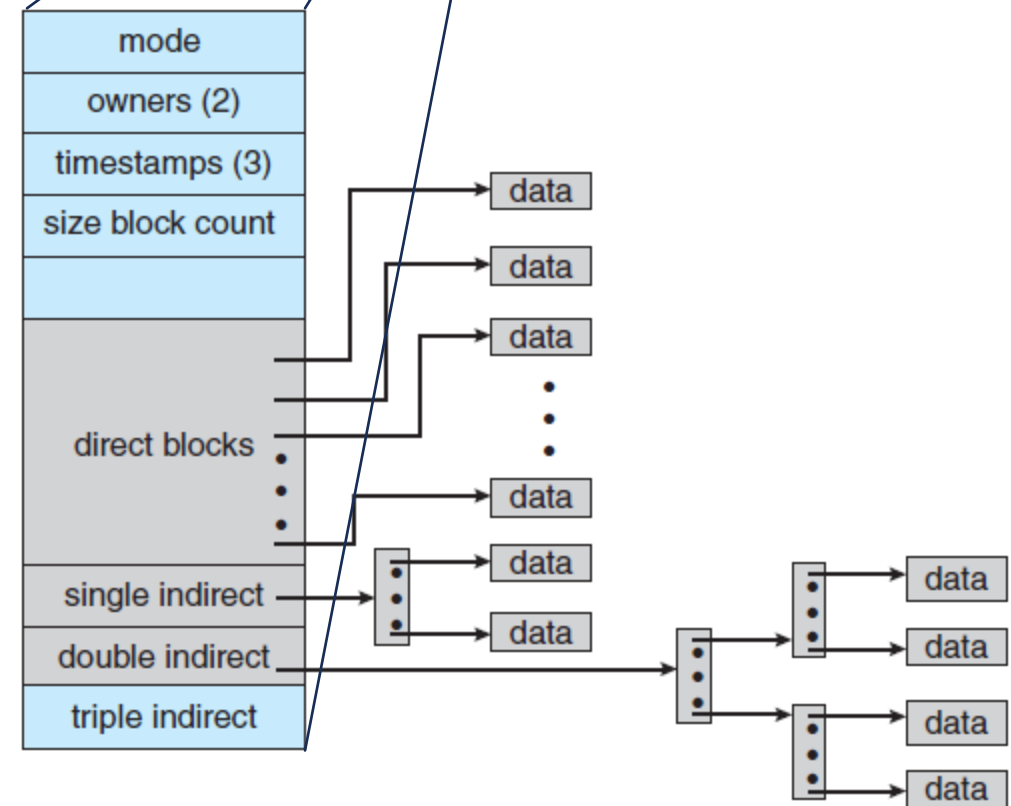
CSCI 509

CSCI 509 - OPERATING SYSTEMS INTERNALS

# LINUX INODE

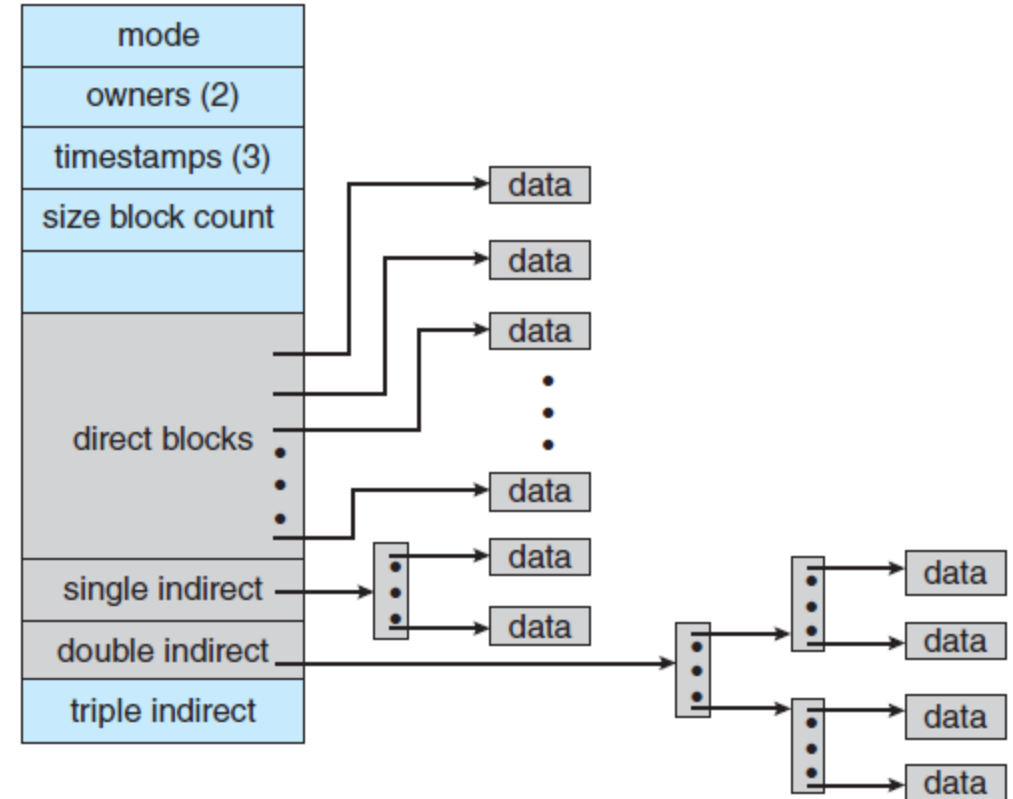


- Contain both metadata and pointers to blocks used.
- Uses various type of indexing.
- First blocks can be addressed directly others could have utilize multilevel indexing.
- Start with using direct block, if that's not enough for the file use indirect.
- Most files are small and usually direct blocks would suffice.




# WORKSHEET

- In an ext2 file system an inode consists of only 15 block pointers.
- The first 12 block pointers are direct block pointers.
- The 13th pointer is an indirect pointer.
- The 14th pointer is a double indirect pointer.
- The 15th pointer is a triple indirect pointer.
- Block size of 4KB
- 32-bit addressing for the blocks
- Which of these pointers will be utilized when the inode represents a file of size 64 KB?
- Which of these pointers will remain unutilized?



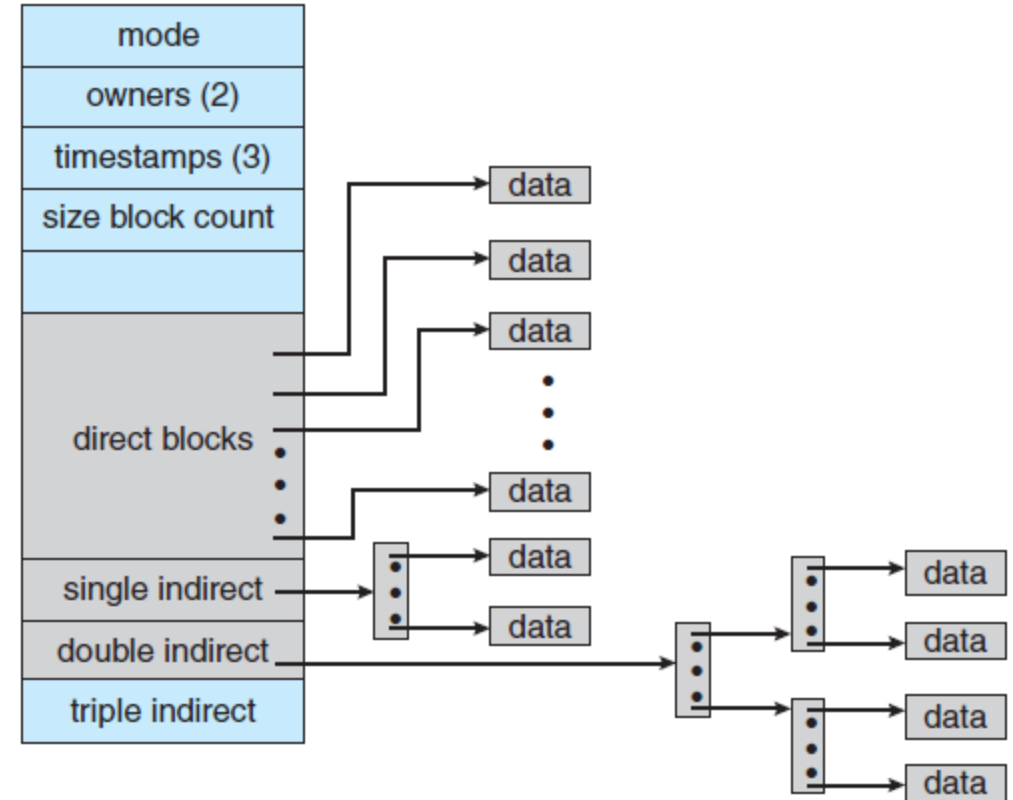
# WORKSHEET

- How many blocks do we need for the file?
- How much “size on disk” does each direct block pointer support?
- How much “size on disk” does a single indirect pointer can support?
  - How many direct pointers can a block on disk hold?



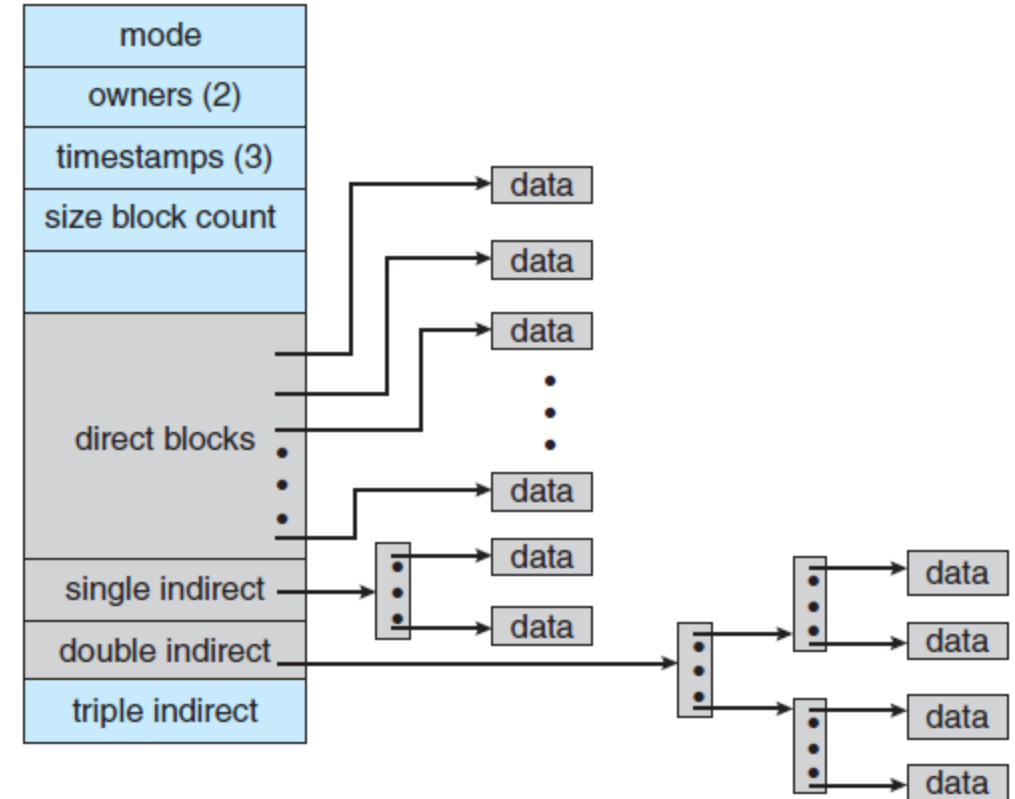
# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.




# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.



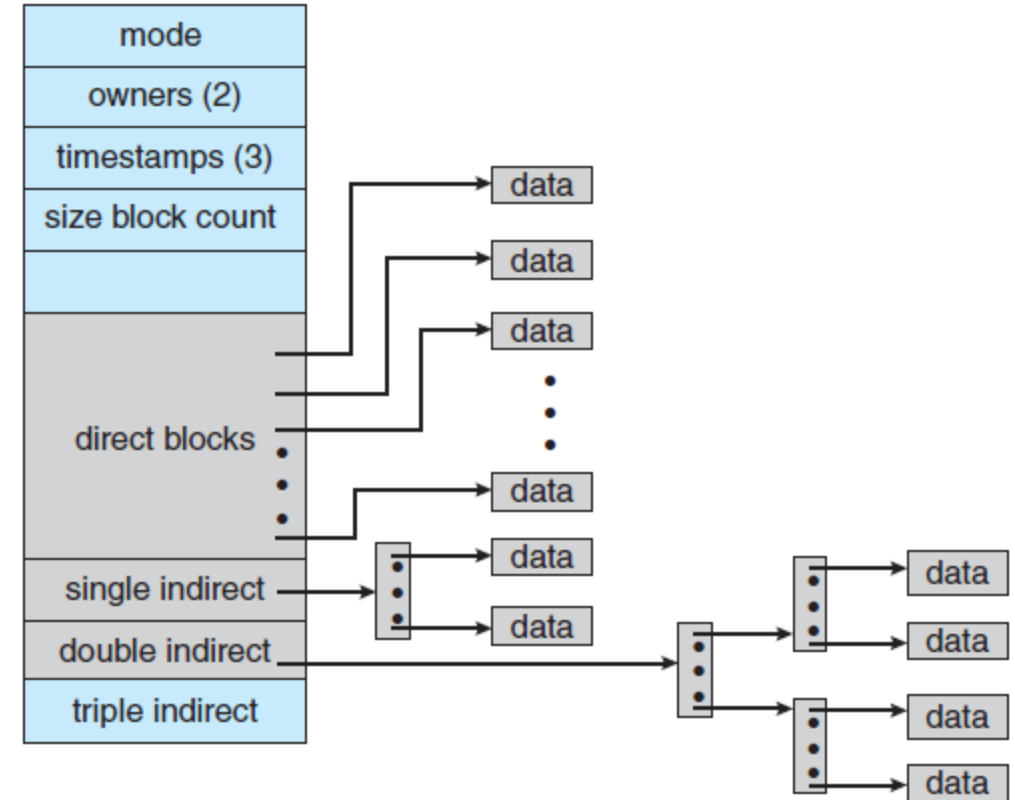
# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.
- How many does single indirect supports?



# INODE EXERCISE


- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.
- How many does single indirect supports?
- 32-bit machine → each direct pointer needs 4 bytes to be stored.






# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.
- How many does single indirect supports?
- 32-bit machine → each direct pointer needs 4 bytes to be stored.
- With a block size of 4KB, one block can store 1024 block addresses or 1024 “direct pointer” ...




# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.
- How many does single indirect supports?
- 32-bit machine → each direct pointer needs 4 bytes to be stored.
- With a block size of 4KB, one block can store 1024 block addresses or 1024 “direct pointer” ...
- We only need 4 more, single indirect provides 1024 more so single indirect should suffice.

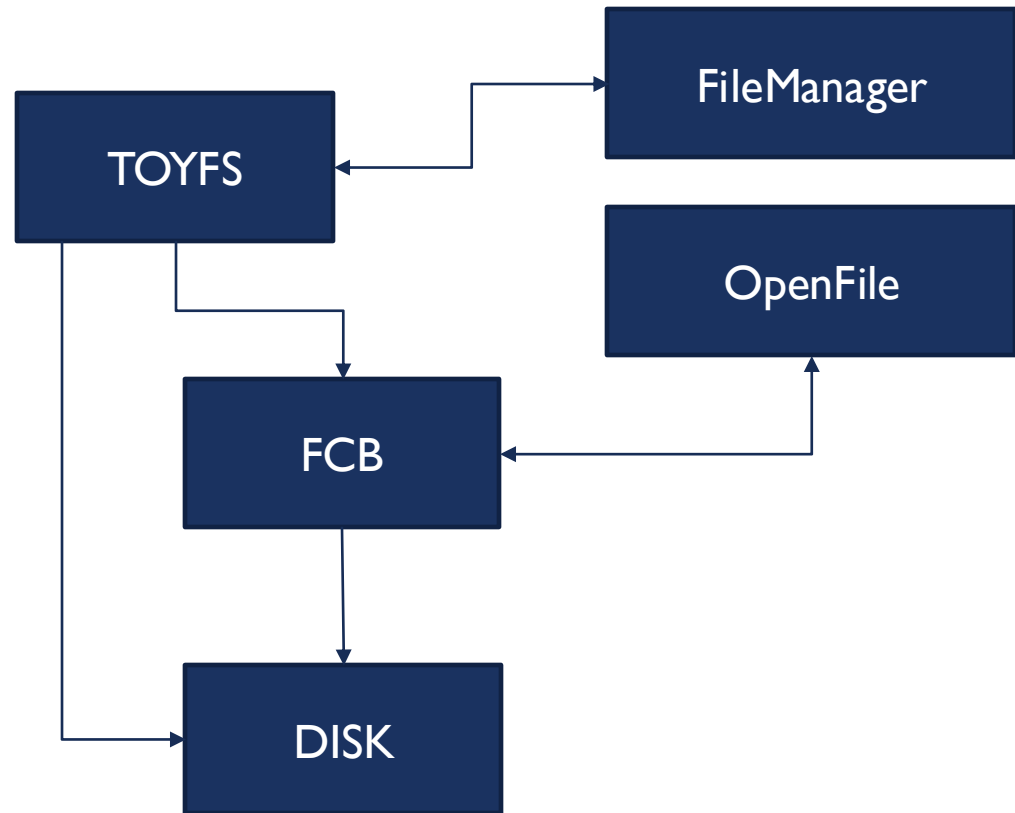


# INODE EXERCISE

- How many blocks do we need for the file?
- 64KB and 4KB per block → 16 blocks.
- We only have 12 direct pointers, we need at least single indirect for the extra 4.
- How many does single indirect supports?
- 32-bit machine → each direct pointer needs 4 bytes to be stored.
- With a block size of 4KB, one block can store 1024 block addresses or 1024 “direct pointer” ...
- We only need 4 more, single indirect provides 1024 more so single indirect should suffice.
- Double and triple indirect are never used.



# TOYFS



There are other classes as well

# FILEMANAGER

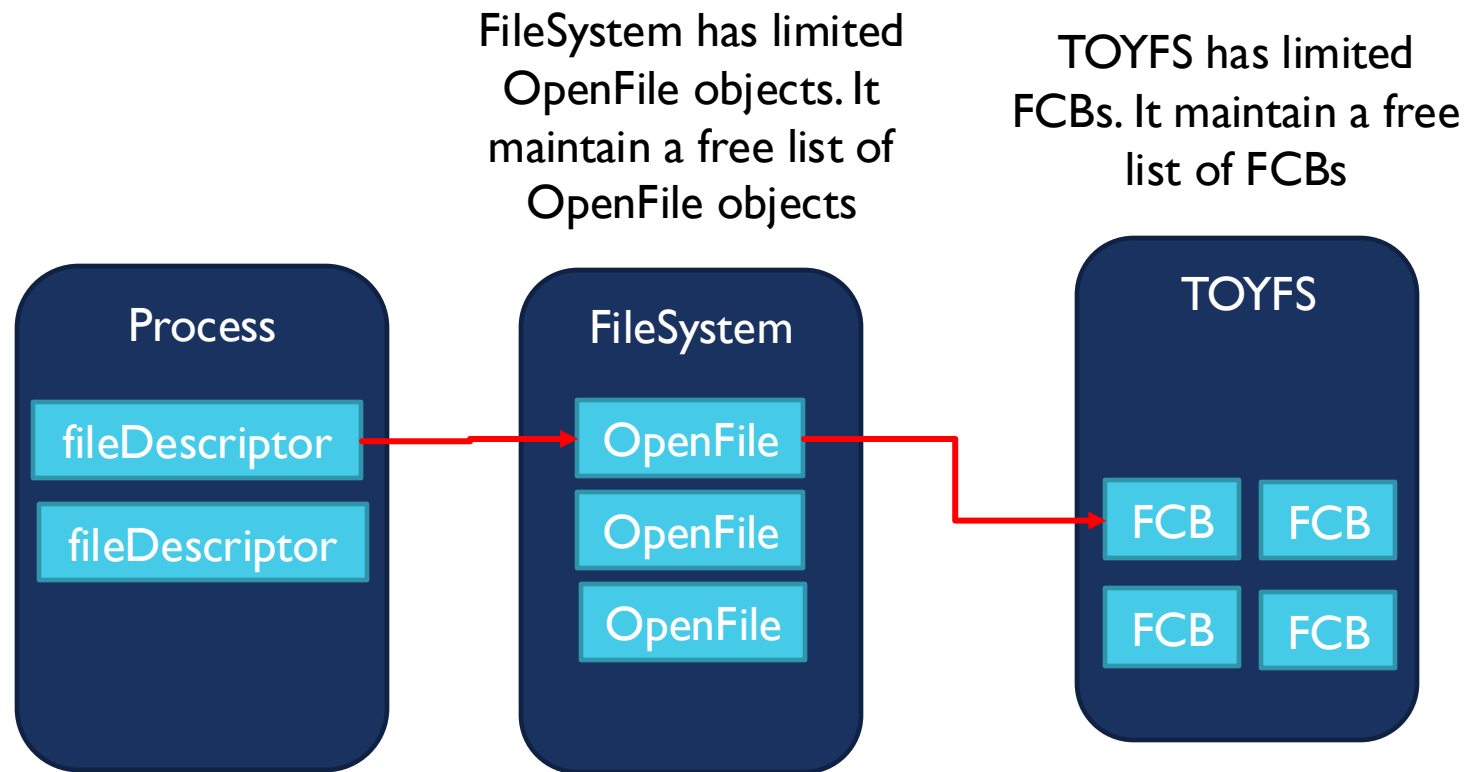
- Manages OpenFile and Pipe Objects
- Methods include
  - GetAnOpenFile
  - Open, Close
  - GetAPipe
  - ...
- Includes data structures
  - OpenFileTable
  - openFileFreeList
  - ...

# FCB

- Records all data associated with a single ToyFS file
- Has an InodeData object that can be used to get file info, like the actual sectors on disk.
- Has methods to read and write to disk

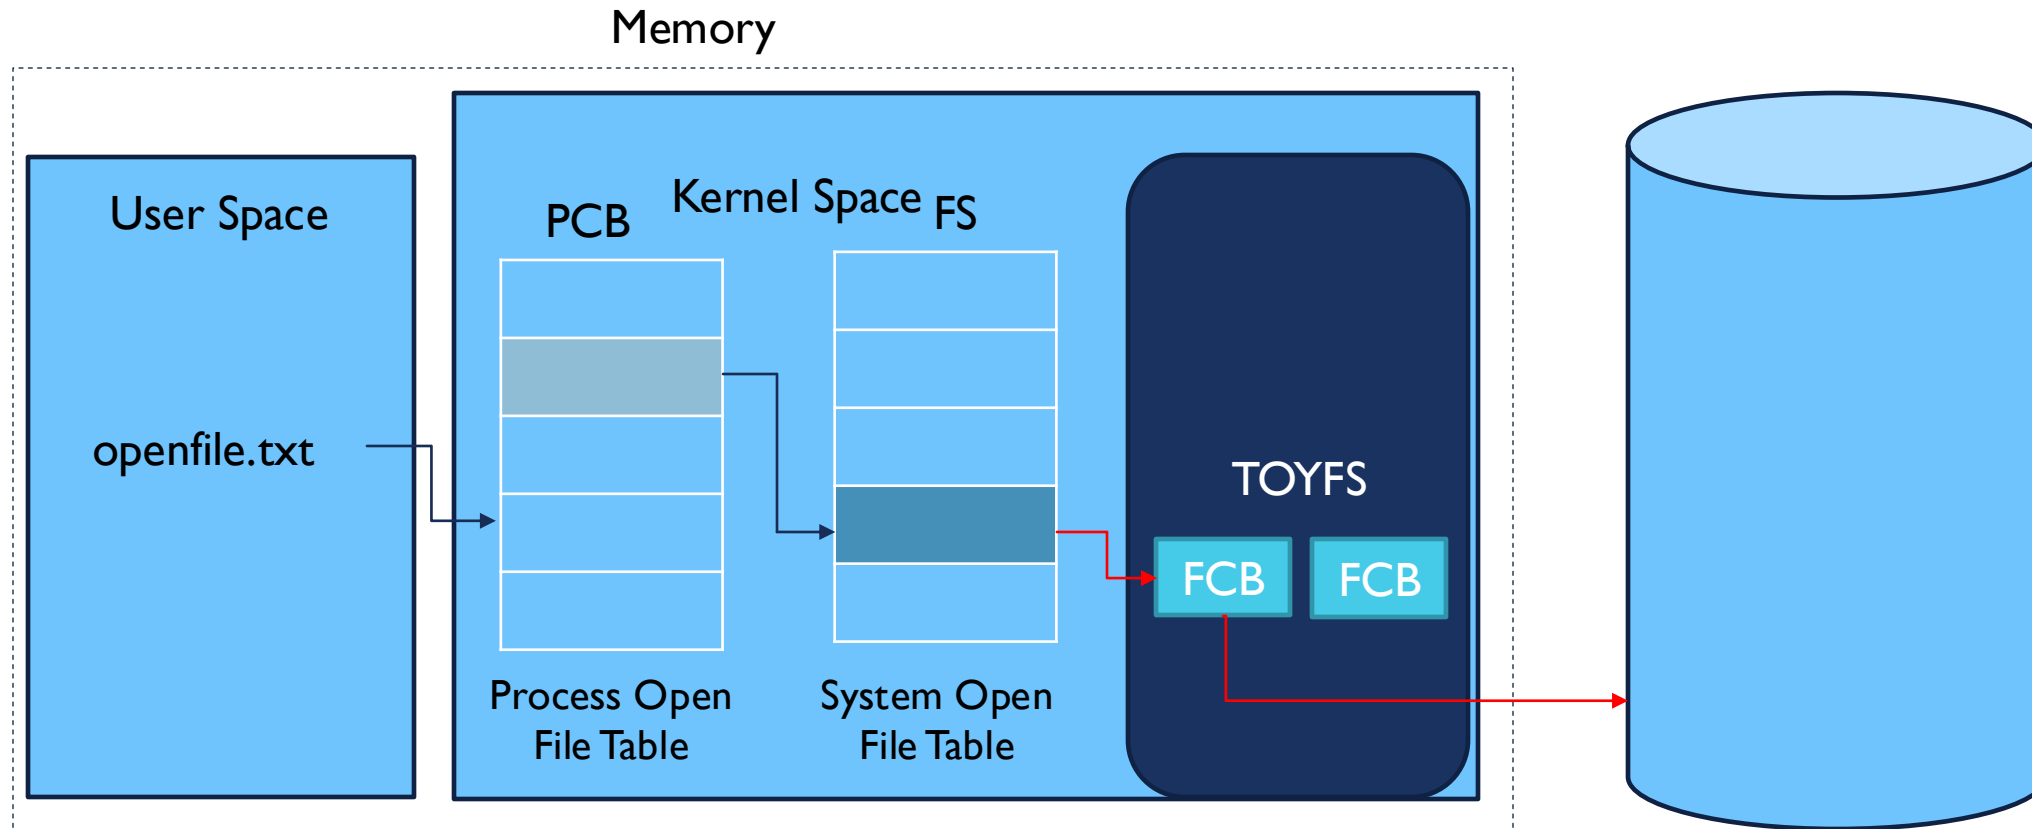
# TOYFS

- Contains the superblock (in memory)
  - This include data and inode bitmaps
- Has a handle to root directory (root always in inode 1 on disk)
- FCB Table, + List of FREE FCB
- Can allocate/free inodes or data blocks
- Methods to look up inode give a file name or an FCB





# IN MEMORY FILE SYSTEM IMPLEMENTATION



# TRACKING FREE DISK SPACE

**Free Space Management** How does the system keep track of available (free) blocks?

- **Bit vector approach** Keep a single bit to specify if the block is free (0) or in use / not free (1).


010001010100010101

**Advantages / Disadvantages :** Easy to implement, but need algorithm to find contiguous space if contiguous allocation is used

You need to search for the '0's ...

# TRACKING FREE DISK SPACE

- Linked list of available block.
- Disadvantage: Complex to implement
- Advantage: Can easily allocate contiguous space.




# FREE SPACE MANAGEMENT

- Counting
  - Because space is frequently contiguously used and freed, with contiguous-allocation allocation, extents, or clustering
  - Keep address of first free block and count of following free blocks
  - Free space list then has entries containing addresses and counts

directory		
File	Start	Length
<i>myFile</i>	1	3
<i>aPic</i>	12	6
<i>song</i>	28	3

**Very similar to contiguous allocation ... we're just keeping track of free blocks instead of file blocks.**




# FREE SPACE MANAGEMENT

- Counting
  - Because space is frequently contiguously used and freed, with contiguous-allocation allocation, extents, or clustering
  - Keep address of first free block and count of following free blocks
  - Free space list then has entries containing addresses and counts

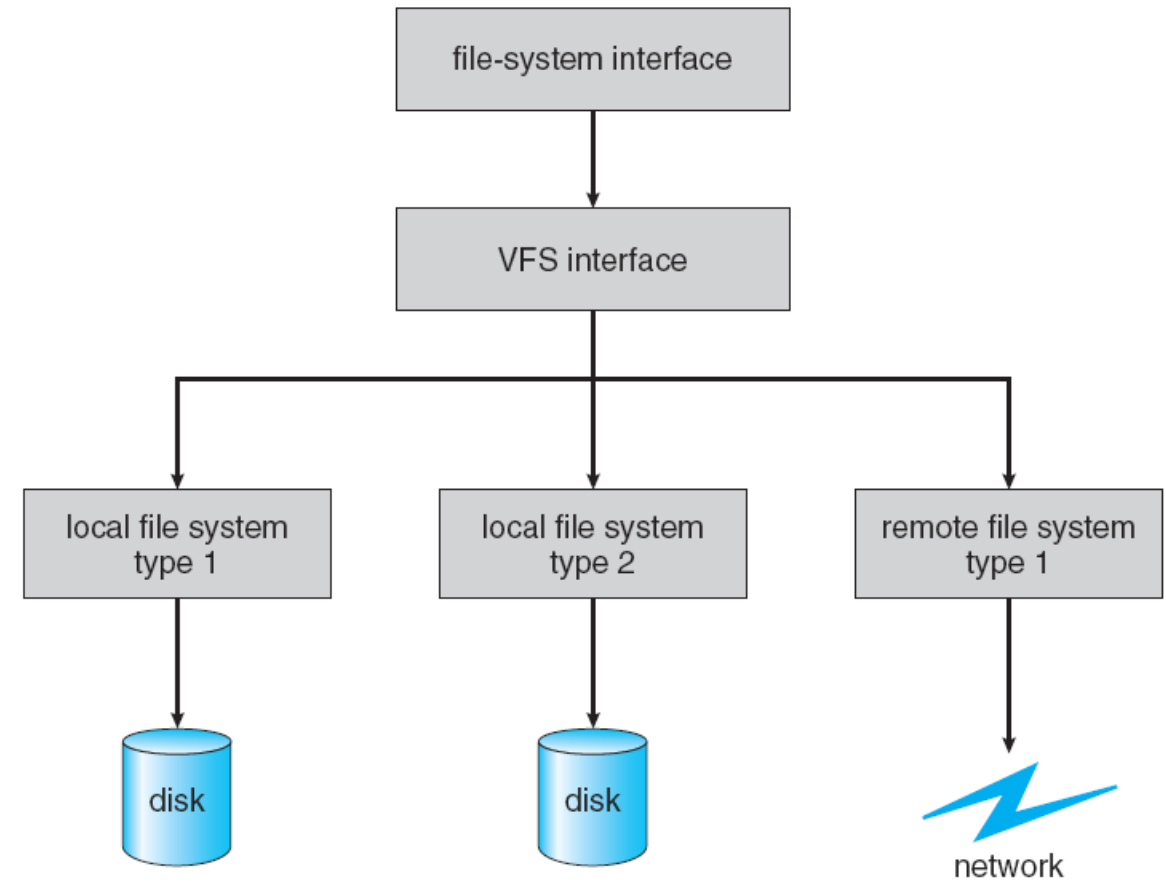
Free Blocks List		
Block	Start	Length
	1	3
	12	6
	28	3

**Very similar to contiguous allocation ... we're just keeping track of free blocks instead of file blocks.**




# VIRTUAL FILE SYSTEM

- Operating systems utilize a virtual file system interface.
- This allows for programs to use the same system calls (such as `open()`, `read()`, `write()`, `close()`, `mkdir()`, etc.) regardless of what file system they are accessing.




# NETWORK FILE SYSTEM

- Advantages?



# NETWORK FILE SYSTEM

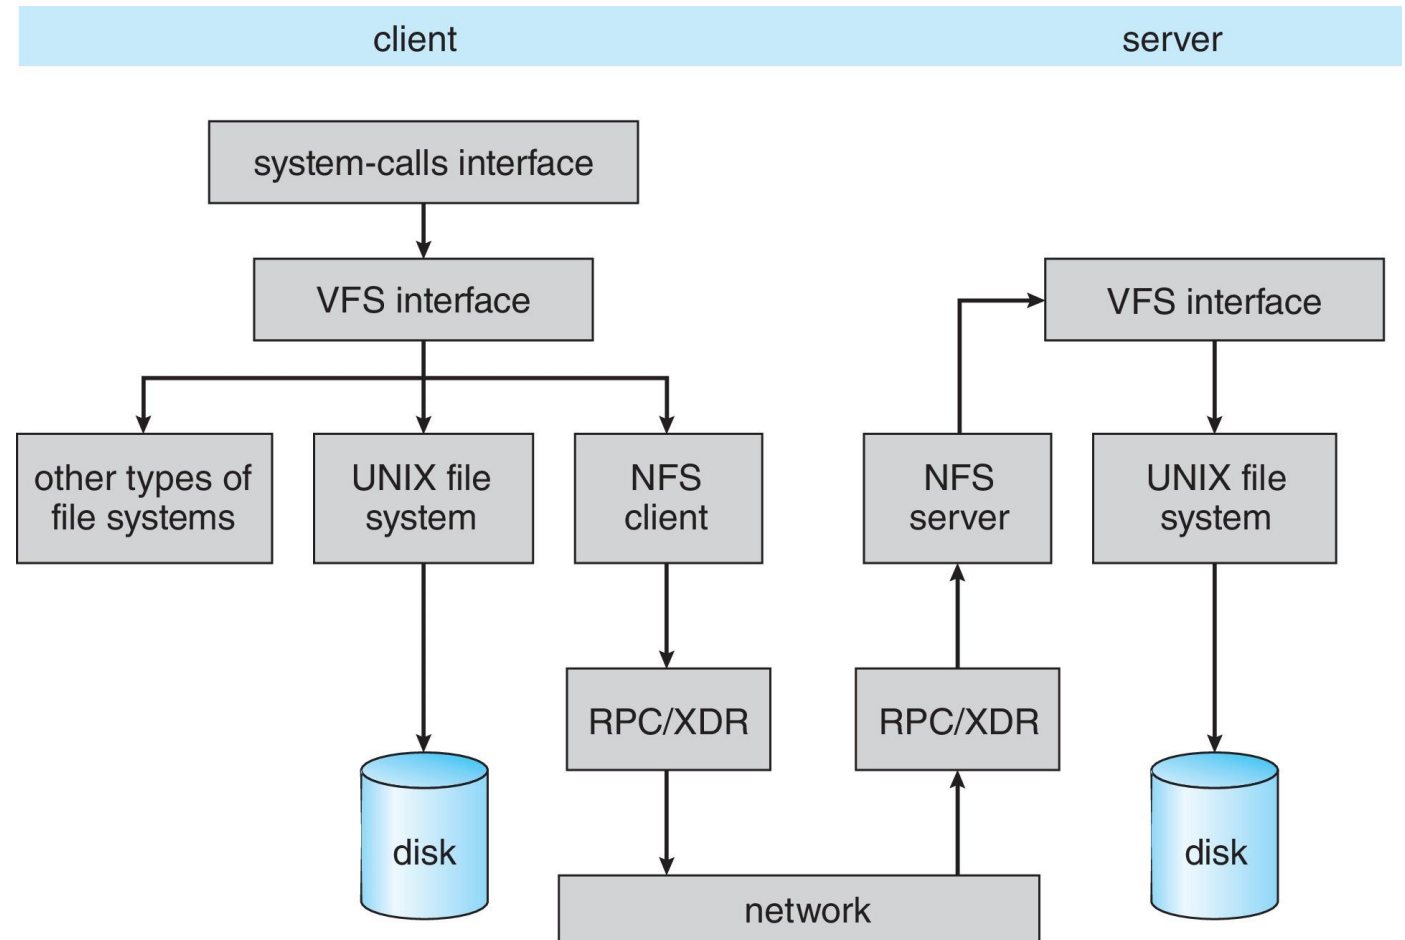
- Sharing
  - Sharing data
  - Sharing free space
- Centralized Administration
  - Backing Up
  - Restoration
- Security





# NETWORK FILE SYSTEM

- Once setup, the network file system is accessed like any local file system.
- Users/Programs utilize the virtual file system where they access remote or local file.




# NETWORK FILE SYSTEM

- NFS was developed by Sun microsystem.
- Latest version: NFSv4, we will look at NFSv2 which was what made NFS popular.
- Design was centered around *Fast Crash Recovery*.
- To achieve this, the NFS used “stateless” design.

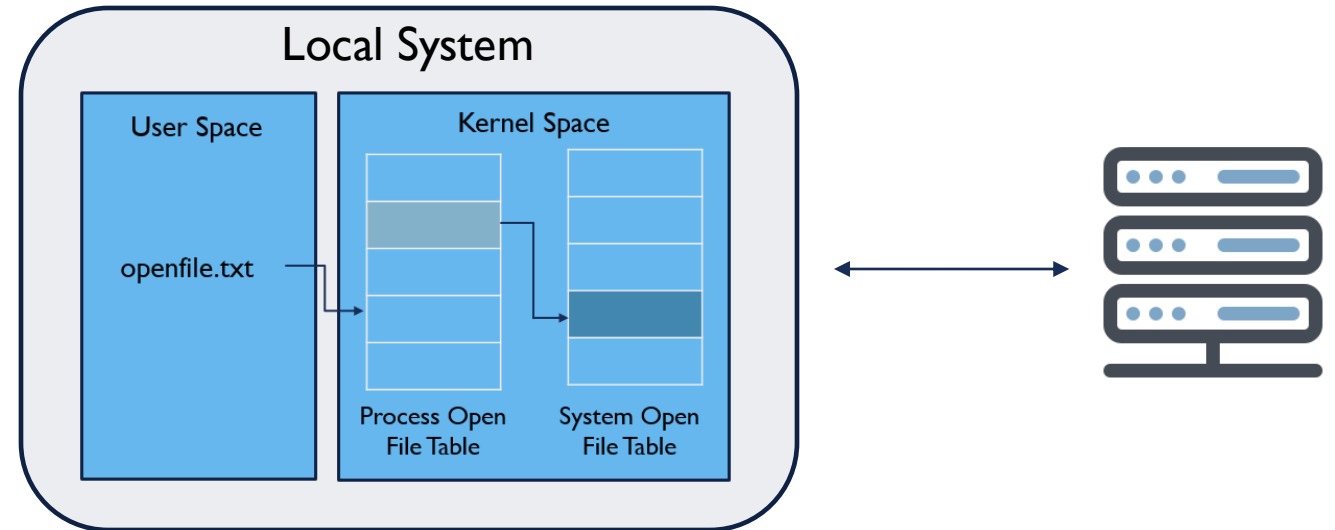
# STATELESS FILE SYSTEM

- “Stateful” file systems, like local file systems keep track of almost all ongoing operations.
- Open file table, file pointers, users accessing the file.
- If a server crash occurs, all this information is lost (which files are open ...)
- If a client crash occurs, it also creates problems.
- Need to implement recovery algorithm and perform recovery with each connecting client.



# HOW TO IMPLEMENT A STATELESS FILE SYSTEM?

- From the client perspective, the NFS is a stateful protocol: Virtual File System.
- The system however, only issues stand-alone commands to the NFS server.



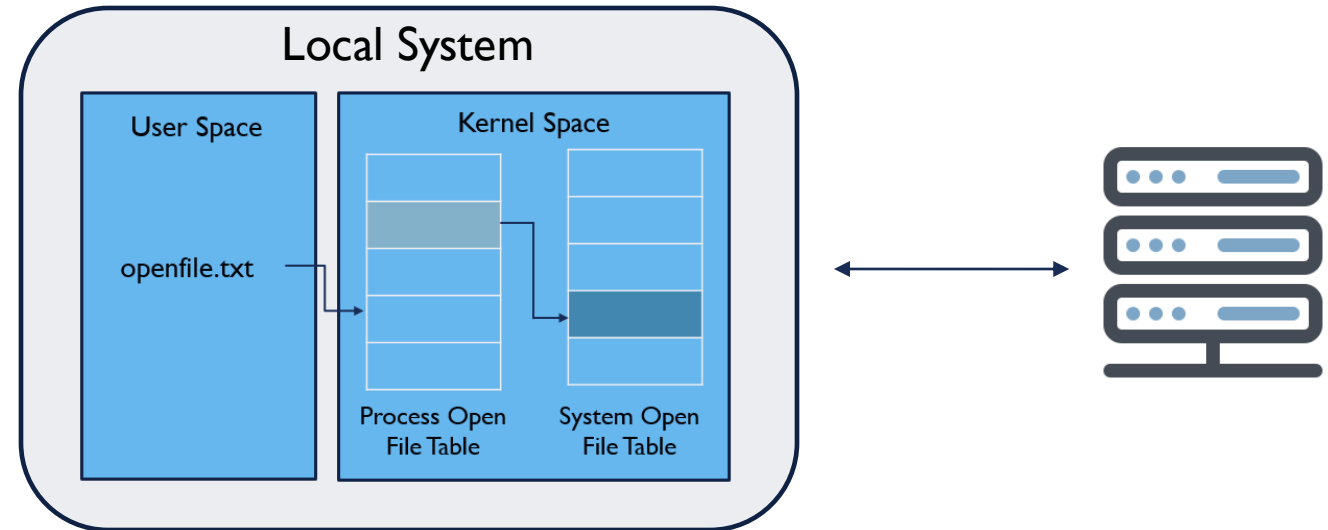
**Q:** What access method should the NFS 2.0 utilize?



**A:** Direct  
**B:** Sequential

# HOW TO IMPLEMENT A STATELESS FILE SYSTEM?

- From the client perspective, the NFS is a stateful protocol: Virtual File System.
- The system however, only issues stand-alone commands to the NFS server.
- Example commands:



**Q:** What access method should the NFS 2.0 utilize?



**A:** Direct  
**B:** Sequential

The server can't keep a record of any state, including file pointers for "read next".

# NFS 4

- The client and server establish a long-lived session
- This session maintains the client's state information, such as file locks and open file handles
- Allows the client to recover its state after a network interruption or server failure without having to re-establish its entire connection.
- With the advancement of computing power, memory capacity and network bandwidth, the overhead of reestablishing a connection became minimal compared to the benefits.

# FILE SYSTEM RECOVERY

- How can the file system recover from crashes?

# FILE SYSTEM RECOVERY

- How can the file system recover from crashes?
- Example: Power failure could occur while writing to disk ...
  - Directory structure could become inaccurate
  - Files could be partially written
  - File meta data could be out of date
  - ...



# FILE SYSTEM RECOVERY

- **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
  - Can be slow and sometimes fails

# FILE SYSTEM RECOVERY

- **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
  - Can be slow and sometimes fails
- Use system programs to **back up** data from disk to another storage device (magnetic tape, other magnetic disk, optical)

# LOG STRUCTURED FILE SYSTEMS

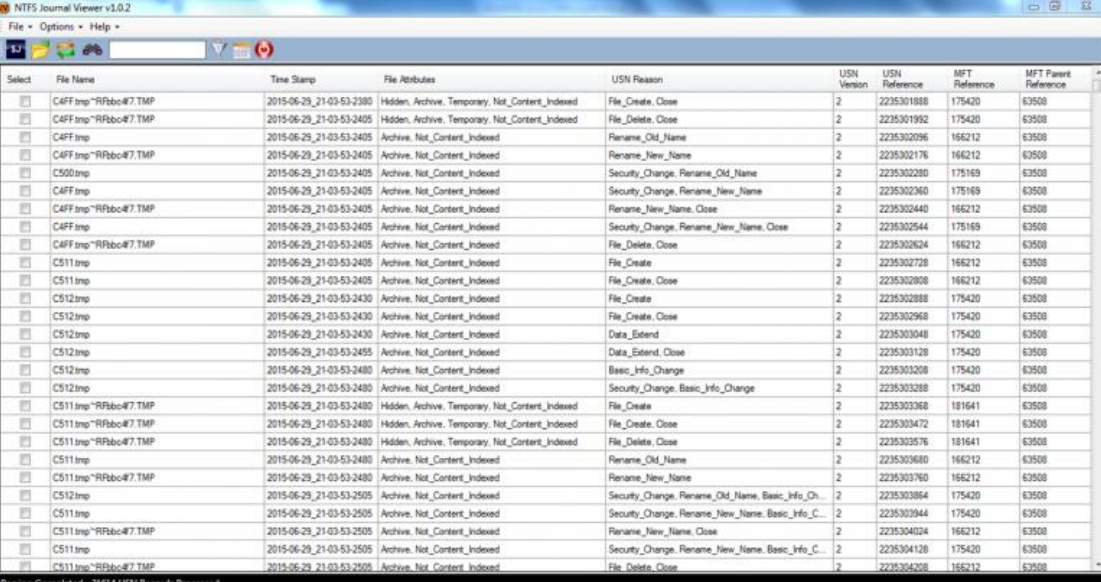
- **Log structured** (or **journaled**) file systems record each metadata update to the file system as a **transaction**

# LOG STRUCTURED FILE SYSTEMS

- **Log structured** (or **journaled**) file systems record each metadata update to the file system as a **transaction**
- All transactions are written to a log
  - A transaction is considered committed once it is written to the log (sequentially)

# LOG STRUCTURED FILE SYSTEMS

- **Log structured** (or **journaled**) file systems record each metadata update to the file system as a **transaction**
- All transactions are written to a log
  - A transaction is considered committed once it is written to the log (sequentially)
- Allows faster recovery from crash, removes chance of inconsistency of metadata



The screenshot displays the NTFS Journal Viewer v1.0.2 application window. It features a menu bar with 'File', 'Options', and 'Help'. Below the menu is a toolbar with icons for file operations. The main area is a table listing file system transactions. The table has columns for 'Select', 'File Name', 'Time Stamp', 'File Attributes', 'USN Reason', 'USN Version', 'USN Reference', 'MFT Reference', and 'MFT Parent Reference'. The data shows a series of transactions for files like 'C:\FF\mp\RRbc\#7.TMP' and 'C:\FF\mp\RRbc\#7.TMP', with reasons such as 'File Create', 'File Delete', 'Rename', 'Security Change', and 'Data Extend'. The status bar at the bottom indicates 'Parsing Completed - 23614 USN Records Processed'.

Select	File Name	Time Stamp	File Attributes	USN Reason	USN Version	USN Reference	MFT Reference	MFT Parent Reference
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2380	Hidden, Archive, Temporary, Not_Content_Indexed	File_Create, Close	2	2235301888	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Hidden, Archive, Temporary, Not_Content_Indexed	File_Delete, Close	2	2235301992	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Rename_Old_Name	2	2235302096	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Rename_New_Name	2	2235302176	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Security_Change, Rename_Old_Name	2	2235302280	175169	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Security_Change, Rename_New_Name	2	2235302360	175169	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Rename_New_Name, Close	2	2235302440	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	Security_Change, Rename_New_Name, Close	2	2235302544	175169	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	File_Delete, Close	2	2235302624	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	File_Create	2	2235302728	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2405	Archive, Not_Content_Indexed	File_Create, Close	2	2235302808	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2430	Archive, Not_Content_Indexed	File_Create	2	2235302888	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2430	Archive, Not_Content_Indexed	File_Create, Close	2	2235302968	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2430	Archive, Not_Content_Indexed	Data_Extend	2	2235303048	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2430	Archive, Not_Content_Indexed	Data_Extend, Close	2	2235303128	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Archive, Not_Content_Indexed	Basic_Info_Change	2	2235303208	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Archive, Not_Content_Indexed	Security_Change, Basic_Info_Change	2	2235303288	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Hidden, Archive, Temporary, Not_Content_Indexed	File_Create	2	2235303368	181641	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Hidden, Archive, Temporary, Not_Content_Indexed	File_Create, Close	2	2235303472	181641	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Hidden, Archive, Temporary, Not_Content_Indexed	File_Delete, Close	2	2235303576	181641	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Archive, Not_Content_Indexed	Rename_Old_Name	2	2235303680	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2480	Archive, Not_Content_Indexed	Rename_New_Name	2	2235303760	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2505	Archive, Not_Content_Indexed	Security_Change, Rename_Old_Name, Basic_Info_Ch...	2	2235303864	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2505	Archive, Not_Content_Indexed	Security_Change, Rename_New_Name, Basic_Info_C...	2	2235303944	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2505	Archive, Not_Content_Indexed	Rename_New_Name, Close	2	2235304024	166212	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2505	Archive, Not_Content_Indexed	Security_Change, Rename_New_Name, Basic_Info_C...	2	2235304128	175420	63508
<input type="checkbox"/>	C:\FF\mp\RRbc\#7.TMP	2015-06-29_21:03:53:2505	Archive, Not_Content_Indexed	File_Delete, Close	2	2235304208	166212	63508

NTFS Logfile

# UNIX FILE SYSTEM

- Initially used “FS” or “File System”, first UNIX file system.
- Has gone through significant changes and iterations
- BFS (Berkley File System)
- ext or Extended File System
- Latest is ext4

# OLD UNIX FILE SYSTEM

- Ken Thompson wrote the first file system.
- Very simple.
- Very poor performance:
  - Not disk aware, treating disk like random memory.
  - High fragmentation.
  - inodes can be allocated very far from their data.
  - Block size too small.
  - Overtime, performance was 2% of actual disk I/O bandwidth.



# BERKLEY FAST FILE SYSTEM

Disk aware:

- Include structure information for each group: file system.
- Directory search, meta data access and modification same group as file access
- Overall, greatly increased performance.

