

OPERATING SYSTEMS INTERNALS

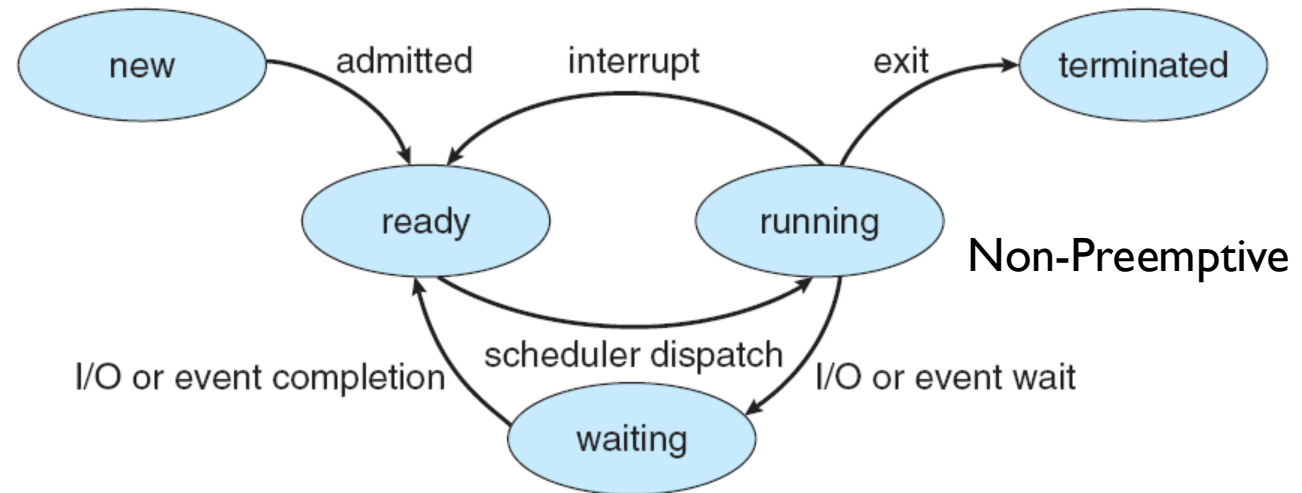
CSCI 509

CHAPTER 5: CPU SCHEDULING



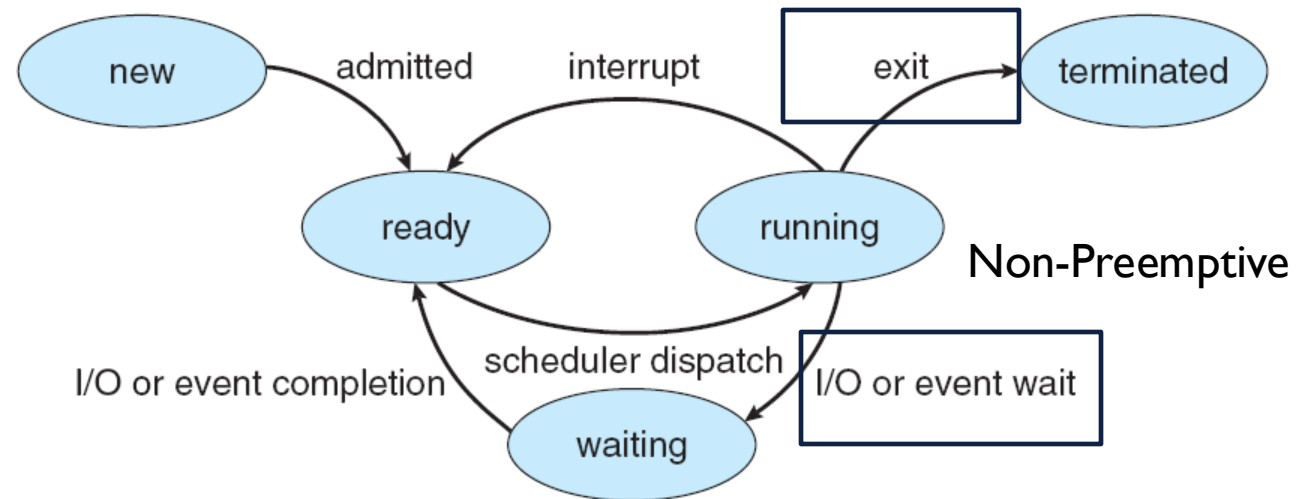
PREEMPTIVE VS NON-PREEMPTIVE

- Preemptive:
Thread/Process
gets switched out
immediately at
scheduler request.



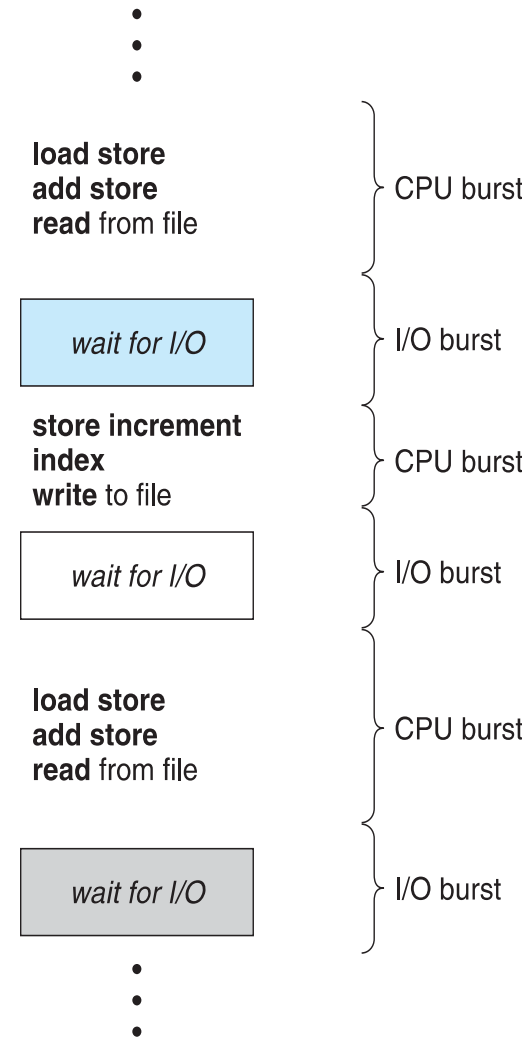
PREEMPTIVE VS NON-PREEMPTIVE

- Preemptive:
Thread/Process gets switched out immediately at scheduler request.
- Non-preemptive:
Thread/Process either decides to switch or is “asked” to switch out.



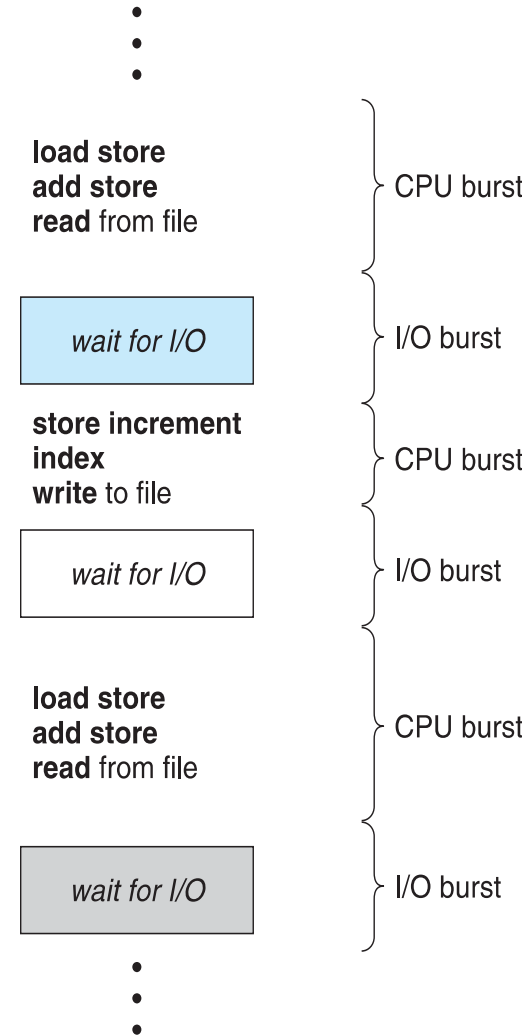
CPU SCHEDULING: I/O BURST VS CPU BURST

- Typical Thread Execution:
- **CPU burst** followed by **I/O burst**



CPU SCHEDULING: I/O BURST VS CPU BURST

- Typical Thread Execution:
- **CPU burst** followed by **I/O burst**
- CPU burst distribution is of main concern.
- Schedulers, in general, want to avoid interrupting CPU bursts. Why?



CPU SCHEDULING CRITERIA

- How to evaluate a Scheduling algorithm?
- How to choose between one algorithm or the other?
- What's the criteria?

CPU SCHEDULING CRITERIA

- How to evaluate a Scheduling algorithm?
- How to choose between one algorithm or the other?
- What's the criteria?
- Worksheet Task: List few scheduling performance criteria.

CPU SCHEDULING CRITERIA

How to evaluate a Scheduling algorithm? How to choose between one algorithm or the other? What's the criteria?

CPU SCHEDULING CRITERIA

How to evaluate a Scheduling algorithm? How to choose between one algorithm or the other? What's the criteria?

Scheduling Criteria

- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less
- **Throughput** : the number of processes completed in some unit of time
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation
- **Response time** : time from submission to the time when the “first” usable data/output is produced

CPU SCHEDULING CRITERIA

How to evaluate a Scheduling algorithm? How to choose between one algorithm or the other? What's the criteria?

Scheduling Criteria

- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less
- **Throughput** : the number of processes completed in some unit of time
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation
- **Response time** : time from submission to the time when the "first" usable data/output is produced

Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

CPU SCHEDULING CRITERIA

Scheduling Criteria

- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less
- **Throughput** : the number of processes completed in some unit of time
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation
- **Response time** : time from submission to the time when the “first” usable data/output is produced



Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

CPU SCHEDULING CRITERIA

Scheduling Criteria

- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less
- **Throughput** : the number of processes completed in some unit of time
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation
- **Response time** : time from submission to the time when the “first” usable data/output is produced






Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

CPU SCHEDULING CRITERIA

Scheduling Criteria





- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less 
- **Throughput** : the number of processes completed in some unit of time 
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting 
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation
- **Response time** : time from submission to the time when the “first” usable data/output is produced

Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

CPU SCHEDULING CRITERIA

Scheduling Criteria






- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less 
- **Throughput** : the number of processes completed in some unit of time 
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting 
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation 
- **Response time** : time from submission to the time when the “first” usable data/output is produced

Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

CPU SCHEDULING CRITERIA

Scheduling Criteria

- **CPU Utilization** : goal ... keep the CPU busy ... ideally idle 10% of the time or less 
- **Throughput** : the number of processes completed in some unit of time 
- **Turnaround** : from the time of submission, to the time of completion, INCLUDING the time spent waiting 
- **Waiting time** : the time spent in the waiting/ready queue ... waiting time does NOT take into account the time a process/thread spends in an I/O queue or time spent on an I/O operation 
- **Response time** : time from submission to the time when the “first” usable data/output is produced 

Q: Which of these
do we want to
minimize?

Q: Which of these
do we want to
maximize?

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

Once a process starts a burst, it won't be interrupted.

Assume 4 processes and their required burst times

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling




FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




**At what time is
P2 allowed to
“begin”?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



Worksheet Q2

**Q: What are the advantages of
this approach?**


**Q: What are the disadvantages
of this approach?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



**Q: What are the advantages of
this approach?**

Easy to implement
No scheduling “calculation” required

**Q: What are the disadvantages
of this approach?**


In a time sharing system, each process
might not get a share of the CPU at some
regular interval → long wait times

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

**P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms**

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



**Q: What are the advantages of
this approach?**

**Q: What are the disadvantages
of this approach?**


**To “see” the disadvantage of this
approach, what if P1 had a burst time
of 300ms. Should it still “go” first?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




**Q: To quantitatively assess the “goodness” of FIFO/FCFS with
nonpreemptive scheduling, what metric (CPU Utilization, Throughput,
Turnaround, Waiting time, Response time) should we calculate?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**




**Q: To quantitatively assess the “goodness” of FIFO/FCFS with
nonpreemptive scheduling, what metric (CPU Utilization, Throughput,
Turnaround, **Waiting time**, Response time) should we calculate?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



Wait times : P1 : 0 ms
 P2 : 24 ms
 P3 : 40 ms
 P4 : 72 ms


Average : 34 ms

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



Wait times : P1 : 0 ms
 P2 : 24 ms
 P3 : 40 ms
 P4 : 72 ms

Average : 34 ms

Q: Can we do better?


**Q: Is a mere reordering of start times
sufficient to reduce average wait time?**

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)

**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms


**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



Wait times : P1 : 0 ms
P2 : 24 ms
P3 : 40 ms
P4 : 72 ms

Average : 34 ms

**Q: How does swapping
the order of P1 and P2
alter the average wait
time for the 4
processes?**




A : reduce
B : increase
C : no change

FIFO : First In First Out == FCFS : First Come First Served
(queue implementation) (behavior description)


**Assume 4 processes and
their required burst times**

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

**If $P1 < P2 < P3 < P4$ and assuming
nonpreemptive scheduling**



**If $P2 < P1 < P3 < P4$ and assuming
nonpreemptive scheduling**



Average Wait times

$P1 < P2 < P3 < P4$: 34 ms
 $P2 < P1 < P3 < P4$: 32 ms

Shortest Job First (SJF)

Assume 4 processes and
their required burst times

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

If all 4 processes are requesting
access to the CPU ... *



* That also means that the ready queue is no
longer a “queue” in the classical sense ...

Shortest Job First (SJF)

Assume 4 processes and
their required burst times

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

If all 4 processes are requesting
access to the CPU ... *



What would be the most
suitable data structure for an
SJF scheduling algorithm?

Shortest Job First (SJF)

Assume 4 processes and
their required burst times

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

If all 4 processes are requesting
access to the CPU ... *



	insert	deleteMin	remove	findMin
ordered array	$O(n)$	$O(1)$	$O(n)$	$O(1)$
ordered list	$O(n)$	$O(1)$	$O(1)$	$O(1)$
unordered array	$O(1)$	$O(n)$	$O(1)$	$O(n)$
unordered list	$O(1)$	$O(n)$	$O(1)$	$O(n)$
binary heap	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(1)$

Shortest Job First (SJF)

Assume 4 processes and
their required burst times

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms

If all 4 processes are requesting
access to the CPU ... *




	insert	deleteMin	remove	findMin
ordered array	$O(n)$	$O(1)$	$O(n)$	$O(1)$
ordered list	$O(n)$	$O(1)$	$O(1)$	$O(1)$
unordered array	$O(1)$	$O(n)$	$O(1)$	$O(n)$
unordered list	$O(1)$	$O(n)$	$O(1)$	$O(n)$
binary heap	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(1)$

Shortest Job First (SJF)

Assume 4 processes and
their required burst times

P1	24 ms
P2	16 ms
P3	32 ms
P4	4 ms

If all 4 processes are requesting
access to the CPU ...



Worksheet Q3

Q: What are the advantages of this approach?
Q: What are the disadvantages of this approach?
(versus FCFS)


Assume 4 processes and
their required burst times

P1 24 ms
P2 16 ms
P3 32 ms
P4 4 ms


Average wait times

FCFS : 34ms
SJF : 17ms


First Come First Serve $P1 < P2 < P3 < P4$



Shortest Job First




ESTIMATING PROCESS BURST DURATION



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

SHORTEST JOB FIRST (SJF): ESTIMATING PROCESS BURST DURATION

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.
- Commonly, α set to $1/2$



IN CLASS EXERCISE

**Assume 5 processes and
their required burst times
and their arrival times**

P11	4 ms	t=0
P37	5 ms	t=5
P7	2 ms	t=2
P16	4 ms	t=3
P25	3 ms	t=7

**Task : If SJF nonpreemptive CPU scheduling is
used, then what is the average wait time for
the 5 processes, and at what time does each
process finish?**

IN CLASS EXERCISE

Task : If SJF nonpreemptive CPU scheduling is used, then what is the average wait time for the 5 processes, and at what time does each process finish?

Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		


Average Wait Time	
-------------------	--

IN CLASS EXERCISE



Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE



Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE



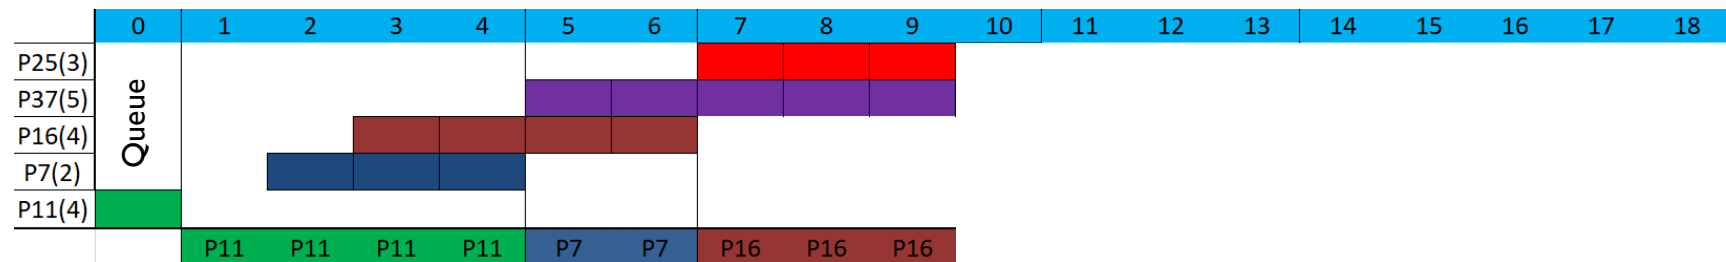
Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE




Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE



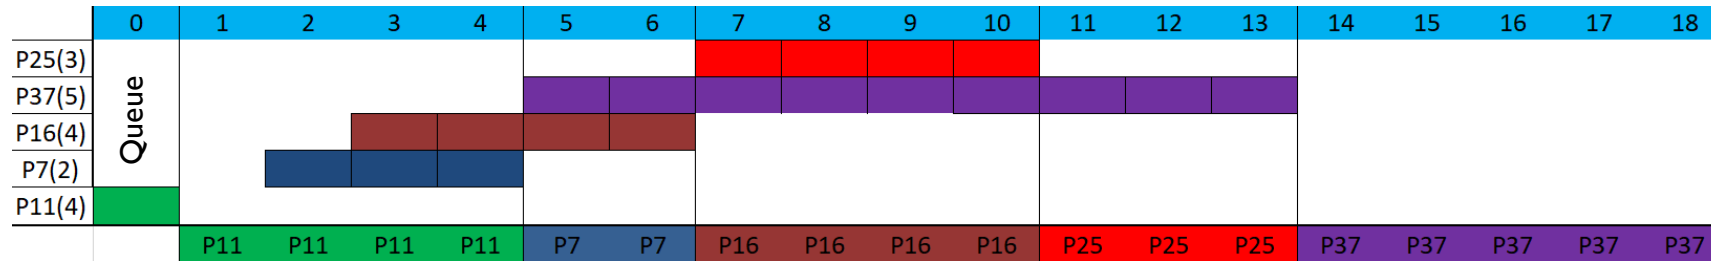
Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE



Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms		
P37	5ms	t=5ms		
P7	2ms	t=2ms		
P16	4ms	t=3ms		
P25	3ms	t=7ms		
Average Wait Time				

IN CLASS EXERCISE



Process	Burst Duration	Arrival Time	Finish Time	Wait Time
P11	4ms	t=0ms	4	0
P37	5ms	t=5ms	18	9
P7	2ms	t=2ms	6	3
P16	4ms	t=3ms	10	4
P25	3ms	t=7ms	13	4
Average Wait Time			4 ms	

PRIORITY SCHEDULING

- Shortest Time First is a special case of priority scheduling:
 - Process priority is the inverse of the process burst time.
- Priority can be established based on other factors.

PRIORITY SCHEDULING

Up until now, all scheduling schemes have assumed a nonpreemptive approach. A more typical approach is for a CPU/OS to employ a preemptive priority scheduling algorithm.

Q: If a preemptive scheduling algorithm is used, how does that affect the order/re-order of Process dispatch?

- Preemptive: CPU can stop the current burst.
- Non-Preemptive: CPU let the process finish its burst.

PRIORITY SCHEDULING

Q: Priorities based on what factor/criteria?

PRIORITY SCHEDULING

Q: Priorities based on what factor/criteria?

Internally defined priority : based on some measurable (computable) quantity ... that the OS can reason about / calculate

PRIORITY SCHEDULING

Q: Priorities based on what factor/criteria?

Internally defined priority : based on some measurable (computable) quantity ... that the OS can reason about / calculate

Externally defined priority : based on criteria outside of the OS

Priority Scheduling

Assume 5 processes, their required burst times, and priorities*

P1	24 ms	1
P2	16 ms	3
P3	2 ms	5
P4	4 ms	2
P5	3 ms	4

If P1-P5 at $t=0$ are all waiting to be dispatched, and using a nonpreemptive scheduling scheme ..

*Low numbers refer to high priorities


Priority Scheduling

Assume 5 processes, their required burst times, and priorities*

P1	24 ms	1
P2	16 ms	3
P3	2 ms	5
P4	4 ms	2
P5	3 ms	4

*Low numbers refer to high priorities

If P1-P5 at $t=0$ are all waiting to be dispatched, and using a nonpreemptive scheduling scheme ..



PRIORITY SCHEDULING

- Preemptive: CPU can stop the current burst.
- Non-Preemptive: CPU let the process finish its burst.

Up until now, all scheduling schemes have assumed a nonpreemptive approach. A more typical approach is for a CPU/OS to employ a preemptive priority scheduling algorithm.

Q: If a preemptive scheduling algorithm is used, how does that affect the order/re-order of Process dispatch?

PRIORITY SCHEDULING

- Preemptive: CPU can stop the current burst.
- Non-Preemptive: CPU let the process finish its burst.

Up until now, all scheduling schemes have assumed a nonpreemptive approach. A more typical approach is for a CPU/OS to employ a preemptive priority scheduling algorithm.

Q: If a preemptive scheduling algorithm is used, how does that affect the order/re-order of Process dispatch?

Processes will be interrupted before they finish their burst.

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

- Assume 5 processes and their required burst times, arrival times, and priorities shown in the table below.
- Q: If **preemptive** (processes can be removed before completion) priority scheduling is used, then at what time does each process finish? What is the total time of each process? How many context switch(es) does each process experience?

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3																	
	P4																	
	P5																	
CPU																		

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3																	
	P4																	
	P5																	
CPU			P2															

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3			5														
	P4																	
	P5																	
CPU			P2	P2														

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3			5														
	P4				2													
	P5																	
CPU			P2	P2	P2													

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3			5														
	P4				2													
	P5																	
CPU			P2	P2	P2	P4												

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1																	
	P2	3																
	P3			5														
	P4				2													
	P5																	
CPU			P2	P2	P2	P4	P4											

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5																	
CPU			P2	P2	P2	P4	P4	P2										

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5																	
CPU			P2	P2	P2	P4	P4	P2	P1									

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5									4								
CPU			P2	P2	P2	P4	P4	P2	P1	P1								

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5									4								
CPU			P2	P2	P2	P4	P4	P2	P1	P1	P1							

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6			
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5									4								
CPU			P2	P2	P2	P4	P4	P2	P1	P1	P1	P2						


Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6	9		
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5									4								
CPU			P2	P2	P2	P4	P4	P2	P1	P1	P1	P2	P2					


Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6	9		
P2	6	3	t=0			
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE



Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6	9		
P2	6	3	t=0	11		
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE



Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6	9		
P2	6	3	t=0	11		
P3	2	5	t=2			
P4	2	2	t=3			
P5	2	4	t=7			

PREEMPTIVE PRIORITY SCHEDULING EXERCISE

		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Queue	P1							1										
	P2	3																
	P3			5														
	P4				2													
	P5								4									
CPU			P2	P2	P2	P4	P4	P2	P1	P1	P1	P2	P2	P5	P5	P3	P3	

Process	Burst Time	Priority	Arrival	Completion Time	Wait time	# of Context Switch(es)
P1	3	1	t=6	9	1	1
P2	6	3	t=0	11	5	3
P3	2	5	t=2	15	12	1
P4	2	2	t=3	5	1	1
P5	2	4	t=7	13	5	1

ROUND ROBIN SCHEDULING

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

ROUND ROBIN SCHEDULING

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Burst time : how much time a process wants

Time Quantum : how much time a process gets (during each “turn” in the CPU)

ROUND ROBIN SCHEDULING

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Burst time : how much time a process wants

Time Quantum : how much time a process gets (during each “turn” in the CPU)

No process can utilize more than the time quantum in the CPU. It has to yield to others.

ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time


Clock tick: 0

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose **5ms**
- Each process has a required burst time


Clock tick: 0

Quantum

FIFO Ready “circular” Queue



Burst times



ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose 5ms
- Each process has a required burst time


Clock tick: 0

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7


P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 0

FIFO Ready “circular” Queue



Q: How long does
P3 have CPU access
for?



ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 0

FIFO Ready “circular” Queue

(tail)



(head)

Q: What is the
turnaround time
for P3?



ROUND ROBIN SCHEDULING

Turnaround time

P3

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 3


FIFO Ready “circular” Queue

(tail)



(head)

Q: What is the
turnaround time
for P3?



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time


Clock tick: 3

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7


P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 3

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7

P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose 5ms
- Each process has a required burst time

Clock tick: 3

FIFO Ready “circular” Queue



Worksheet Q1: At what time does P32
Leave the CPU? What would be its
remaining CPU time?

ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7


P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 3

FIFO Ready “circular” Queue



Worksheet Q1: At what time does P32
Leave the CPU? What would be its
remaining CPU time?

ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 8

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7


P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 8

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7


P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick: 8

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1

Round Robin Scheduling


- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

10

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

10

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:


15

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

15

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:


15

FIFO Ready “circular” Queue

(tail)



(head)



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

20

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

20

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3 3ms

P32

P7 10ms

P1


Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:

21

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3	3ms
P32	
P7	10ms
P1	21ms

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time

Clock tick:
21

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

Turnaround time

P3	3ms
P32	23ms
P7	10ms
P1	21ms

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let’s choose 5ms
- Each process has a required burst time


Clock tick:

23

FIFO Ready “circular” Queue



Q2: What is the disadvantage of Round Robin Scheduling? How to mitigate it?



ROUND ROBIN SCHEDULING

Turnaround time


P3	3ms
P32	23ms
P7	10ms
P1	21ms

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose **5ms** → Quantum
- Each process has a required burst time

Clock tick: 0

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

- The time quantum should be large compared with the context switch time
- However, if the time quantum is too large, RR scheduling degenerates to an FCFS policy.
- A rule of thumb is that 80 percent of the CPU bursts should be shorter than the time quantum.

Turnaround time


P3	3ms
P32	23ms
P7	10ms
P1	21ms

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose **5ms** → Quantum
- Each process has a required burst time

Clock tick: 0

FIFO Ready “circular” Queue



ROUND ROBIN SCHEDULING

- The time quantum should be large compared with the context switch time
- However, if the time quantum is too large, RR scheduling degenerates to an FCFS policy.
- A rule of thumb is that 80 percent of the CPU bursts should be shorter than the time quantum.

Turnaround time


P3	3ms
P32	23ms
P7	10ms
P1	21ms

Round Robin Scheduling

- “Circular” Queue
- Time Quantum .. Here let's choose **5ms**
- Each process has a required burst time

Clock tick: 0


FIFO Ready “circular” Queue



Round Robin is simply trying to cut off “long bursts” while avoiding cutting off any “regular burst” processes.

MULTILEVEL QUEUE

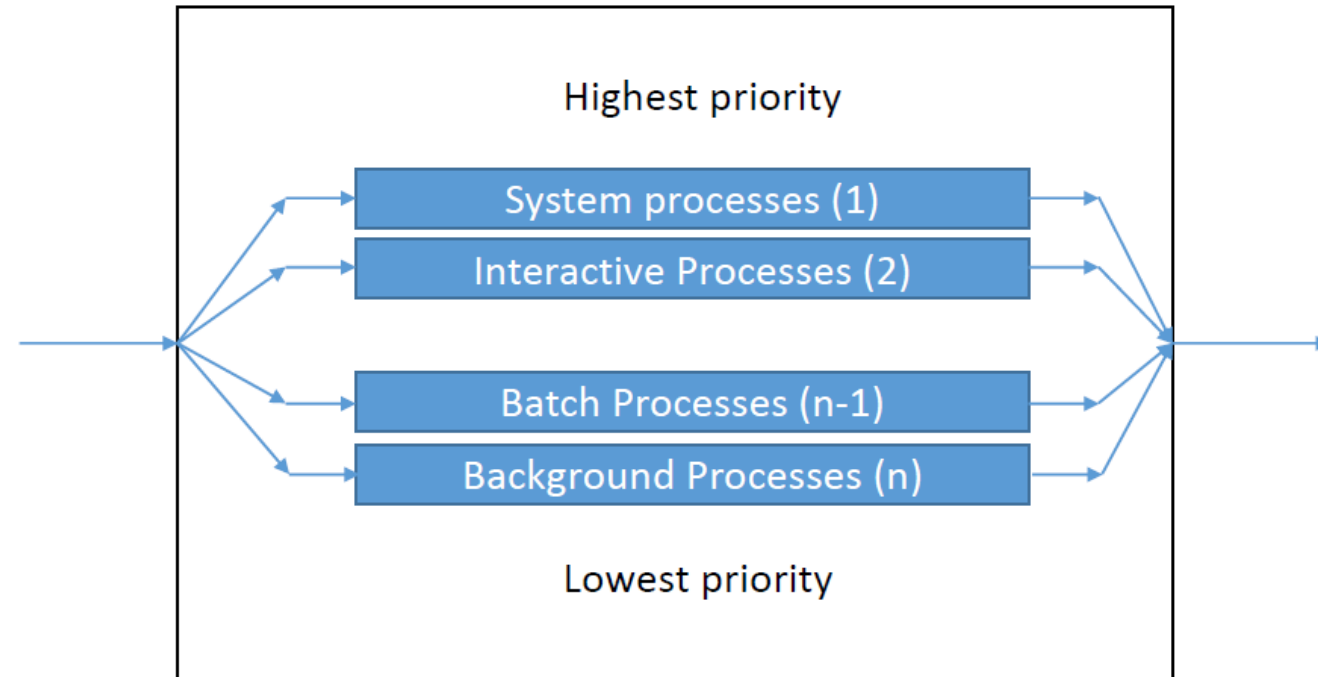
Multilevel Queue The ready queue is partitioned into several separate queues



MULTILEVEL QUEUE

Multilevel Queue The ready queue is partitioned into several separate queues

Creating different queues for different priorities.

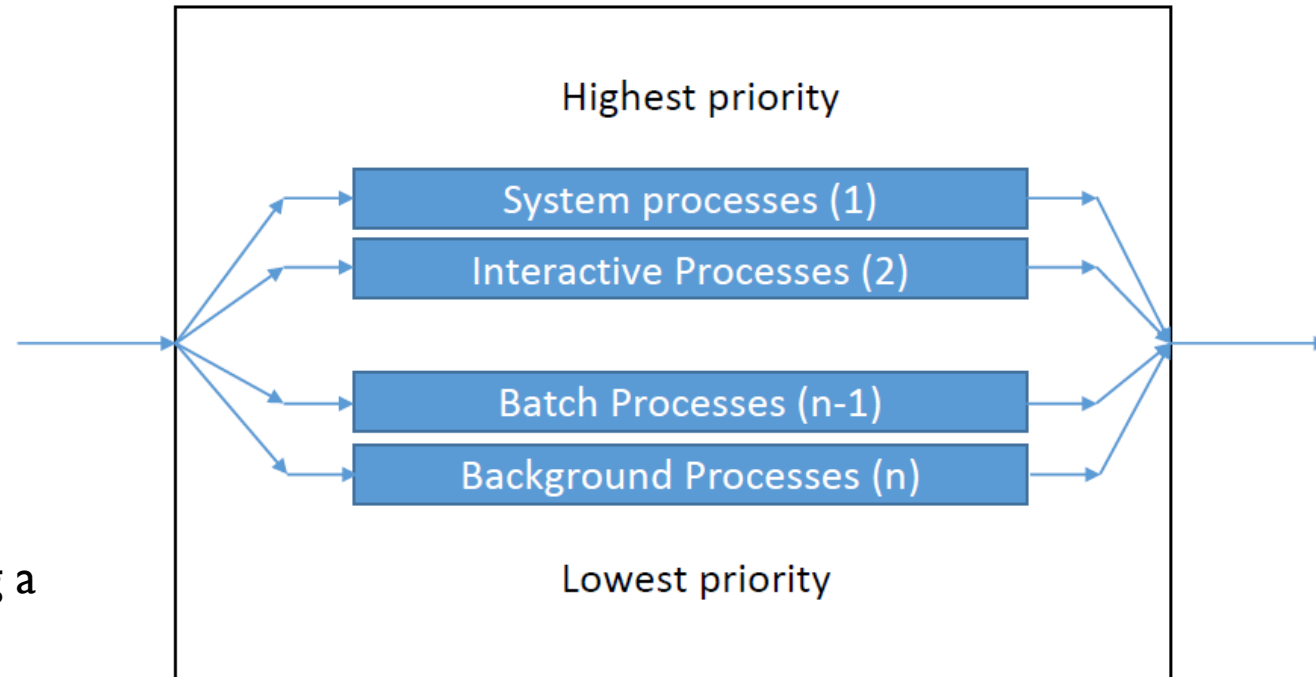


MULTILEVEL QUEUE

Multilevel Queue The ready queue is partitioned into several separate queues

Creating different queues for different priorities.

Q3: What is the advantage of this approach over having a single queue for all priorities?



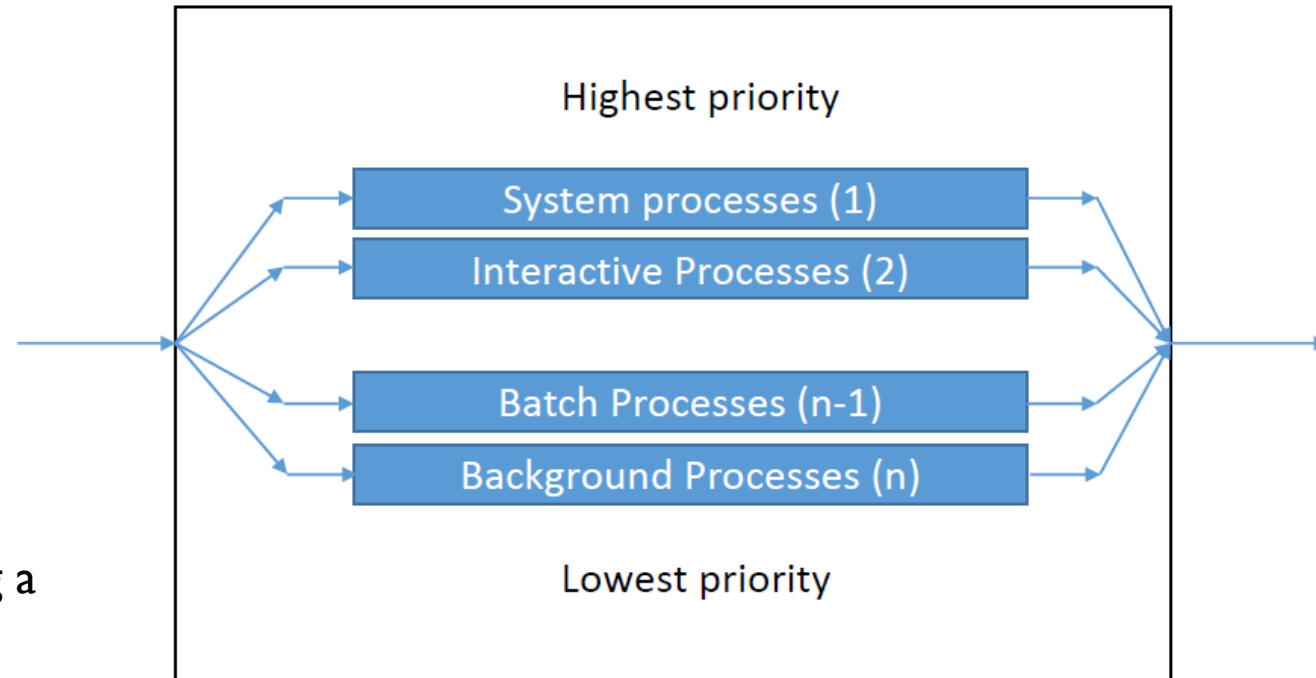
MULTILEVEL QUEUE

Multilevel Queue The ready queue is partitioned into several separate queues

Creating different queues for different priorities.

Q3: What is the advantage of this approach over having a single queue for all priorities?

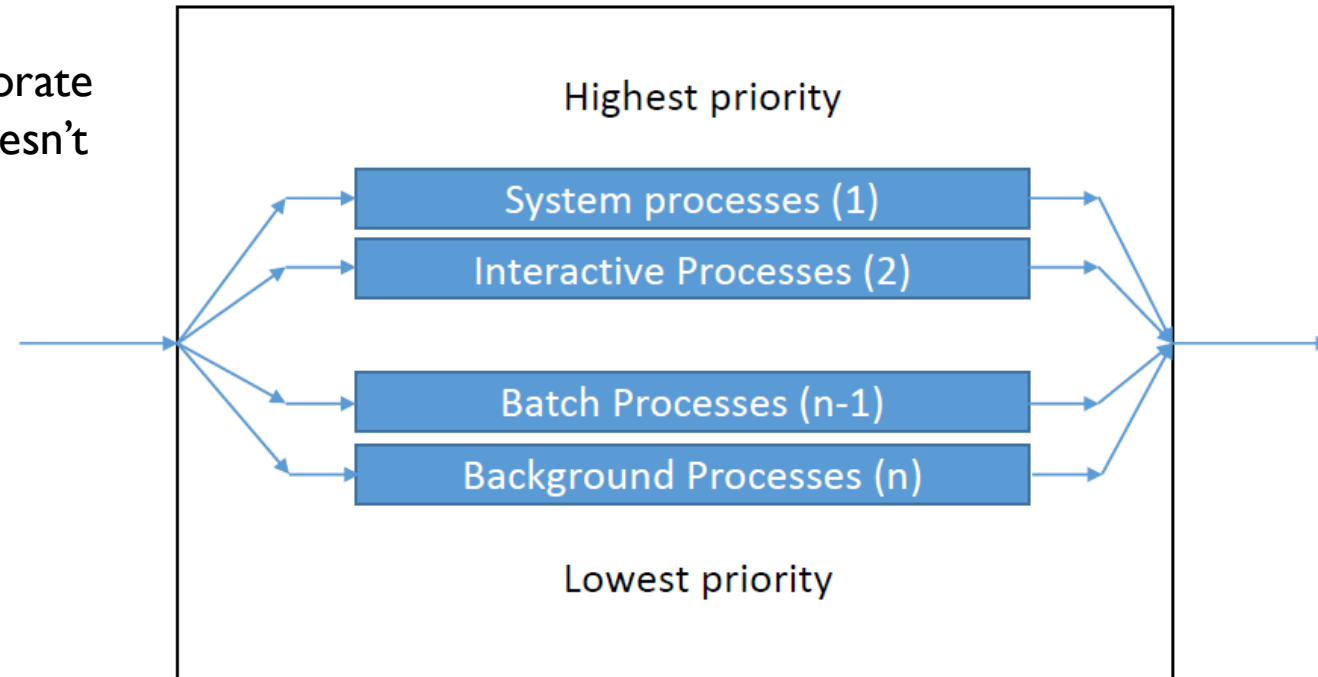
No need for sorting!



MULTILEVEL QUEUE

Multilevel Queue The ready queue is partitioned into several separate queues

We can do more elaborate scheduling. Priority doesn't have to be absolute.

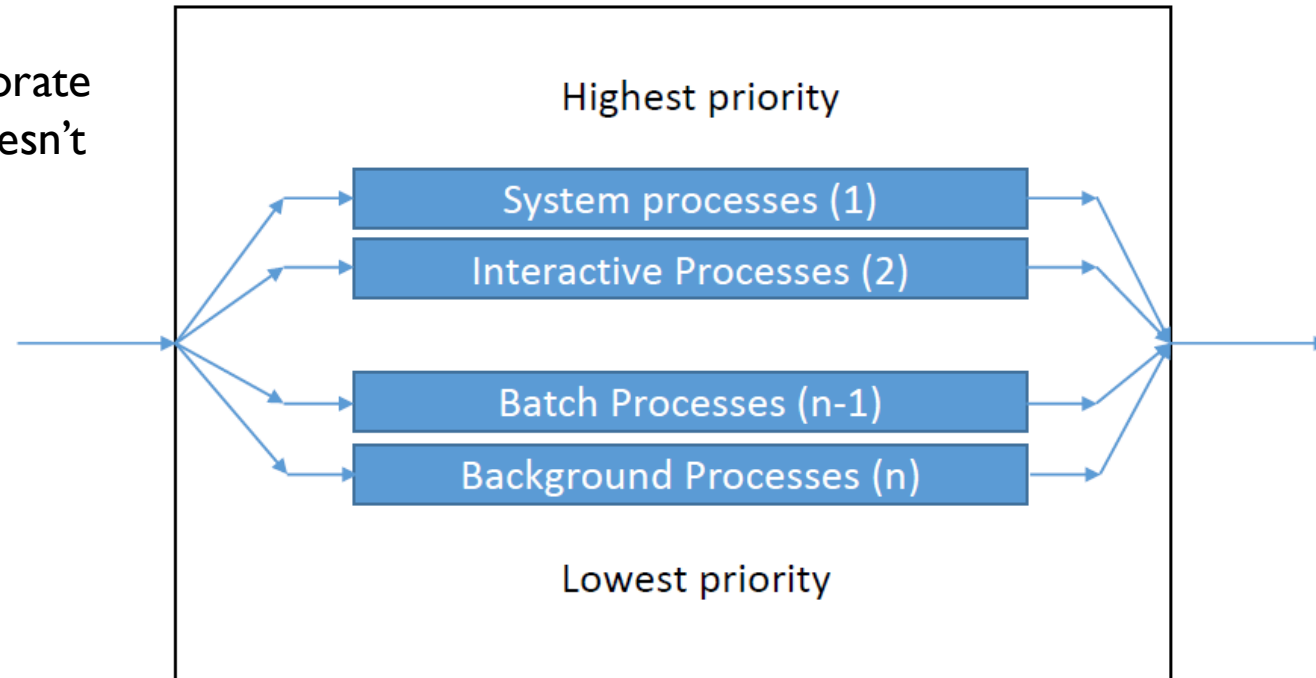


MULTILEVEL QUEUE

Multilevel Queue The ready queue is partitioned into several separate queues


We can do more elaborate scheduling. Priority doesn't have to be absolute.

Highest priority queue might have 70% of CPU time, while the others 30%



MULTILEVEL QUEUE WITH FEEDBACK

Processes can migrate from one queue to another.




MULTILEVEL QUEUE WITH FEEDBACK

Multilevel Feedback Queue

Possible design

- If a dispatched process from the quantum=5 queue does not finish, a context switch occurs, and process is placed at tail end of quantum=10 queue.




MULTILEVEL QUEUE WITH FEEDBACK

Multilevel Feedback Queue


Possible design

- If a dispatched process from the quantum=5 queue does not finish, a context switch occurs, and process is placed at tail end of quantum=10 queue.
- If a process from the quantum=10 queue does not finish, a context switch occurs, and process is placed at the tail end of the quantum=20 queue, etc. etc.




MULTILEVEL QUEUE WITH AGING

- Can a low priority process be starved?



MULTILEVEL QUEUE WITH AGING


- Can a low priority process be starved?
- If the processor is always busy with high priority processes, low priority ones can become starved.



MULTILEVEL QUEUE WITH AGING

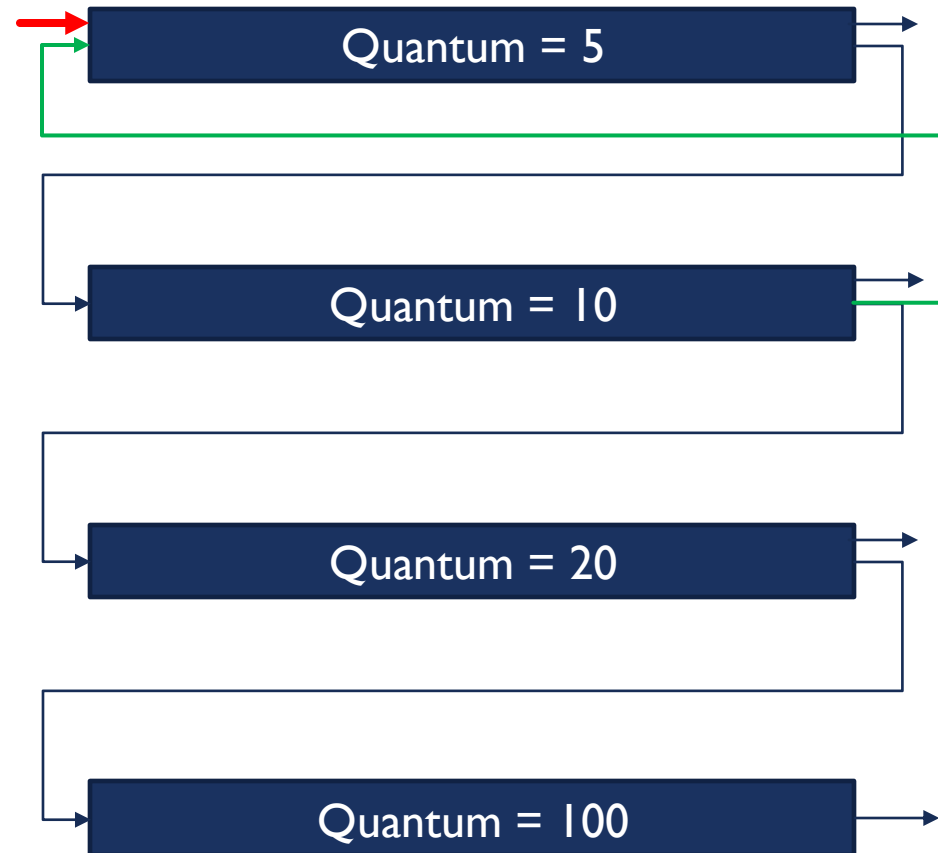
- Can a low priority process be starved?
- If the processor is always busy with high priority processes, low priority ones can become starved.

Q4: How would you fix the starvation problem?



MULTILEVEL QUEUE WITH AGING

- Can a low priority process be starved?
- If the processor is always busy with high priority processes, low priority ones can become starved.
- Aging can solve this.
- At predefined intervals, all processes in queue have their priority increased.
- After execution, a process priority goes back to its default value.



REAL-TIME SCHEDULING

- Time sensitive application:
 - Transportation: smart cars, jets, trains ...
 - Manufacturing pipelines
 - Playing audio/video

REAL-TIME SCHEDULING

- Time sensitive application:
 - Transportation: smart cars, jets, trains ...
 - Manufacturing pipelines
 - Playing audio/video
- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
 - But only guarantees soft real-time

REAL-TIME SCHEDULING

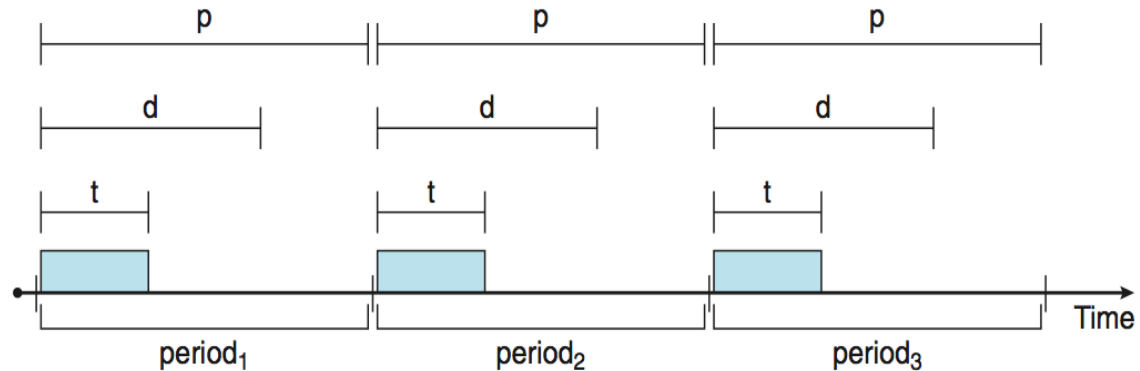
- Time sensitive application:
 - Transportation: smart cars, jets, trains ...
 - Manufacturing pipelines
 - Playing audio/video
- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
 - But only guarantees soft real-time
- For **hard** real-time must also provide ability to meet deadlines. If it can't, the OS would refuse to run the program.

REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - processing time t
 - deadline d
 - period p
 - $0 \leq t \leq d \leq p$
 - **Rate** of periodic task is $1/p$

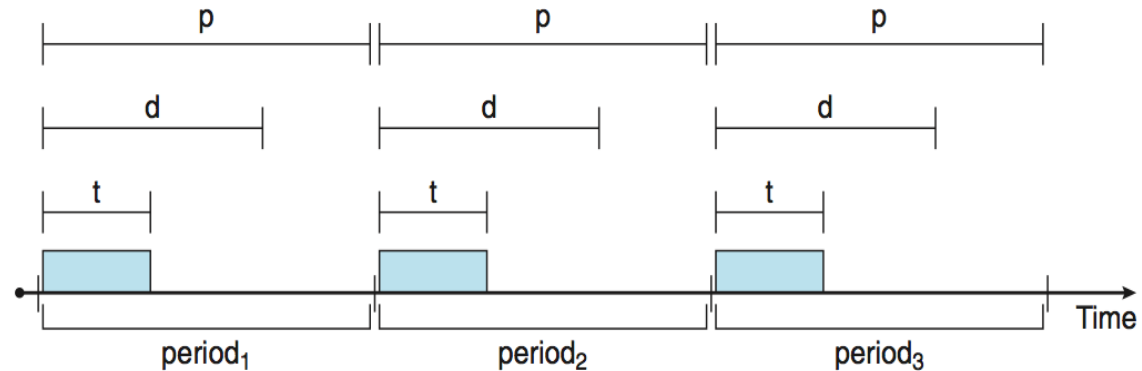
REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - processing time t
 - deadline d
 - period p
 - $0 \leq t \leq d \leq p$
 - Rate** of periodic task is $1/p$



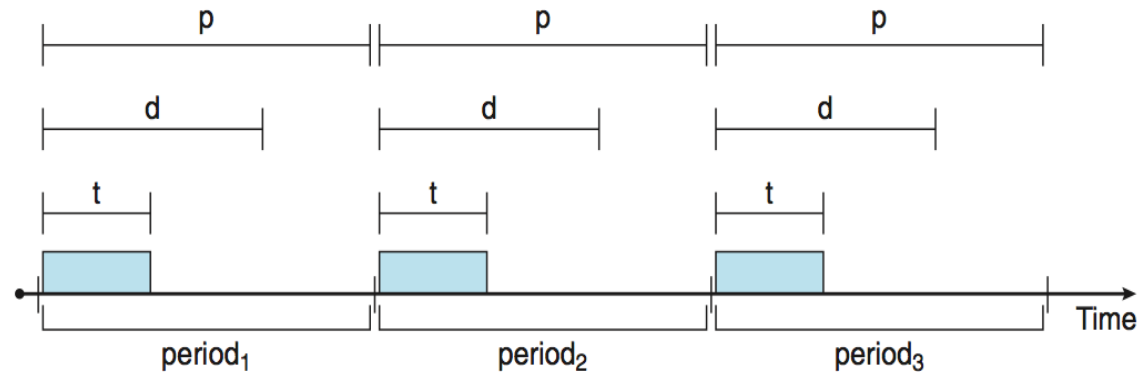
REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - processing time t
 - deadline d
 - period p
 - $0 \leq t \leq d \leq p$
 - Rate** of periodic task is $1/p$
- If $p < t$... what does that mean?
Processor is not fast enough!




REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - Has processing time t , deadline d , period p
 - $0 \leq t \leq d \leq p$
 - Rate** of periodic task is $1/p$



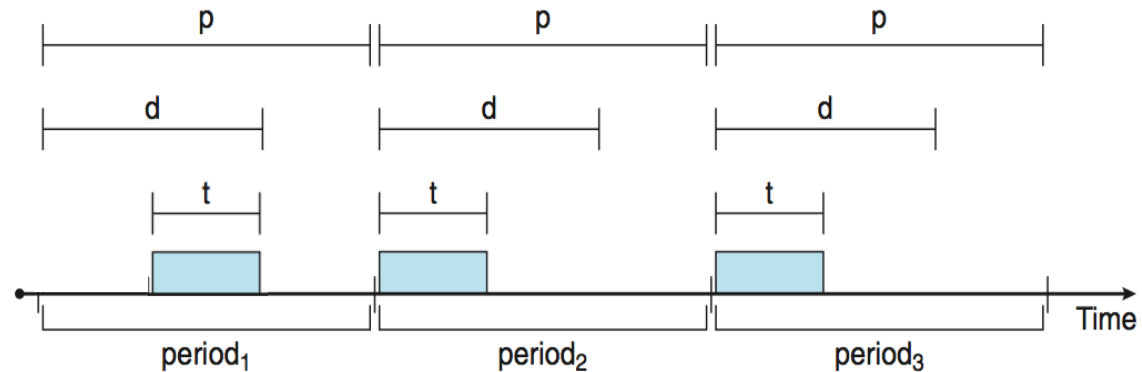
REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - Has processing time t , deadline d , period p
 - $0 \leq t \leq d \leq p$
 - Rate** of periodic task is $1/p$



REAL-TIME SCHEDULING

- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - Has processing time t , deadline d , period p
 - $0 \leq t \leq d \leq p$
 - Rate** of periodic task is $1/p$



REAL-TIME SCHEDULING

Periodic : occurring at a constant interval, or period (p)

Processing time : time needed to execute (burst time, t)

Deadline : time by which a process must finish

All units
are ms

Q: if a requesting process does not “satisfy”
the $0 \leq t \leq d \leq p$
requirement, the CPU scheduler
{might/should/should not} reject the process



A : might
B : should
C : should not

REAL-TIME SCHEDULING

Periodic : occurring at a constant interval, or period (p)

Processing time : time needed to execute (burst time, t)

Deadline : time by which a process must finish

All units
are ms

Q: if a requesting process does not “satisfy”
the $0 \leq t \leq d \leq p$
requirement, the CPU scheduler
{might/should/should not} reject the process



A : might
B : should
C : should not

The processor should **reject** the process
as it can never satisfy its runtime
requirement.


REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period

P1


P2



REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




If P2 is scheduled before P1,
what is the timeline?

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




If P2 is scheduled before P1,
what is the timeline?

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




If P2 is scheduled before P1,
what is the timeline?

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




If P2 is scheduled before P1,
what is the timeline?

P1 misses its deadline!

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




What if P1 goes first?

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




What if P1 goes first?

REAL-TIME SCHEDULING

Process 1 p:50ms, t:20ms, d:once each period

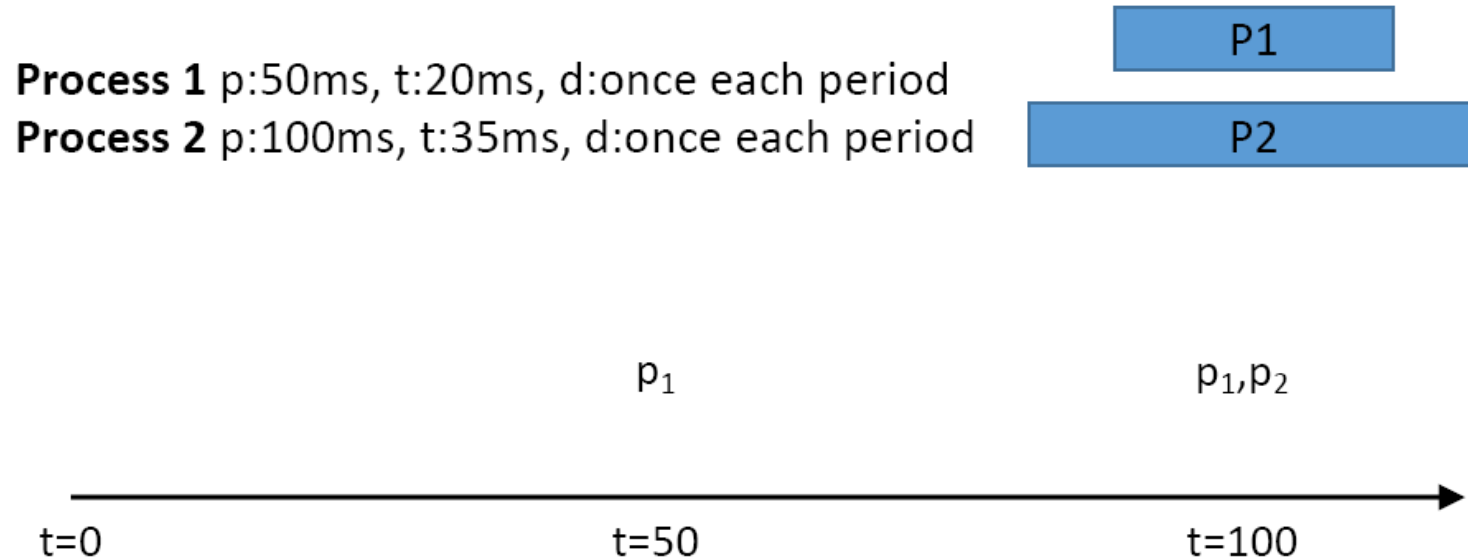
Process 2 p:100ms, t:35ms, d:once each period



What if P1 goes first?

Both processes can meet their deadline.

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

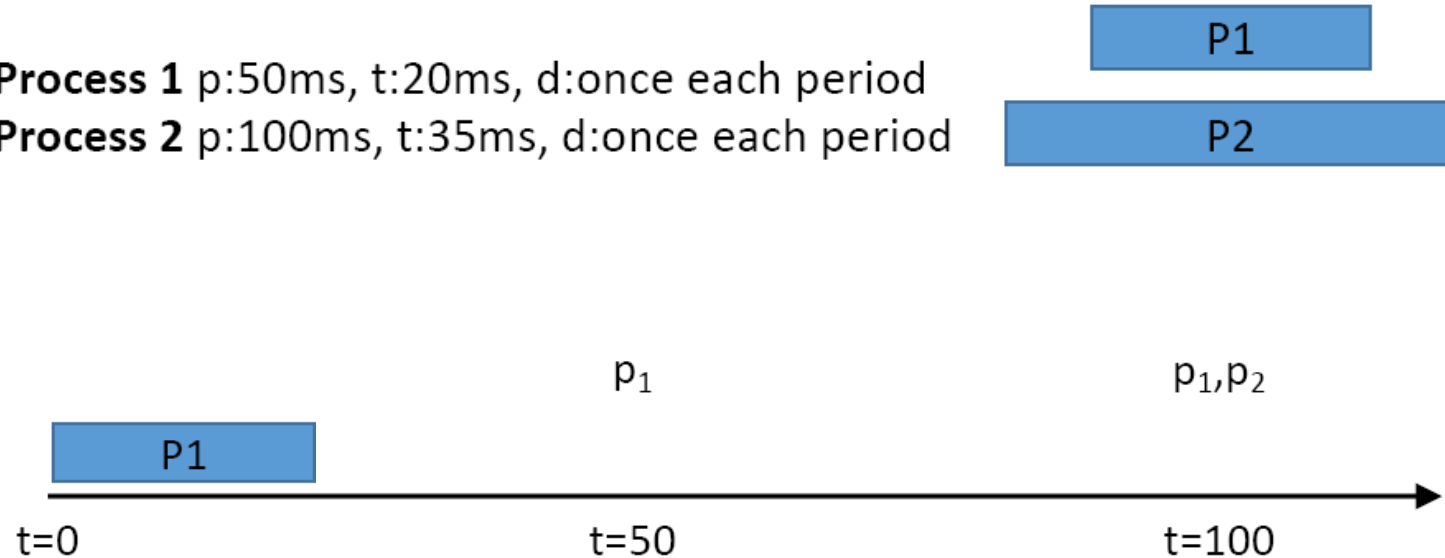


RMS: Assigns higher priority to shorter period with preemptive scheduling.

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




Priority of P1 : $1/50$

Priority of P2 : $1/100$

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period



Priority of P1 : $1/50$

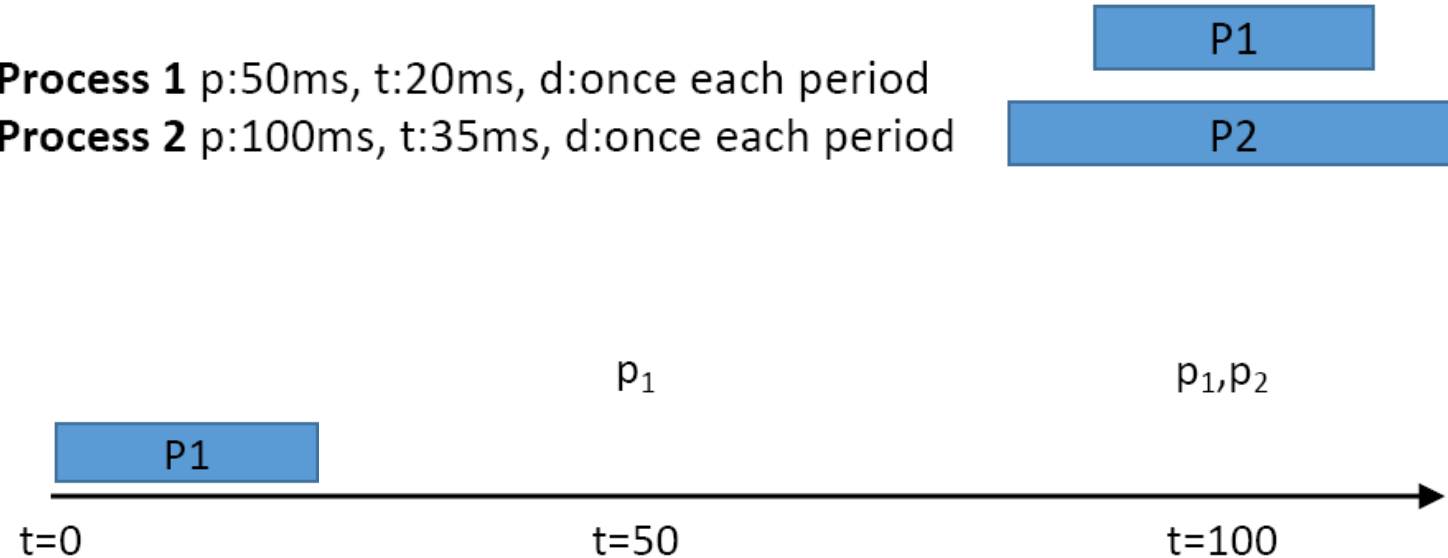
Priority of P2 : $1/100$

- **P1 will have higher priority**
- **Priority is static**

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period



Priority of P1 : $1/50$

Priority of P2 : $1/100$


- P1 will have higher priority
- Priority is static

- Scheduling is preemptive: lower priority process will be removed as soon as higher priority ones are available

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




Priority of P1 : $1/50$
Priority of P2 : $1/100$

Scheduling is preemptive:
lower priority process will be
removed as soon as higher
priority ones are available

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period



Priority of P1 : $1/50$


Priority of P2 : $1/100$

Scheduling is preemptive:
lower priority process will be
removed as soon as higher
priority ones are available

REAL-TIME SCHEDULING: RATE MONOTONIC SCHEDULING (RMS)

Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period



Priority of P1 : $1/50$


Priority of P2 : $1/100$

Scheduling is preemptive:
lower priority process will be
removed as soon as higher
priority ones are available

REAL-TIME SCHEDULING

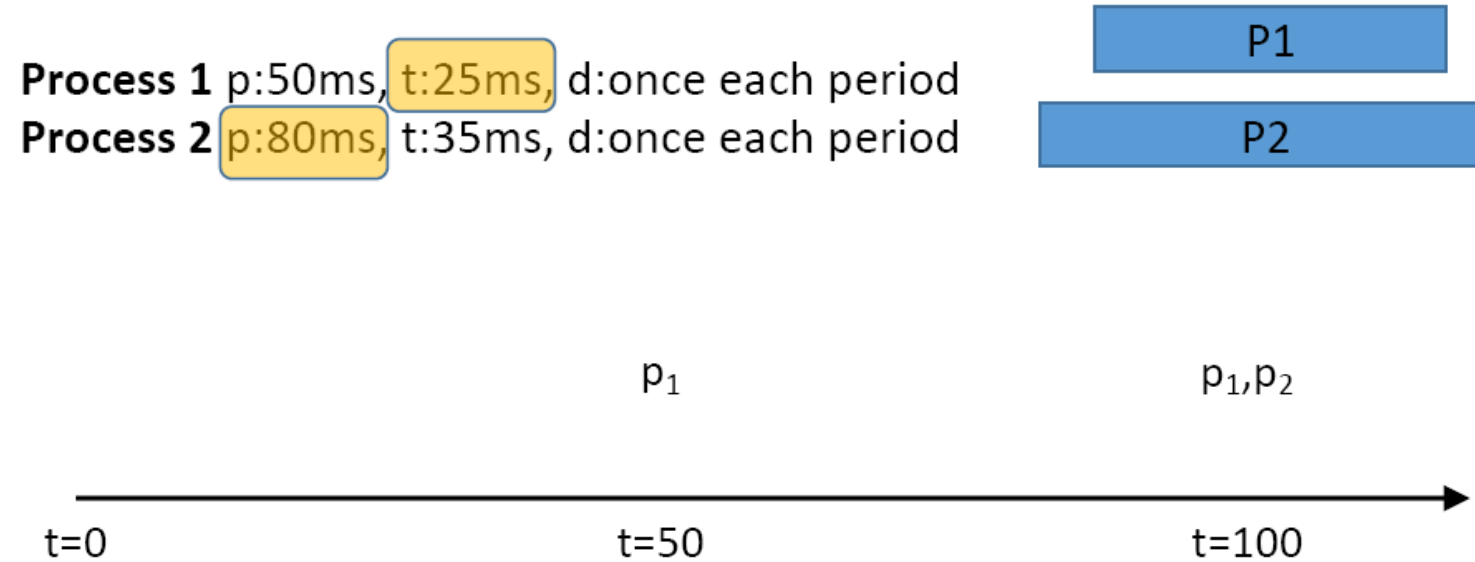
Process 1 p:50ms, t:20ms, d:once each period

Process 2 p:100ms, t:35ms, d:once each period




Cycle will repeat

REAL-TIME SCHEDULING



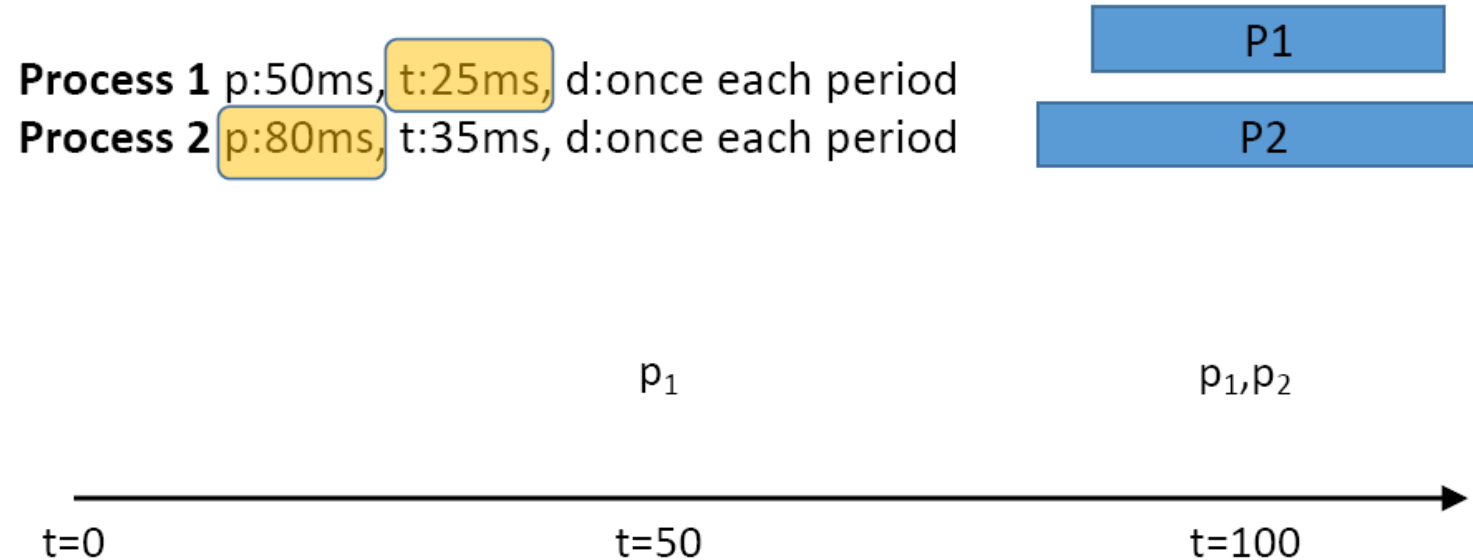
REAL-TIME SCHEDULING



**Q: Do we have enough
CPU time for both?**

**P1: $25/50 = 0.5$
P2: $35/80 = 0.4375$
Total = $0.94375 < 1$**


REAL-TIME SCHEDULING



Q: Do we have enough CPU time for both?
Yes, we should have.

P1: $25/50 = 0.5$
P2: $35/80 = 0.4375$
Total = $0.94375 < 1$

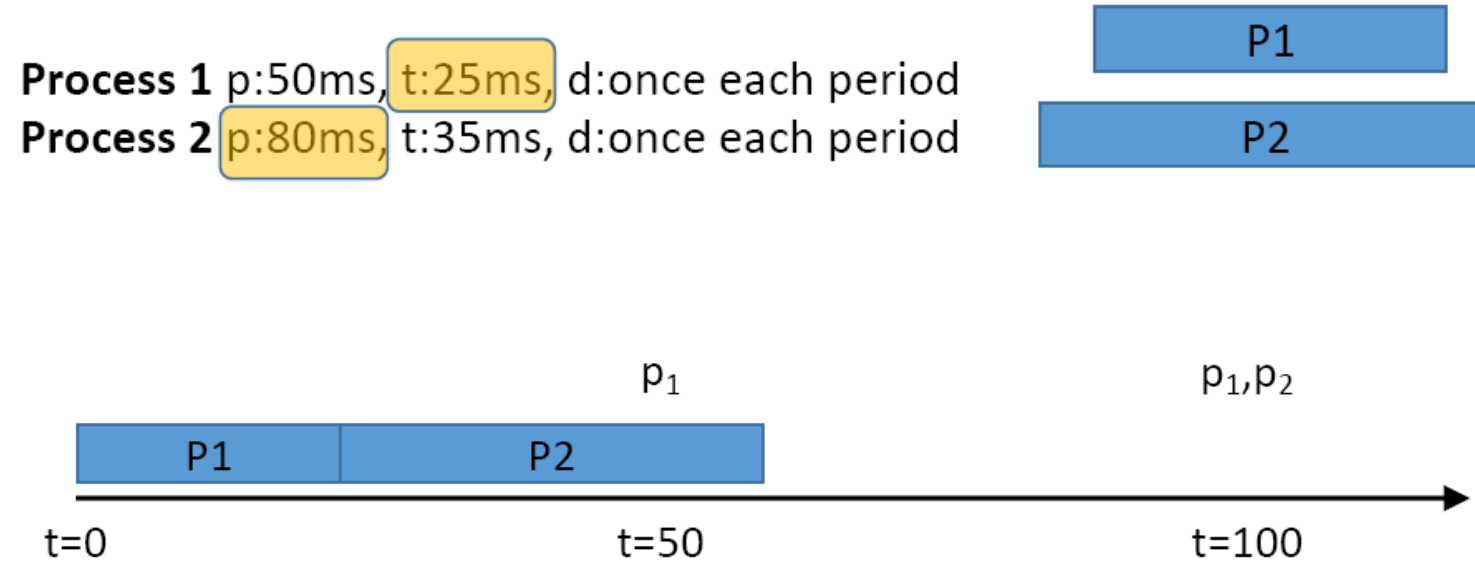
REAL-TIME SCHEDULING



P1 Priority = $1/50$

P2 Priority = $1/80$

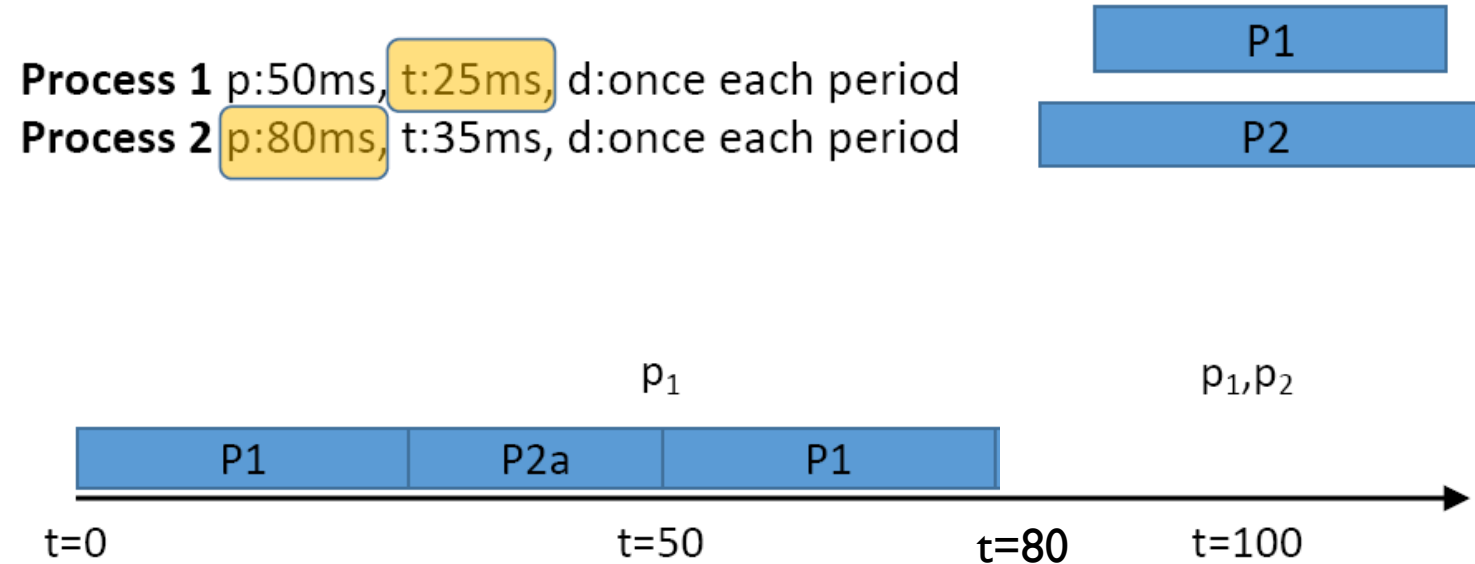
REAL-TIME SCHEDULING



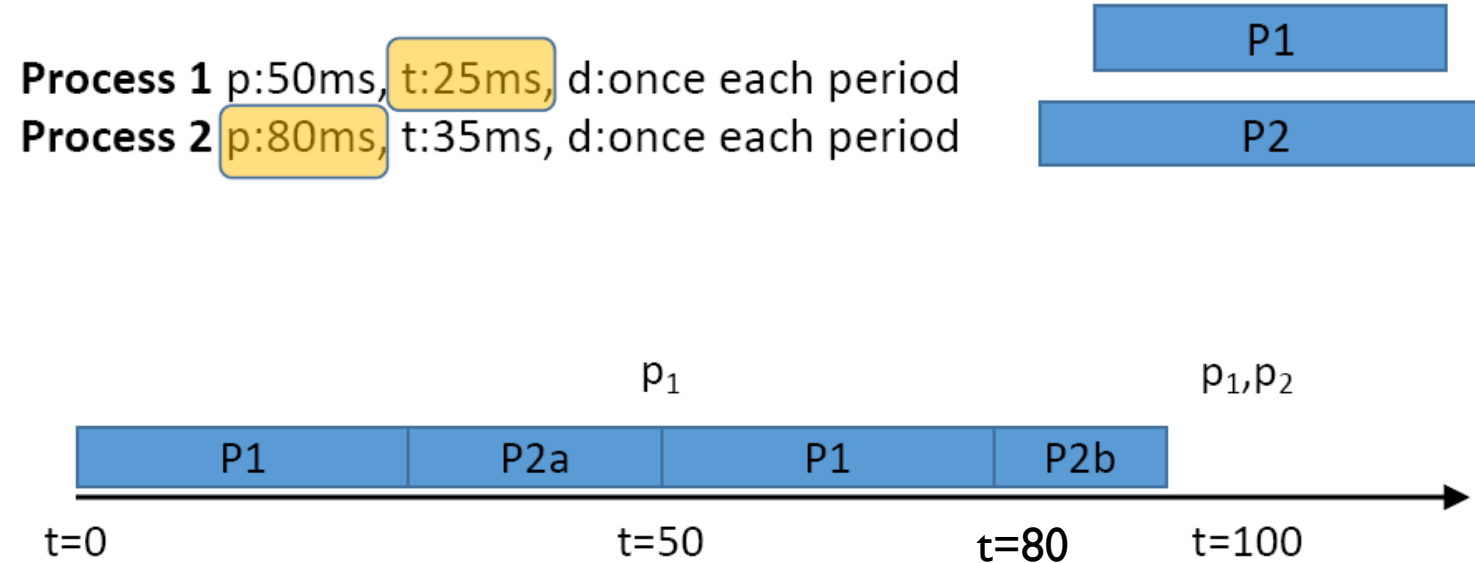
P1 Priority = $1/50$

P2 Priority = $1/80$

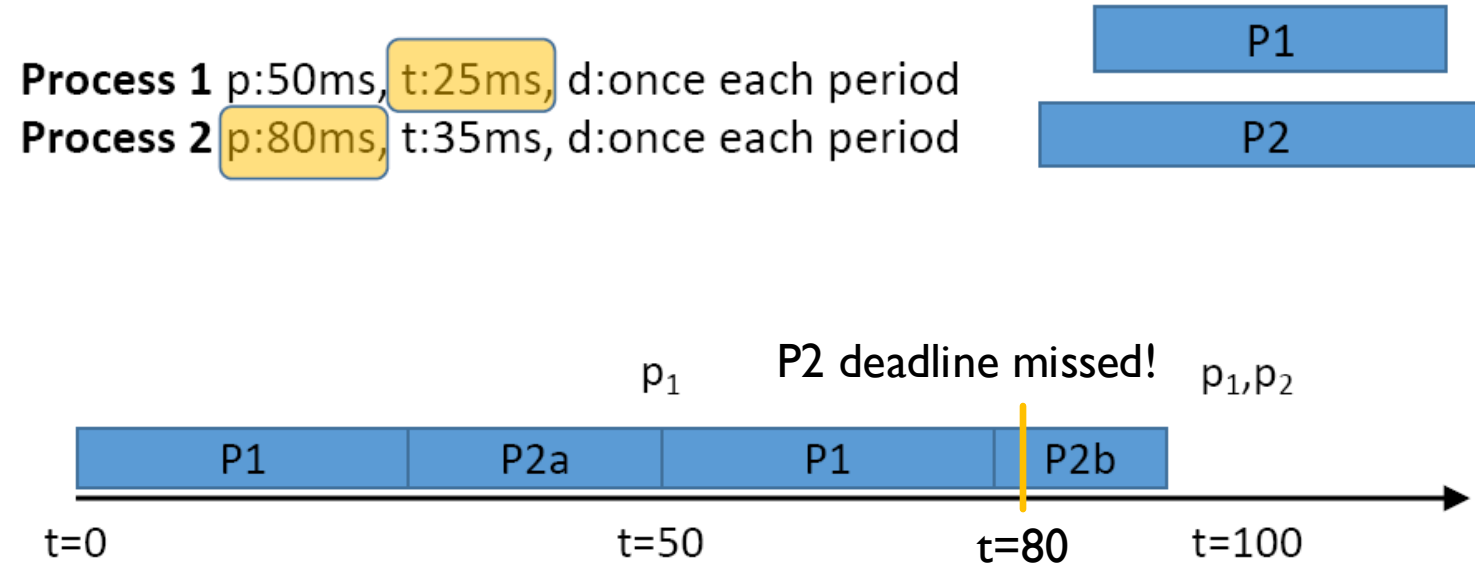
REAL-TIME SCHEDULING



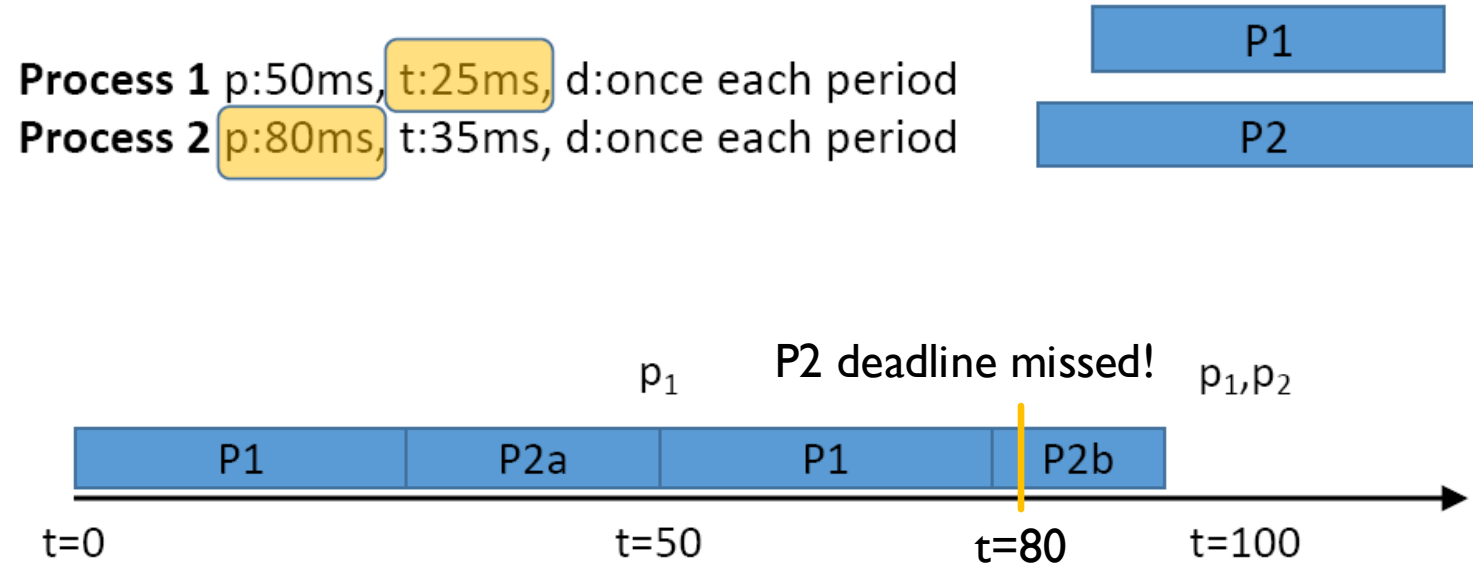
REAL-TIME SCHEDULING



REAL-TIME SCHEDULING


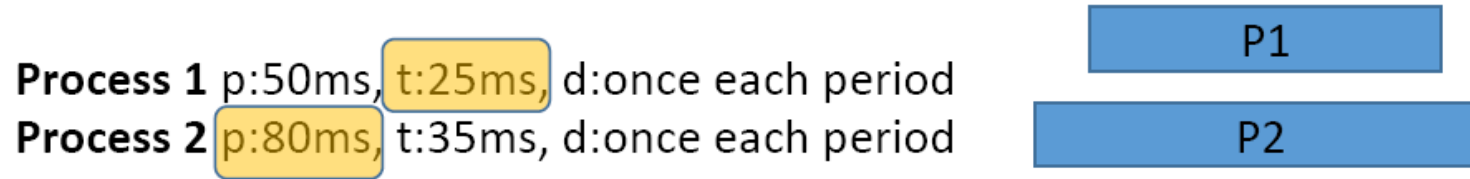


REAL-TIME SCHEDULING



For RMS CPU utilization is actually capped at $N(2^{1/N} - 1)$ where N is the number of processes.


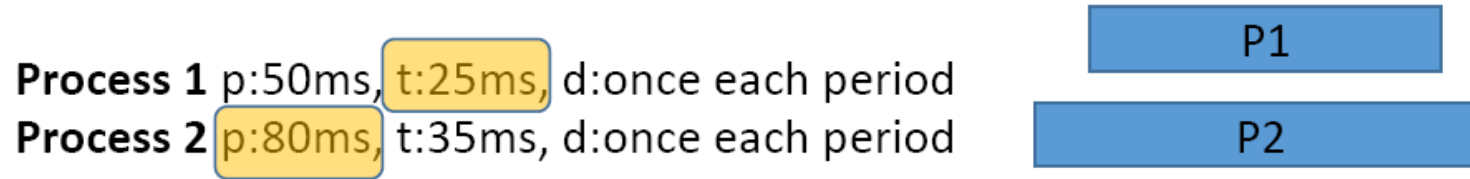
REAL-TIME SCHEDULING



For RMS CPU utilization is actually capped at $N(2^{1/N} - 1)$ where N is the number of processes.

$$\text{CPU Cap: } N(2^{1/N} - 1) = 2(2^{1/2} - 1) = 0.83$$


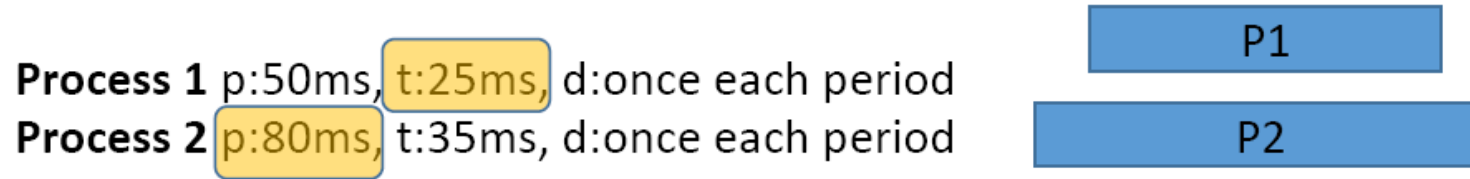
REAL-TIME SCHEDULING



For RMS CPU utilization is actually capped at $N(2^{1/N} - 1)$ where N is the number of processes.

CPU Cap: $N(2^{1/N} - 1) = 2(2^{1/2} - 1) = 0.83$
CPU Utilization required: 0.94


REAL-TIME SCHEDULING



For RMS CPU utilization is actually capped at $N(2^{1/N} - 1)$ where N is the number of processes.

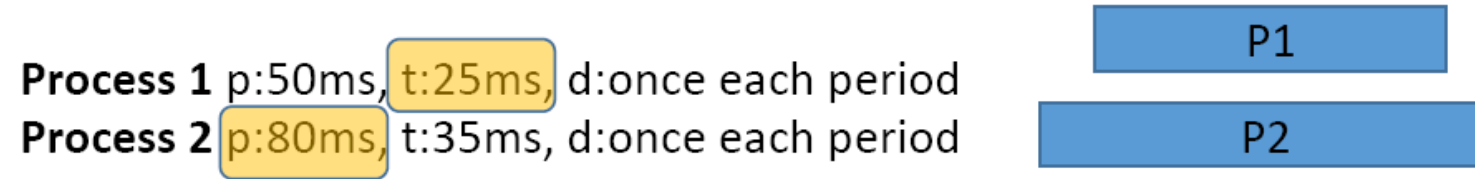
$$\begin{aligned} \text{CPU Cap: } N(2^{1/N} - 1) &= 2(2^{1/2} - 1) = 0.83 \\ \text{CPU Utilization required: } 0.94 &= \\ &= 35/80 + 25/50 \end{aligned}$$

REAL-TIME SCHEDULING




Q: How to fix this?

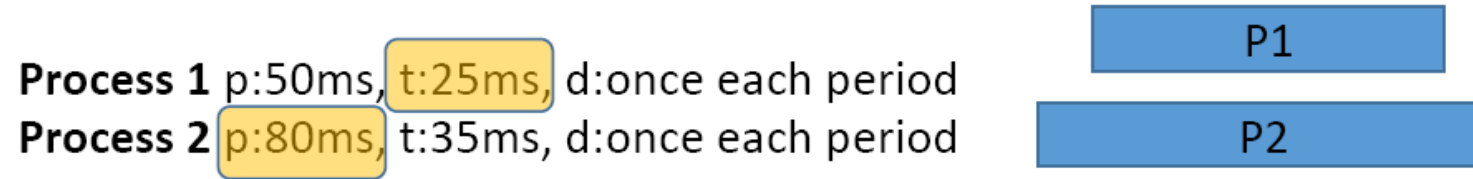
REAL-TIME SCHEDULING




P1 lets P2 complete



REAL-TIME SCHEDULING

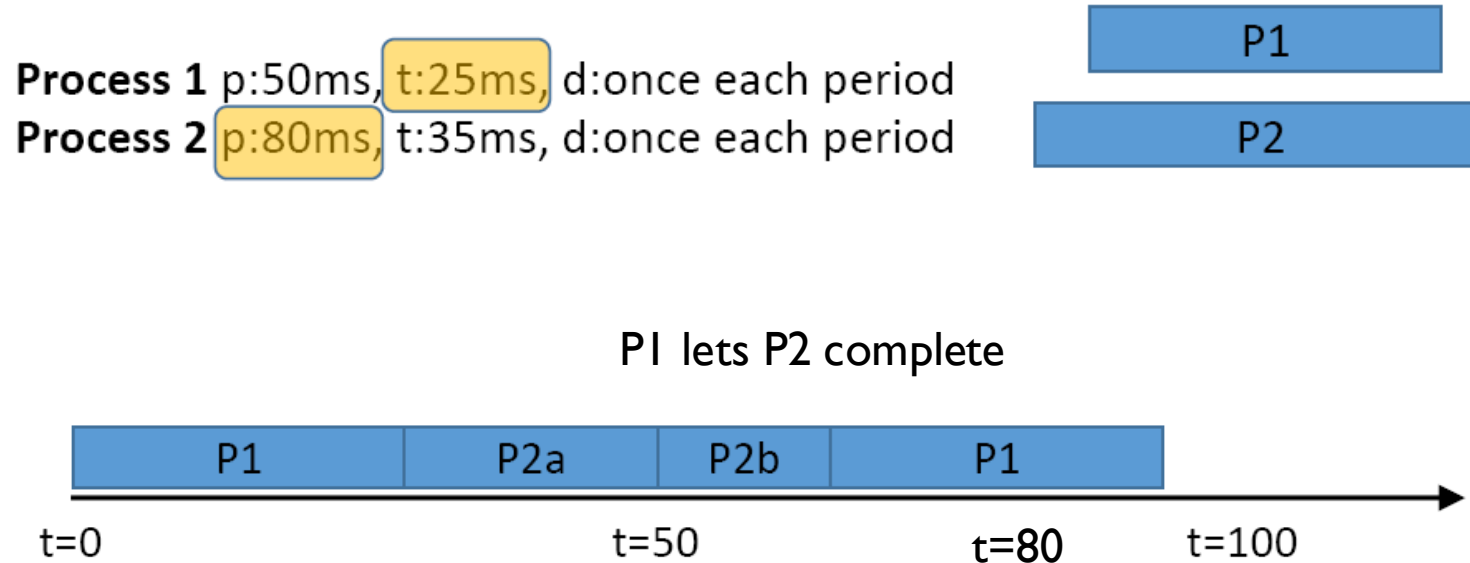


P1 lets P2 complete



- This works but this requires dynamic priority.
- First period: priority was to P1, second period it was to P2.


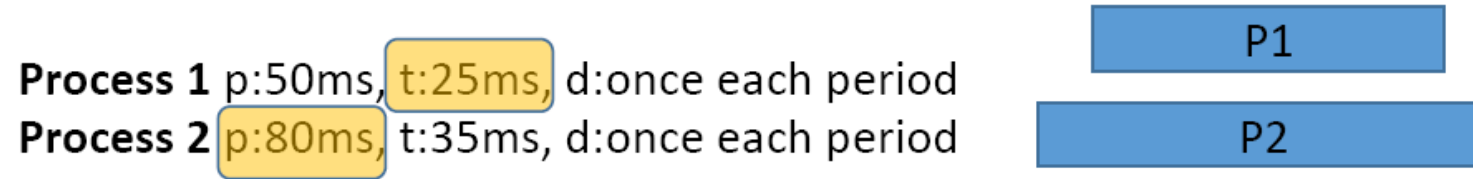
REAL-TIME SCHEDULING



- This works but this requires dynamic priority.
- First period: priority was to P1, second period it was to P2.

Need a new algorithm with dynamic priority.


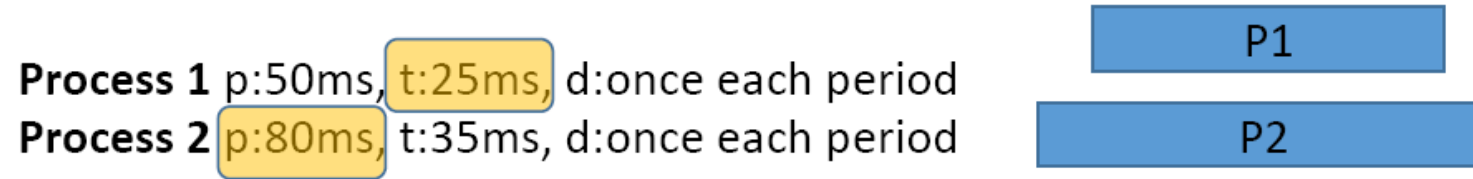
REAL-TIME SCHEDULING: EARLIEST DEADLINE FIRST



Earliest Deadline First Algorithm:

At $t=0$, earliest deadline is P1's so priority goes to P1


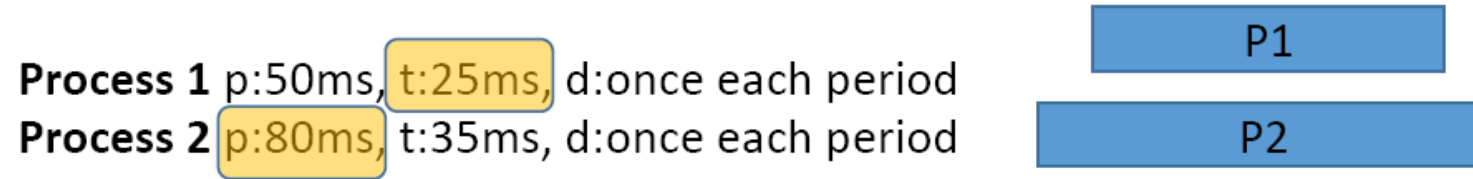
REAL-TIME SCHEDULING: EARLIEST DEADLINE FIRST



At t=50:

- Deadline For P1: 100ms
- Deadline for P2: 80ms
- Earliest Deadline First: P2 keeps executing


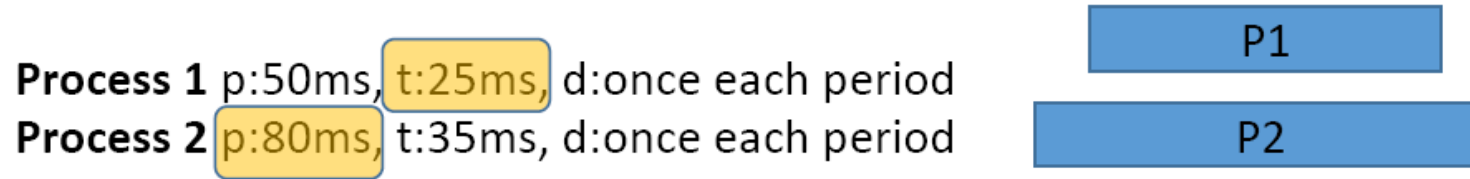
REAL-TIME SCHEDULING



At t=50:

- Deadline For P1: 100ms
- Deadline for P2: 80ms
- Earliest Deadline First: P2 keeps executing

REAL-TIME SCHEDULING




At t=50:

- Deadline For P1: 100ms
- Deadline for P2: 80ms
- Earliest Deadline First: P2 keeps executing
- P1 follows


MULTITHREADED MULTICORE SYSTEM

- CPU scheduling more complex when multiple CPUs are available




MULTITHREADED MULTICORE SYSTEM

- **Homogeneous processors** within a multiprocessor: all processors/cores are the same.
- **Heterogeneous processors** within a multiprocessor: more than one core/processor model.



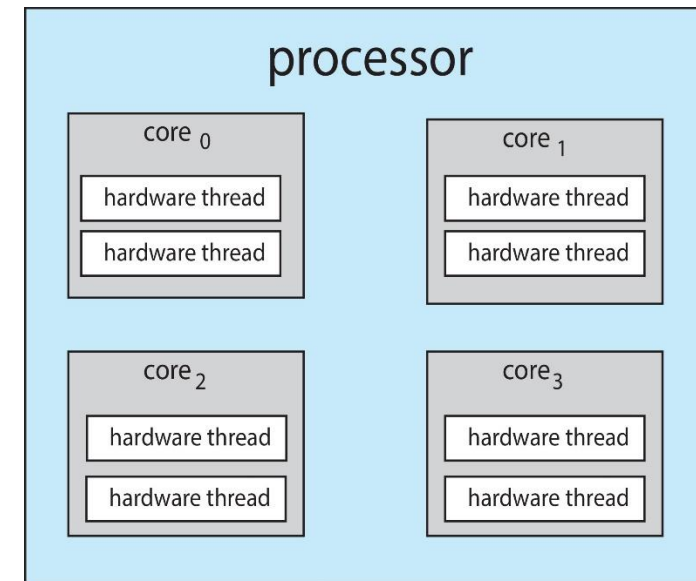
MULTITHREADED MULTICORE SYSTEM

- **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
 - Currently, most common




PROCESSOR AFFINITY

- **Processor affinity** – process has affinity for processor on which it is currently running
 - **soft affinity**
 - **hard affinity**
- **Soft Affinity**– process will resist switching to another core. CPU has to be highly overloaded for it to migrate.
- **Hard affinity**– process will never switch to another core.



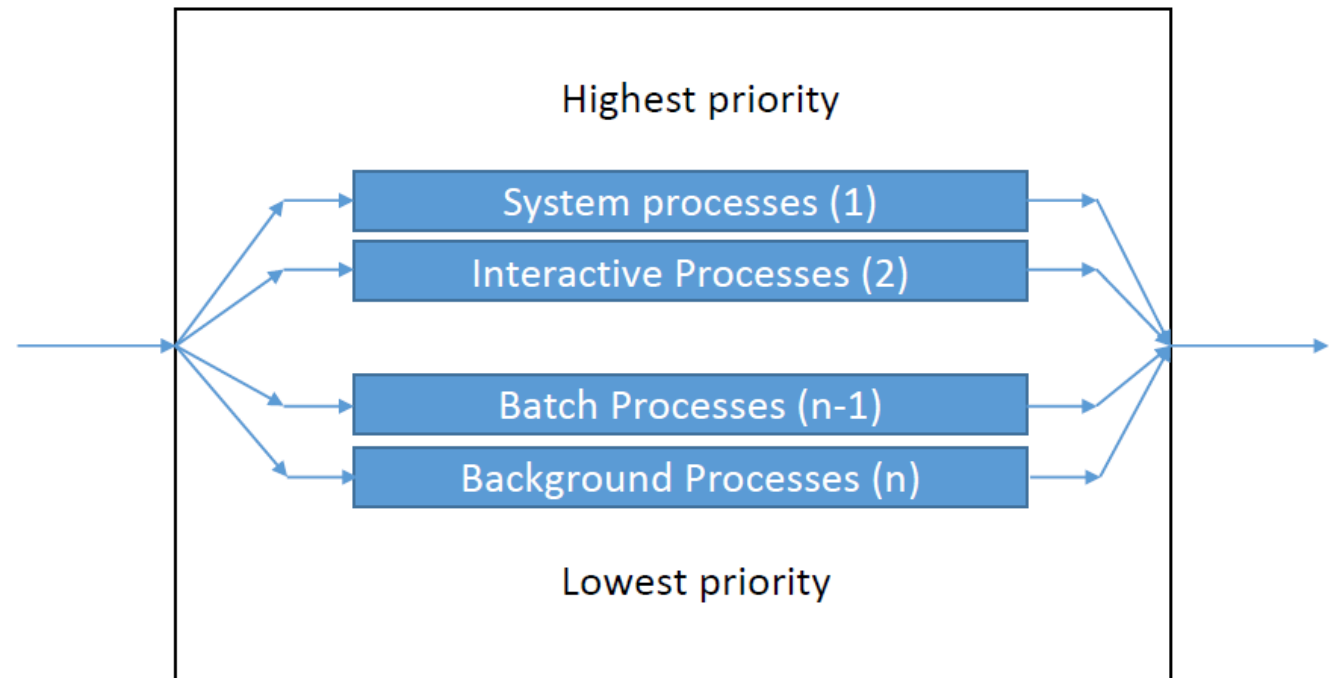
LOAD BALANCING

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
- **Pull migration** – idle processors pulls waiting task from busy processor



WINDOWS SCHEDULING

- Windows uses priority-based preemptive scheduling
- Highest-priority thread runs next
- Thread runs until (1) blocks, (2) uses time slice, (3) preempted by higher-priority thread
- Real-time threads can preempt non-real-time
- 32-level priority scheme
- **Variable class** is 1-15, **real-time class** is 16-31
- Queue for each priority
- If no run-able thread, runs **idle thread**



LINUX SCHEDULER

Completely Fair Scheduler (CFS)

- **Scheduling classes**
 - Each has specific priority
 - Scheduler picks highest priority task in highest scheduling class



LINUX SCHEDULER

Completely Fair Scheduler (CFS)

■ **Scheduling classes**

- Each has specific priority
- Scheduler picks highest priority task in highest scheduling class
- Rather than quantum based on fixed time allotments, based on proportion of CPU time
- Quantum is longer when the processor have less load.

LINUX SCHEDULER

Completely Fair Scheduler (CFS)

- Quantum calculated based on **nice value** from -20 to +19
 - Lower value is higher priority
- CFS scheduler maintains per task **virtual run time** in variable **vruntime**
 - Associated with decay factor based on priority of task – lower nice value is higher decay rate
 - Normal default priority yields virtual run time = actual run time

LINUX SCHEDULER

Completely Fair Scheduler (CFS)

- Quantum calculated based on **nice value** from -20 to +19
 - Lower value is higher priority
- CFS scheduler maintains per task **virtual run time** in variable **vruntime**
 - Associated with decay factor based on priority of task – lower nice value is higher decay rate
 - Normal default priority yields virtual run time = actual run time
- Virtual runtime does not affect quantum duration. It simply keeps track of how often a task was scheduled.
- To decide next task to run, scheduler picks task with lowest virtual run time

LINUX SCHEDULER

Completely Fair Scheduler (CFS)

- Quantum calculated based on **nice value** from -20 to +19
 - Lower value is higher priority
- CFS scheduler maintains per task **virtual run time** in variable **vruntime**
 - Associated with decay factor based on priority of task – lower nice value is higher decay rate
 - Normal default priority yields virtual run time = actual run time
- Virtual runtime does not affect quantum duration. It simply keeps track of how often a task was scheduled.
- To decide next task to run, scheduler picks task with lowest virtual run time

Aging

DEADLOCKS

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource

DEADLOCKS

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes

DEADLOCKS

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task

DEADLOCKS

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

FINAL EXAM


- Comprehensive
- 2 Hours
- Tuesday, March 14th from 1:00 to 3:00 pm on Canvas (online exam)
- Two pages of notes (single sheet).

FINAL EXAM


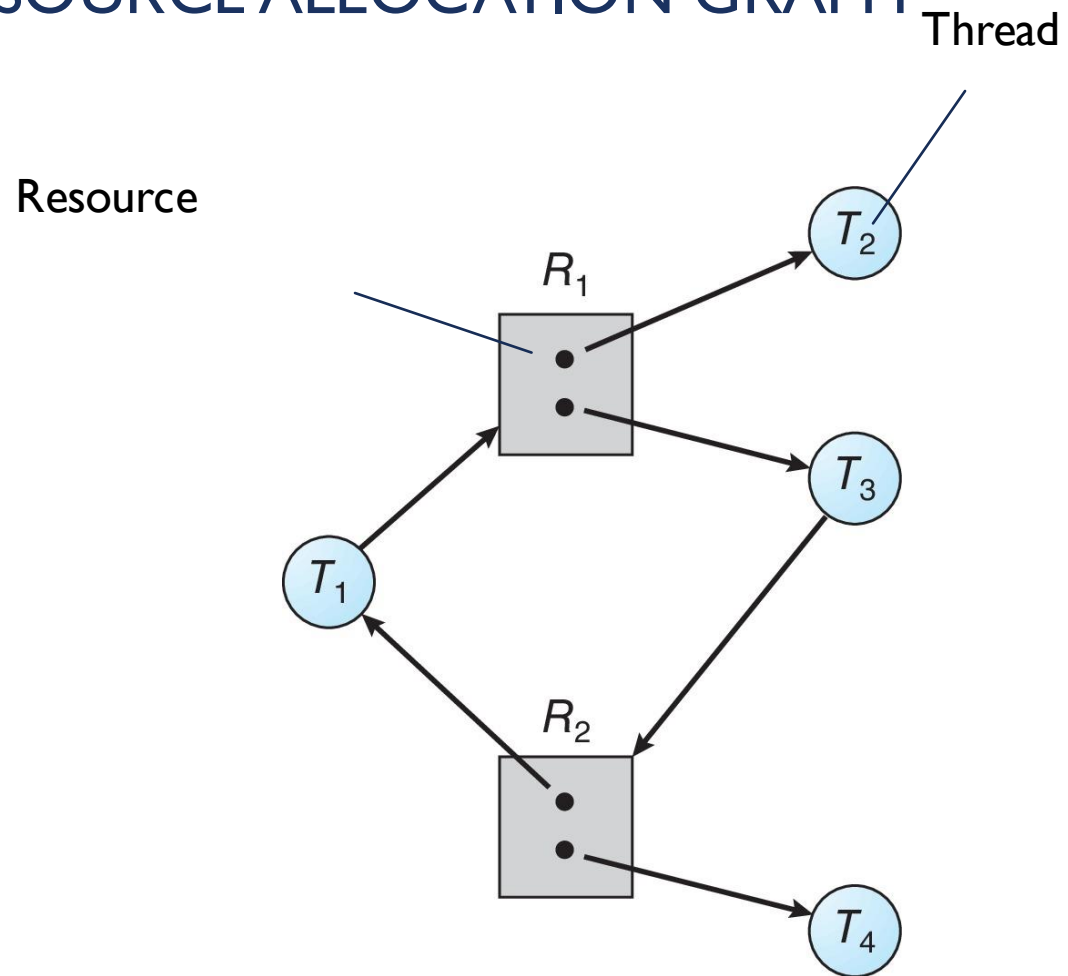
- Multiple Choice
- True/False
- Essay questions.
- Questions will include a wide variety of topics.
- Questions guaranteed to be there:
 - FS and their implementation
 - Segmentation and TLB/Paging question: similar to worksheet
 - Scheduling (be ready for real-time or non-real time scheduling question)
- Review quizzes and worksheets: some questions will be very similar.
- This is not a comprehensive list of questions!

DEADLOCKS IN MULTITHREADED APPLICATION



- Mutual Exclusion can cause deadlocks when not well planned.
- Can be illustrated with a **resource allocation graph**:




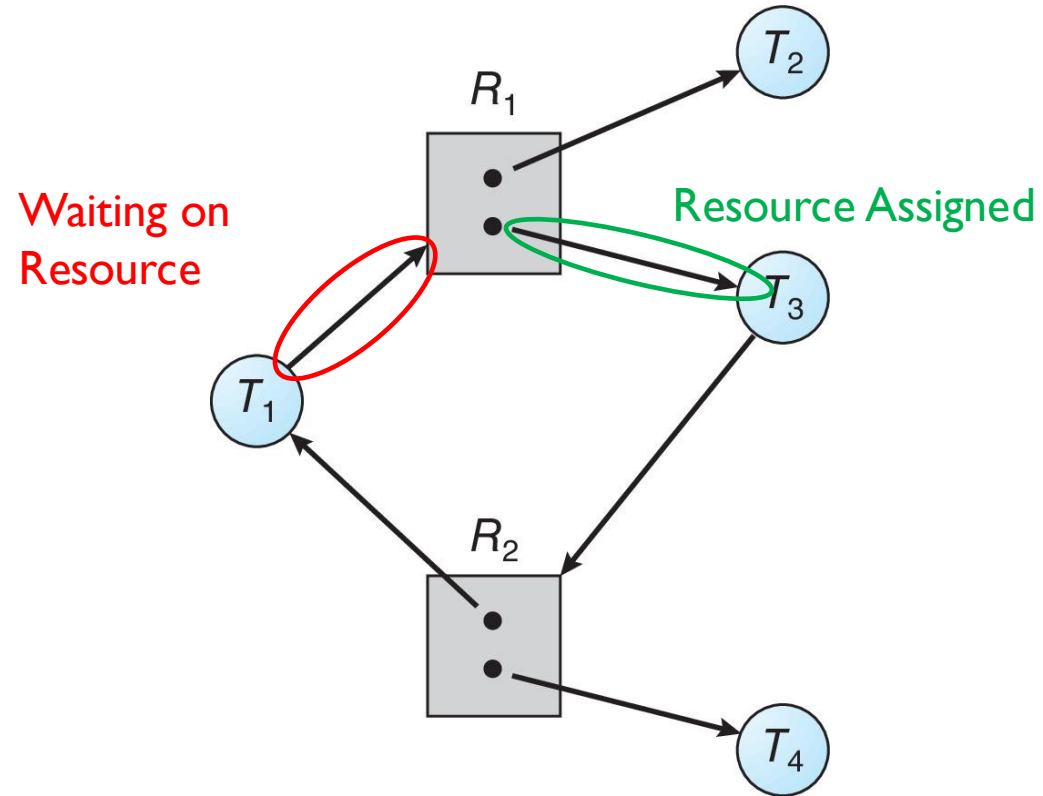
RESOURCE ALLOCATION GRAPH




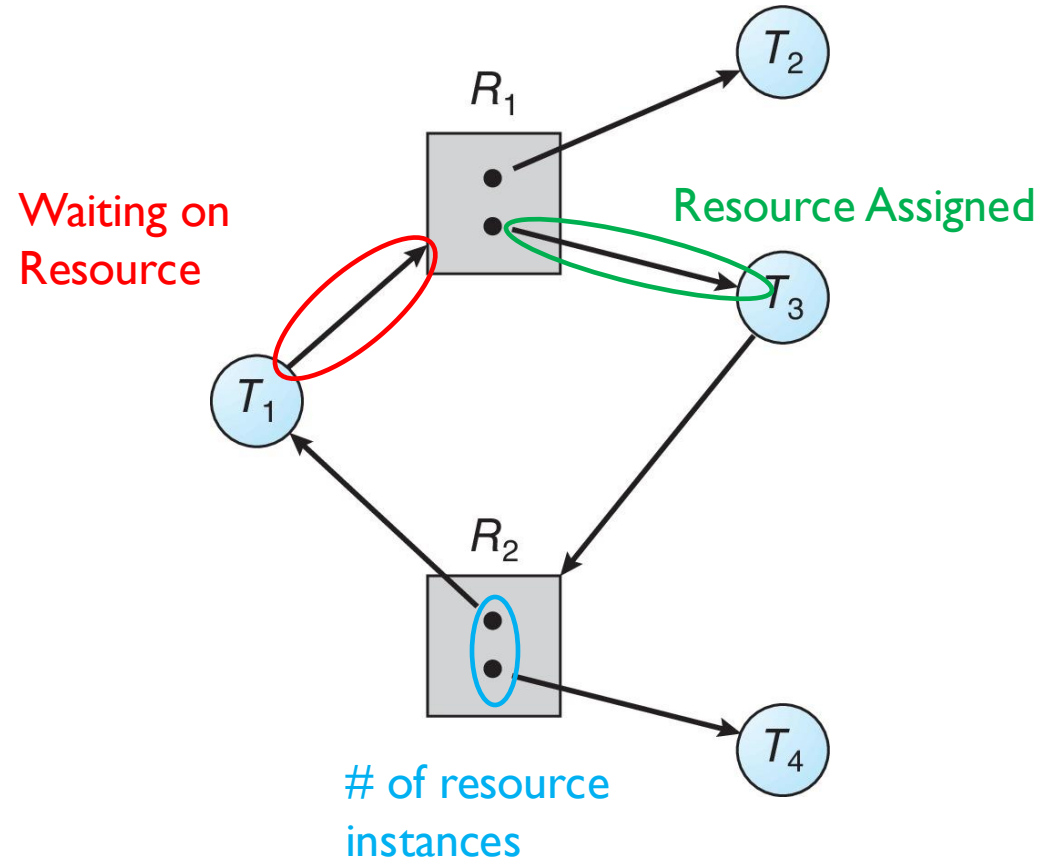
RESOURCE ALLOCATION GRAPH





RESOURCE ALLOCATION GRAPH





RESOURCE ALLOCATION GRAPH





RESOURCE ALLOCATION GRAPH: CLOSED LOOP





RESOURCE ALLOCATION GRAPH: CLOSED LOOP




RESOURCE ALLOCATION GRAPH: CLOSED LOOP




RESOURCE ALLOCATION GRAPH: CLOSED LOOP



RESOURCE ALLOCATION GRAPH: CLOSED LOOP



A



B

RESOURCE ALLOCATION GRAPH: CLOSED I

A

B

C


D

A: Deadlock only at A


B: Deadlock only at B

C: Both Deadlocked

D: None




A




B

RESOURCE ALLOCATION GRAPH: CLOSED LOOP




A: No Deadlock




B: Deadlocked

RESOURCE ALLOCATION GRAPH: CLOSED LOOP



Some resources are assigned to threads outside the loop.

A: No Deadlock



B: Deadlocked

All resources instances are part of the loop threads.

HANDLING DEADLOCKS

Three approaches for deadlocks:

- Prevent/Avoid Deadlocks
- Detect and Recover
- Ignore Deadlocks

HANDLING DEADLOCKS

Three approaches for deadlocks:

- Prevent/Avoid Deadlocks
 - Detect and Recover
 - Ignore Deadlocks
-
- Operating Systems ignore deadlocks!

HANDLING DEADLOCKS

Three approaches for deadlocks:

- Prevent/Avoid Deadlocks
 - Detect and Recover
 - Ignore Deadlocks
-
- Operating Systems ignore deadlocks!
 - As a programmer, you have to make sure to use mutual exclusion properly.